



# Computer Systems for AI-inspired Cloud: Theory & Lab.

Fall 2018

Dr. JongWon Kim

[\(jongwon@gist.ac.kr\)](mailto:jongwon@gist.ac.kr)

Networked Computing Systems Lab.,  
School of Electrical Engineering and Computer Science, GIST

- 기존 국문제목: 클라우드, AI기반 소프트웨어 중심의 정보통신
- 수정 영문제목:

Computer Systems for AI-inspired Cloud: Theory & Lab.

- 수정 영문제목:

AI 기반 클라우드를 위한 컴퓨터 시스템: 이론 및 실습

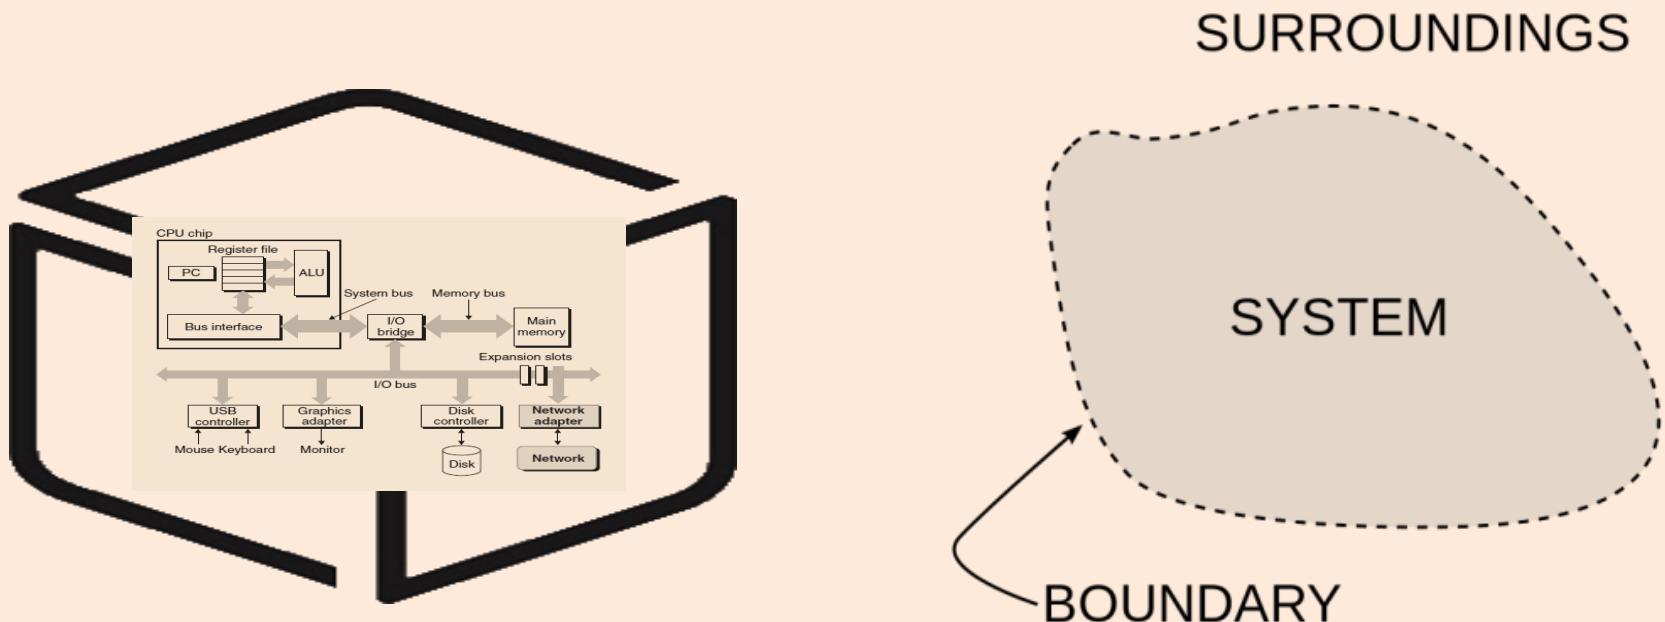
# Contents: Theory

- Chapter 1. What is Computer System?
- Chapter 2. Computer System & OS Software
- Chapter 3. Computer System: Virtualization & Containers
- Chapter 4. Computer System: Cloud-native Computing
- Chapter 5. Computer System: Clusters for HPC/BigData & AI
- Chapter 6. Computer Systems for Cloud Data Center and SDI
- Chapter 7. Computer System & SmartX Abstraction

# Contents: Labs

- Lab #0: Overview Lab
- Lab #1: Box Lab
- Lab #2: Inter-Connect Lab
- Lab #3: Tower Lab
- Lab #4: IoT Lab
- Lab #5: Cluster Lab
- Lab #6: Analytics Lab

# Chapter 1: What is Computer System?



BOUNDARY

SURROUNDINGS

SYSTEM

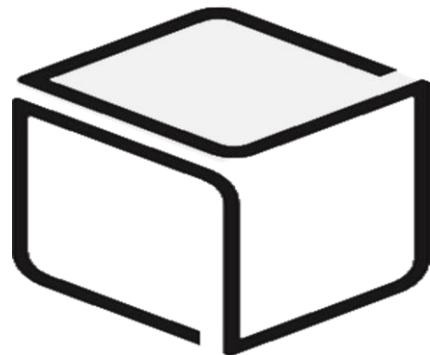
# Chapter Keywords

Storage

Compute

Box

Networking

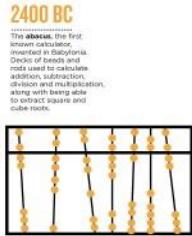
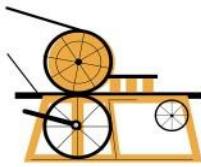


Check <https://www.livescience.com/20718-computer-history.html>

# THE EVOLUTION OF COMPUTERS

PC NINJA TRAVELS THROUGH TIME, revealing the history of how computers became our sidekicks. From sliding pebbles on a simple machine to swiping your fingers across a touchscreen, technology has transformed radically!

## HISTORIC



1439  
Johannes Gutenberg invents the printing press

1642  
Blaise Pascal invents the *Pascaline* (Pascal's calculator), a mechanical adding machine



KEY

PC HISTORY

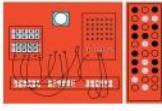
MAC HISTORY

INTERNET HISTORY

OTHER DEVICES HISTORY

## MODERN

1943-44



The ENIAC (Electronic Numerical Integrator and Computer) was the first true digital computer. It used vacuum tubes, switches, and relays through a plugboard and switches to solve a full range of scientific computing problems. ENIAC occupied 170 square feet, 30 by 50 feet in size and weighed 30 tons!

1949



First e-mails sent by Ray Tomlinson on the Arpanet, the precursor to the Internet, which were then stored on different hosts. He used the @ sign to separate people from machines, which became the first part of the e-mail address.

1952



The first mass-produced computer terminal and host by the IBM Company.

1955



Charles Babbage designs the Difference Engine, a mechanical calculating machine that is considered one of the first programmable computers, though the technology was not advanced to be built in his time.

1950



Mauchly and Eckert receive funding from the Census Bureau to build the UNIVAC (Universal Automatic Computer), the first commercially successful computer for business and government applications.

1952



The first computer game, "Tic-Tac-Toe," designed and created by A.S. Douglas.

1955



Standards under a team of mathematicians. As the first fully functional stored-program electronic computer, it could keep programmed instructions close to technology more accessible to the general public.

1964



Computer systems that use icons, windows, and a mouse developed by Douglas Engelbart, which will later become the personal computer with a GUI that features a drop-down menu with icons that can be clicked using a little box with button called a "mouse."

1971



The Apple II, the final personal computer with color graphics is determined.

1973



Word processing becomes a reality as Microsoft's first Microsoft releases WordStar.

1979



The "Personal Computer" released by the IBM Corporation, under the name PC. It uses Microsoft's MS-DOS operating system, has an Intel chip, two floppy disks and an optional color monitor.

## MICRO SOFT

1977



The Apple II, the first personal computer with color graphics is determined.

1979



Word processing becomes a reality as Microsoft's first Microsoft releases WordStar.

1981



The "Personal Computer" released by the IBM Corporation, under the name PC.

1985



Bill Gates and Microsoft release the first graphical desktop operating system, *Microsoft Windows*, desktop computing.

1989



The "World Wide Web" invented by Tim Berners-Lee, uses HyperText Markup Language (HTML) to design websites.

1990



Release of the first Microsoft Office 1.0, featuring Word 1.1, Excel 2.0, and Powerpoint 2.0.

1995



Amazon Bookstore launches, and online shopping becomes popular.

1999



The term Wi-Fi becomes part of everyday language and users begin connecting to the internet without wires.

2001



Microsoft unveils the Mac OS X operating system, featuring protected memory security, a new pre-emptive multitasking engine, and support for 3D graphics, making Microsoft roll out Windows XP, which has been largely redesigned GUI.

2007



Apple introduces the iPhone, both a phone and handheld computer that runs "iOS."

2008



The first 3D printer created by Chuck Hull at 3D Systems Corp.

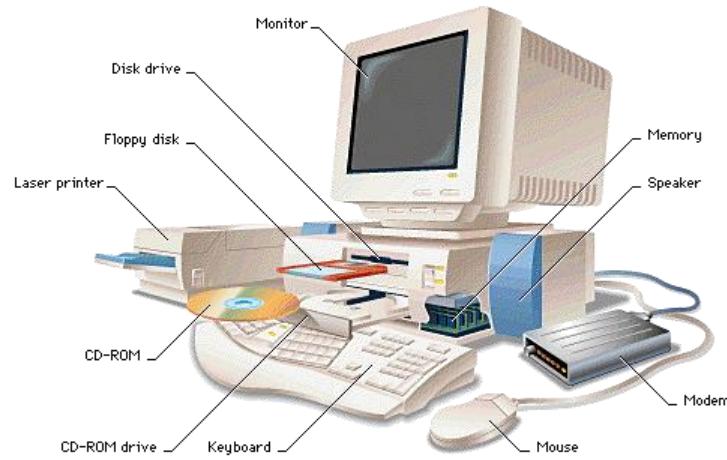
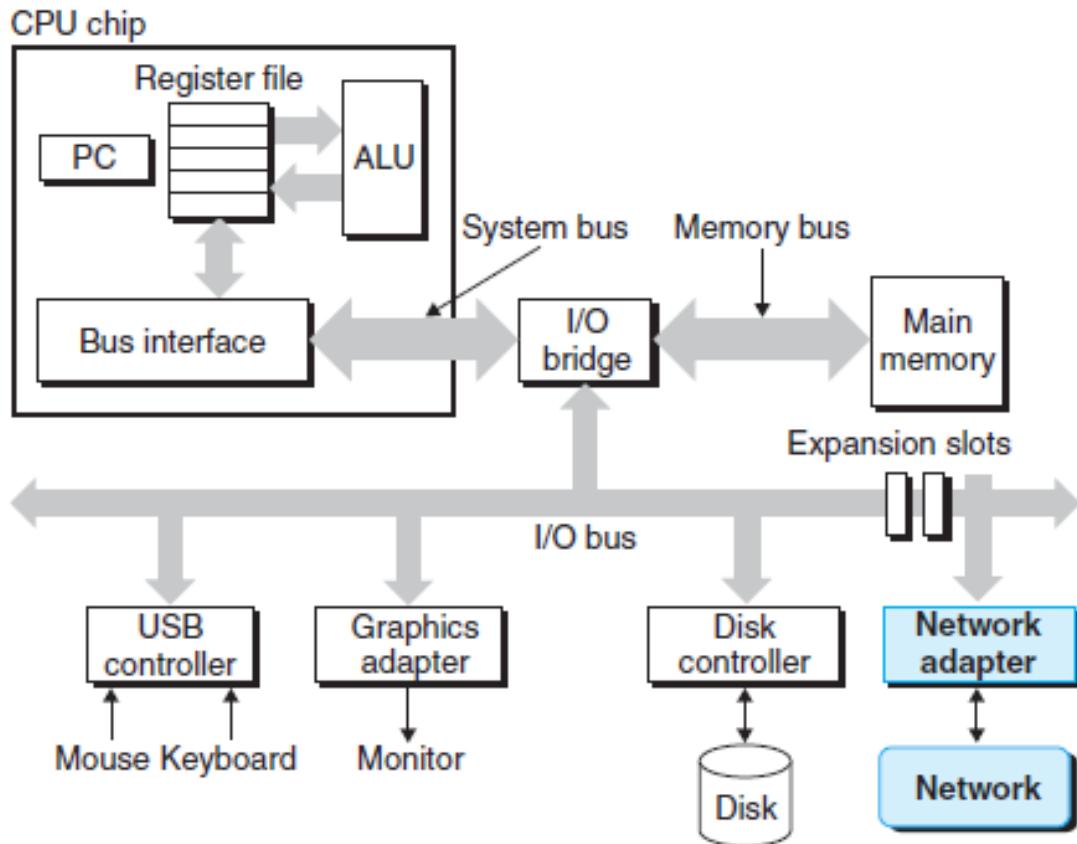
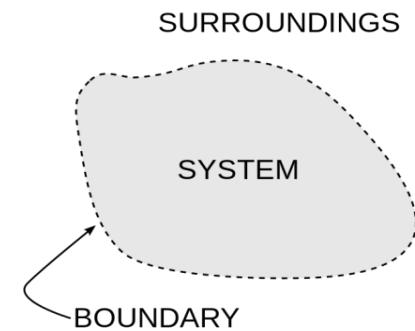
2013



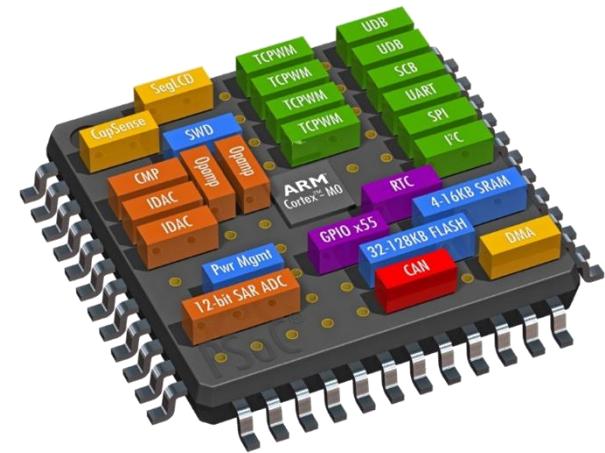
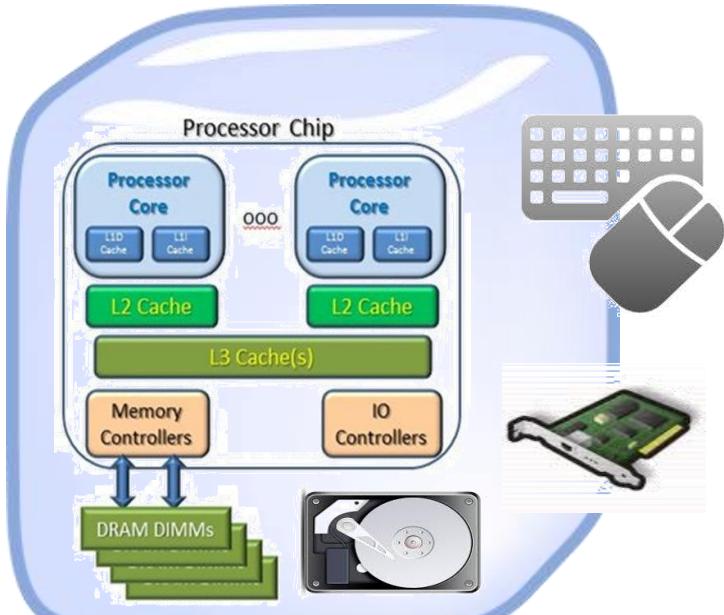
Microsoft releases Windows 8, a completely revamped operating system designed toward touchscreens, tablets, and cloud computing.

# Computer System: Definition

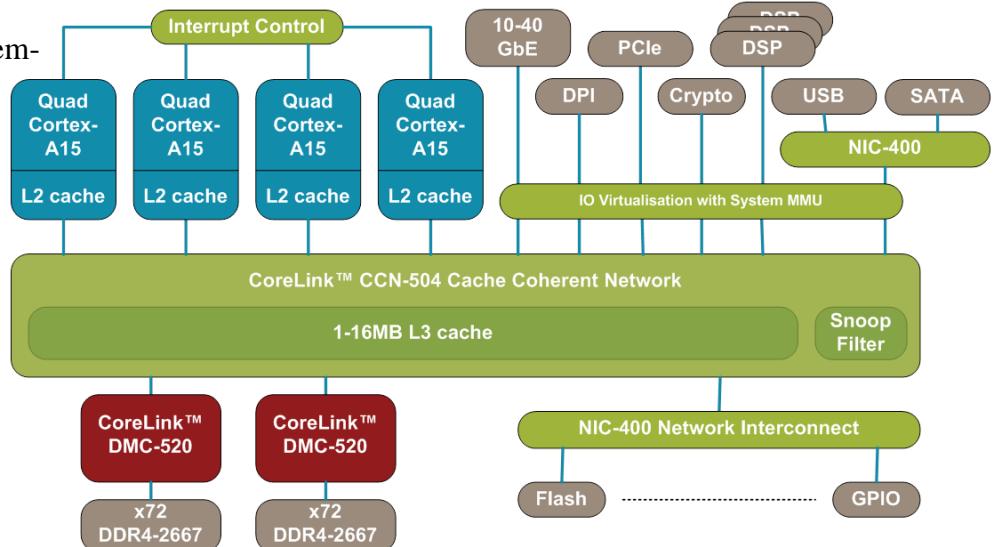
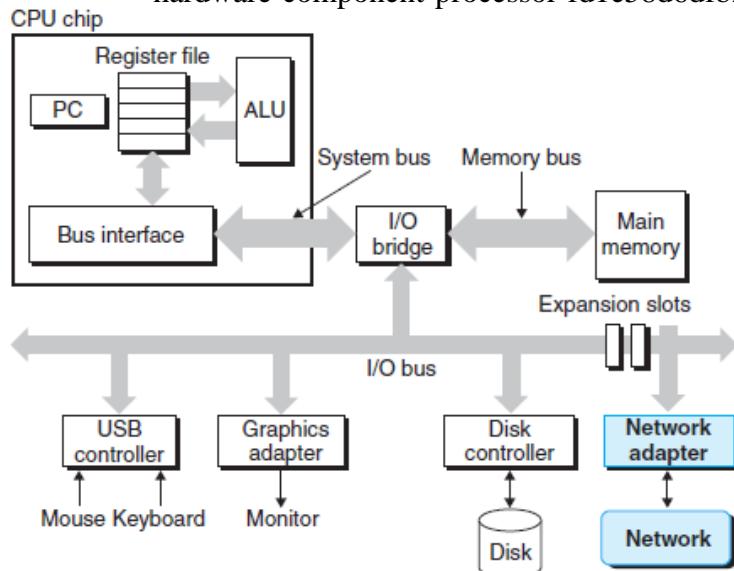
**Computer:** an electronic device for storing and processing data, typically in binary form, according to instructions given to it in a variable program.



# Evolution of Computer System (Chip)



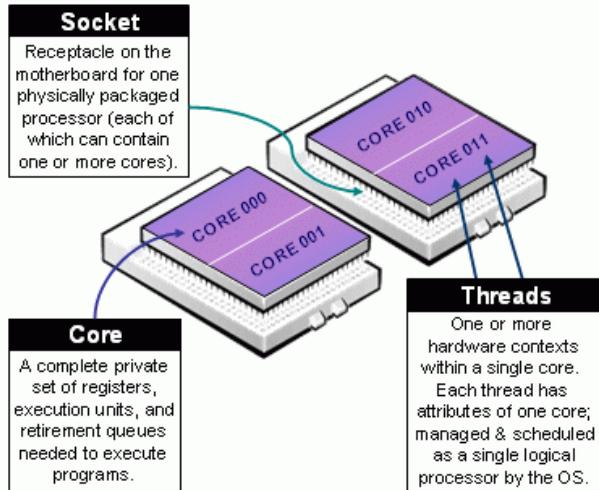
<https://medium.com/computing-technology-with-it-fundamentals/system-hardware-component-processor-fd1e58d6dfb2>



SOC Chip (ARM Cortex + Corelink)

<https://community.arm.com/processors/b/blog/posts/coherent-interconnect-technology-supports-exponential-data-flow-growth>

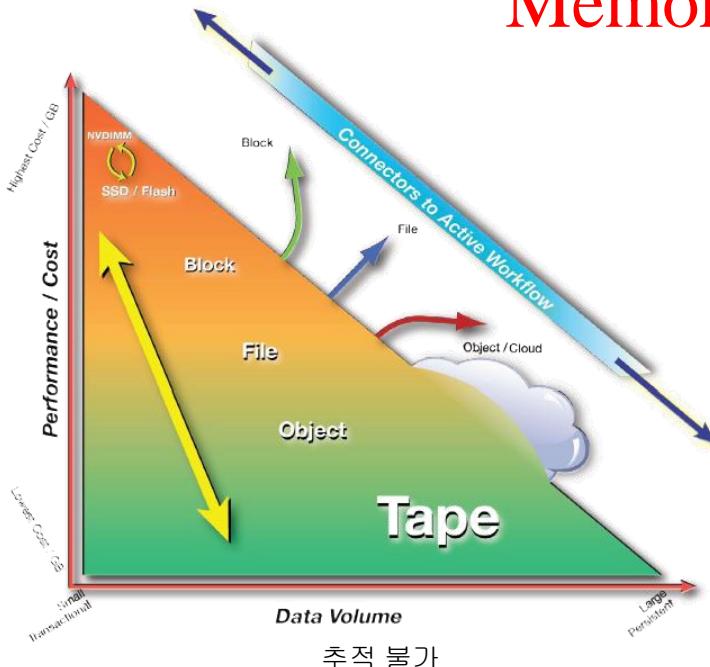
# Major Components of Computer Systems



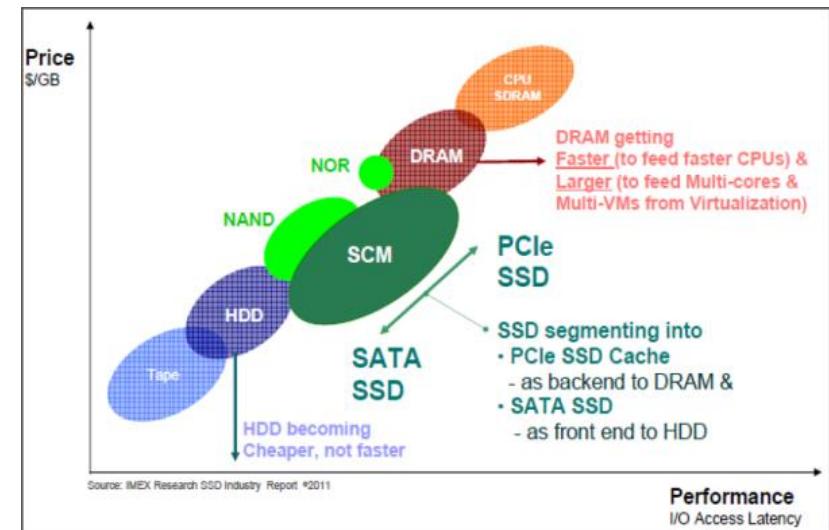
[https://slurm.schedmd.com/mc\\_support.html](https://slurm.schedmd.com/mc_support.html)

- Compute
  - CPU & Memory
- Storage
- Network

## Memory/Storage Hierarchy



주적 불가

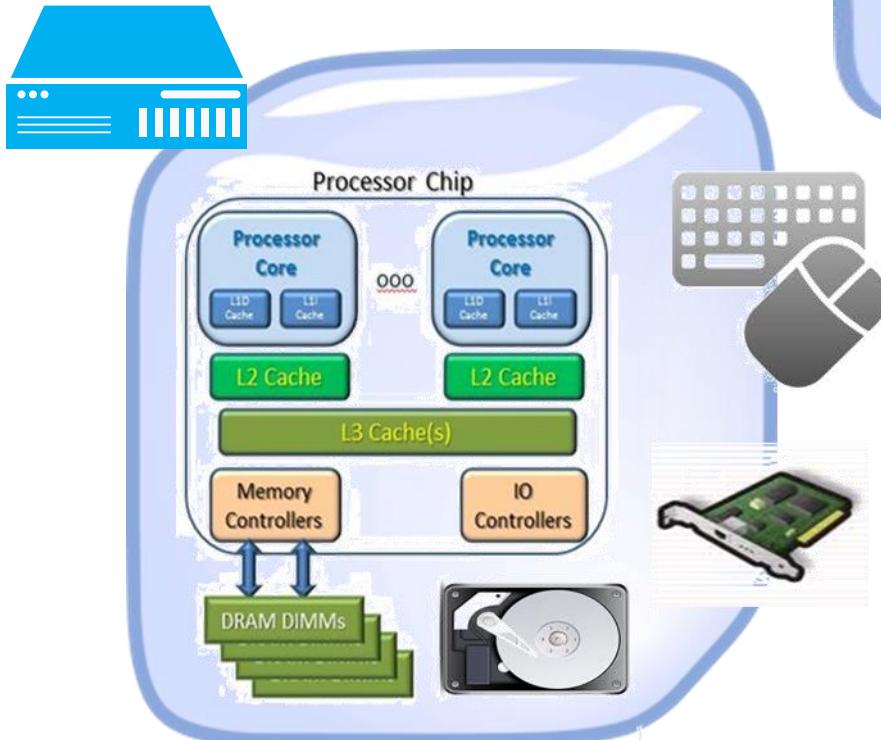


<https://www.semanticscholar.org/paper/Optimizations-of-Management-Algorithms-for-Memory-Oren/2799edf40be793a6c7b86b391270014442191cd7>

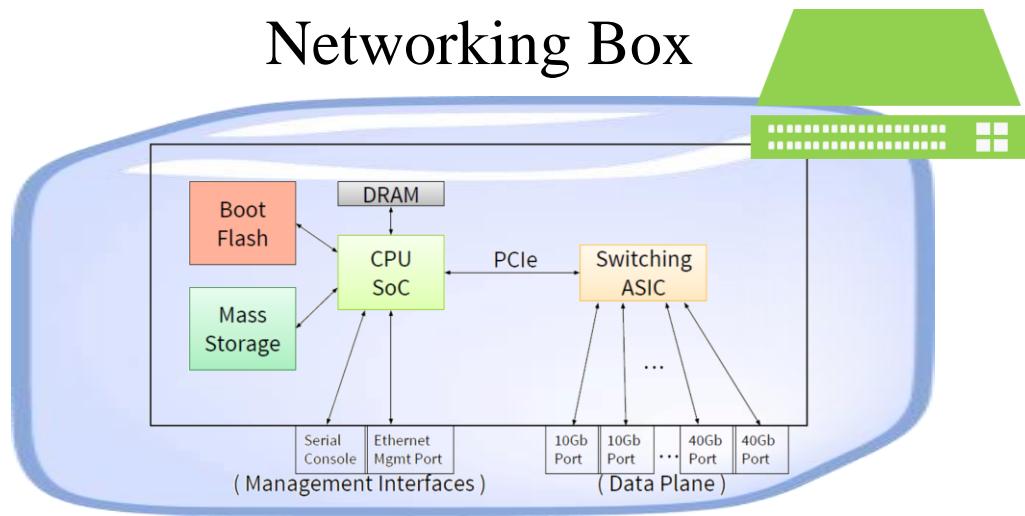
# Diversity of Computer Systems →→ Box



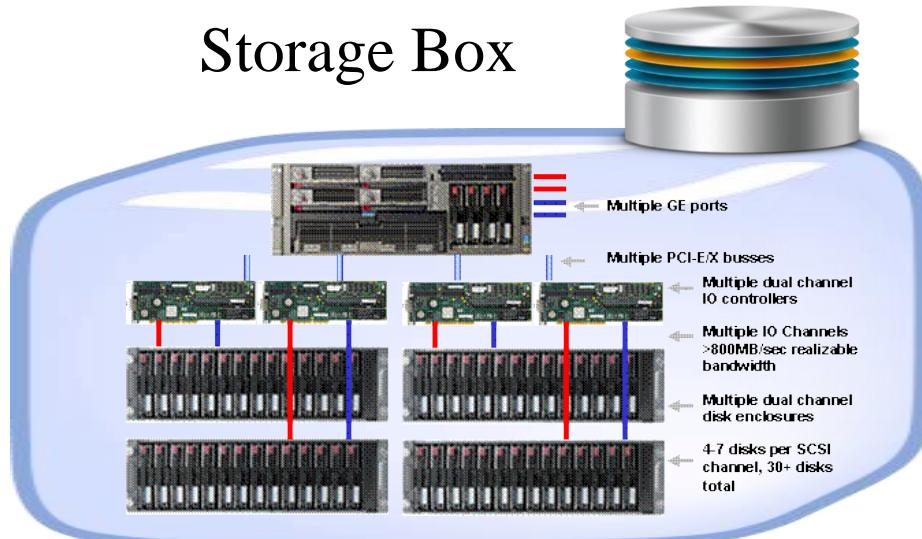
## Compute Box



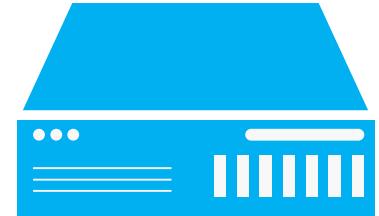
## Networking Box



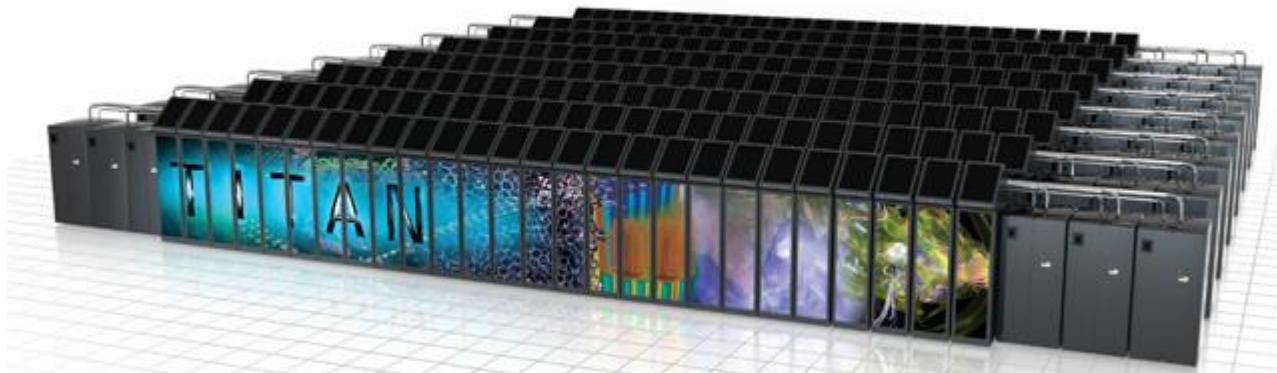
## Storage Box



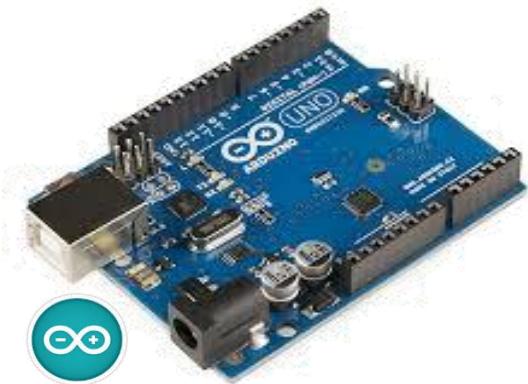
# Compute Box: Computers & Computer (Server) Clusters



<https://www.engadget.com/2014/03/31/intel-sd-card-sized-edison-change/>



<https://www.clubic.com/pro/it-business/calcul-distribue/actualite-601866-supercalculateur-tianhe-2-premier-top500.html>



[https://commons.wikimedia.org/wiki/File:Arduino\\_Ethernet\\_Board.jpg](https://commons.wikimedia.org/wiki/File:Arduino_Ethernet_Board.jpg)

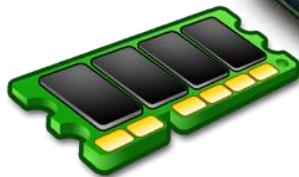
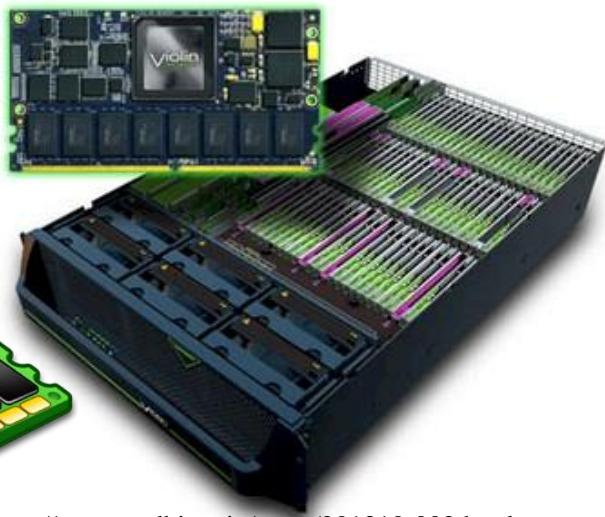


[https://gl.wikipedia.org/wiki/Raspberry\\_Pi](https://gl.wikipedia.org/wiki/Raspberry_Pi)

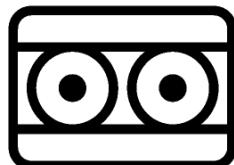


[https://hatewait.ru/product/samsung\\_galaxy\\_a8\\_2018\\_black.aspx](https://hatewait.ru/product/samsung_galaxy_a8_2018_black.aspx)

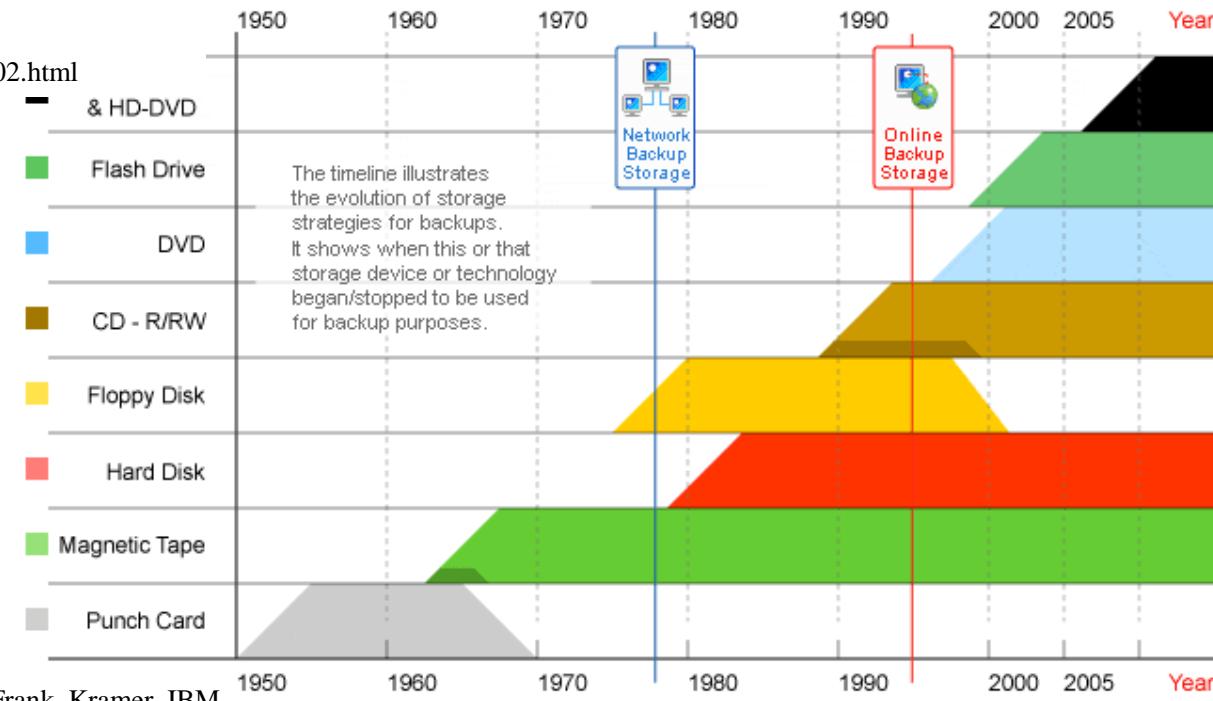
# Storage Box: Storage Servers



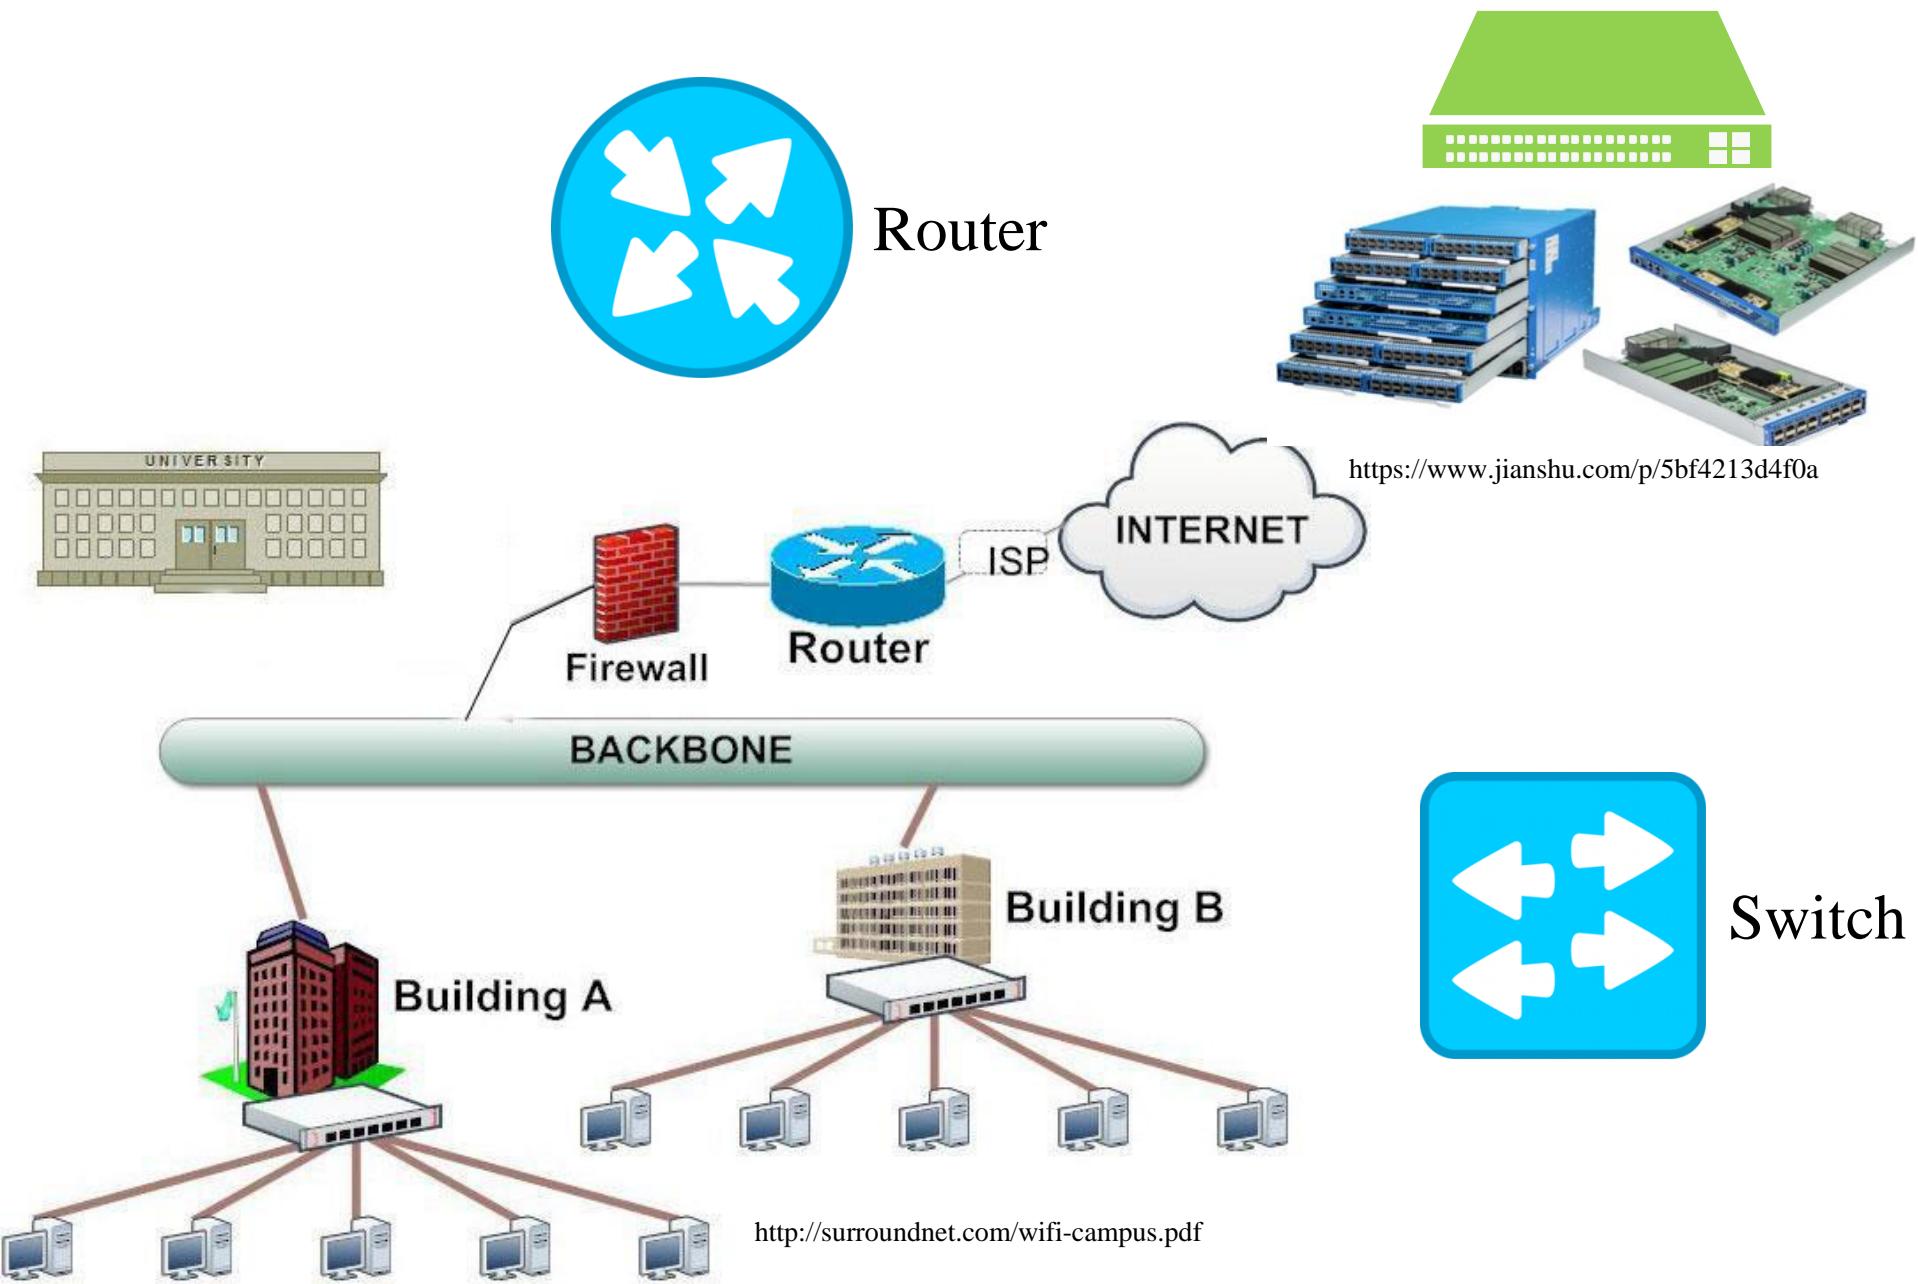
<http://www.nsdbi.co.jp/news/201310-002.html>



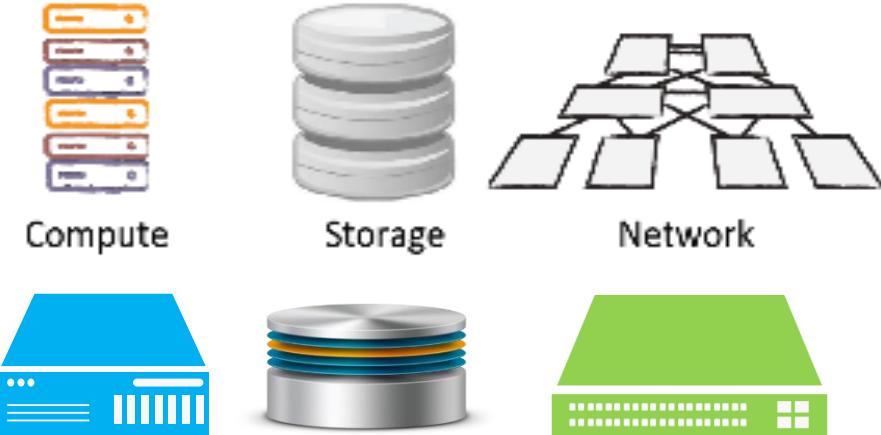
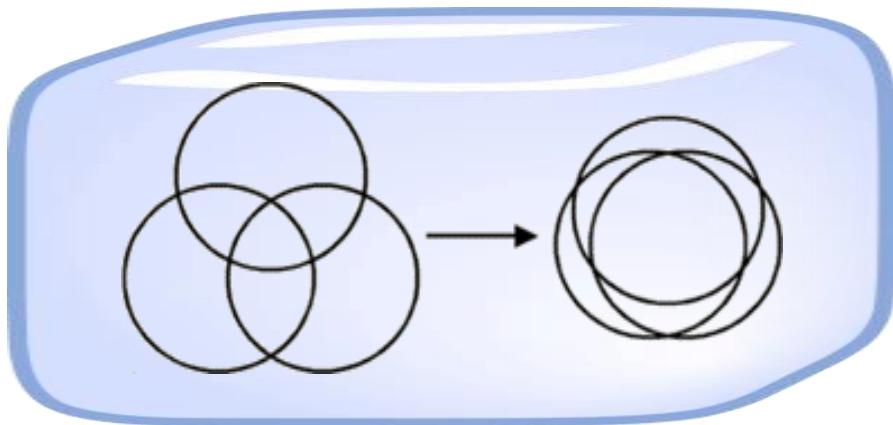
<http://www.bigfixtech.com/>  
**Timeline: Data Backup Storage**



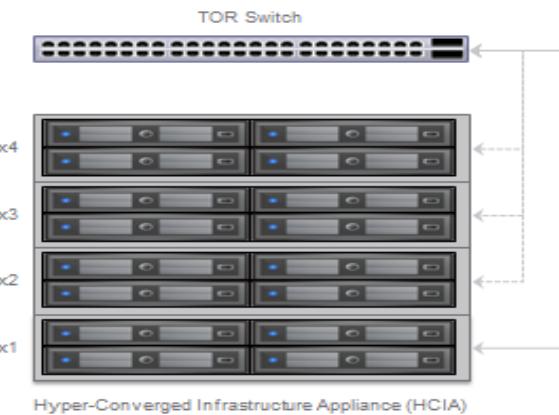
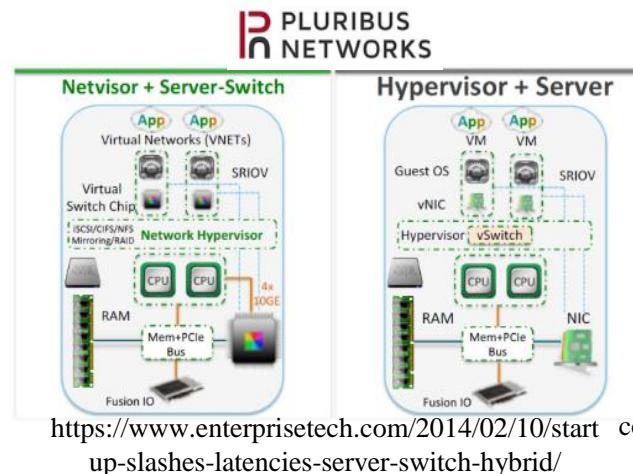
# Networking Box: Switches & Routers



# Futuristic Box: Hyper-Converged Infrastructure

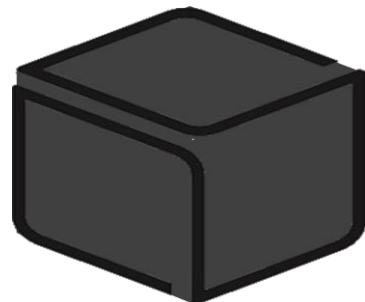
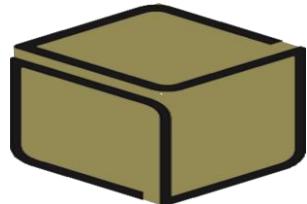


<http://www.bytelife.com/blog/software-defined-and-mekeskus-pariselt/>



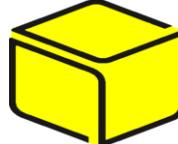
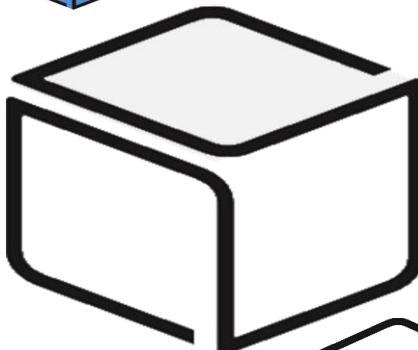
# Multiple Box Abstractions: Black/Brown/White and SmartX Boxes

Brown/Brite Box

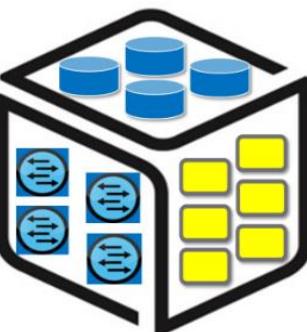
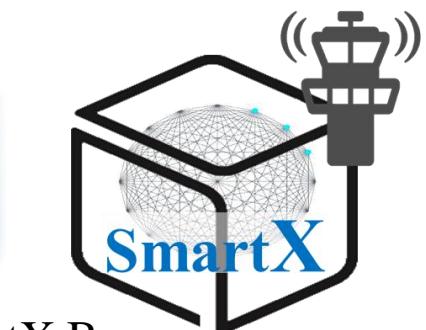


Black Box

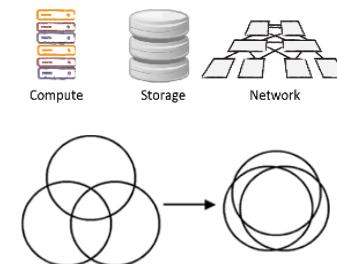
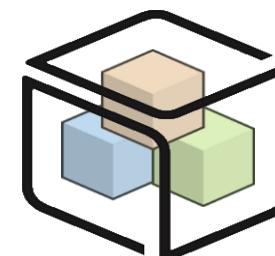
White Box



SmartX Box



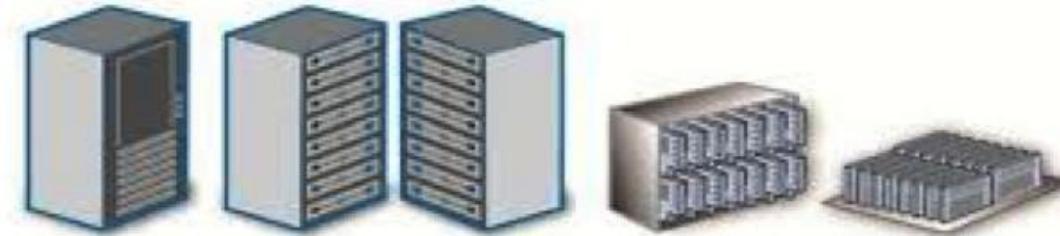
Virtualized Box



Hyper-converged Box

# Converged/Disaggregated Boxes in Data Centers

<https://internetofthingsagenda.techtarget.com/definition/microserver>



Tower Server

Rack Server

Blade Server

Micro Server



2000

2004

2008

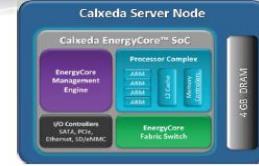


AMD  
SBrick

2010



hp

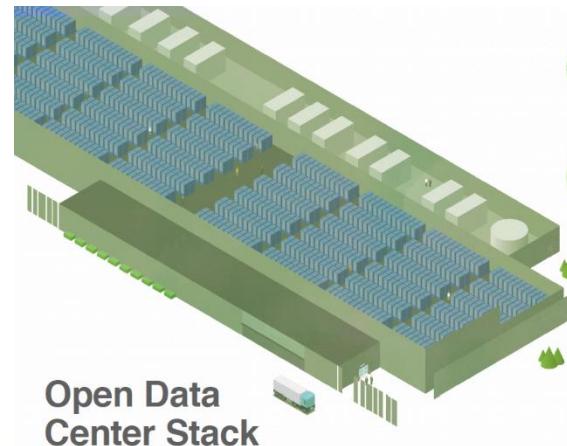


2012

2014



Optical PCIe via  
Intel Silicon Photonics



## Open-Source Hardware

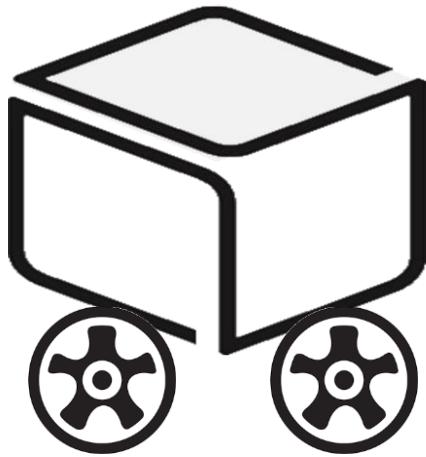
## Rack Scale Computing

Silicon Photonics

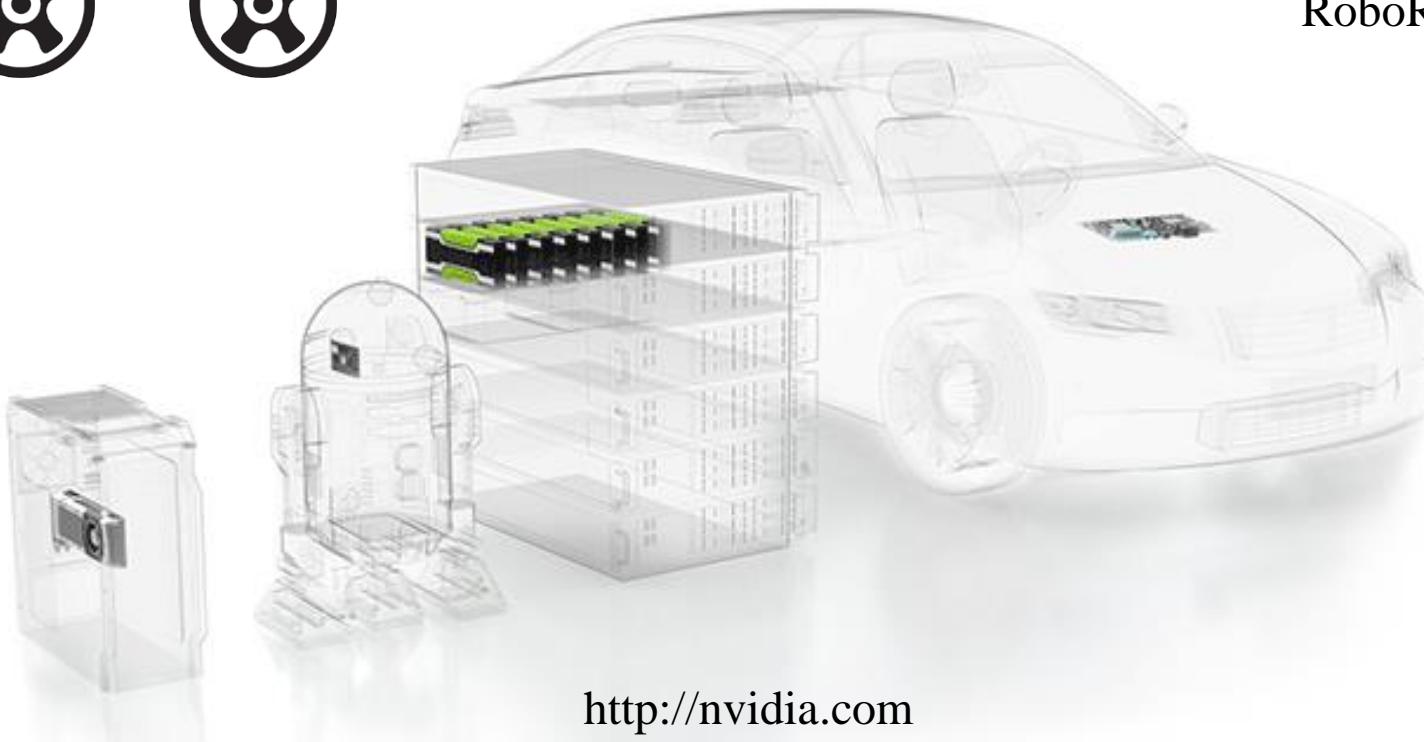
High-bandwidth / low-latency interconnect (resource disaggregation)

# Computer Systems Everywhere

<http://poe.goha.ru/news/99/2016-04-01%2014:00:18>



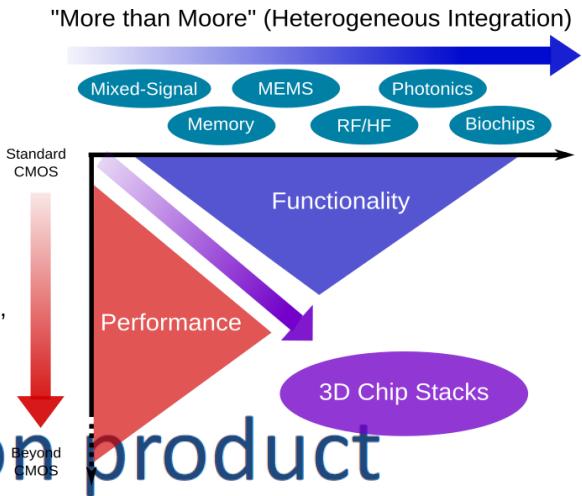
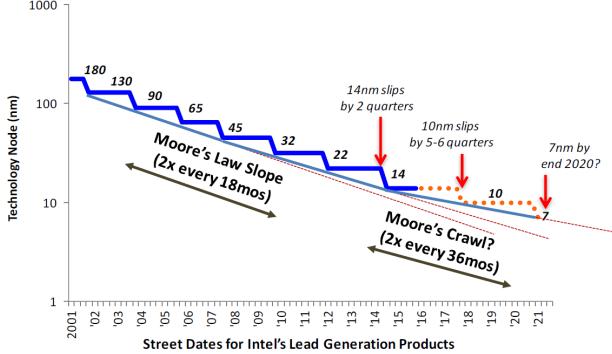
RoboRace



<http://nvidia.com>

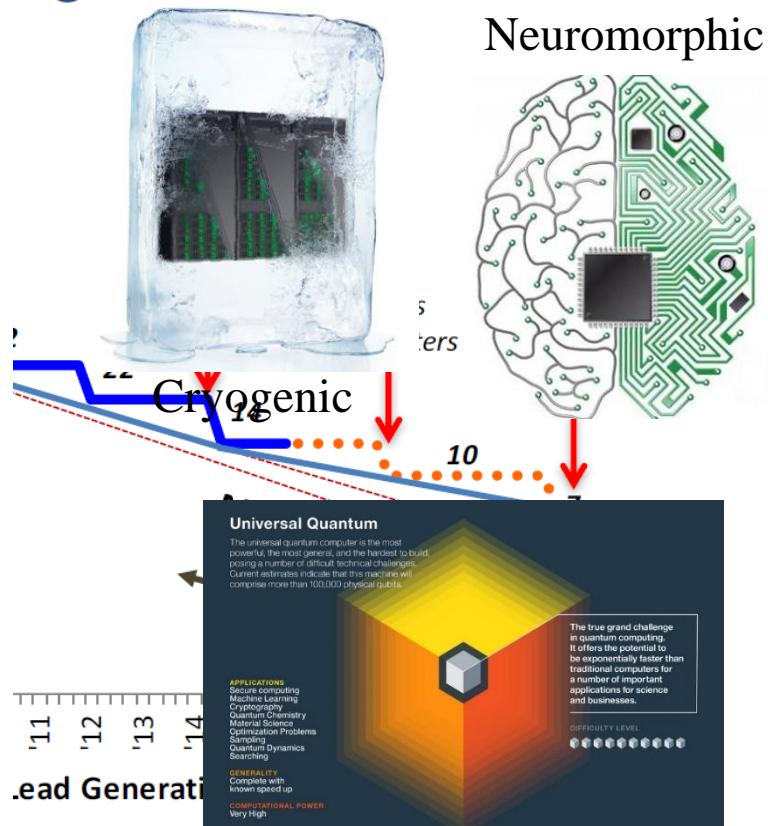
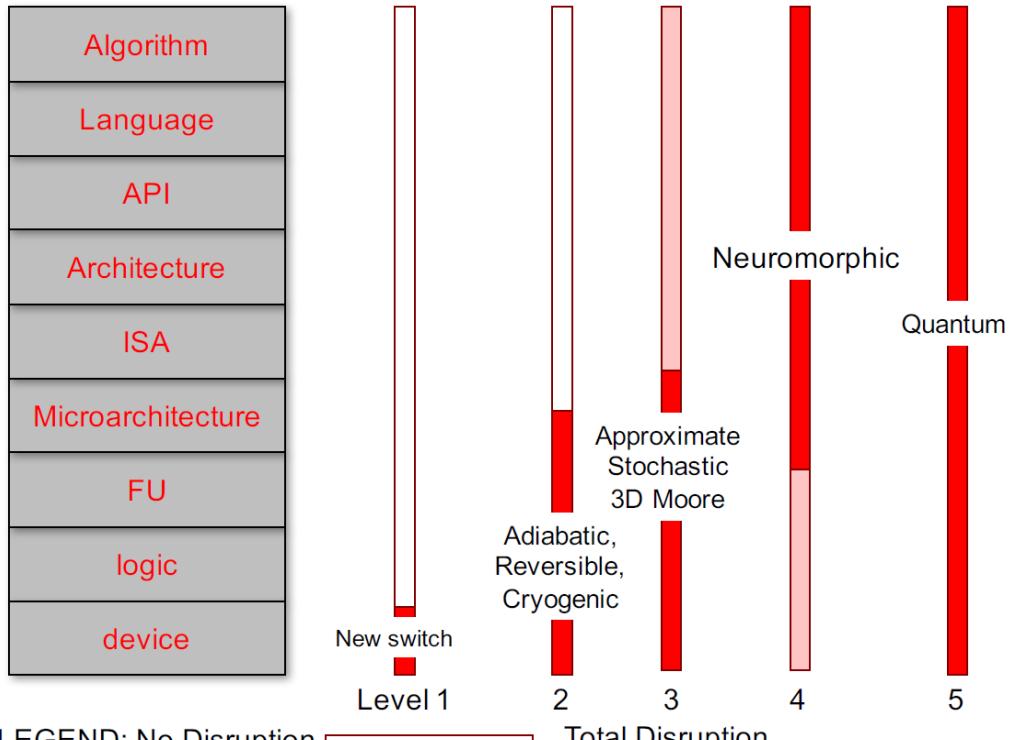
# Future of Compute Boxes

## Moore's Law impact on product release cycles



act on product  
release cycles

## Differing Levels of Disruption in Computing Stack





Gwangju Institute of  
Science & Technology



Thank you!

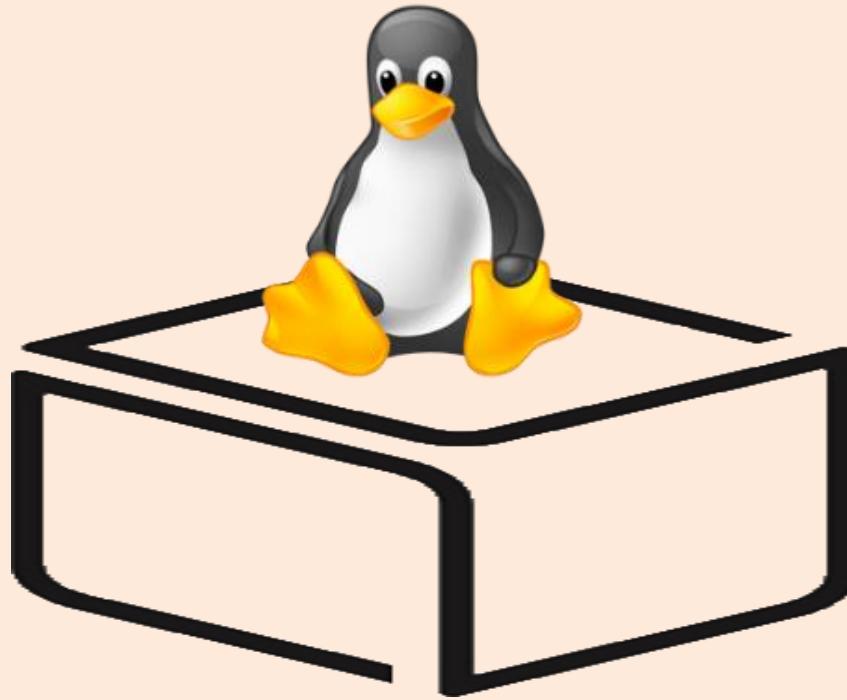
Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>

# Chapter 2:

# Computer System

# & OS Software



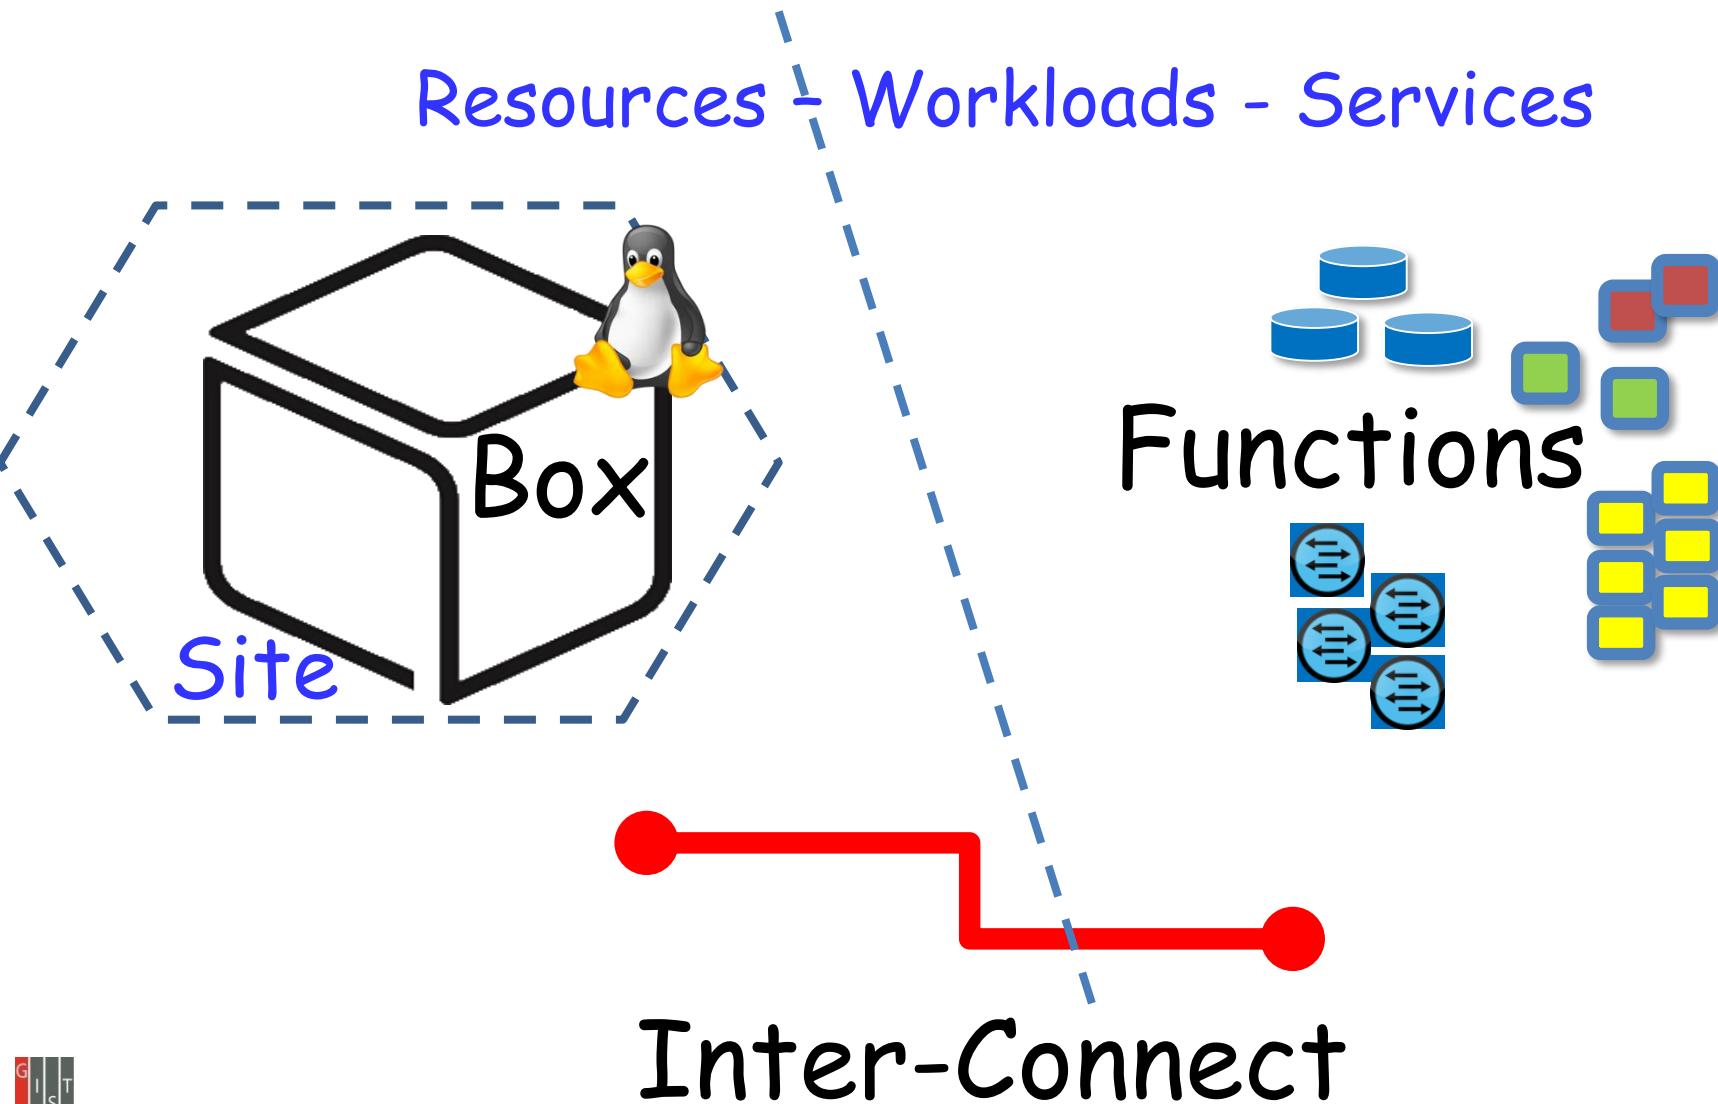
# Chapter Keywords

Operating System (OS)

Kernel

Linux

# Computer System: Inter-Connected Functions inside/across Boxes/Sites



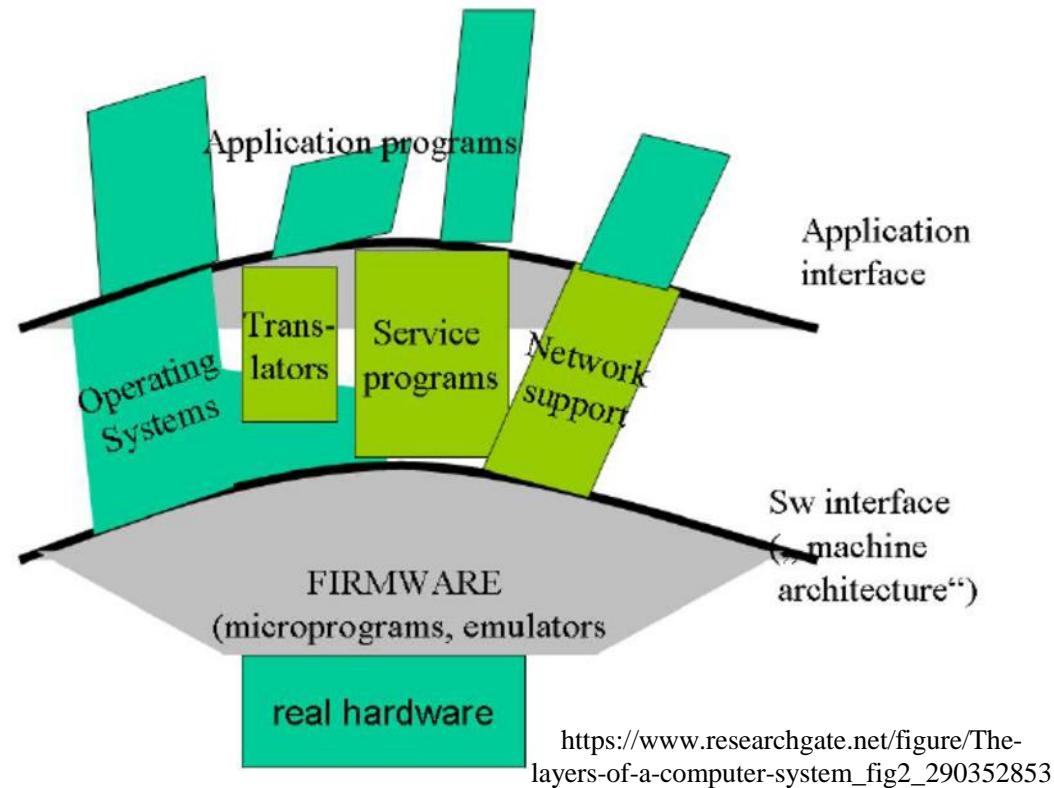
# Computer System: Hardware & Software

## Hardware

VS

## Software

(**Firmware**, **Operating System (OS)**,  
Middleware, Application, ...)

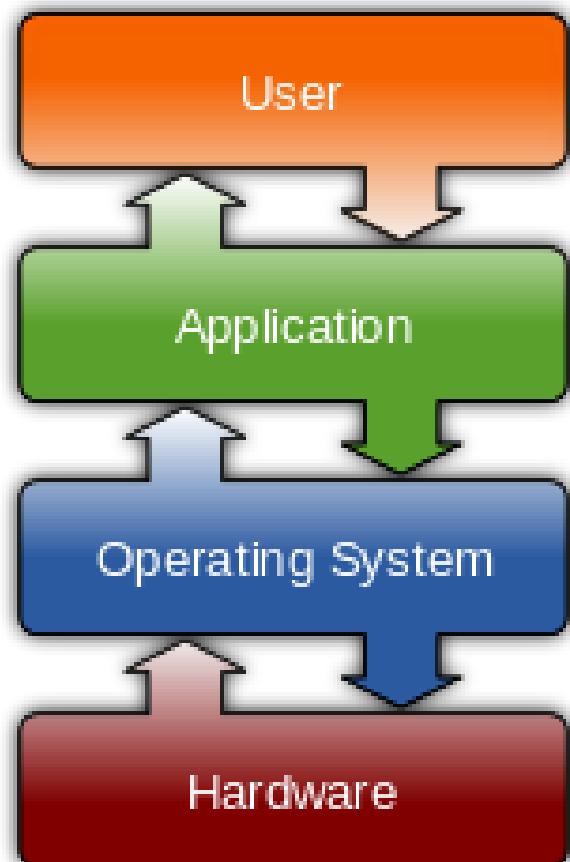


**Middleware** ("software glue"): Computer software that provides services to software applications beyond those available from OS; Helps to communicate and manage data of *distributed systems*.

**Firmware**: A block of program instructions for specific purposes, recorded in non-volatile memory, which establishes the lowest level logic that controls the electronic circuits of a device of any type; Helps to control the functionality of *the hardware or electronic devices*.

# Computer System: Operating System (OS)

Operating System (OS): Software that manages computer hardware and software resources and provides common services for computer programs.

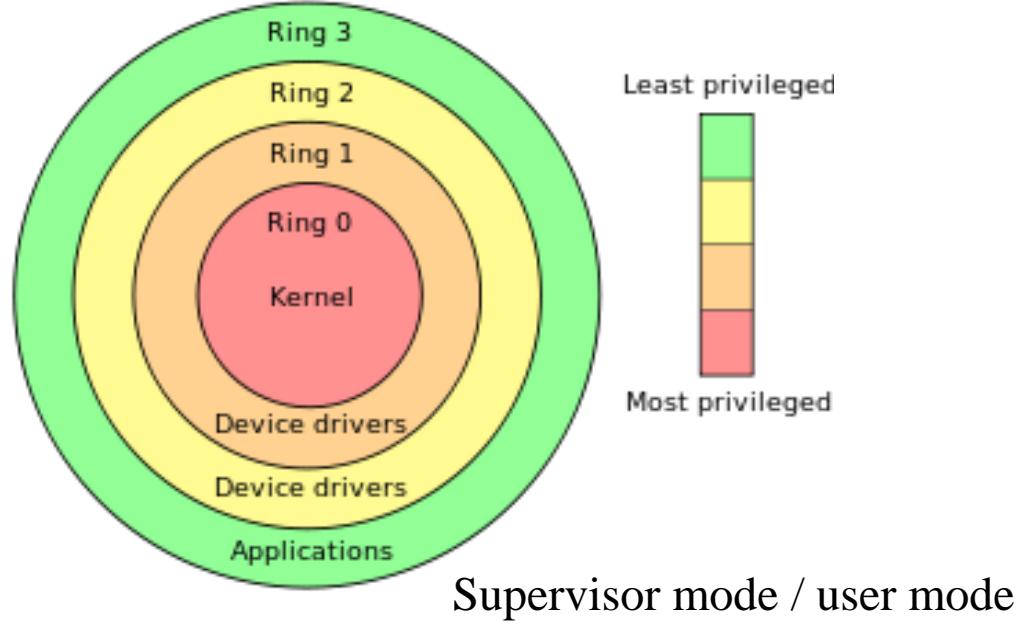


<https://www.topsimages.com/images/utility-system-fe.html>

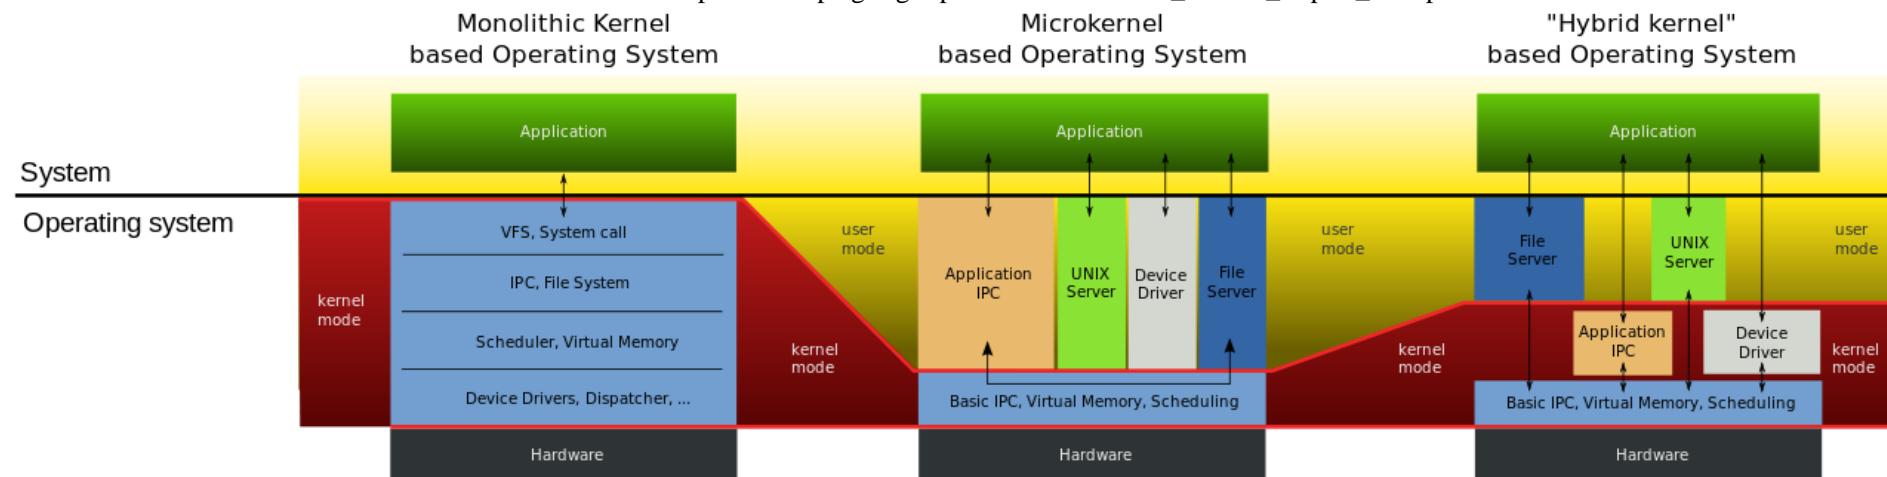
# Computer System: OS & Kernel

Kernel: a computer program that is the core of a computer's operating system, with complete control over everything in the system.

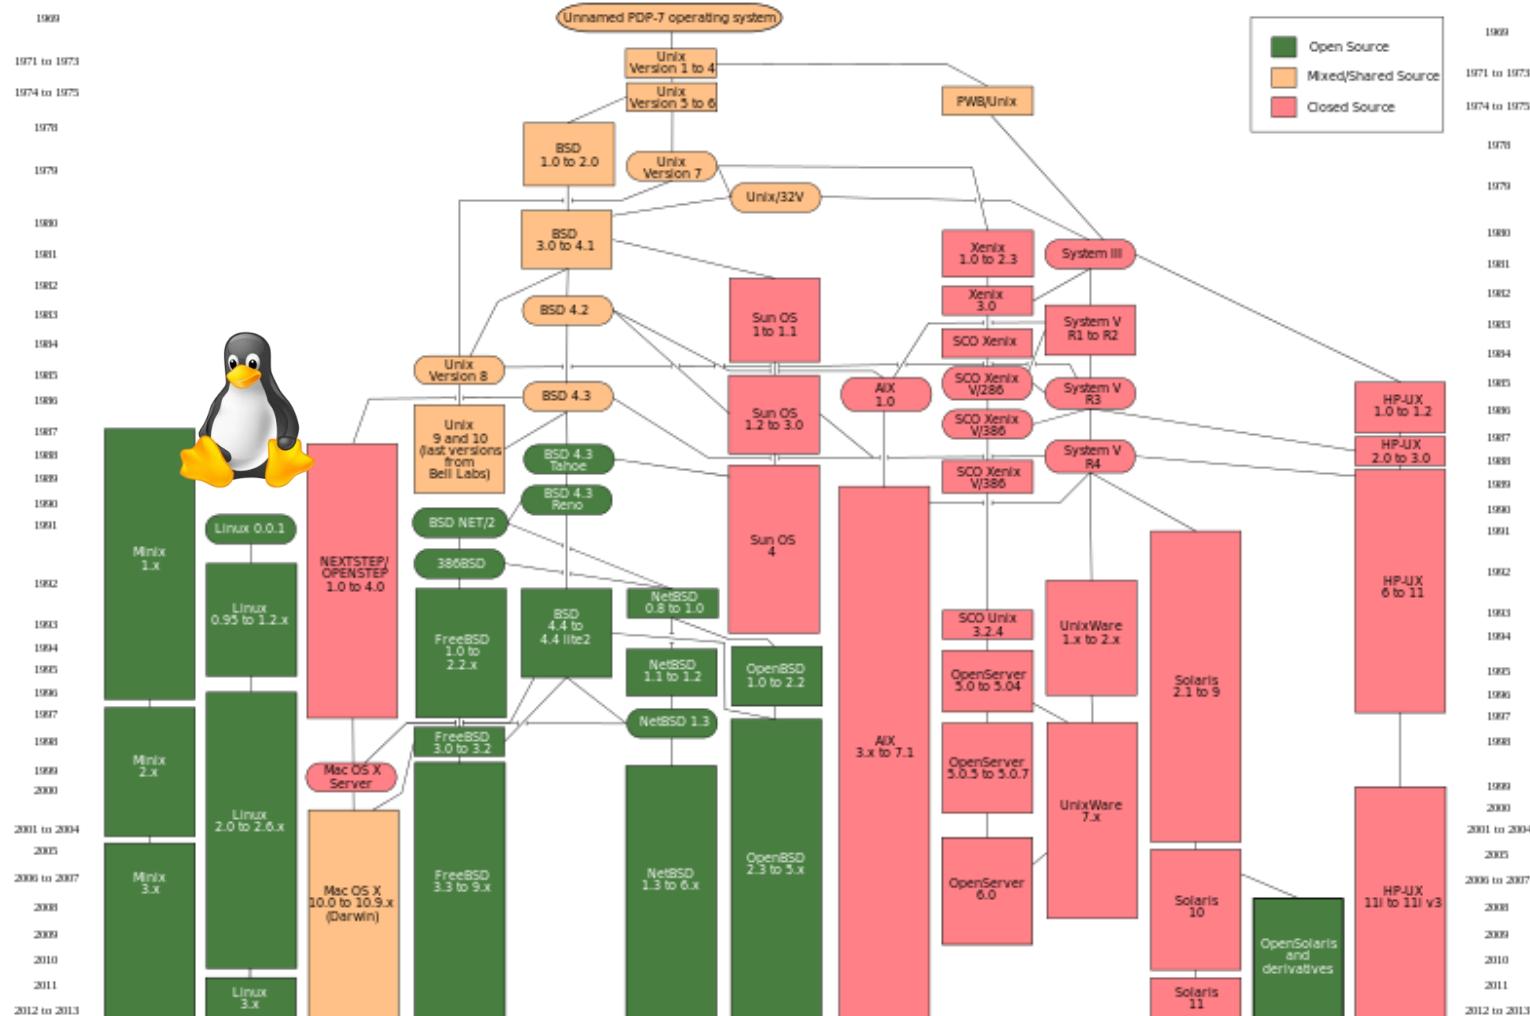
[https://de.wikipedia.org/wiki/Datei:Priv\\_rings.svg](https://de.wikipedia.org/wiki/Datei:Priv_rings.svg)



[https://docs.apwg.org/reports/mobile/APWG\\_Mobile\\_Report\\_v1.9.pdf](https://docs.apwg.org/reports/mobile/APWG_Mobile_Report_v1.9.pdf)



# History of Unix-like Operating Systems



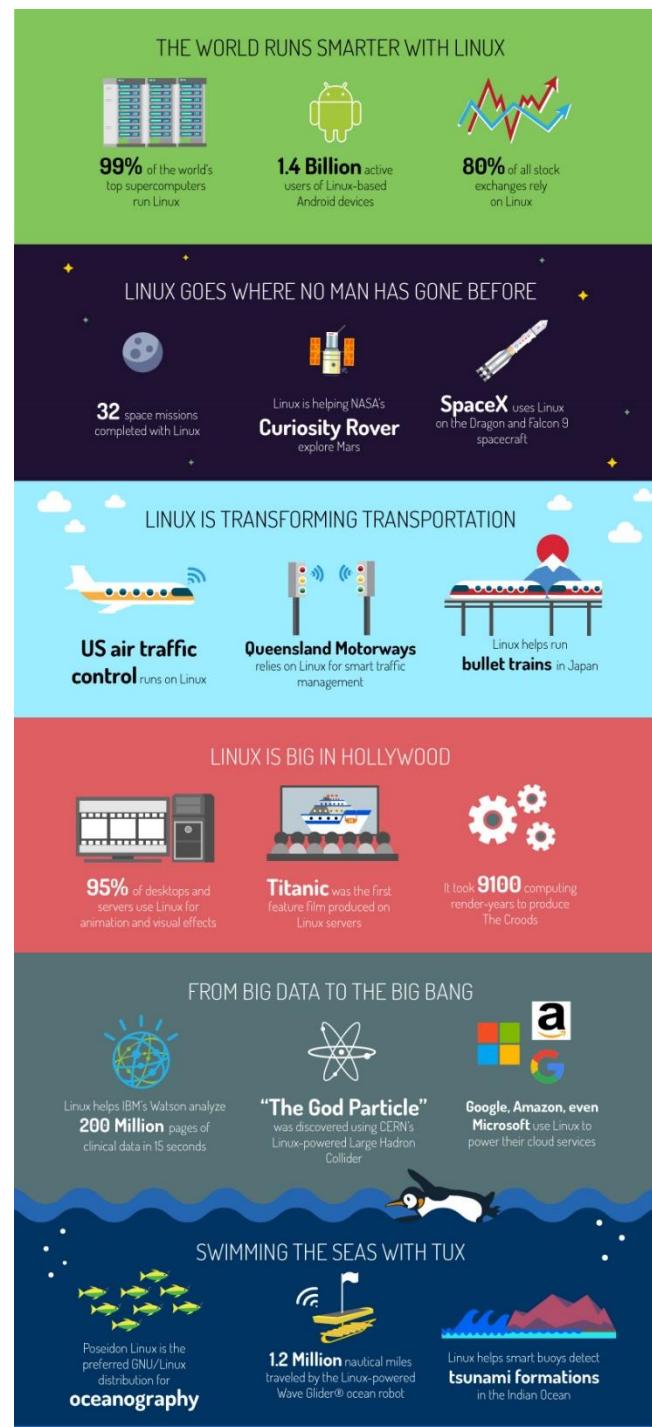
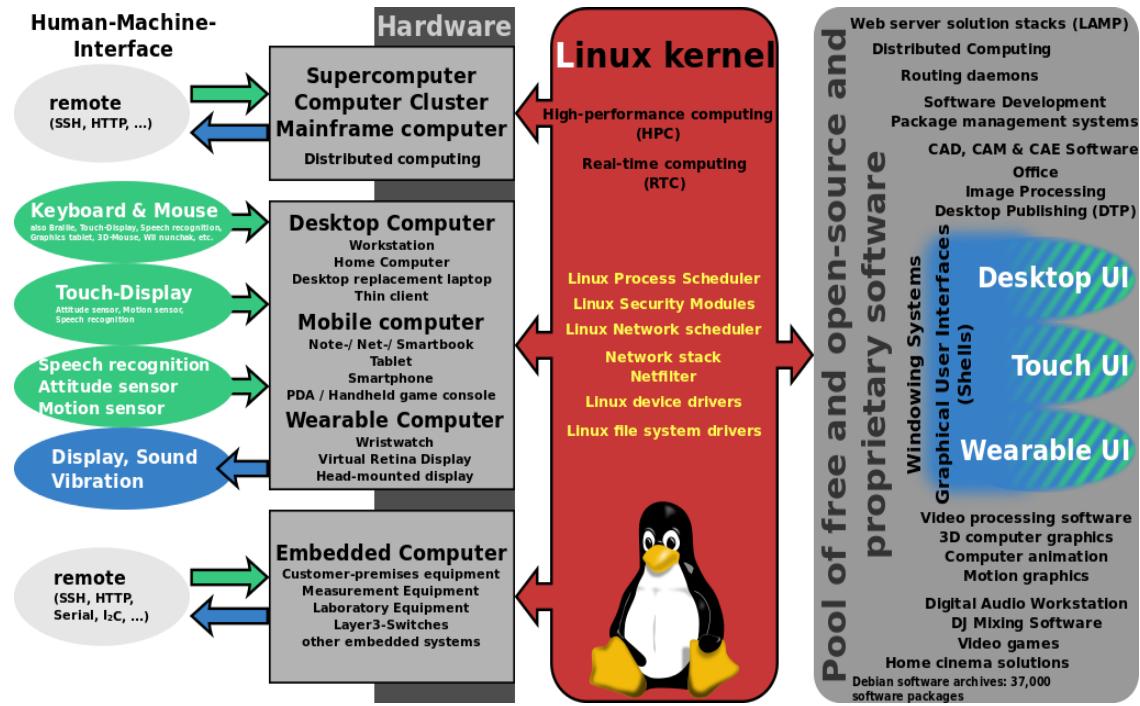
[Video] How Linux is built [https://youtu.be/yVpbFMhOAwE?list=PLAbMs\\_G4N5IWpQC2fDpv-hQhbCf-MPPQL](https://youtu.be/yVpbFMhOAwE?list=PLAbMs_G4N5IWpQC2fDpv-hQhbCf-MPPQL)

[Video] Linux Kernel Development, 1991-2015 <https://youtu.be/5iFnzr73XXk>

# Linux OS Everywhere



<http://www.linuxsc.net/linux/windows-domina-el-escritorio-pero-linux-gana-el-mundo>



# Linux Kernel Diagram



functions

system

networking

storage

memory

processing

human interface

layers

user space interfaces

System calls

Sockets

files and directories

memory access

Processes

char devices

virtual subsystems

proc, sysfs  
file systems

protocol families

Virtual file system

Virtual memory

Tasks

input subsystem

bridges

Device Model

system run,  
modules,  
generic  
HW access

NFS

page cache

memory  
mapping

Swap

logical

protocols:  
TCP, UDP, IP

logical  
filesystems  
ext2, ext3

logical  
memory

Scheduler

HI class  
drivers

bus drivers

network  
devices  
and drivers

Block  
devices  
and drivers

Page  
Allocator

interrupts  
core

HI  
peripherals  
drivers

bus  
controllers:  
PCI, USB

network  
cards:  
Ethernet, WiFi

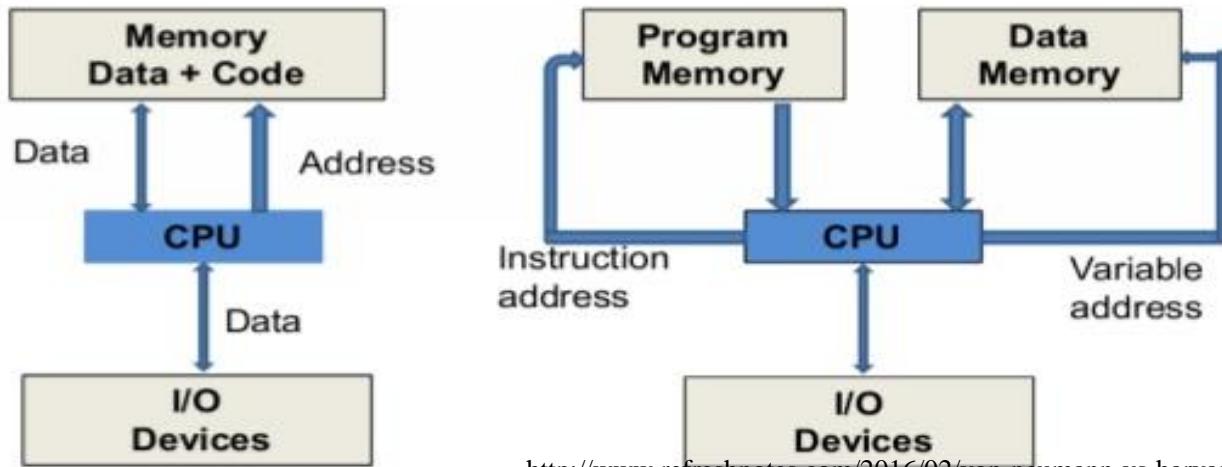
disk  
controllers:  
IDE, SCSI

MMU, RAM

CPU

display  
keyboard  
mouse, audio

# Computer System: Microcontroller Architecture



<http://www.refreshnotes.com/2016/02/von-neumann-vs-harvard-architecture.html>

## Von Neumann Model

Named after the mathematician and early computer scientist John Von Neumann.

Required only one memory for their instruction and data.

Design of the von Neumann architecture is simple.

Required only one bus for instruction and data.

Processor needs two clock cycles to complete an instruction.

Low performance as compared to Harvard architecture.

Cheaper.

## Harvard Model

Originated from "Harvard Mark I" a relay based computer.

Required two memories for their instruction and data.

Design of Harvard architecture is complicated.

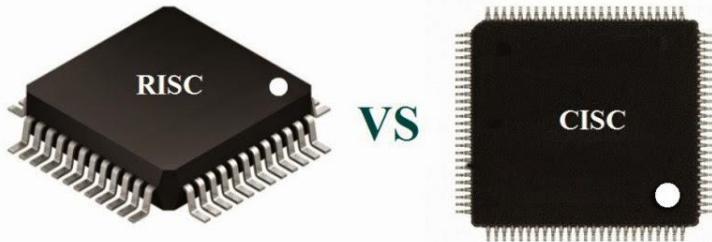
Required separate bus for instruction and data.

Processor can complete an instruction in one cycle

Easier to pipeline, so high performance can be achieved.

Comparatively high cost.

# Computer: CPU – RISC vs CISC



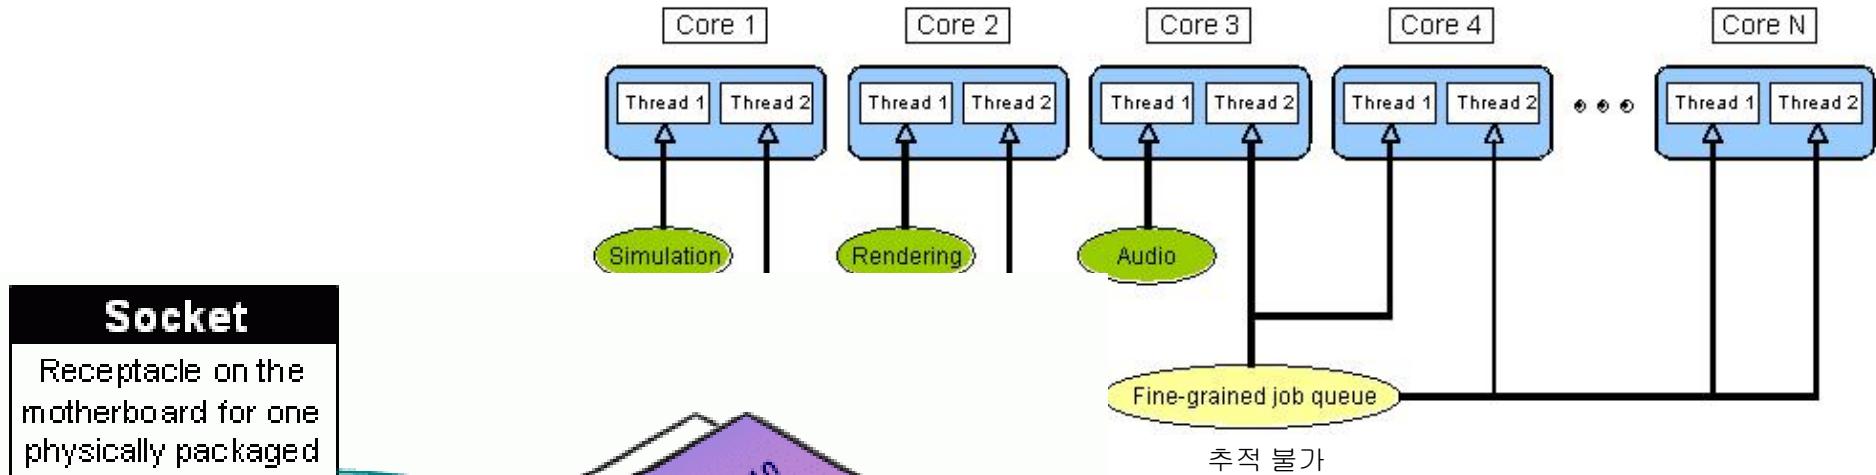
CISC	RISC
Emphasis on hardware	Emphasis on software
Multiple instruction sizes and formats	Instructions of same set with few formats
Less registers	Uses more registers
More addressing modes	Fewer addressing modes
Extensive use of microprogramming	Complexity in compiler
Instructions take a varying amount of cycle time	Instructions take one cycle time
Pipelining is difficult	Pipelining is easy

- Intel x86: complex instruction set computing (CISC) architecture
- ARM, MIPS: reduced instruction set computing (RISC) → low power consumption

<https://csarassignment.wordpress.com/2013/05/02/risc-vs-cisc/>

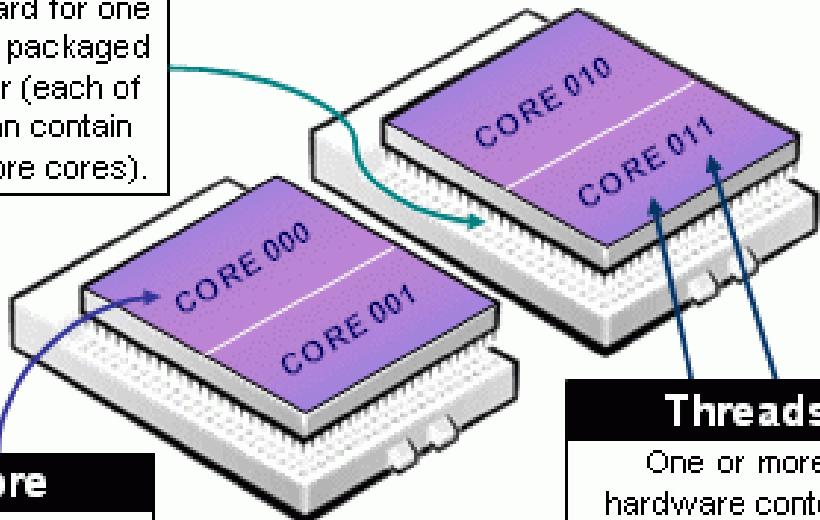
ARM:  
A chips for high-performance application;  
R chips for real-time processing;  
M chips for micro-controllers.

# OS & Compute: Socket/Core & Threads



## Socket

Receptacle on the motherboard for one physically packaged processor (each of which can contain one or more cores).



## Core

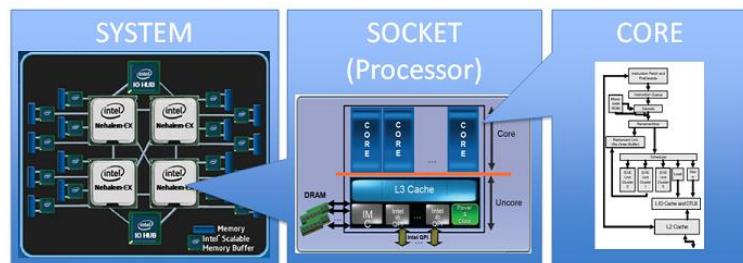
A complete private set of registers, execution units, and retirement queues needed to execute programs.

## Threads

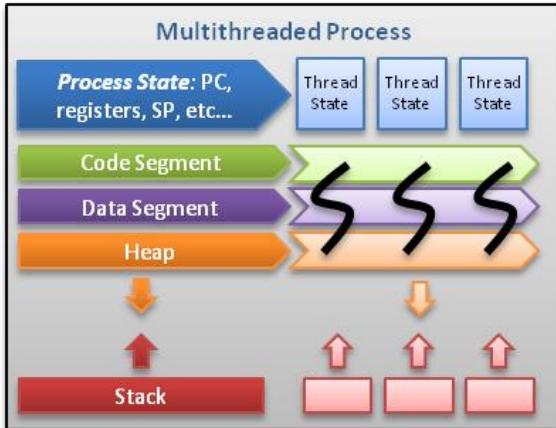
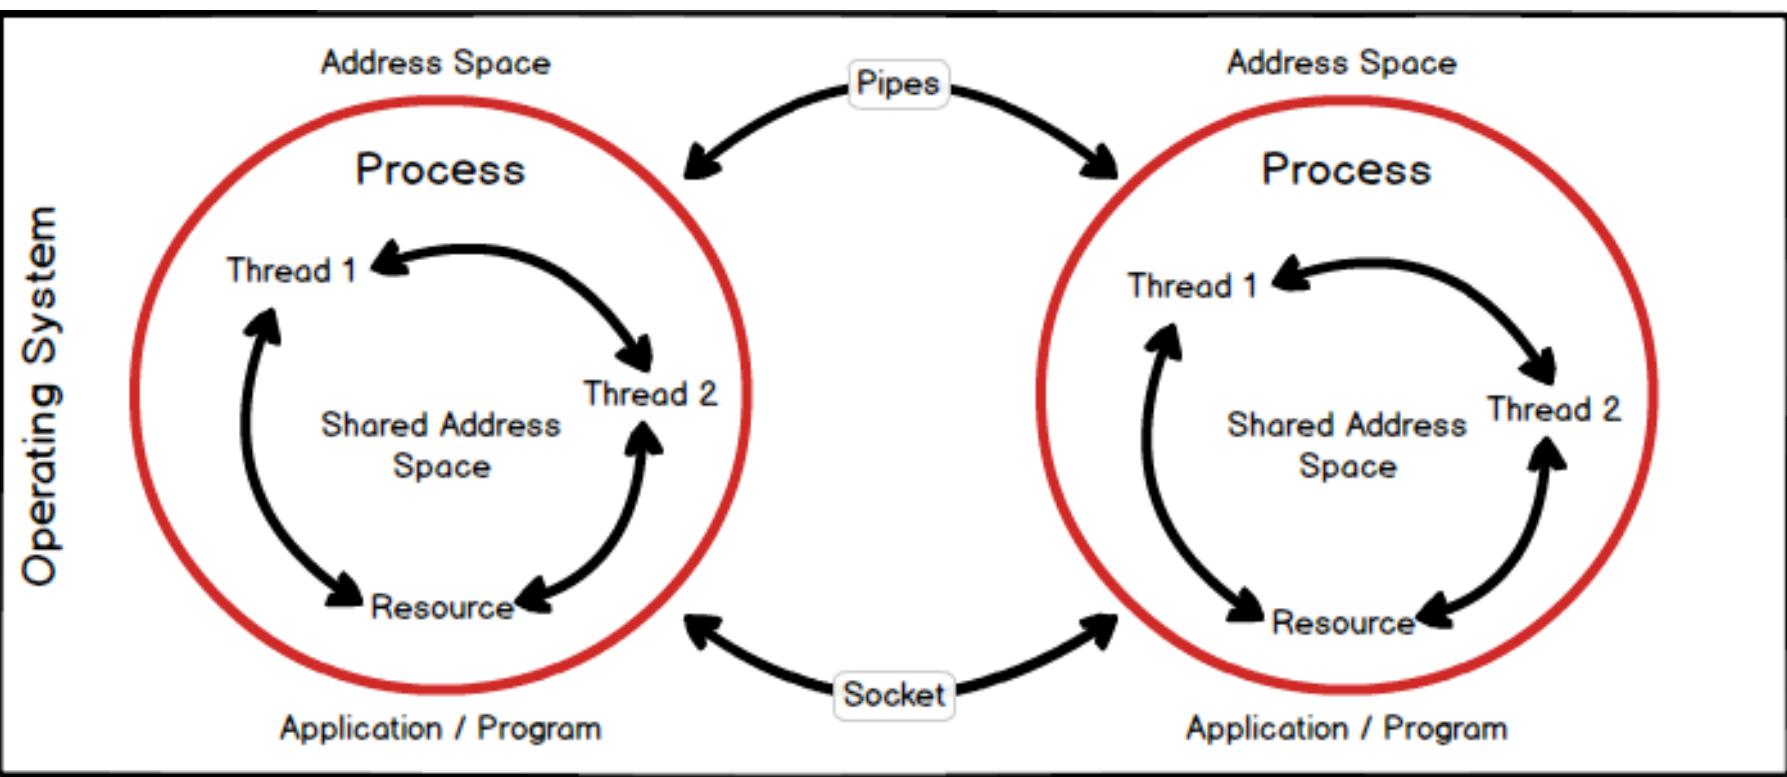
One or more hardware contexts within a single core.

Each thread has attributes of one core; managed & scheduled as a single logical processor by the OS.

[http://www.hpc.lsu.edu/training/weekly-materials/2018-Fall/HPC\\_UserEnv\\_Fall\\_2018\\_session\\_1.pdf](http://www.hpc.lsu.edu/training/weekly-materials/2018-Fall/HPC_UserEnv_Fall_2018_session_1.pdf)

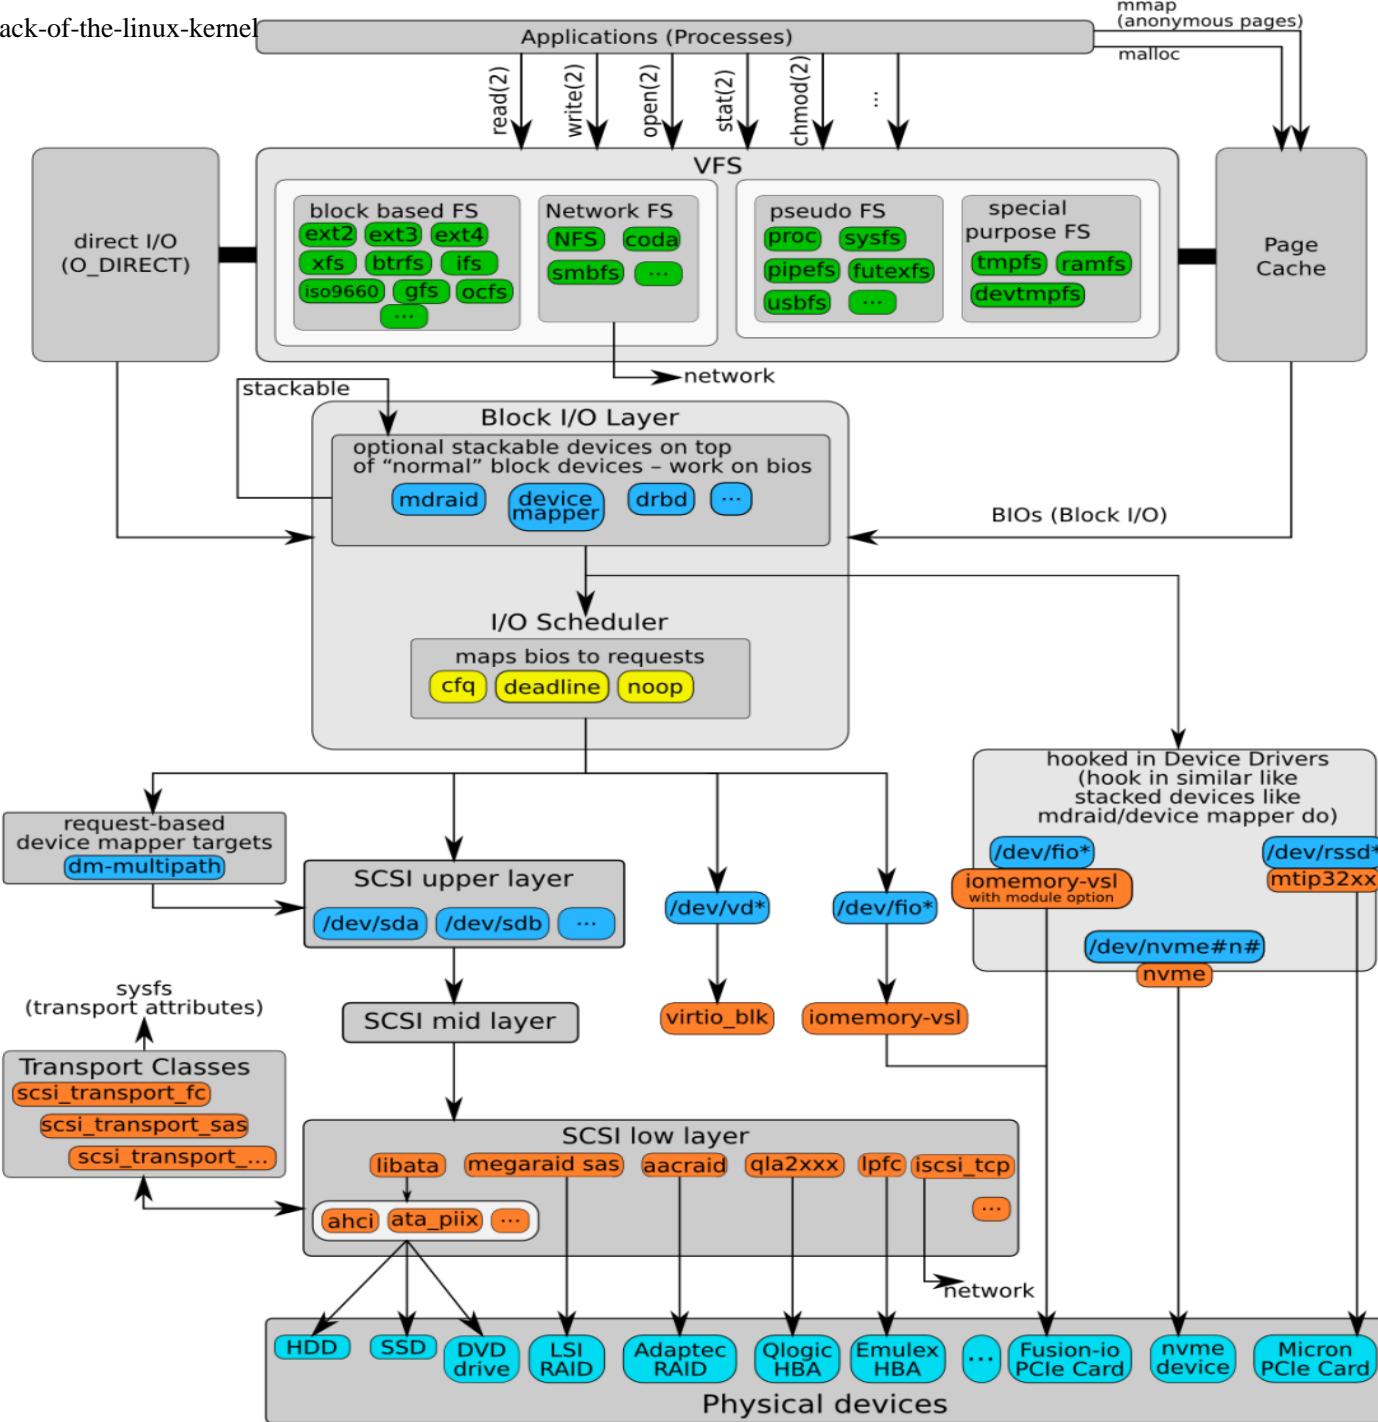


# OS & Compute: Threads/Processes



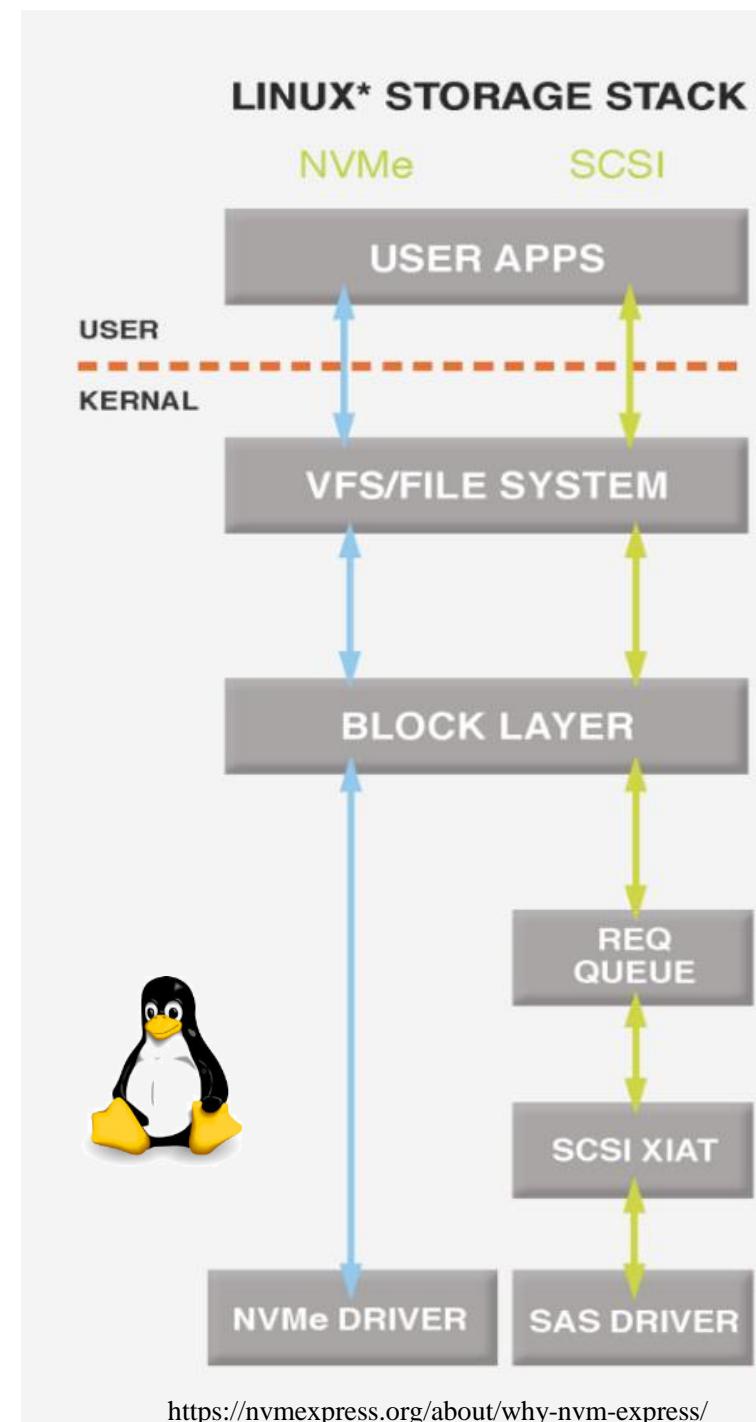
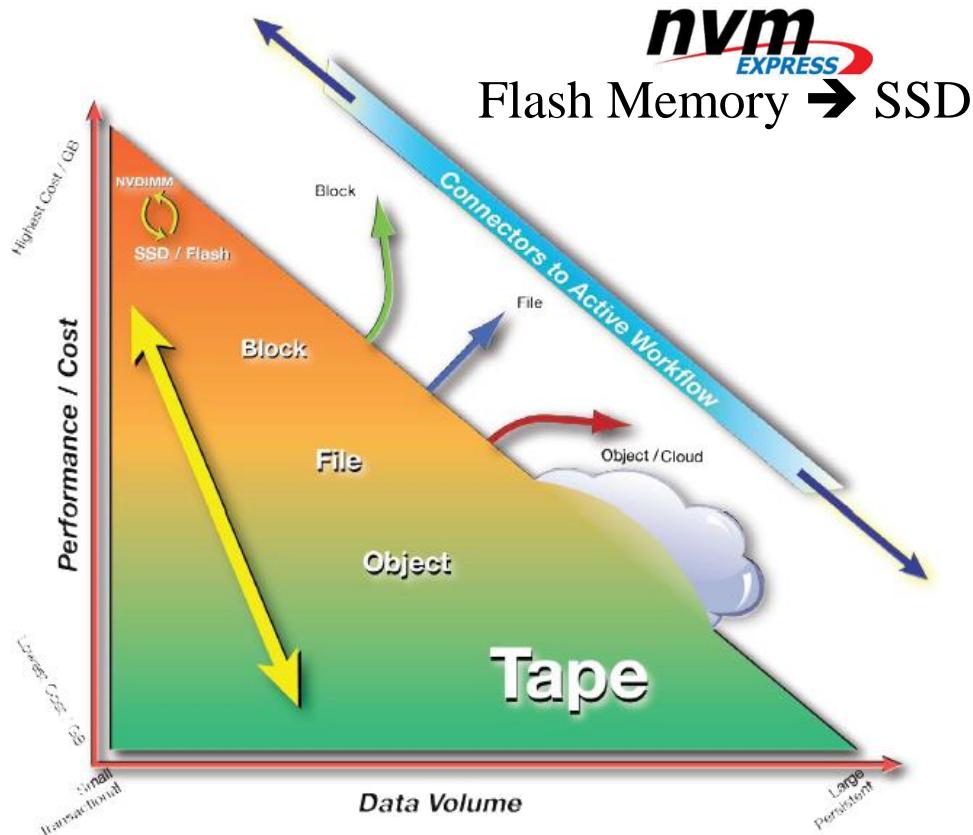
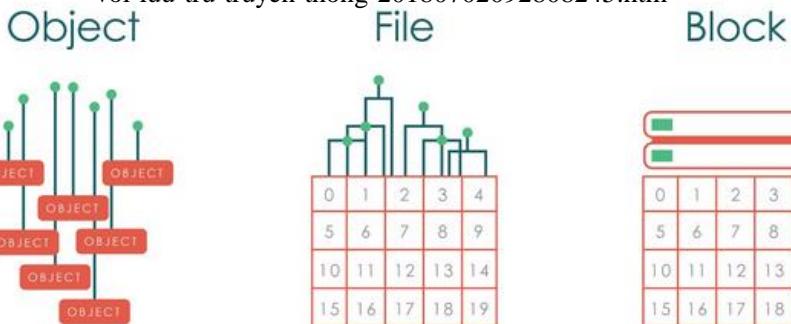
**Threads** contain only necessary information, such as a stack, a copy of the registers, program counter and any thread-specific data to allow them to be scheduled individually.

# Linux OS: I/O Stack



# OS & Memory/Storage

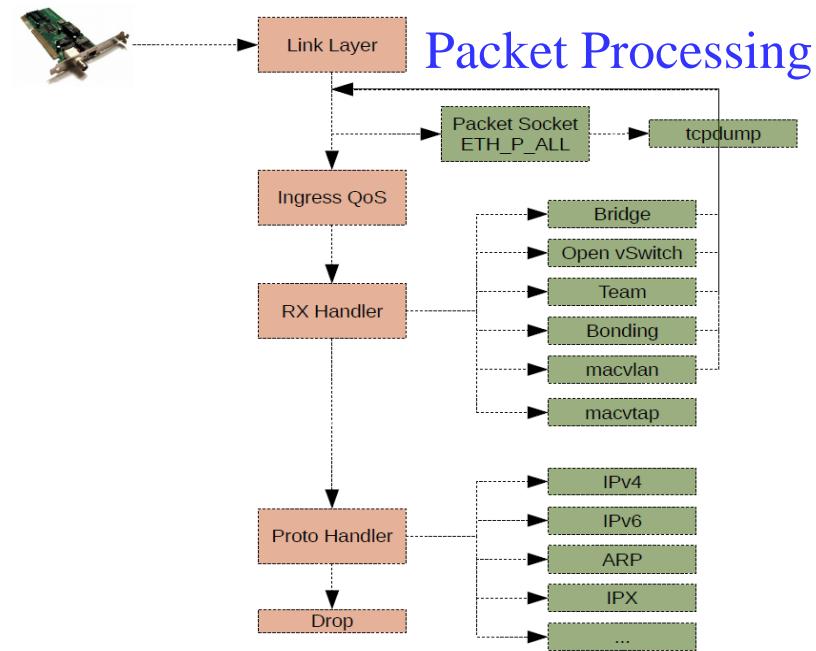
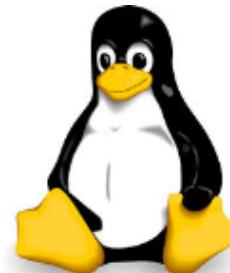
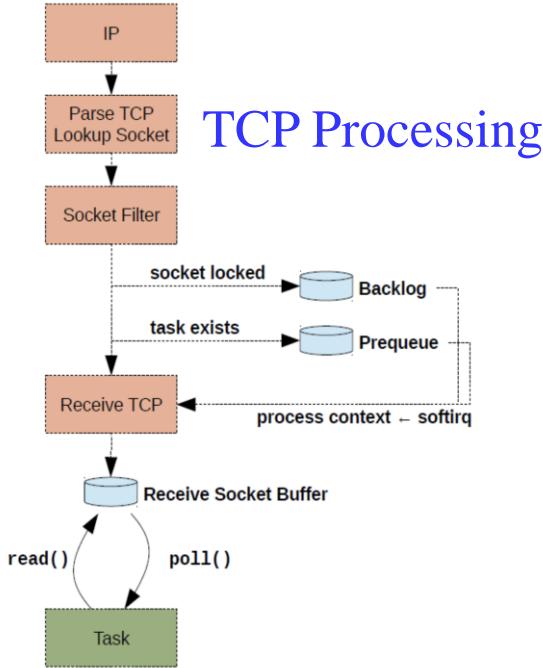
<https://tech.vccloud.vn/object-storage-luu-tru-doi-tuong-la-gi-khac-gi-voi-luu-tru-truyen-thong-20180702092808245.htm>



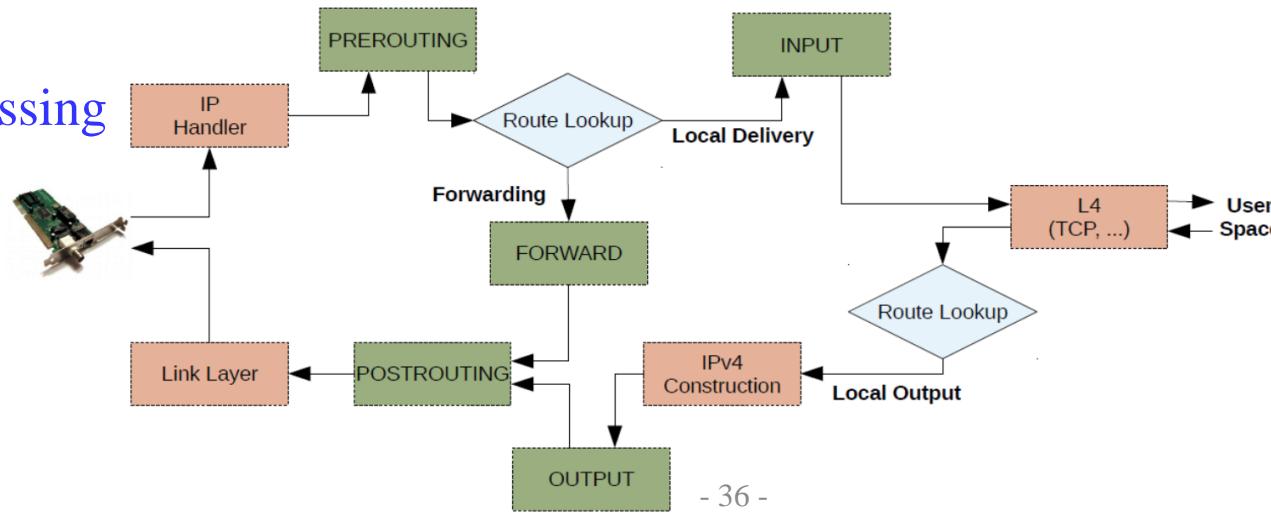
<https://nvmexpress.org/about/why-nvm-express/>

# Linux Kernel Networking: L2/L3/L4...

<https://www.slideshare.net/ThomasGraf5/devconf-2014-kernel-networking-walkthrough>



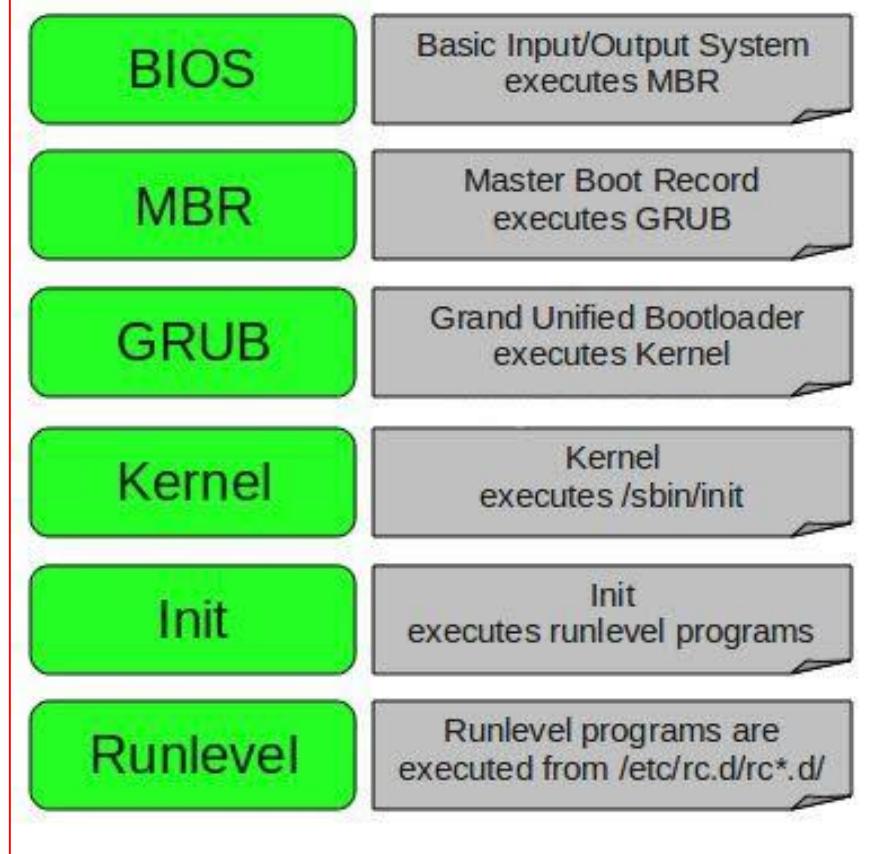
## IP Processing



# Linux OS: Boot Procedure



<http://www.matsuura.com.br/2014/04/processo-de-inicializacao-do-linux.html>



<https://www.thegeekdiary.com/centos-rhel-7-booting-process/>

Run Level	Name	Description
0	Halt	Shuts down all services when the system will not be rebooted.
1	Single User	Used for system maintenance. No Networking capabilities.
2	MultiUser No Network Support	Used for maintenance and system testing.
3	MultiUser Network Support	Non-Graphical Text Mode operations for server systems.
4	-	Custom Mode, used by SysAdmin
5	Graphical <small>X11</small>	Graphical login with same usability of Run Level 3.
6	Reboot	Shuts down all services when the system is being rebooted.

<https://docplayer.es/4085195-Systemd-software-freedom-day-2015-alex-callejas.html>

## Lilo → GRUB (Grand Unified)

Illustration 1: an MBR-partitioned harddisc with sector size of 512 or 4096Bytes

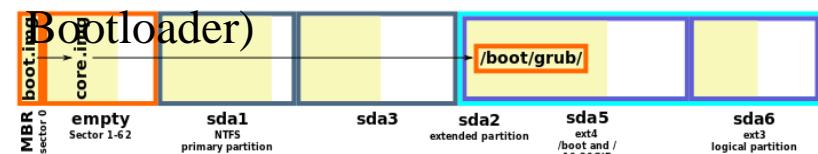
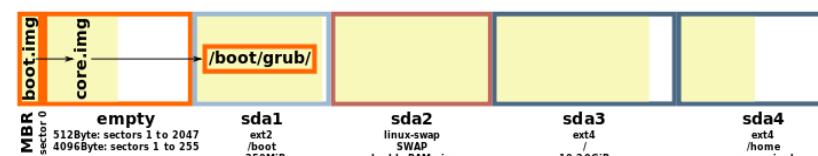


Illustration 2: recommended partitioning



<https://unix.stackexchange.com/questions/111237/how-does-linux-deal-with-a-separate-boot-partition>



Gwangju Institute of  
Science & Technology



Thank you!

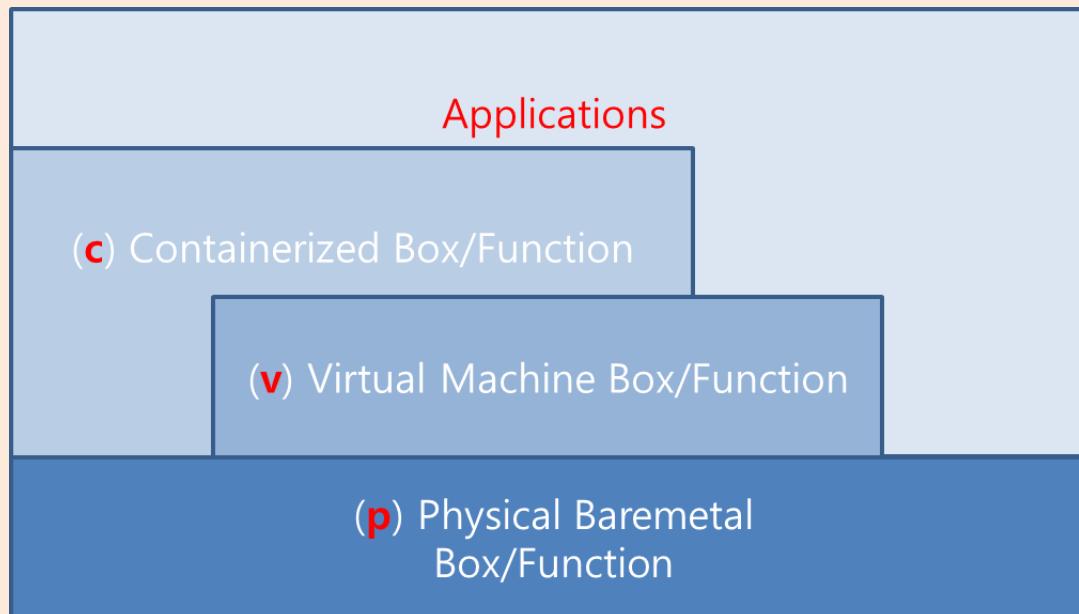
Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>

# Chapter 3:

# Computer System

## Virtualization & Containers



# Chapter Keywords

Virtualization

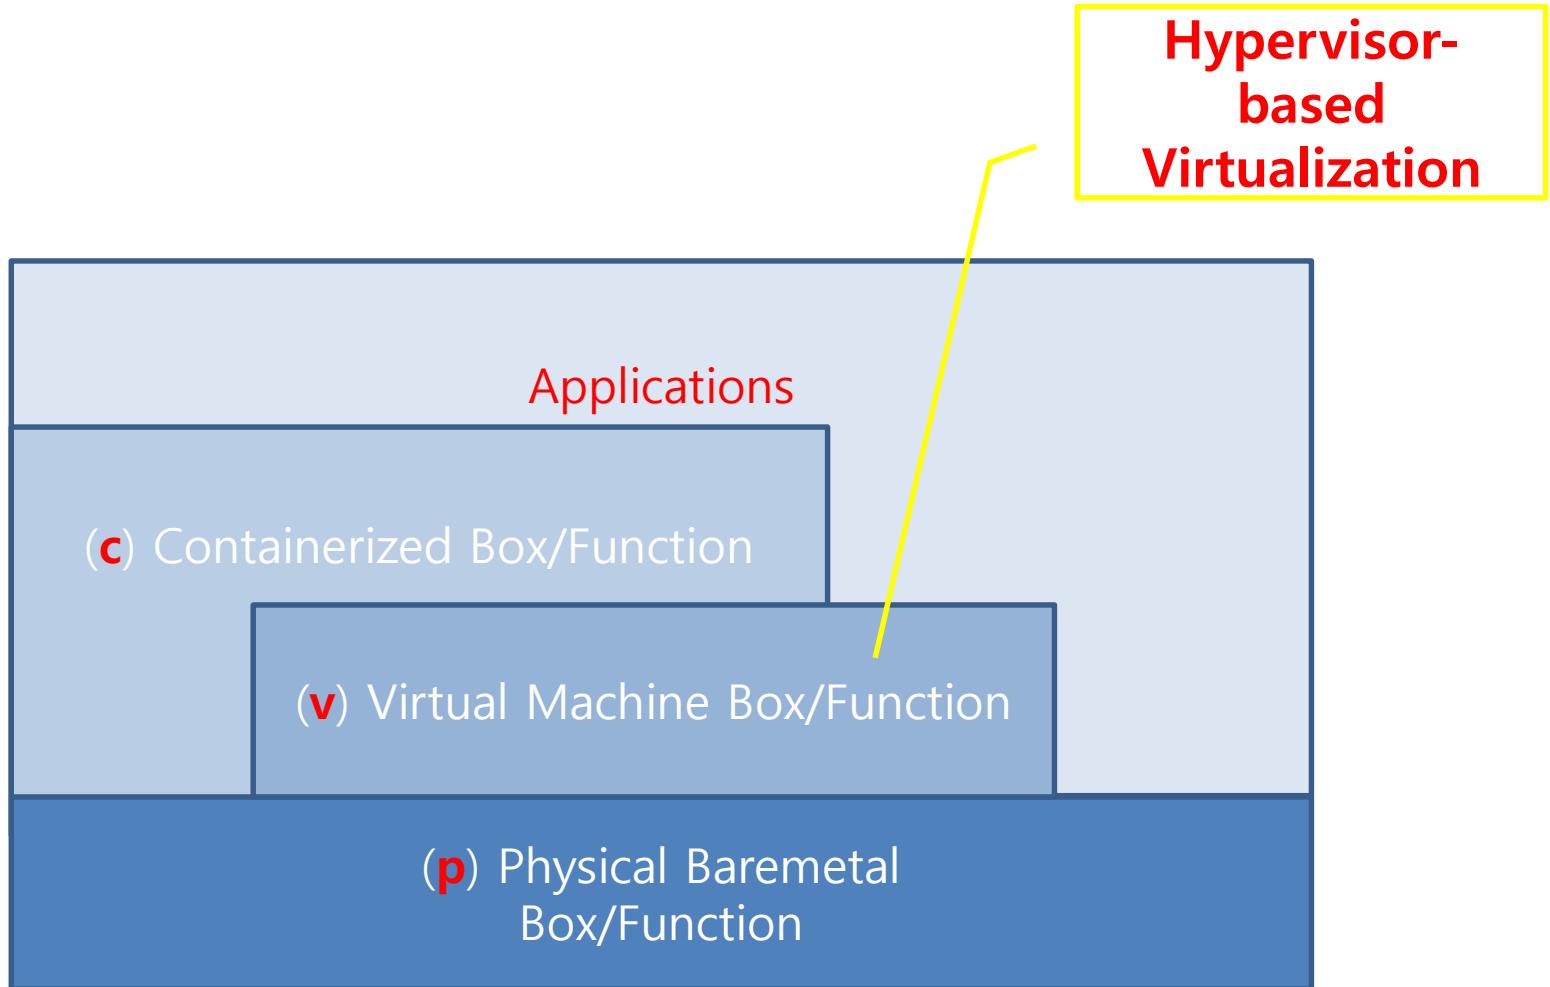
Hypervisor

Virtual Machine

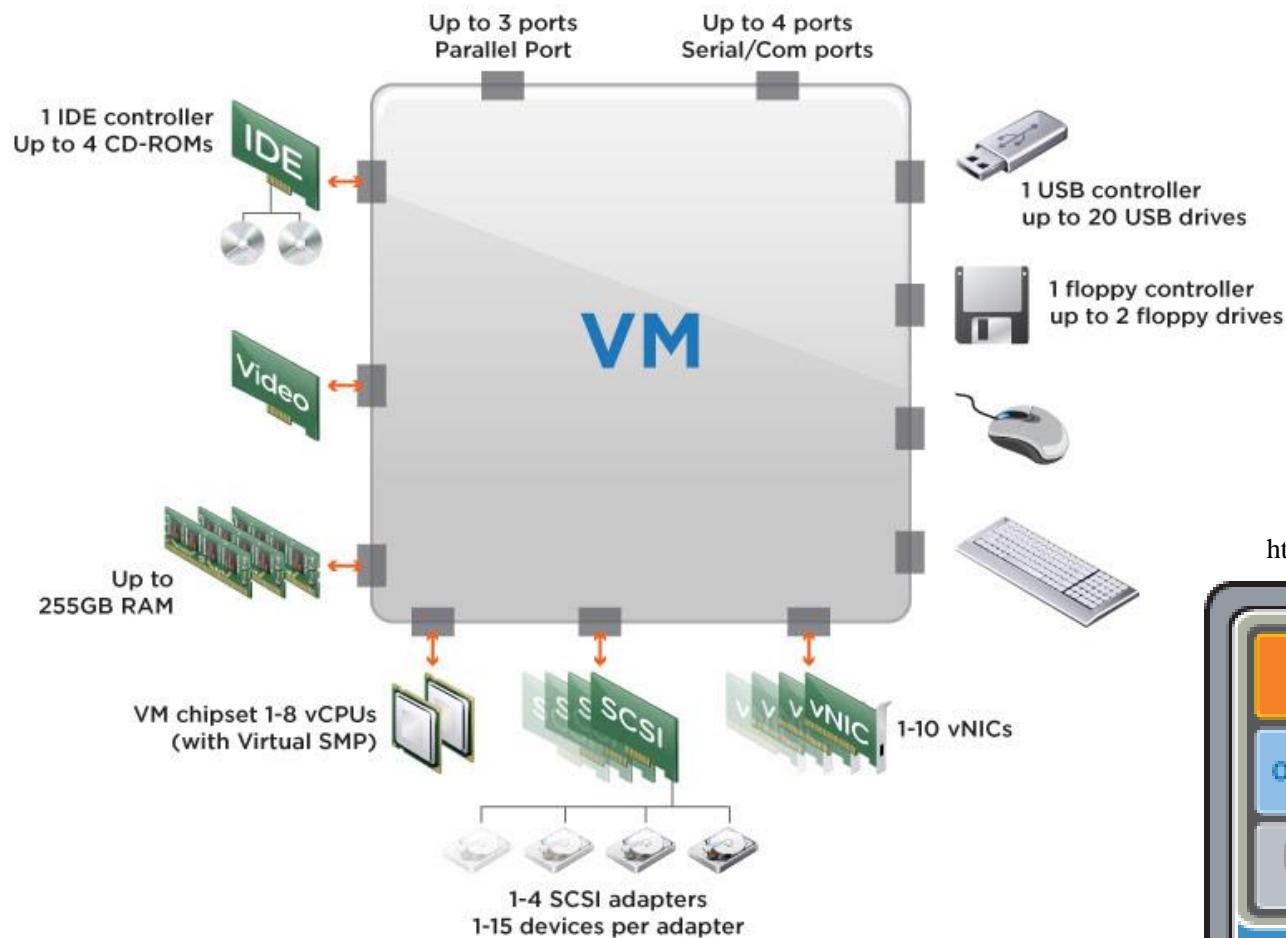
Containers

**P**(Baremetal) + **V**(Virtual Machine) + **C**(Container)

# Harmonization

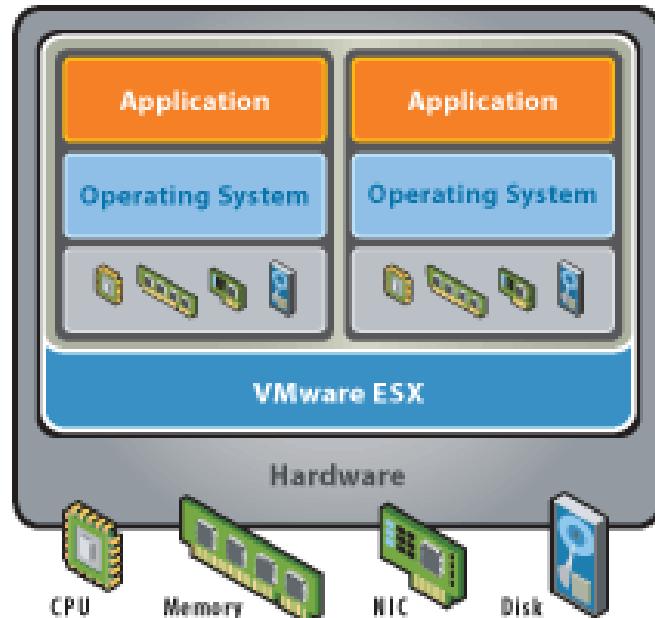


# Virtual Machines (VMs)

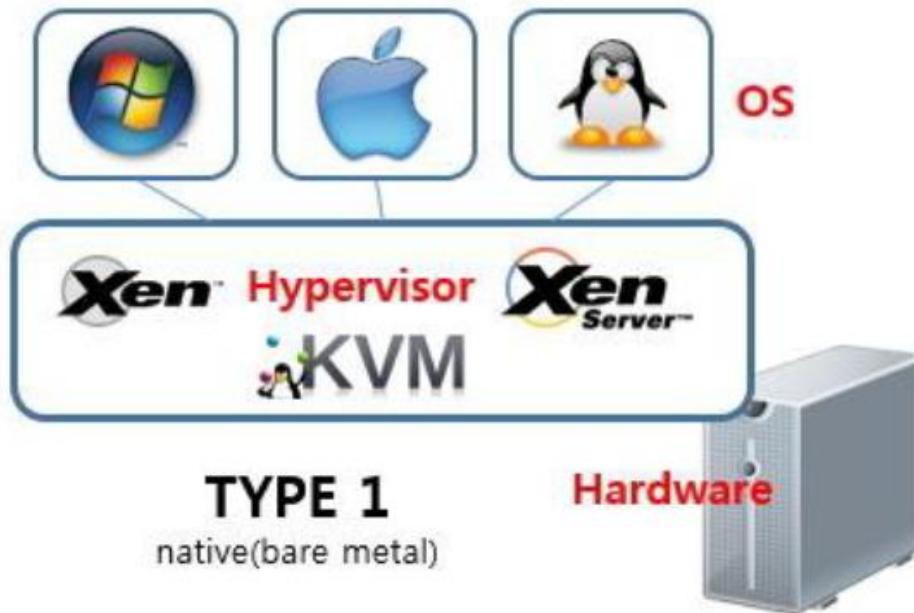


<http://kuas1095108129.blogspot.com/>

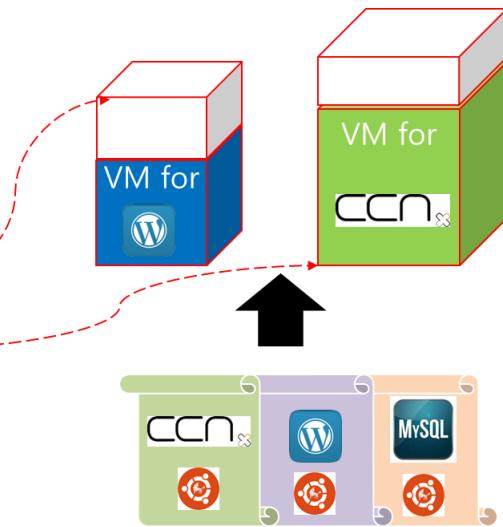
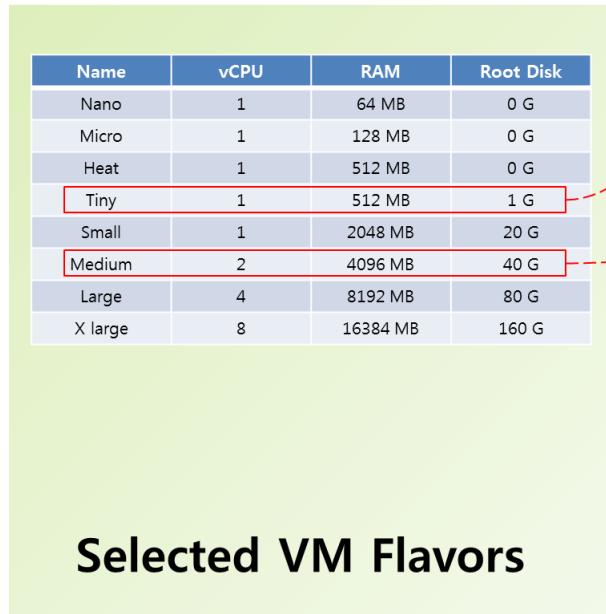
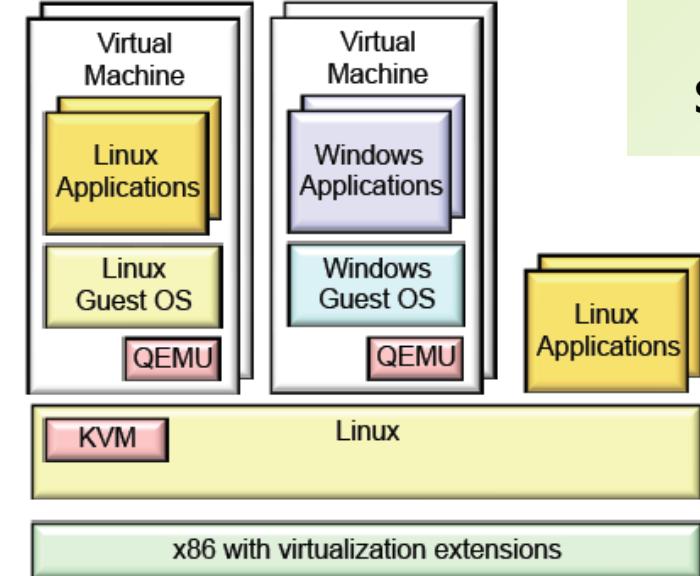
<http://hesed.info/blog/virtuelle-maschine.abp>



# Hypervisor Type 1 & 2 for Virtualization



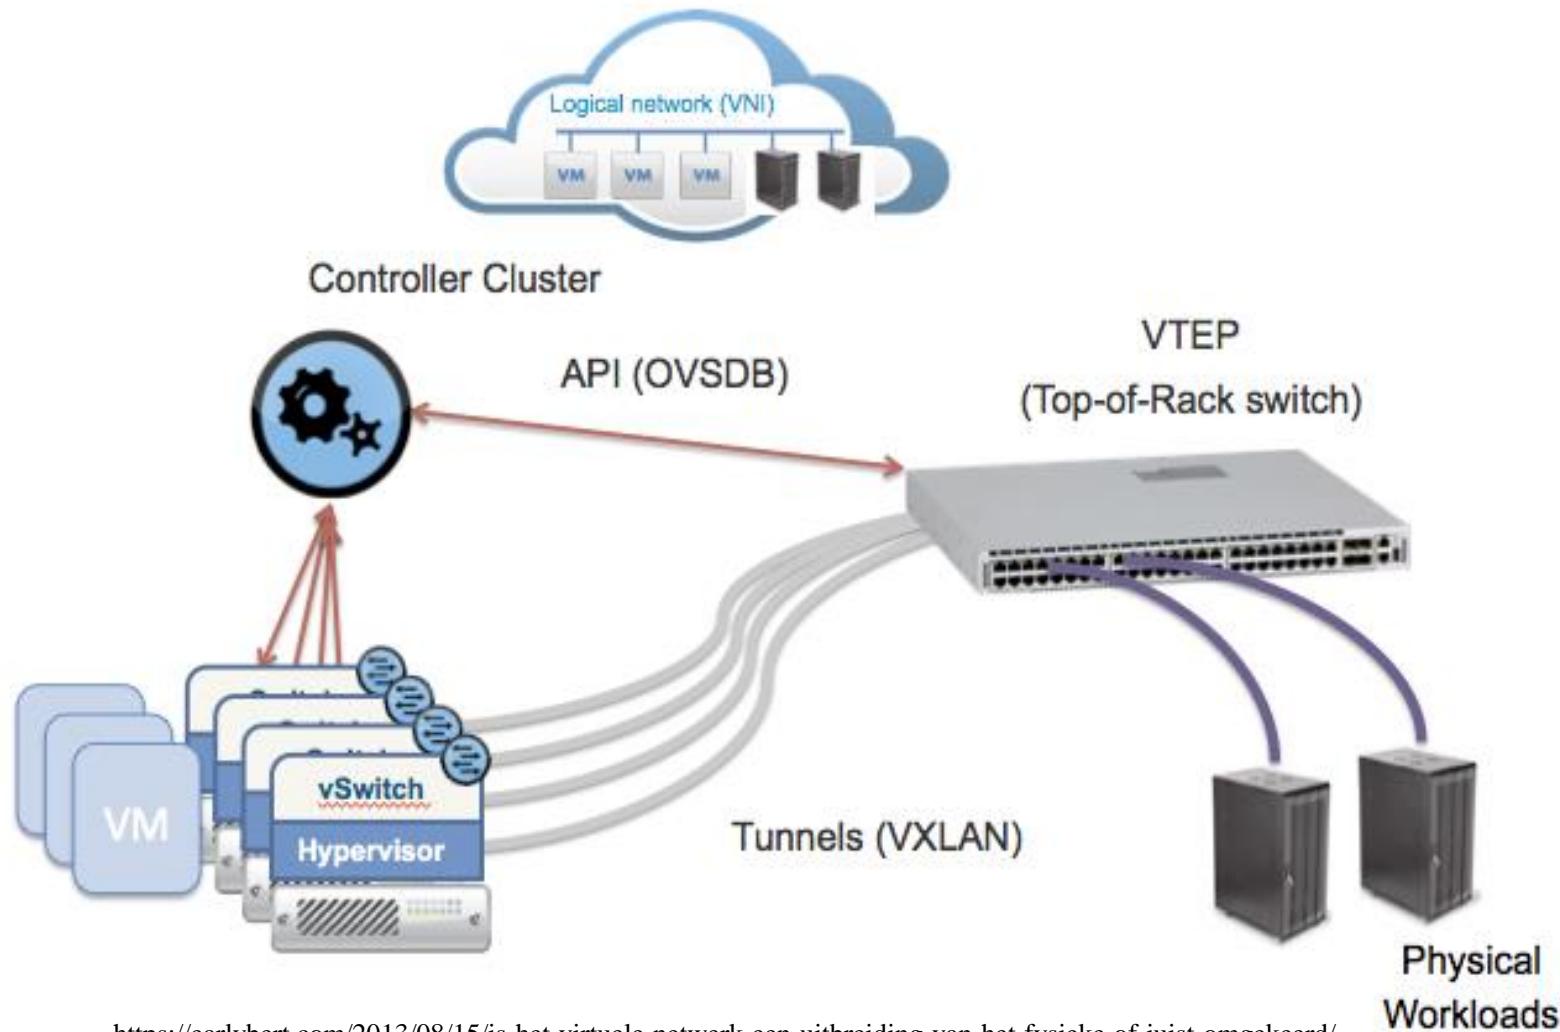
# VMs (Functions) with Hypervisor-based Virtualization



VM images with  
preinstalled Functions

추적 불가

# Physical and Virtual Machines: Inter-connection

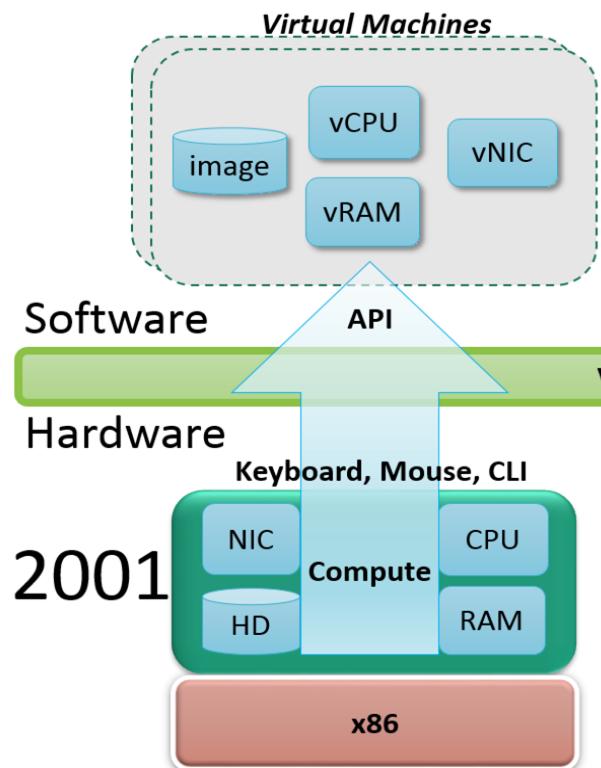


<https://earlybert.com/2013/08/15/is-het-virtuele-netwerk-een-uitbreiding-van-het-fysieke-of-juist-omgekeerd/>

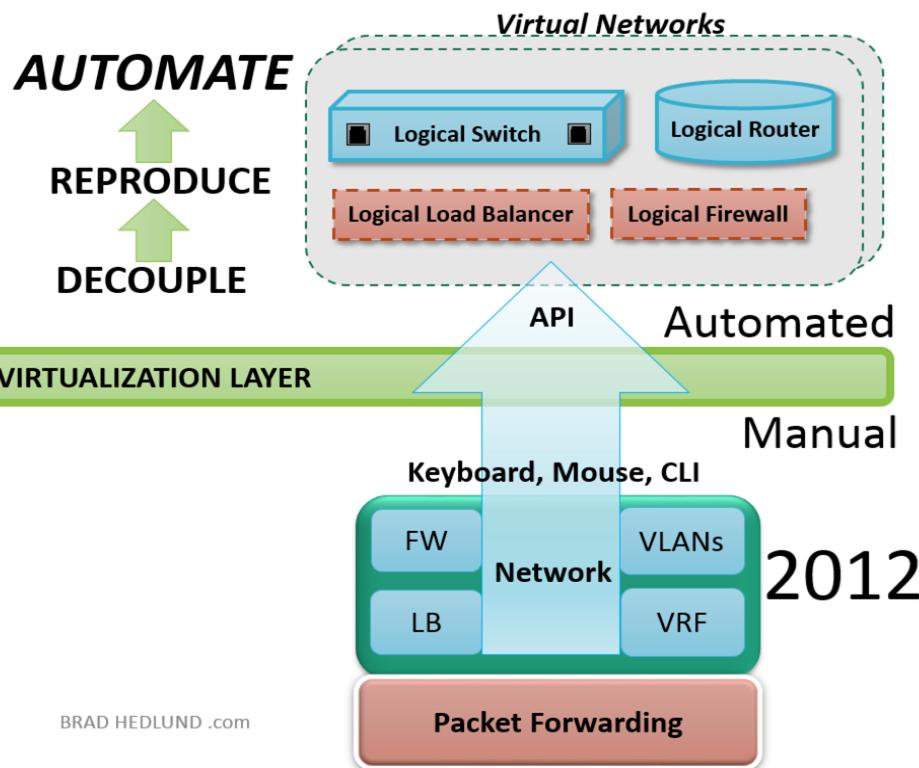
# Virtualization for Compute/Network

Virtualization: Basic act of **decoupling** an infrastructure service from the physical assets on which that service operates.

## Compute Virtualization



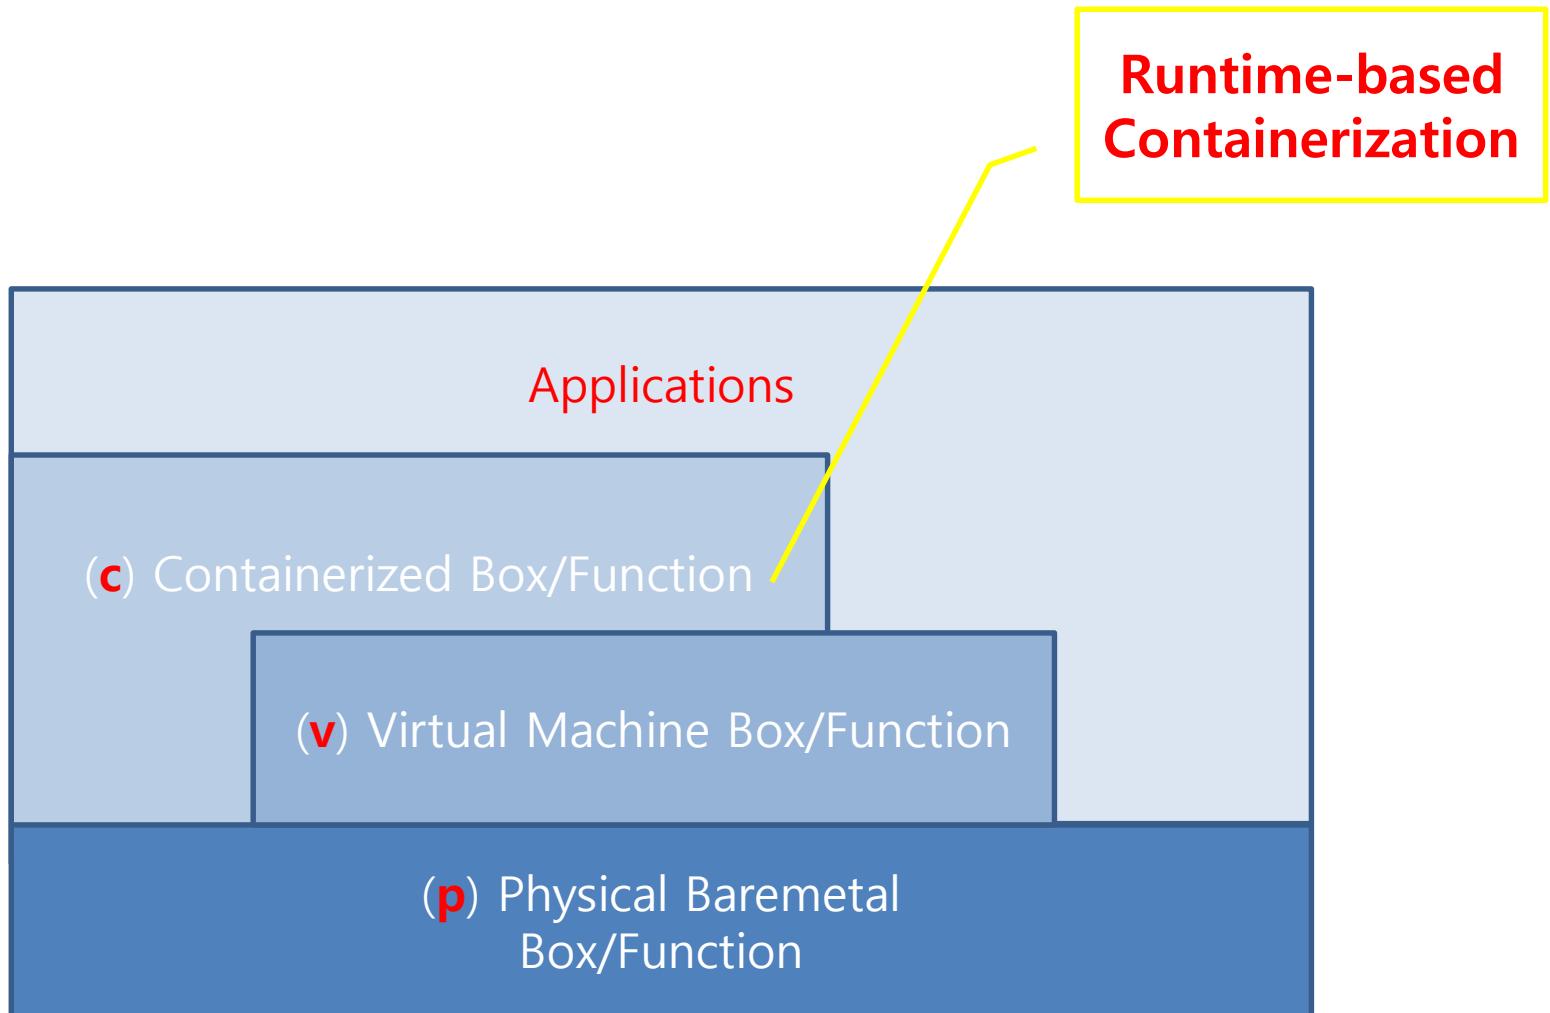
## Network Virtualization



BRAD HEDLUND .com

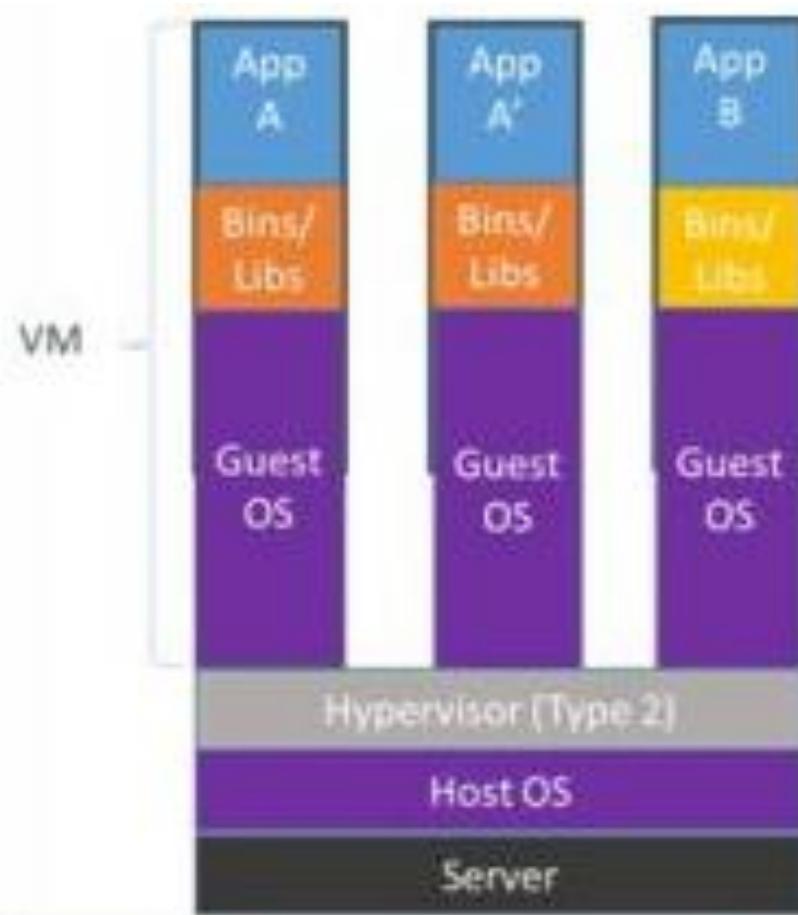
**P**(Baremetal) + **V**(Virtual Machine) + **C**(Container)

# Harmonization

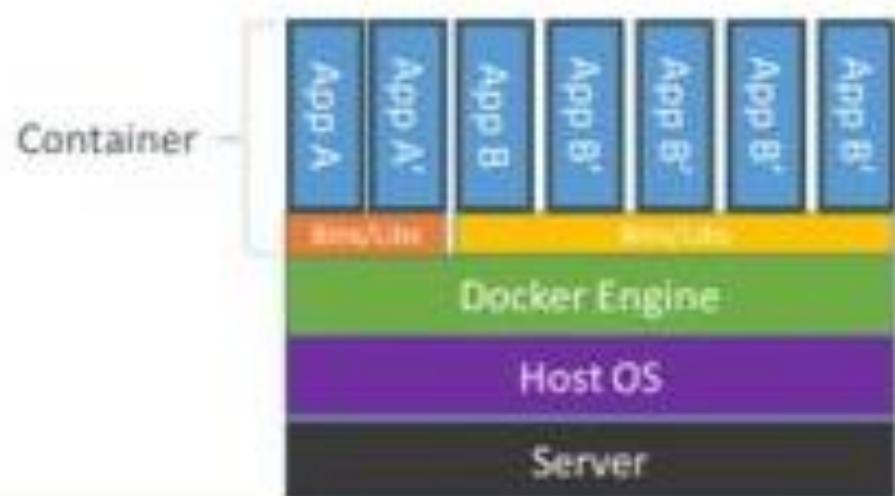


# Containers vs VMs

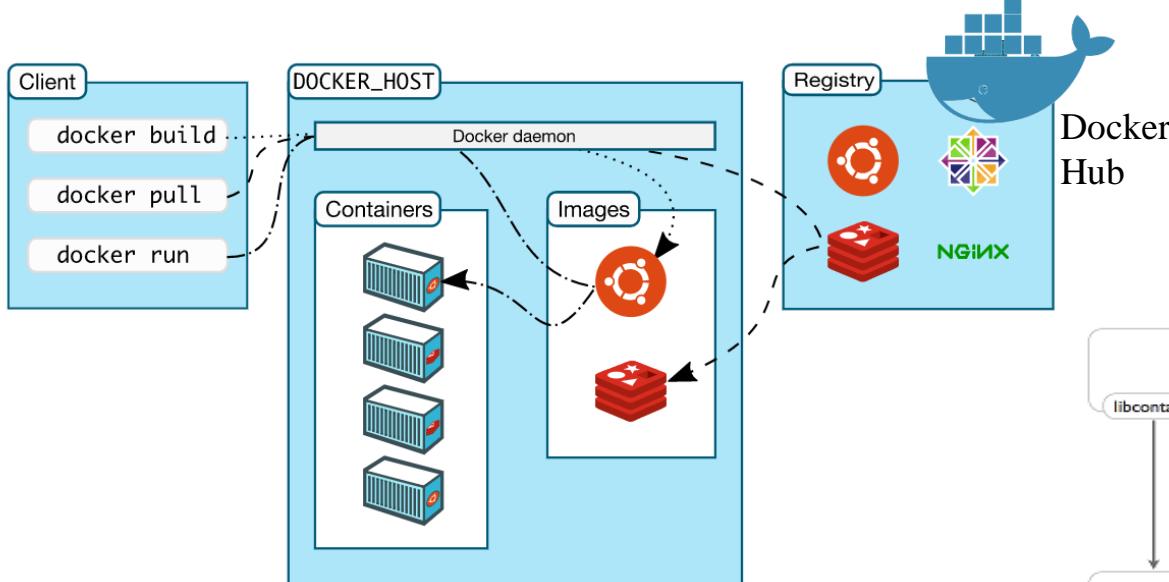
## Packaging workload tasks and Scaling ...



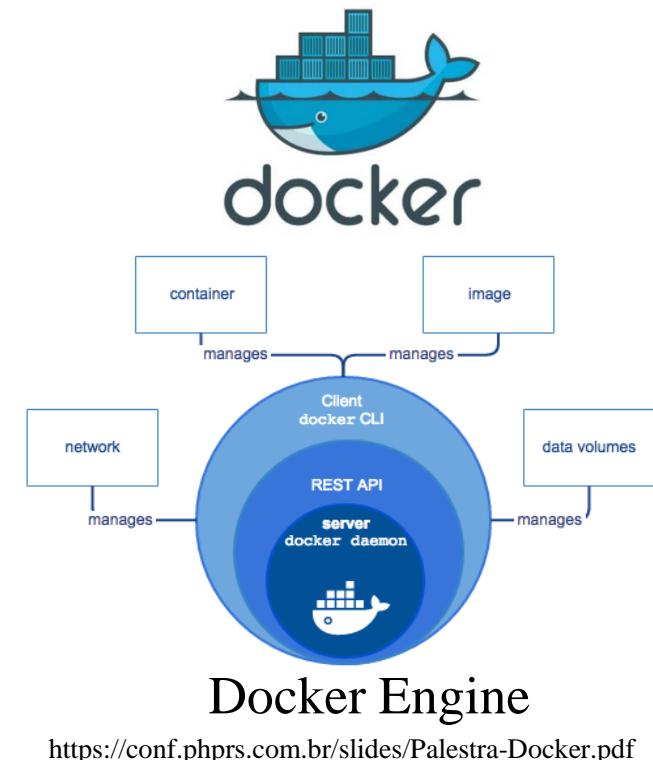
Containers are isolated, but share OS and, where appropriate, bins/libraries



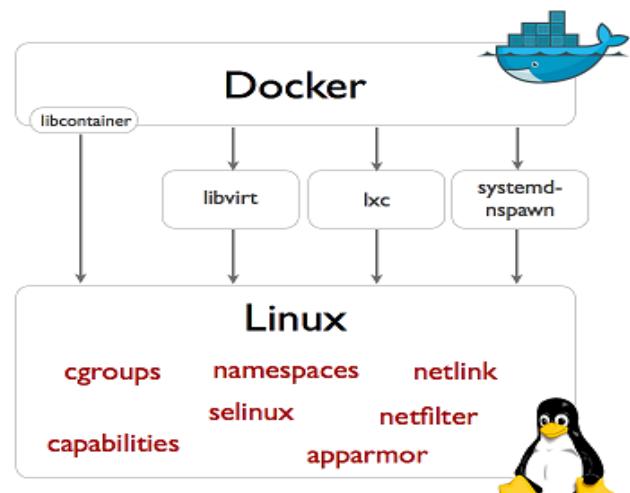
# Docker: Light-weight Process (Application) Container



[https://ecs.victoria.ac.nz/foswiki/pub/Courses/NWEN406\\_2018T1/LectureSchedule/NWEN406-presentation-summary-10.pdf](https://ecs.victoria.ac.nz/foswiki/pub/Courses/NWEN406_2018T1/LectureSchedule/NWEN406-presentation-summary-10.pdf)



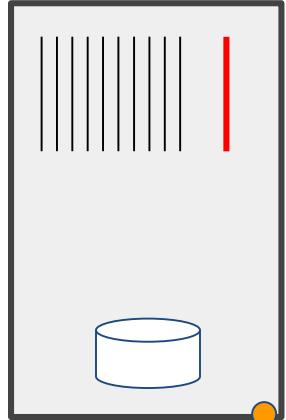
<https://conf.phprs.com.br/slides/Palestra-Docker.pdf>



<https://www.cnblogs.com/lijuanhu321/archive/2018/07/15/9313159.html>

# Containers

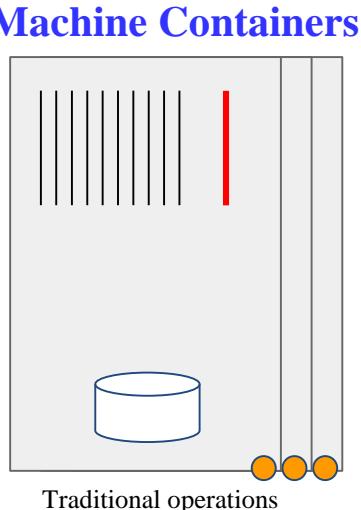
## Virtual Machines



Traditional operations



KVM

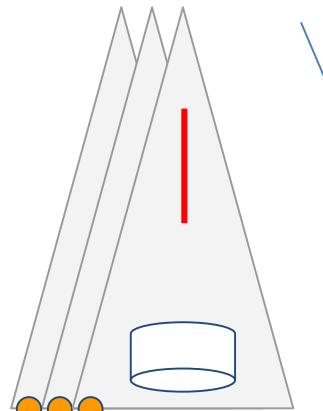


Traditional operations

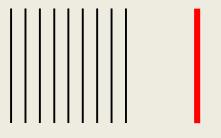


LXD

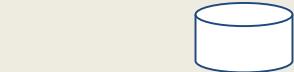
## Process Containers



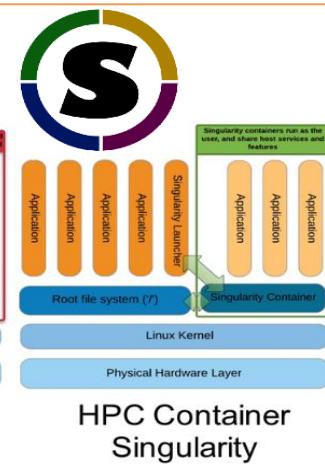
Mesos / Kubernetes / Swarm



Host Linux Filesystem

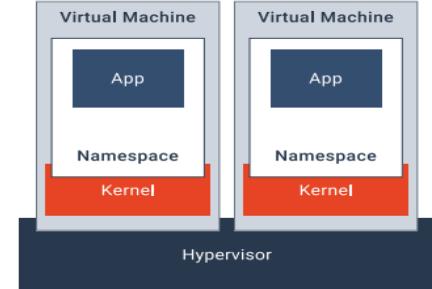


## HPC-targeted Containers



<http://technodocbox.com/71051449-Windows/Singularity-for-gpu-and-deep-learning.html>

<http://ju.outofmemory.cn/entry/337394>

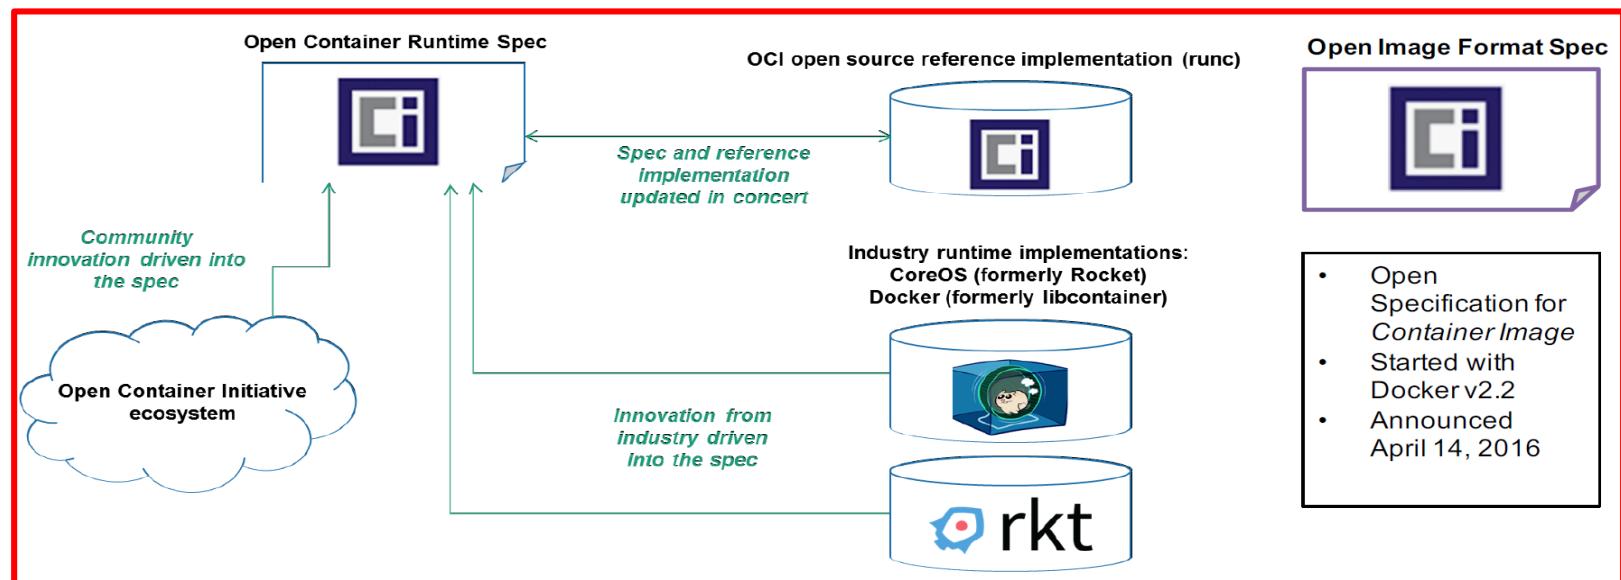
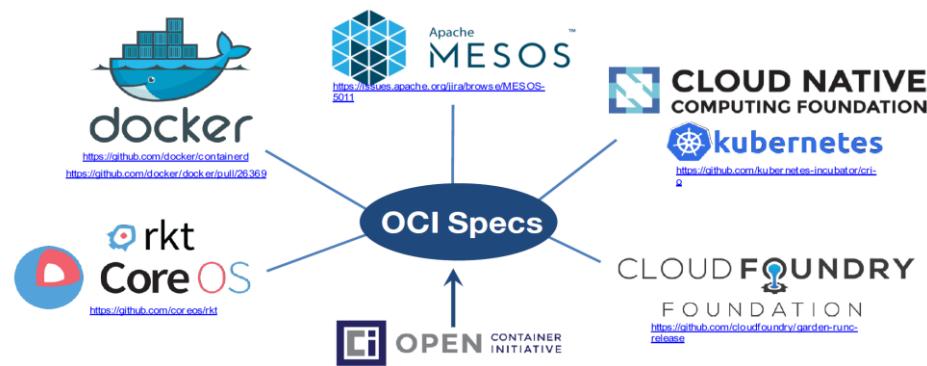


## Sandbox-isolated Containers

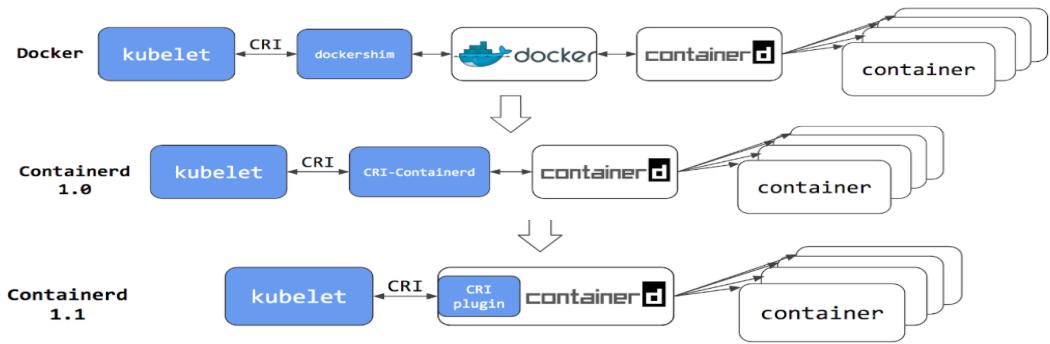
# OCI (Open Container Initiative)



To encapsulate a software component and all of its dependencies in a format that is self-describing and portable, so that any compliant runtime can run it without extra dependencies, regardless of the underlying machine and contents of the container.



# CRI(Container Runtime Interface): Containerd, rkt, CRI-O, ...



**Containerd**: An industry-standard container runtime with an emphasis on simplicity, robustness and portability.



**rkt**: CRI designed for composability, security, and speed.

**CRI-O**: OCI-based implementation of Kubernetes CRI; vs cri-containerd.

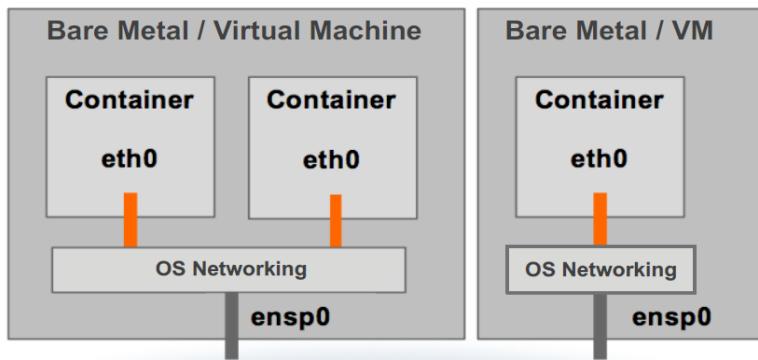


CRI-O: OCI-based  
Kubernetes Runtime

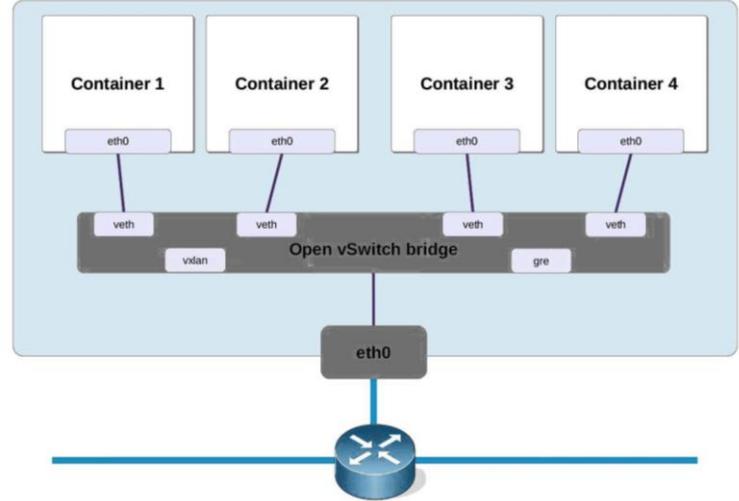
<https://kubernetes.io/blog/2018/05/24/kubernetes-containerd-integration-goes-ga/>

# p+v+c Harmonized Networking

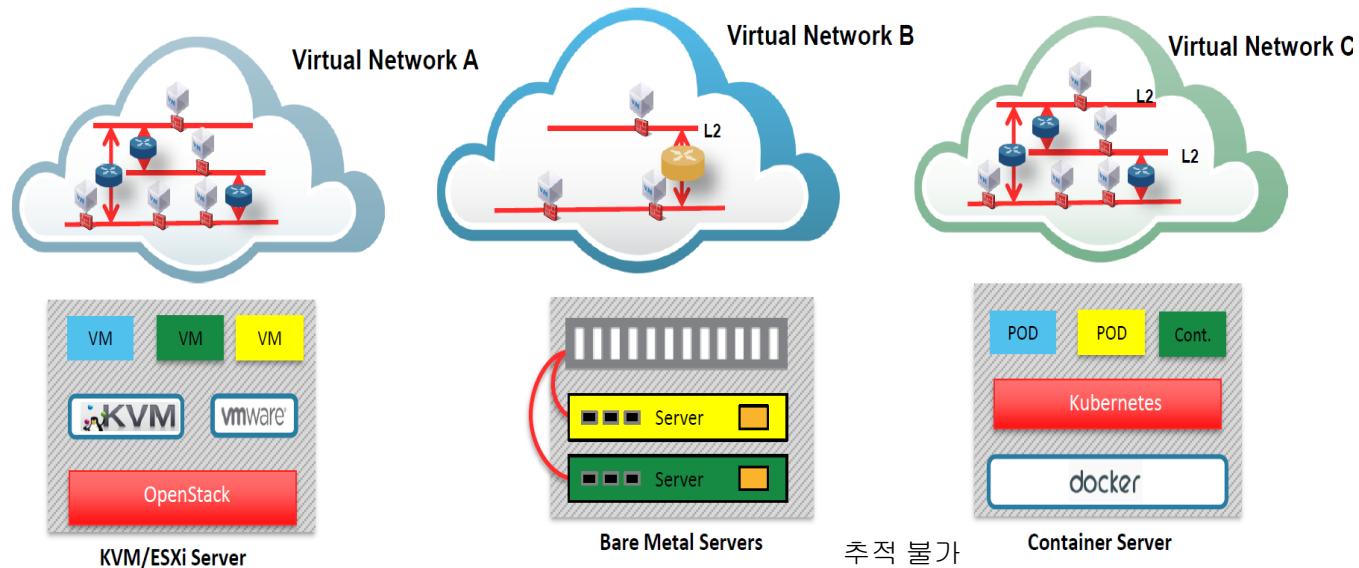
추적 불가



추적 불가



**p(Baremetal) + v(Virtual Machine) + c(Container) Harmonization**



추적 불가

추적 불가



Gwangju Institute of  
Science & Technology



Thank you!

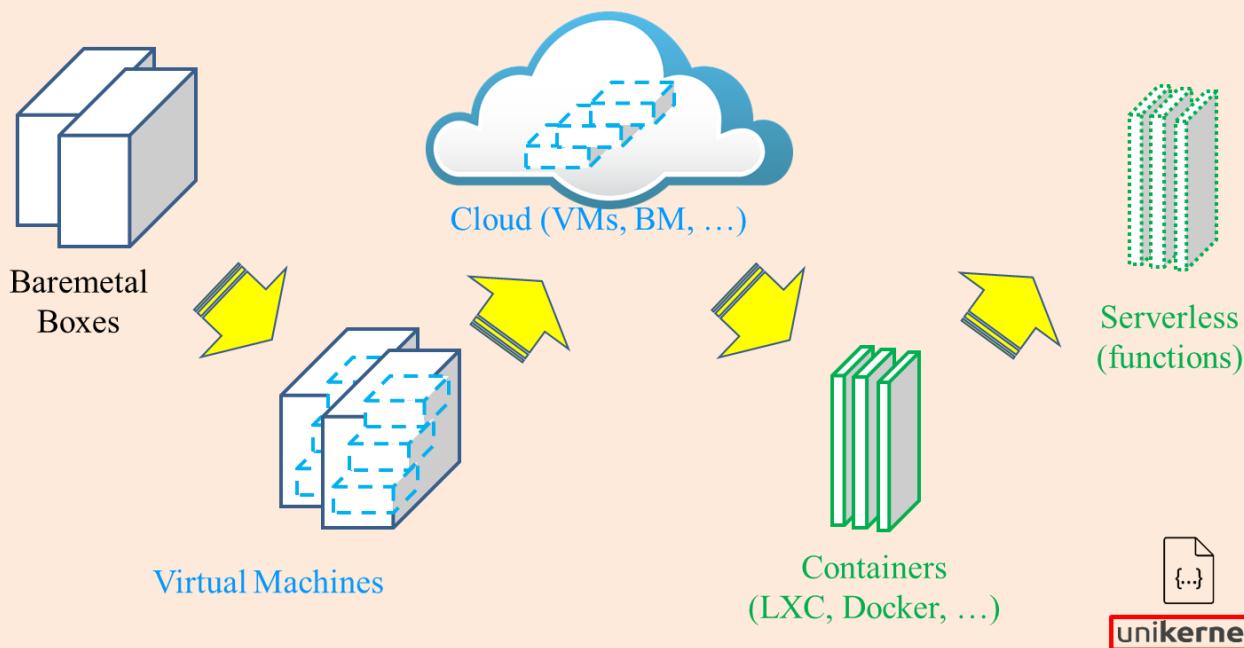
Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>

# Chapter 4:

# Computer System

# Cloud-native Computing



# Chapter Keywords

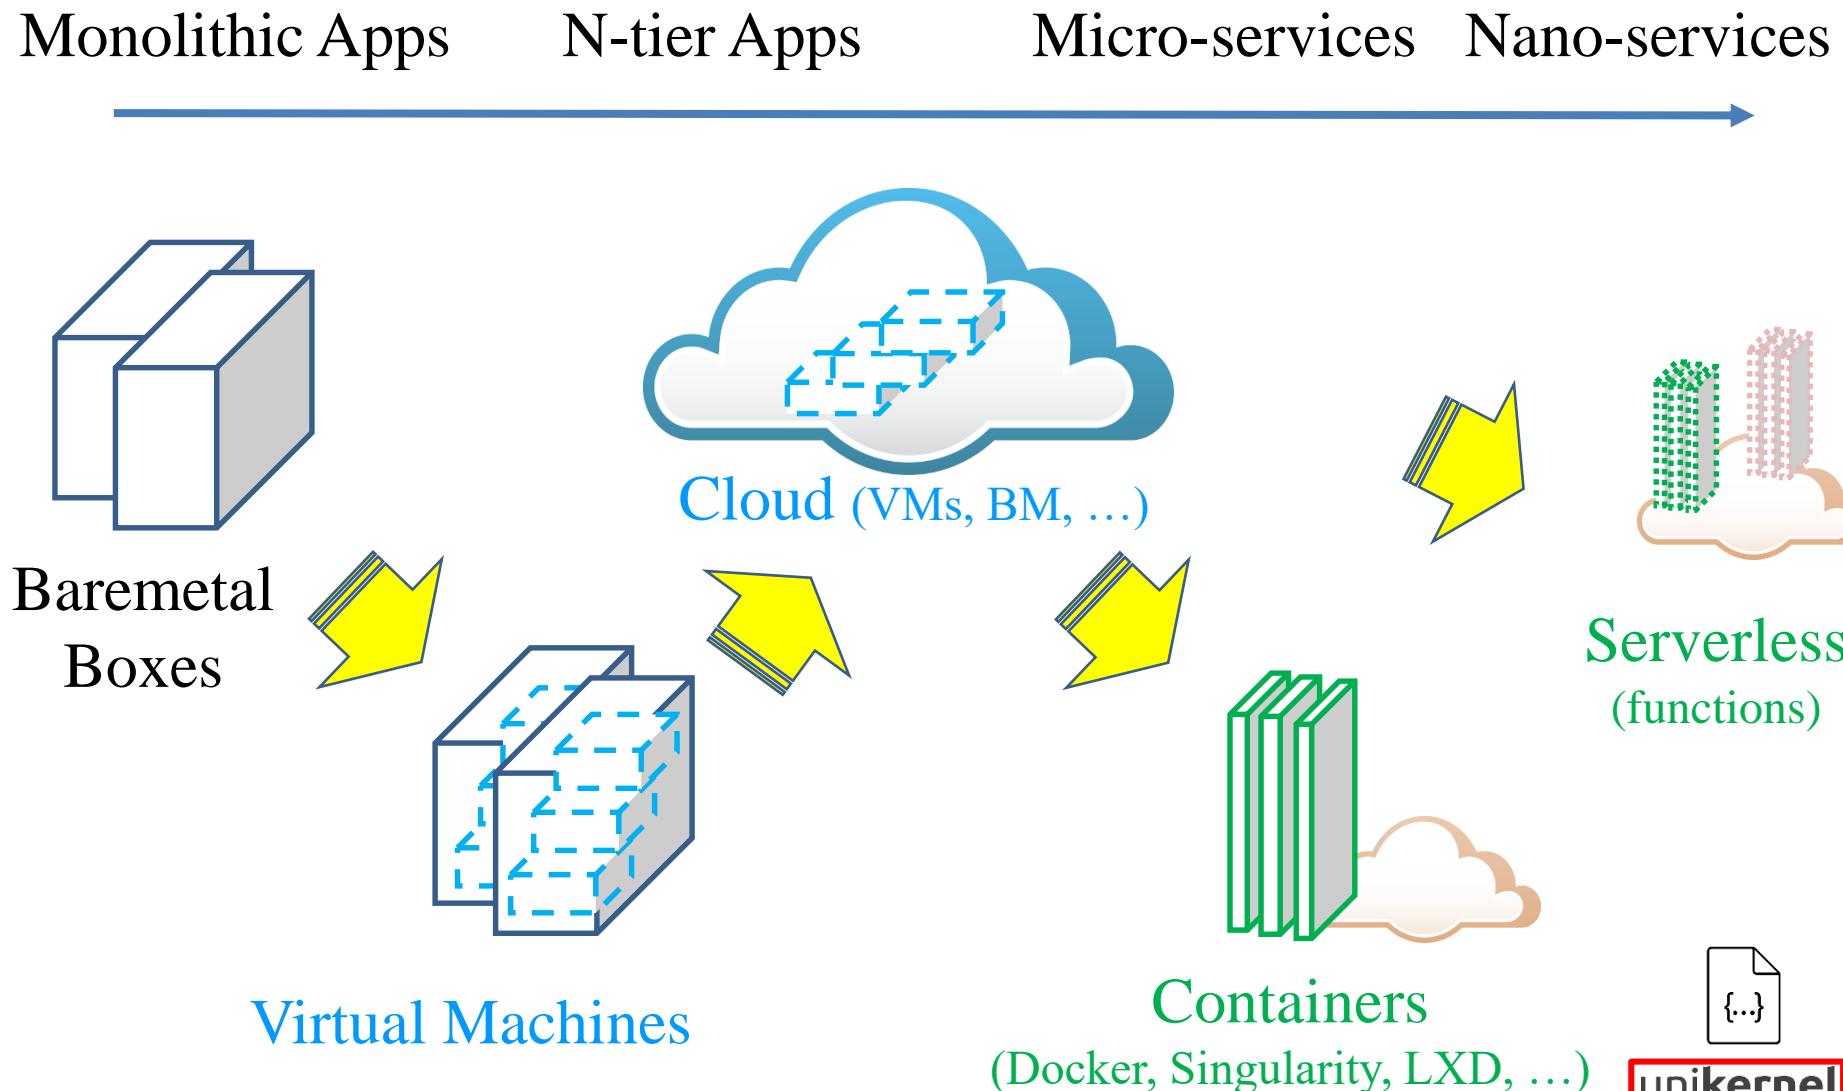
Cloud-native  
Computing

Serverless  
Computing

Orchestration

Open Interface

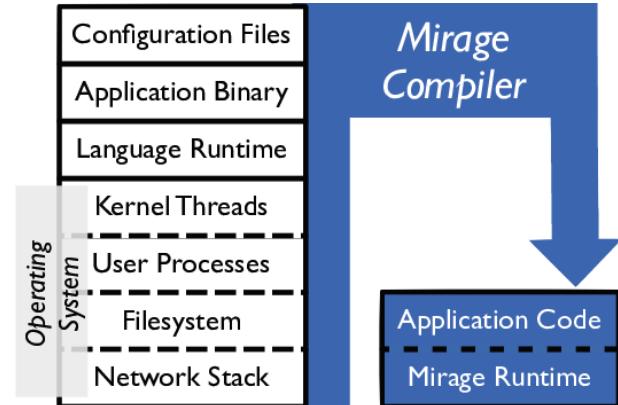
# From Baremetal to Containers / Serverless



{...}

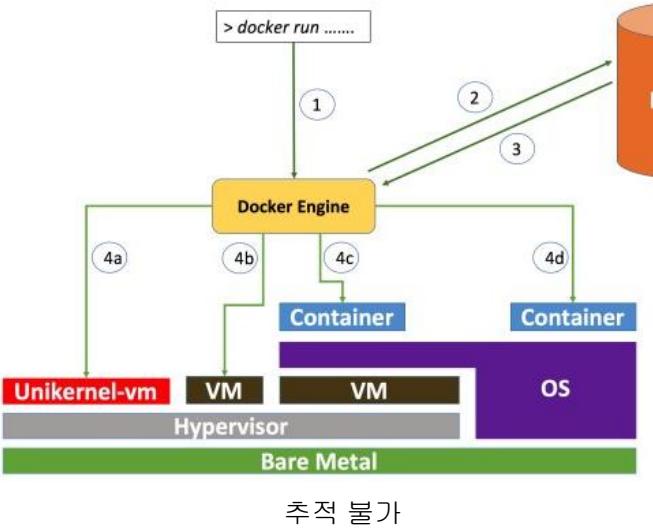
# unikernel

A specialized, single address space machine image constructed by using **library operating systems**.

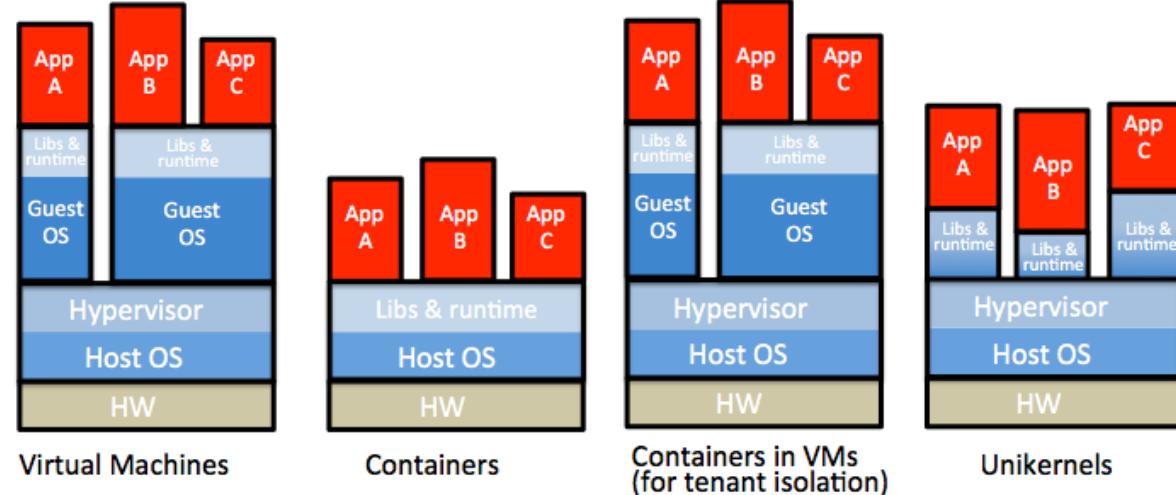


<https://manuzhang.github.io/2016/01/28/unikernel.html>

## unikernel



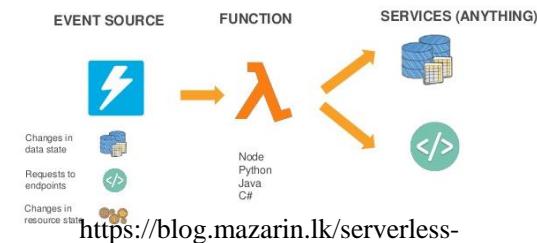
<https://www.industrynetwork.nl/tag/security/>



# Serverless (FaaS) computing



AWS Lambda



Event-driven,  
Continuous-scaled,  
Pay-by-use

<https://medium.com/horizon-four/lets-talk-about-serverless-que-tal-na-azure-bc307f11952a>



Azure Functions



Apache  
OpenWhisk

<https://www.smoothterminal.com/articles/exploding-rails-knative-k8s-serverless-elixir-qol-js-fundamentals-stable-apis>



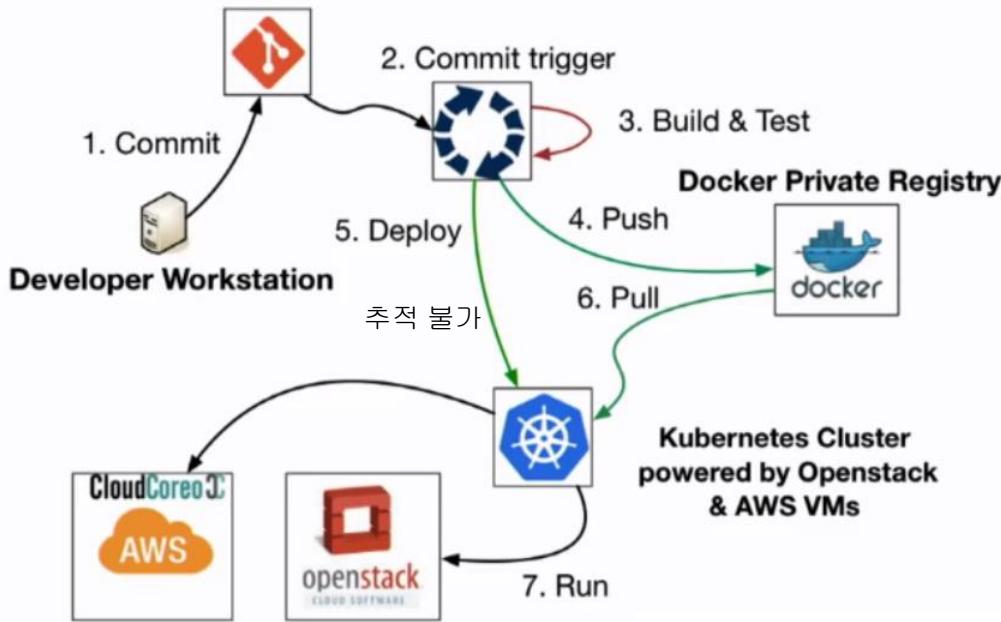
Knative



Google Cloud Functions

# Container-leveraged Microservices Architecture for Agile & Automated Service Realization

## Microservice Pipeline



**DEV OPS**

Automation Service

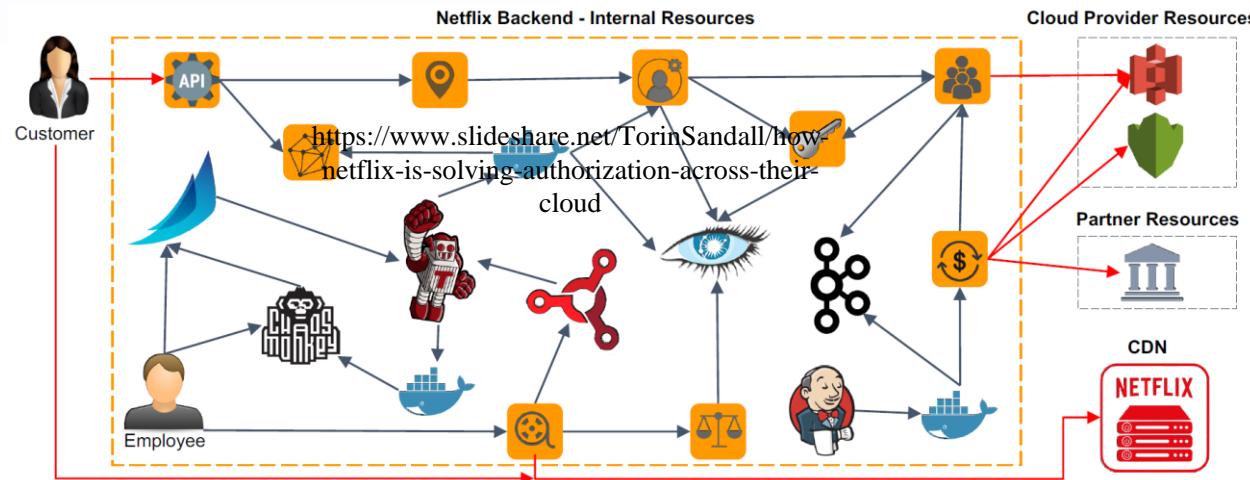
[https://static.rainfocus.com/oracle/pow18/sess/1525990149968001qZyNPF/DEV5564\\_15405732841760016eb0.pdf](https://static.rainfocus.com/oracle/pow18/sess/1525990149968001qZyNPF/DEV5564_15405732841760016eb0.pdf)

GIST

Container-leveraged  
Orchestration of  
Service Composition

&

Dynamically-scaled  
Resource Pooling



# Container Orchestration Early View (2015): Docker Containers + Mesos Scheduler + Kubernetes Orchestration

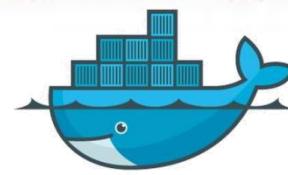
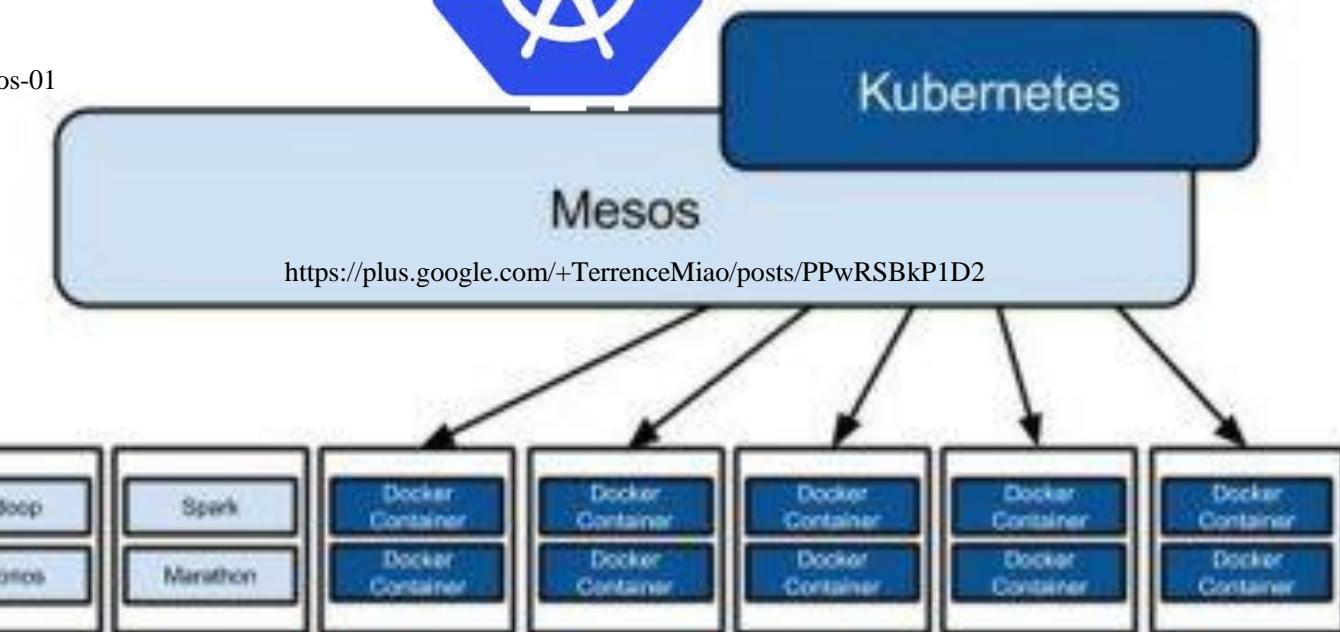
<https://en.wikipedia.org/wiki/Kubernetes>



<https://www.brandeps.com/logo/M/Mesos-01>

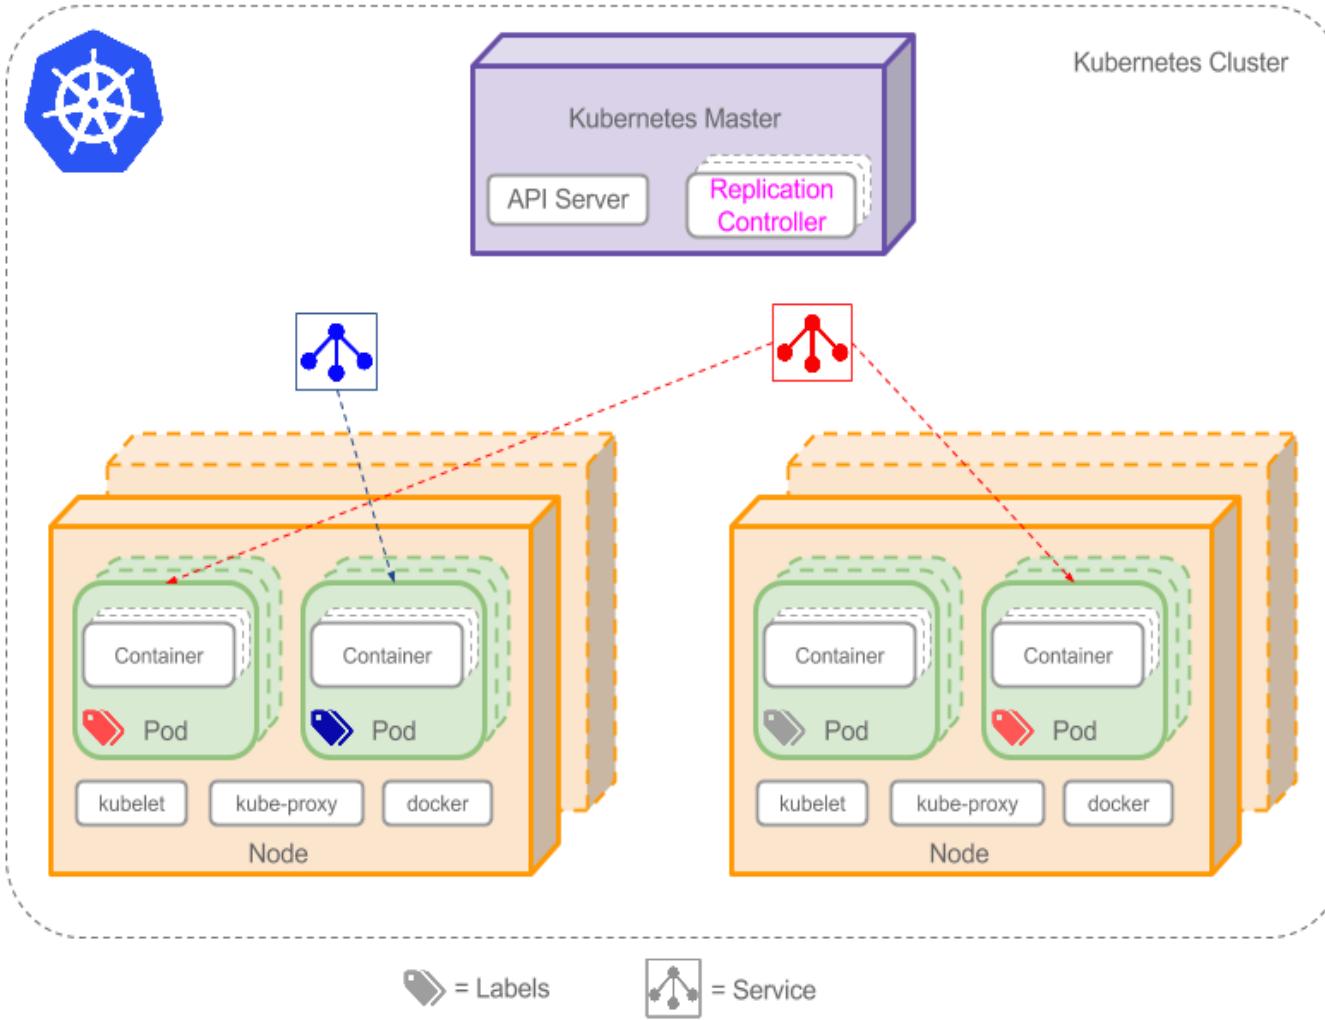


**compute  
nodes**  
(machines or VMs)



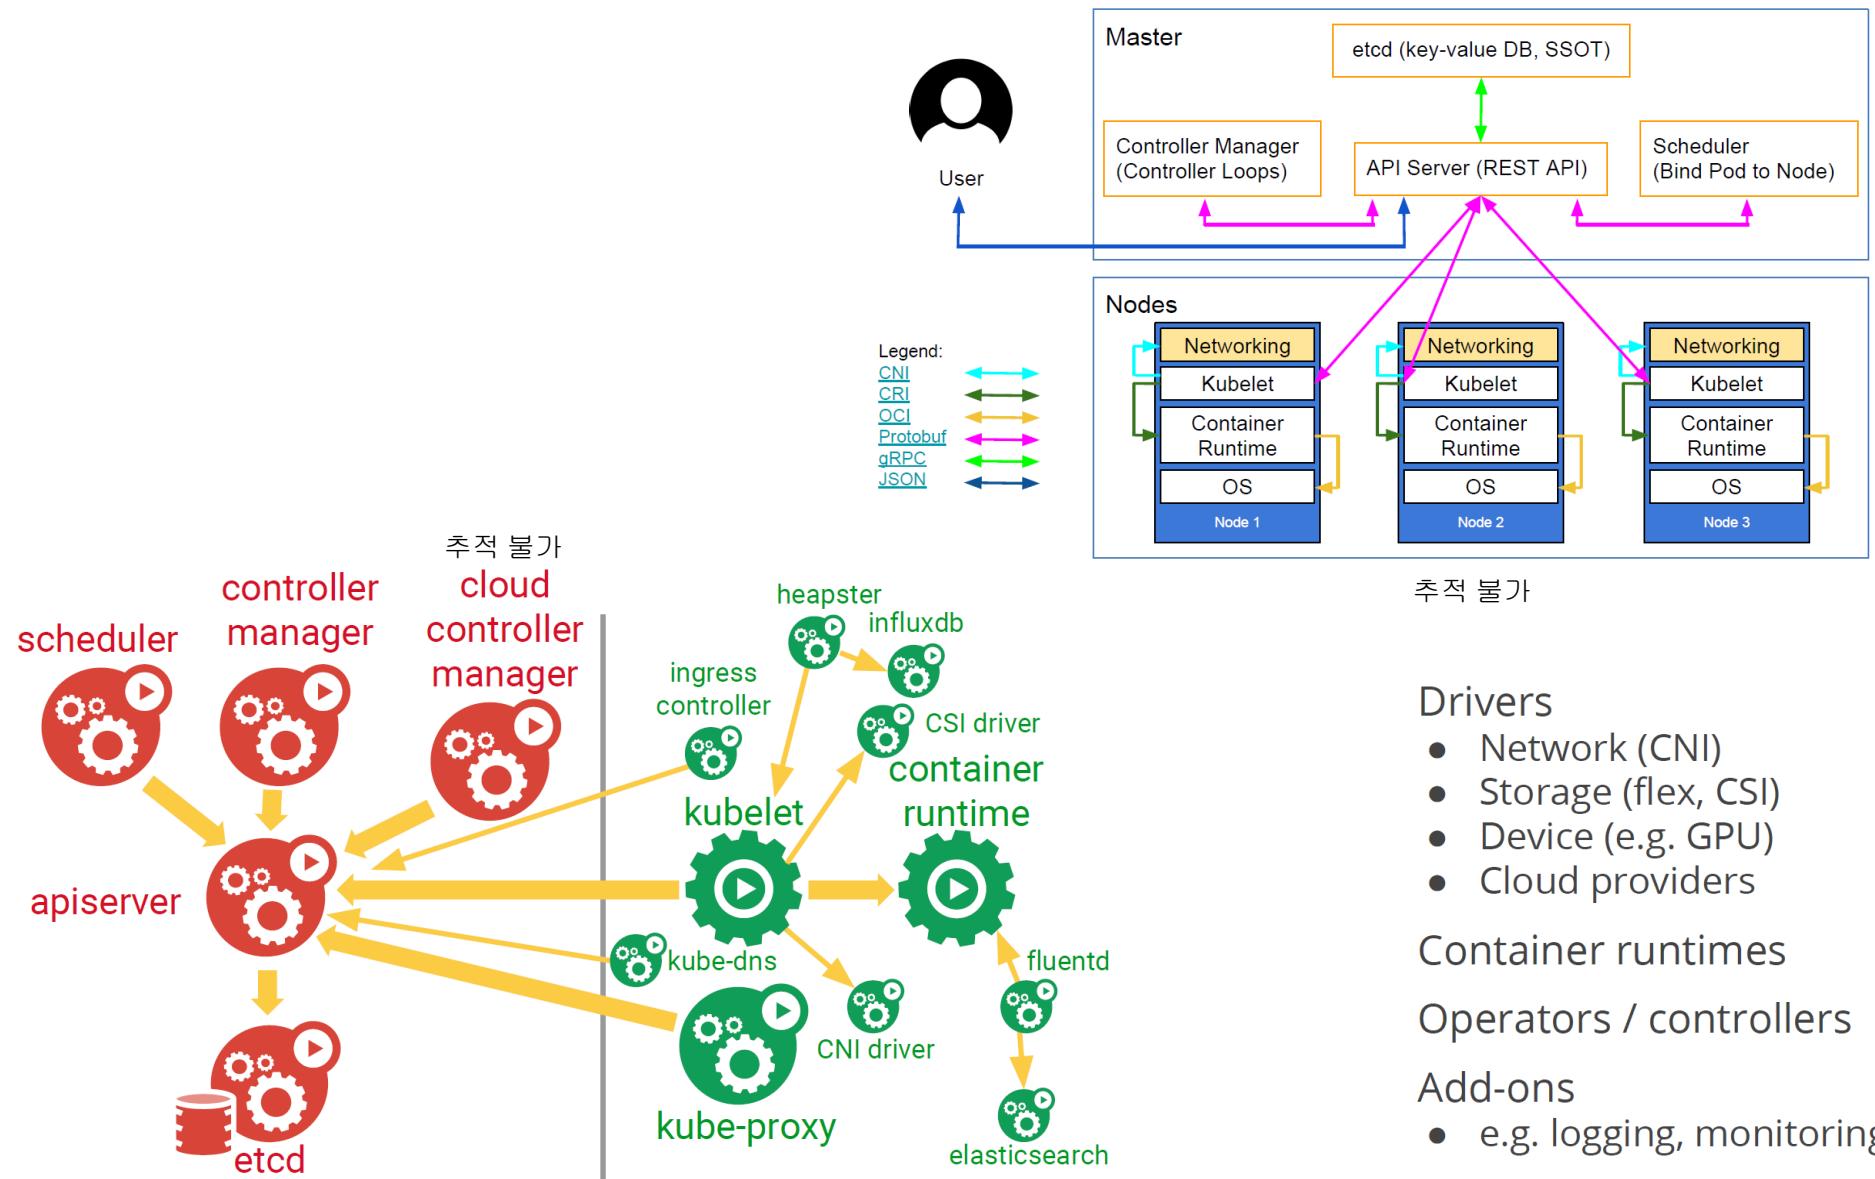


# Kubernetes for Large-scale Container Orchestration: Architecture



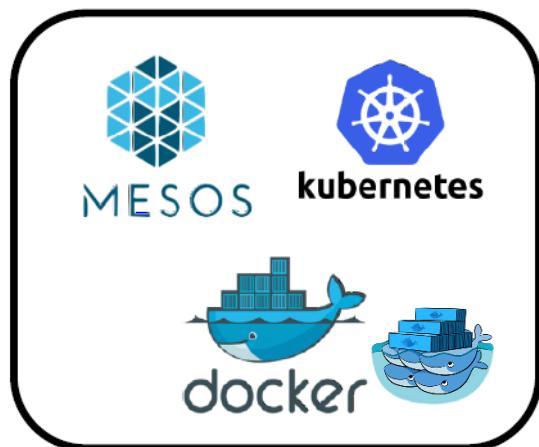
**Kubernetes (k8s)**  
is an open-source system for automating deployment, scaling, and management of containerized applications.

# Kubernetes for Large-scale Container Orchestration: Operation



# Container Platforms & Orchestration Tools (2017)

## Container Orchestration Engine



<https://docs.docker.com/samples/library/swarm/>

## Platform-as-a-Service (PaaS)



## Container-as-a-Service (CaaS)





Kubernetes  
Orchestration



Prometheus  
Monitoring



OpenTracing  
Distributed Tracing API



Fluentd  
Logging



linkerd  
Service Mesh



gRPC  
Remote Procedure Call



CoreDNS  
Service Discovery



containerd  
Container Runtime



rkt  
Container Runtime



CNI  
Networking API



Envoy  
Service Mesh



Jaeger  
Distributed Tracing



Notary  
Security



TUF  
Software Update Spec



CLOUD NATIVE  
COMPUTING FOUNDATION



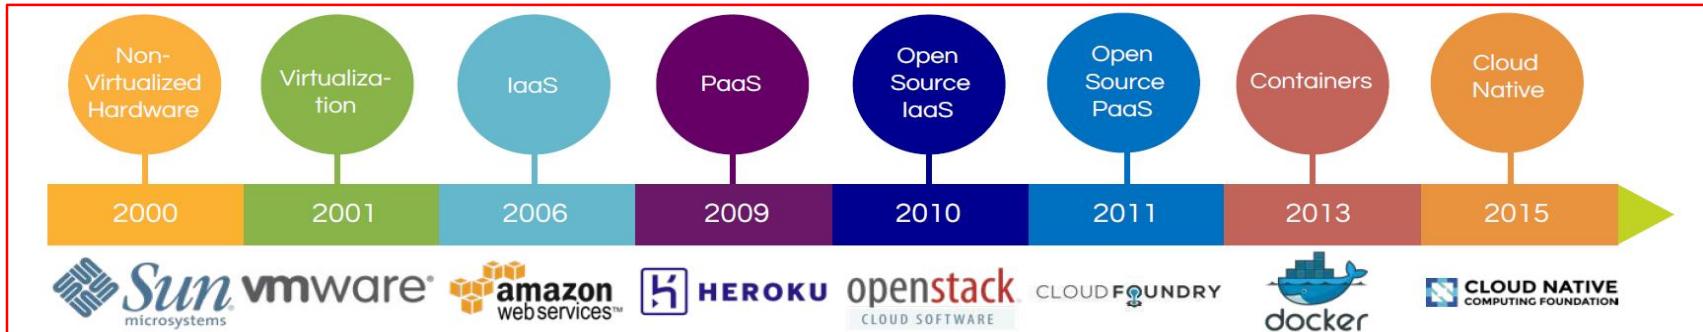
# CNCF & K8S



<https://blog.netsil.com/kubernetes-kontainers-and-kubernetes-native-in-2018-f5d6fb851ad0>



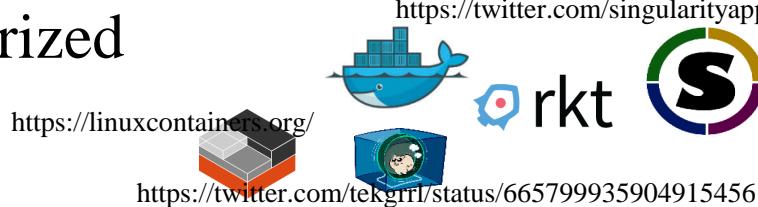
# Cloud-native Computing



<https://opensource.com/article/18/2/how-kubernetes-became-solution-migrating-legacy-applications>

Cloud native computing uses an open source software stack to be:

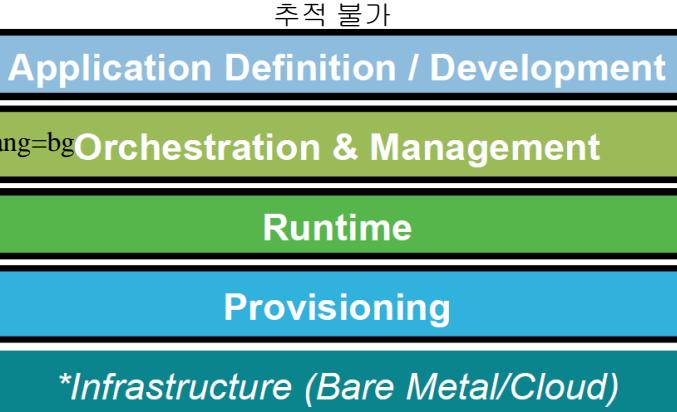
- Containerized



- Dynamically orchestrated



- Microservices oriented



microservices + automation +  
DevOps



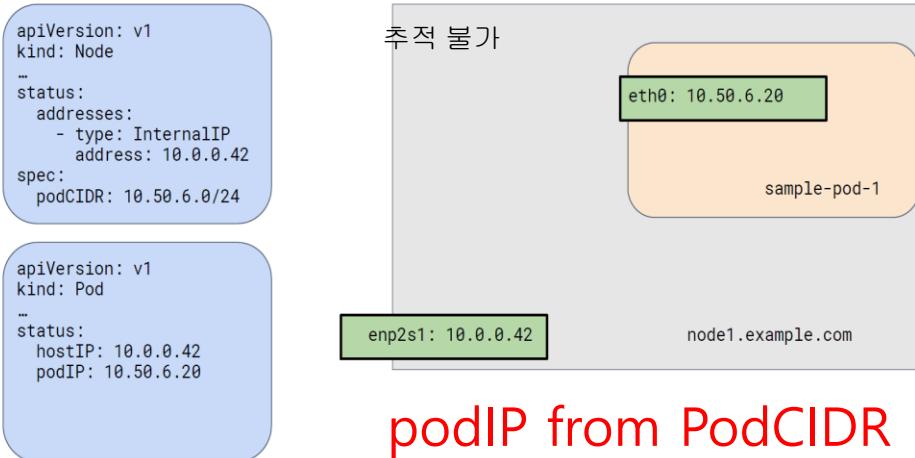
Agility of software teams +  
Resilience of applications.

# CNI (Container Network Interface)



CNI is at least three different things:

- A vendor-neutral protocol used by container runtimes to make request to networking providers (Not just for kubernetes!)
- A set of commonly used network plugins maintained by the community
- A “kubelet network plugin” (e.g. “--network-plugin=cni” on the command line)



## CNI Plugins for

- **Connectivity:** Create a “veth” (a point-to-point virtual tunnel) pair → Move one end of the pair in to the container’s namespace → Configure an ip and route in the container’s namespace
- **Reachability:** Make every PodCIDR reachable from every Node (by announcing dynamic routes to some peers) with a per-host daemon that programs the network with learned routes; Programming via 1) Overlay networks (Flannel, Weave, Calico), 2) Cloud provider APIs (cp-azure, Flannel), 3) Routing protocols (Calico, Romana)

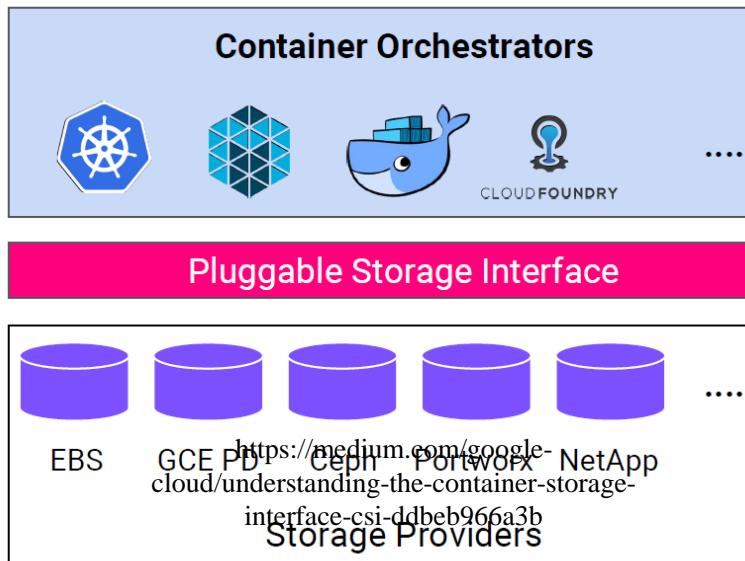
# CSI (Container Storage Interface)



Problems with existing storage interfaces: CLI-based interface, Lack of idempotency on APIs, In-tree interface, Tightly coupled with an implementation, Too heavyweight

→ **CSI** for Interoperability + Vendor neutral + Focus on specification + Control plane only + Keep it simple and boring

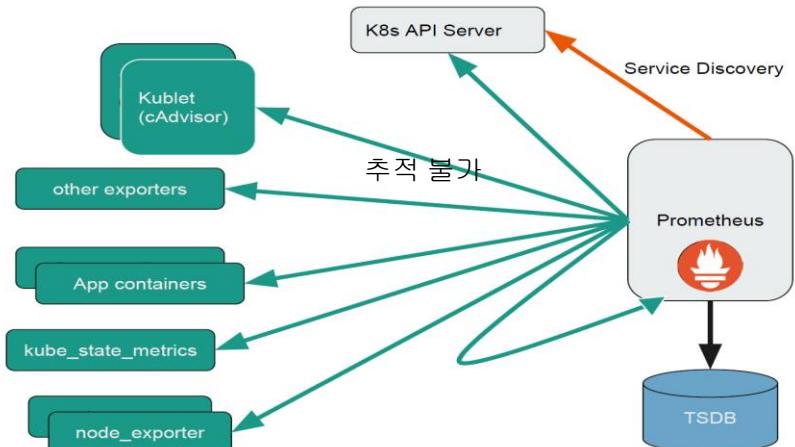
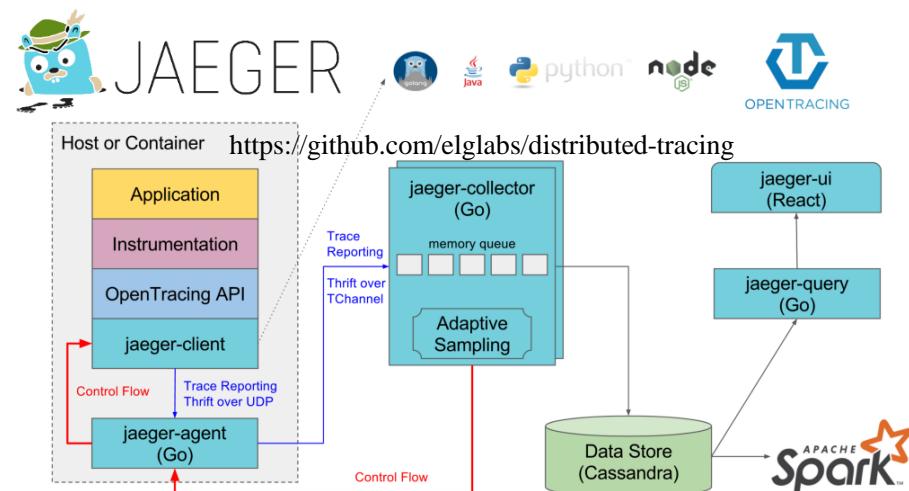
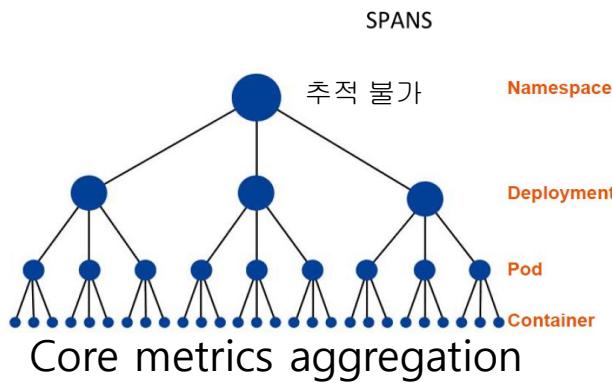
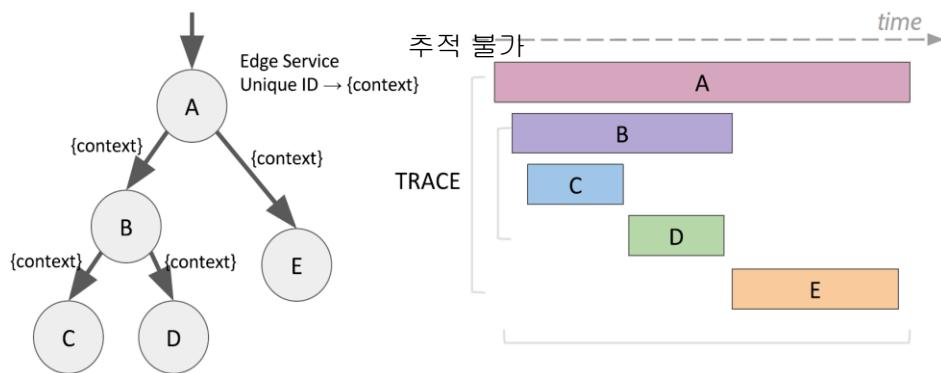
- Docker volume plugins [[link](#)]
- K8s FlexVolume [[link](#)]
- K8s In-tree volume plugins [[link](#)]
- Libstorage storage drivers [[link](#)]
- OpenSDS volume drivers [[link](#)]



# Unified Visibility (Monitoring/Tracing) for Kubernetes-orchestrated Services



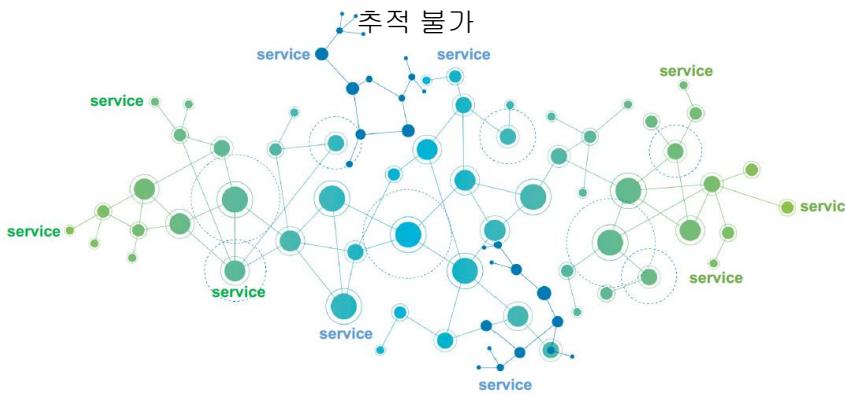
**OpenTracing**: Vendor-neutral / Cross-language Instrumentation API for Context Propagation + Distributed tracing + Contextualized logging/metrics



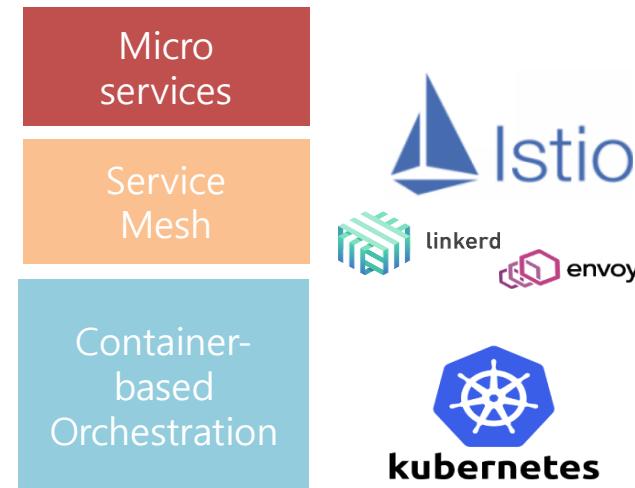
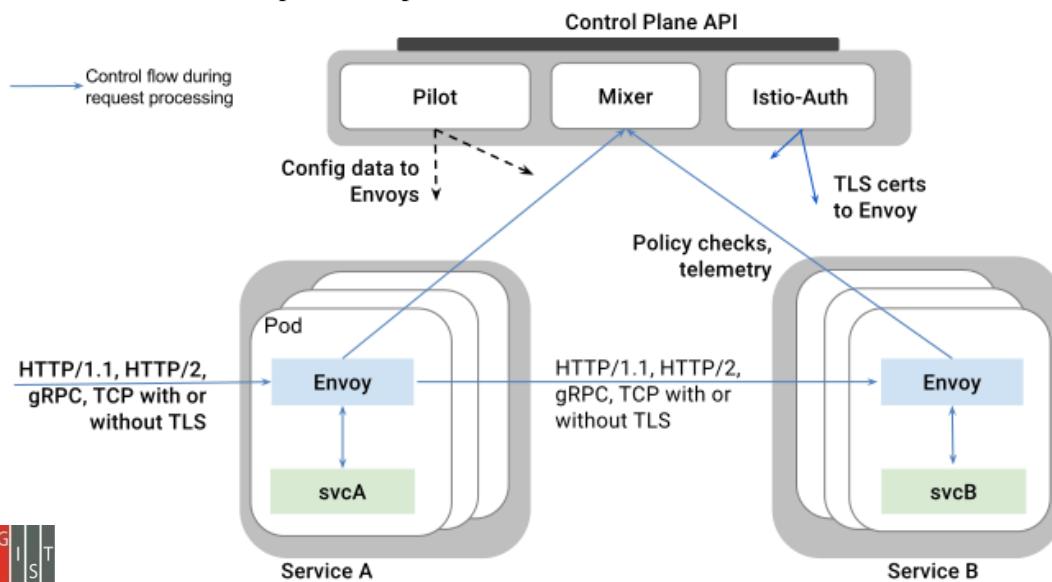
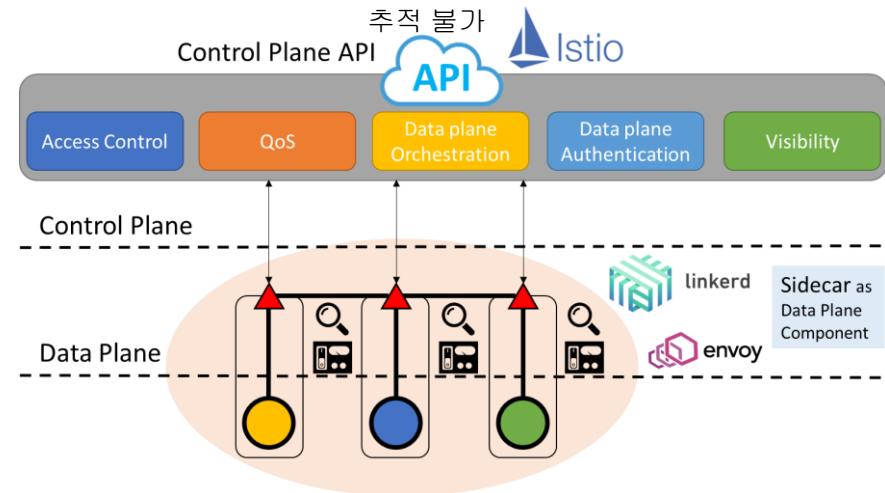
Visibility sources in Kubernetes: Node via the node\_exporter + Container metrics via the kubelet and cAdvisor + Kubernetes API server + etcd + Derived metrics via kube-state-metrics

# Service Mesh: Thin-layer for Reliable Microservices with Traffic Monitoring & fine-grained QoS/Access Control

To solve common issues with routing/re-routing for graceful degradation as services fail, secure inter-service communication and rate limiting between services.



<https://developer.ibm.com/courses/how-istio-works-dwc024/>





Gwangju Institute of  
Science & Technology



Thank you!

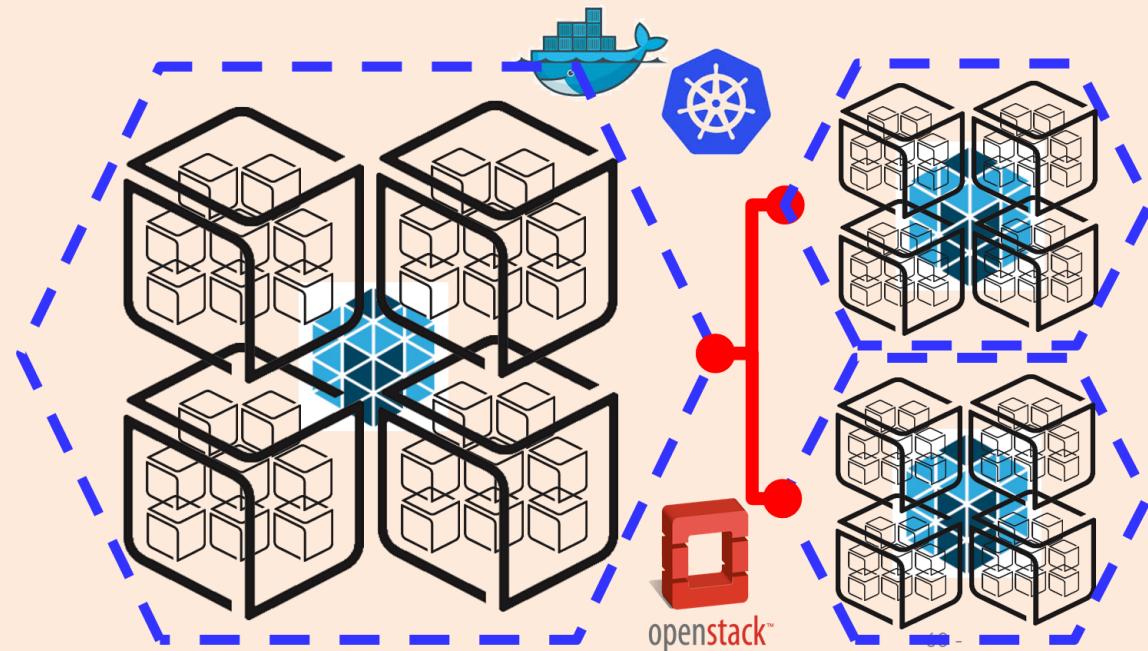
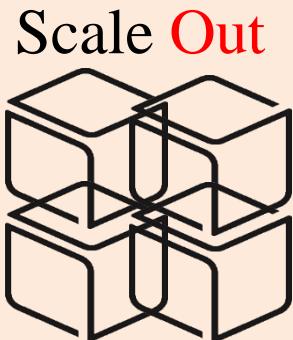
Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>

# Chapter 5:

# Computer System

## Clusters for HPC/BigData & AI



# Chapter Keywords

Scale-Out/Scale-Up

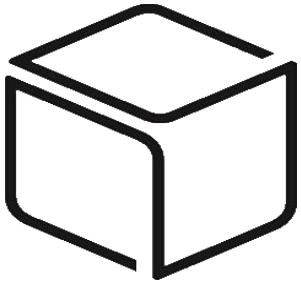
Cluster / Grid / Cloud

HPC Cluster

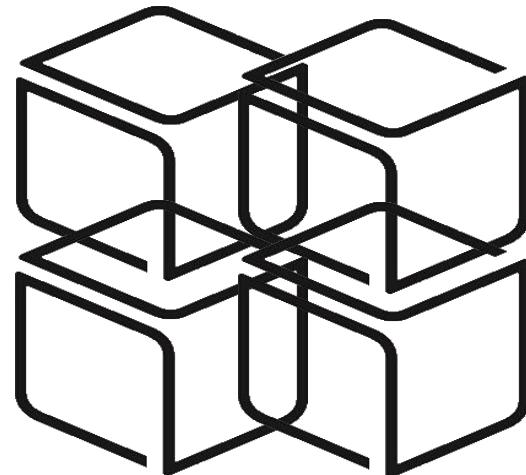
Storage/BigData Cluster

Data & AI Cluster

# Computer Systems: Scaling Options

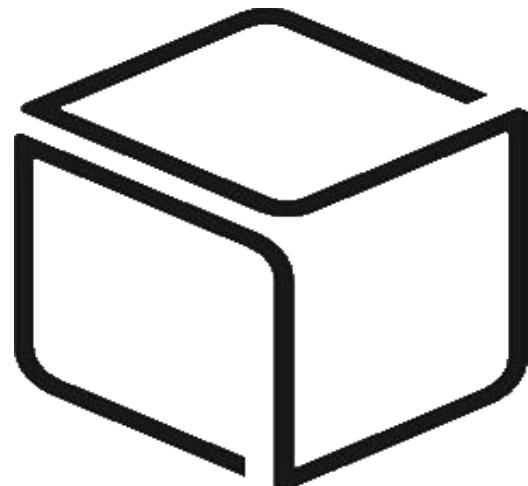


Scale Out



vs

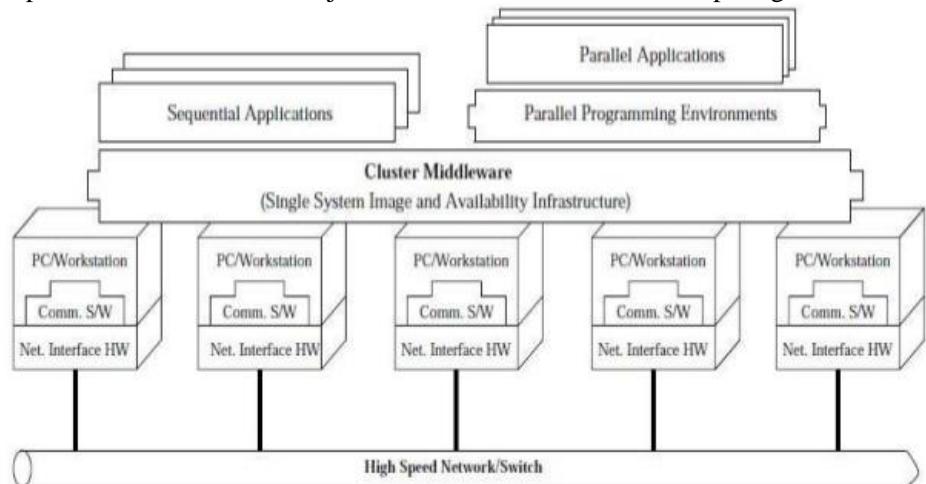
Scale Up



# Computer Systems: Cluster Computing

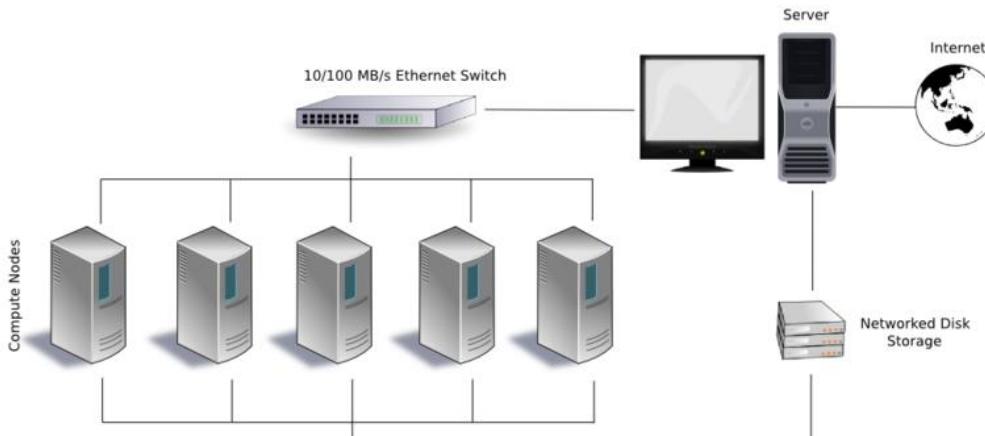
[Computer Cluster](#) (Wikipedia): A set of loosely or tightly connected computers that work together so that, in many respects, they can be viewed as a single system. Unlike grid computers, computer clusters have each node set to perform the same task, controlled and scheduled by software.

<https://www.slideshare.net/anjalibhandari11011995/cluster-computing-53768699>



Cluster architecture

[http://www.ppcr.tj/documents/ADBTAA8090Inception\\_Report.pdf](http://www.ppcr.tj/documents/ADBTAA8090Inception_Report.pdf)

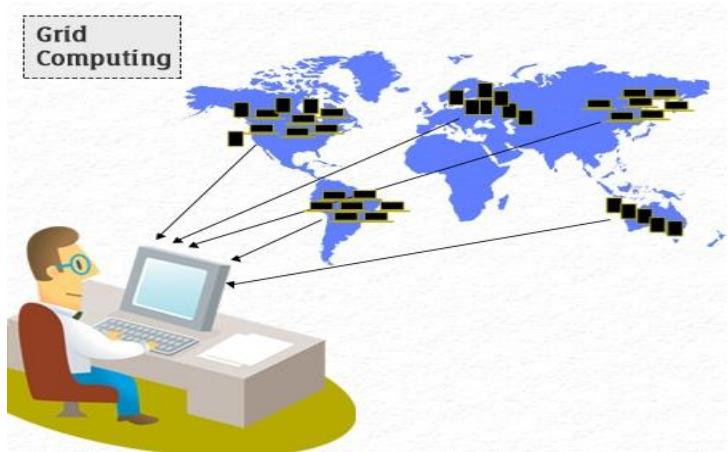


[Beowulf cluster](#) (1994 ~): A computer cluster with identical/commodity computers for high-performance processing, which are networked into a small LAN with shared libraries and programs.

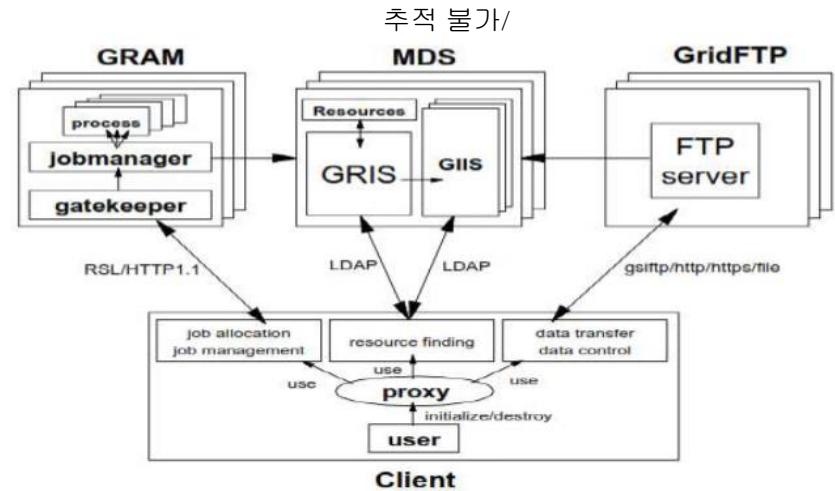
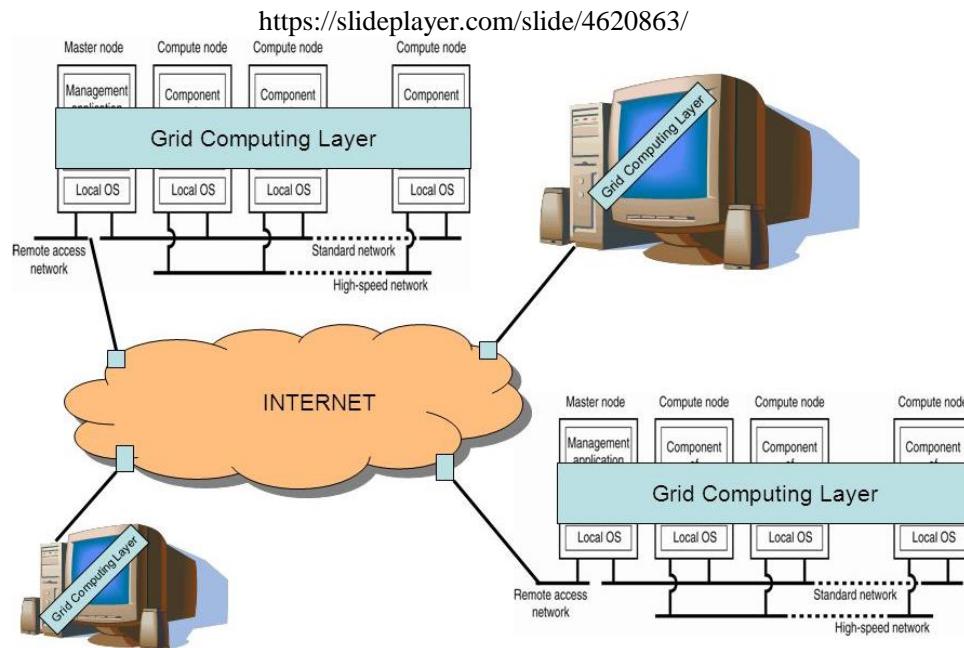
# Computer Systems: Grid Computing

**Grid computing:** The use of geographically distributed and heterogeneous computer resources to reach a common goal.

- Using distributed computers and resources collectively.
- Usually associated with geographically distributed computers and resources on a high-speed network.



<http://rhodyzekiel.blogspot.com/2017/05/grid-computing.html>

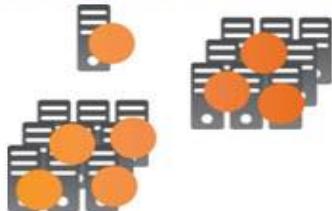


# Computer Systems - Evolution: Cluster / Grid / Cloud Computing

Scope of sharing

- HPC Cluster**
- Commodity HW
  - Compute / data intensive apps

Distributed Clusters



1992

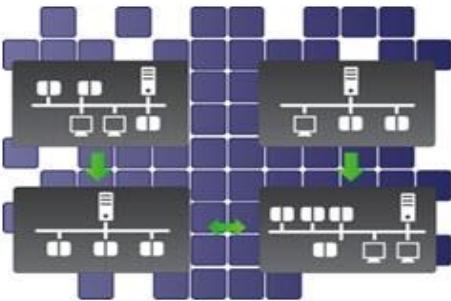
2002

2012

## Enterprise Grid

- Dynamic workload using static resources
- Policy-based scheduling

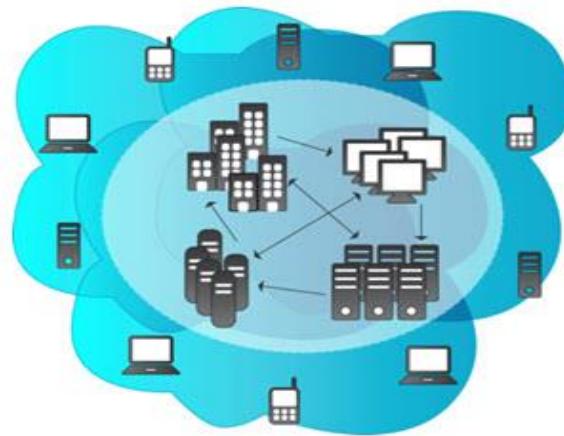
## Enterprise Grid



## Public & Private Cloud

- Web 2.0, HPC & bus apps
- Dynamic resource provisioning
- On demand, utility, self-service

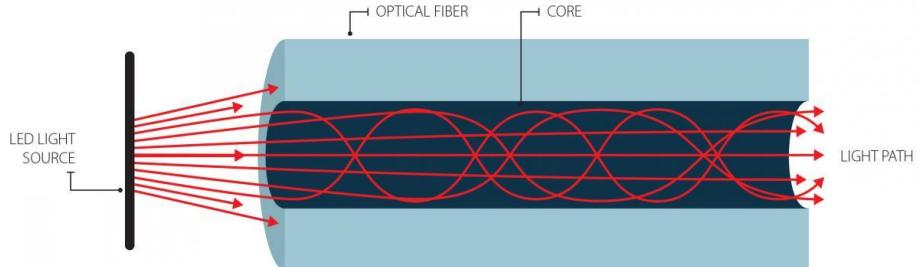
## Public & Private Cloud



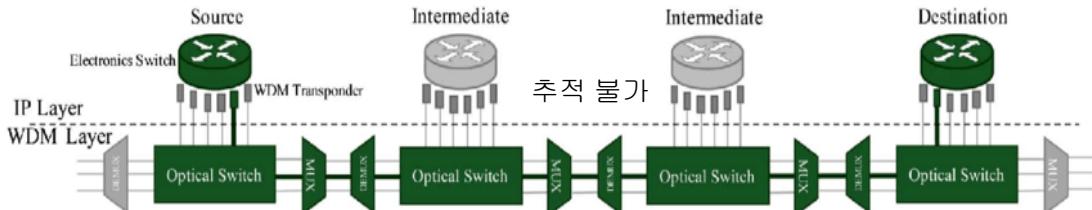
# Computer Systems - Evolution: Enabled by High-Performance Optical Networking

Cluster/Grid/Cloud  
Computing Enabled by  
Global-scale Lambda (Optical  
Light Path) Networking

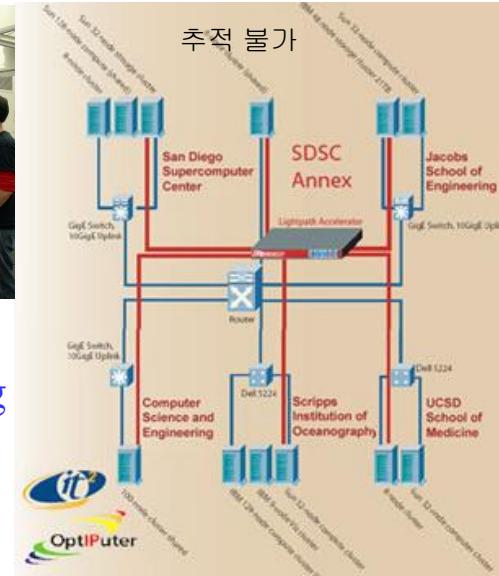
<http://www.fibre-systems.com/viewpoint/reducing-measurement-uncertainty-multimode-optical-fiber>



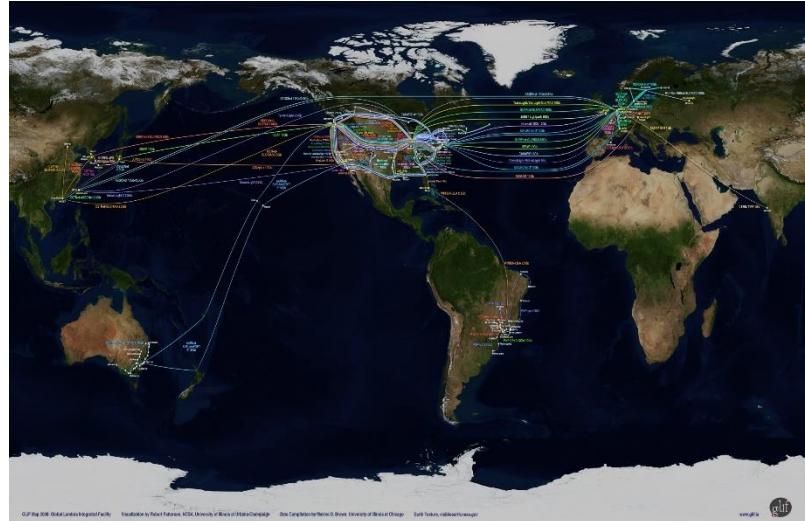
Light Path & Optical Networking



OptIPuter &  
SAGE Visual Sharing



<http://generic.wordpress.soton.ac.uk/sites/79/>

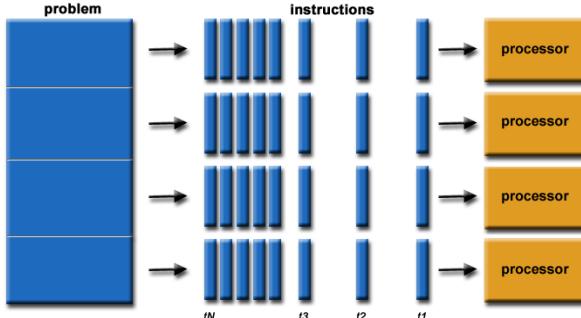
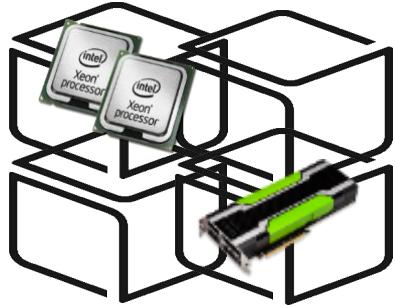




# Computer Systems:

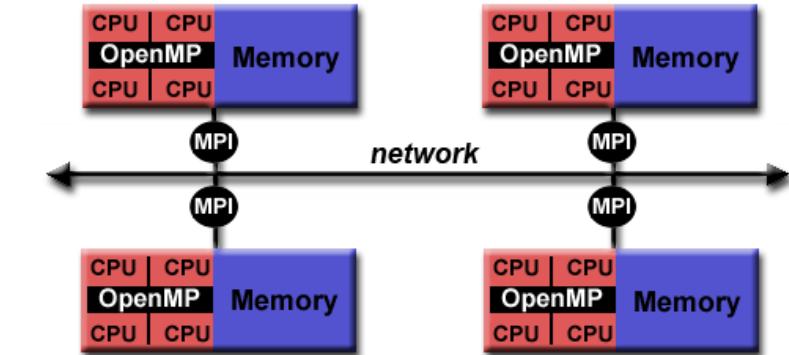
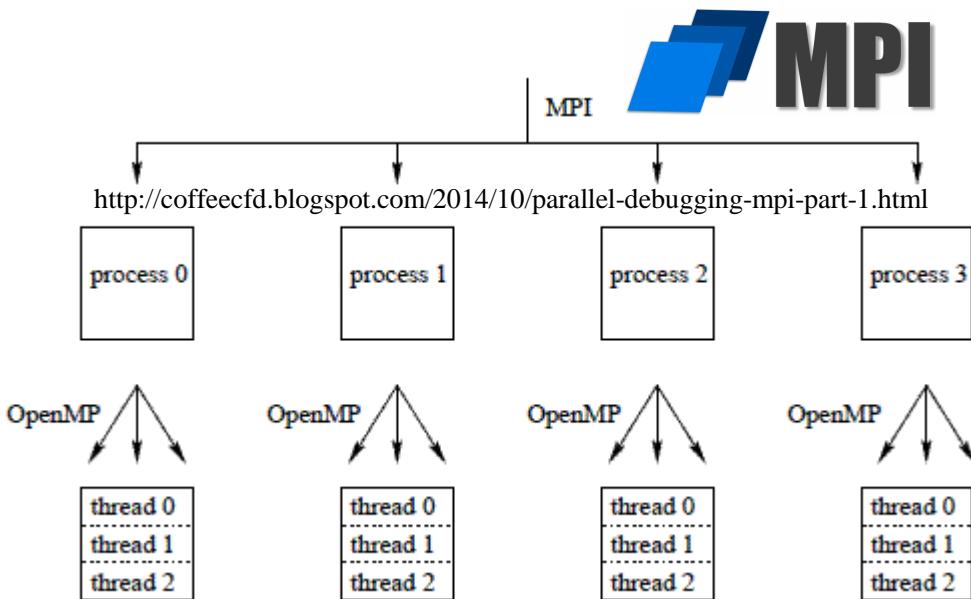
# HPC (High Performance Computing) Clustering

<https://eds-uga.github.io/cbio4835-fa18/slides/Lecture5.slides.pdf>



HPC with Parallel Computing

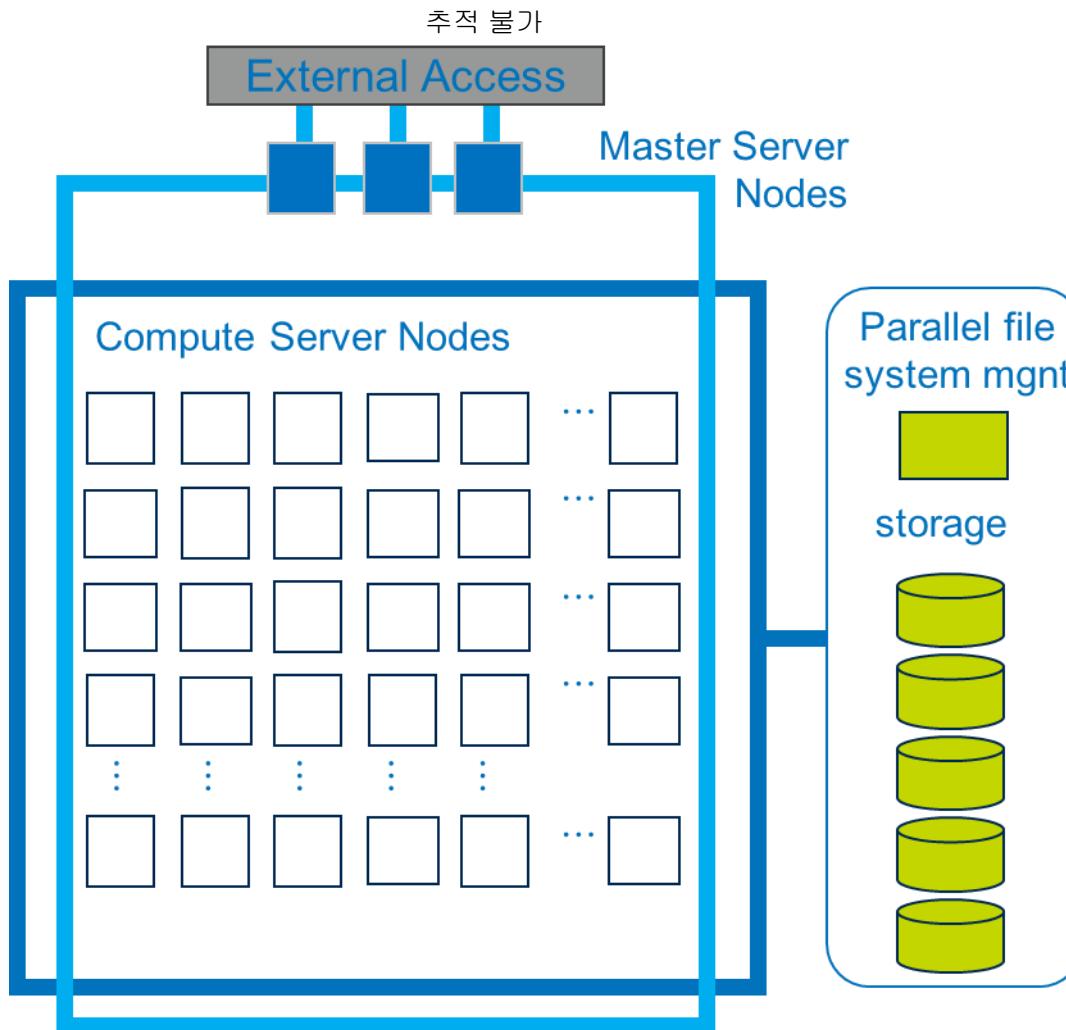
<https://events.prace-ri.eu/event/729/contribution/3/material/slides/0.pdf>



OpenMP

Parallel computing: MPI & OpenMP

# Computer Systems: HPC Clustering Requirement



**Powerful master nodes**

- manage the compute nodes
- optimize overall capacity



**Many fast compute nodes**

- number crunching, rendering, compiling, or other manipulation



**Low latency / high BW fabric**

- master to compute node comms
- inter-compute node communications



**Fast storage access**

- high performance data network
- high performance **parallel file system**

-mix of HDD and SSD (more SSDs with Big Data)

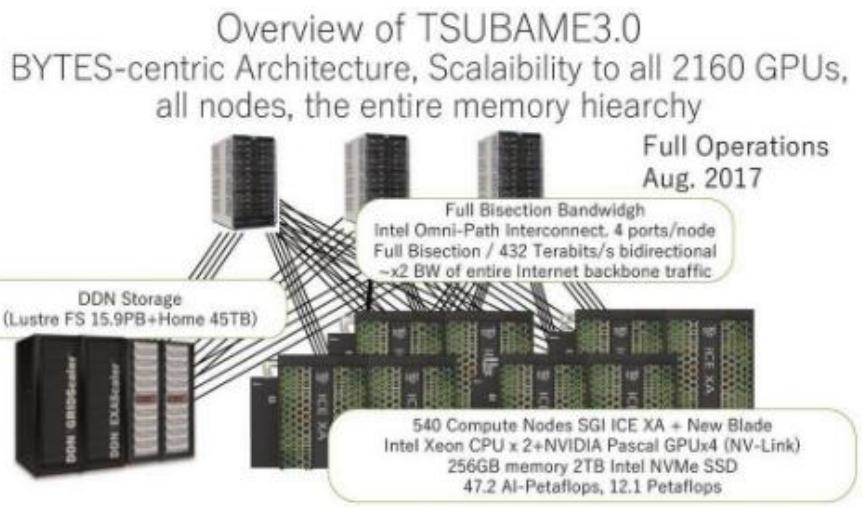
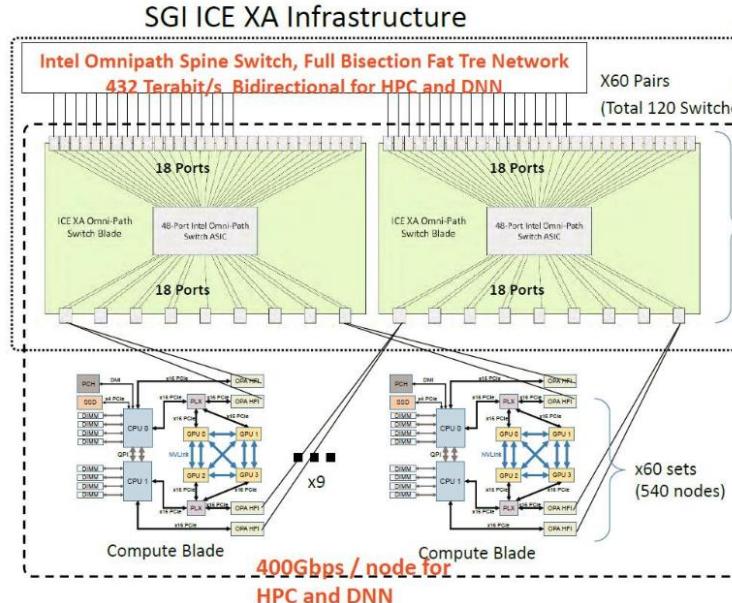
*AND...easy to deploy, manage, and use (e.g. apps compatibility)*

# Computer Systems: HPC Clustering for Super Computer



# TSUBAME3.0

**400Gbps / Box**  
(with Intel Omni Path Adapter)



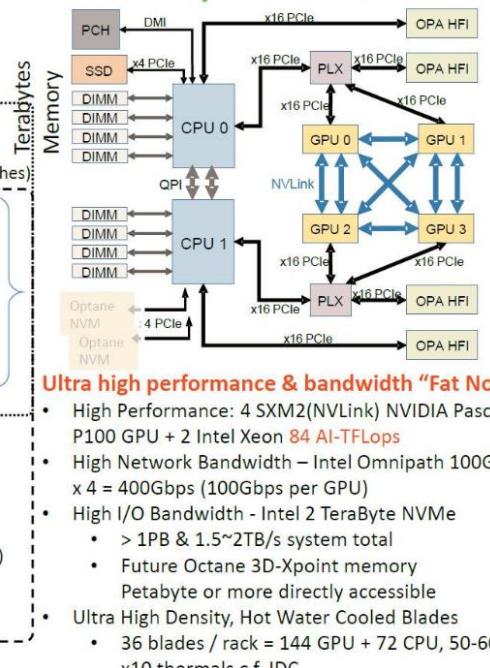
<https://www.forbes.com/sites/marcochiappetta/2017/02/20/nvidias-accelerated-computing-platform-to-power-japans-fastest-ai-supercomputer/#112589b97708>

추적 불가  
**TSUBAME3.0 Compute Node SGI ICE-XA, a New GPU Compute Blade Co-  
 Designed by SGI and Tokyo Tech GSIC**



```

graph LR
    PCH[PCH] -- DMI --> Bus[x16 PCIe]
    Bus --> OP[OP]
  
```

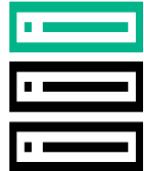


# Storage Challenges: Latency & Cost

## Data Storage Challenges

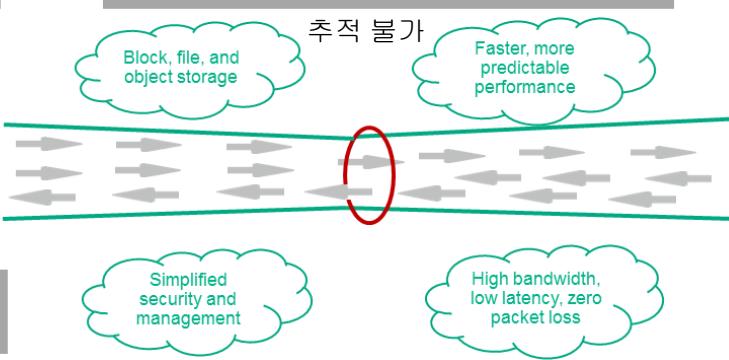
Predictable Performance, Deterministic & Secure Fabrics

Servers

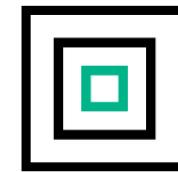


- Higher processing capability
- High-density virtualization

Fabrics

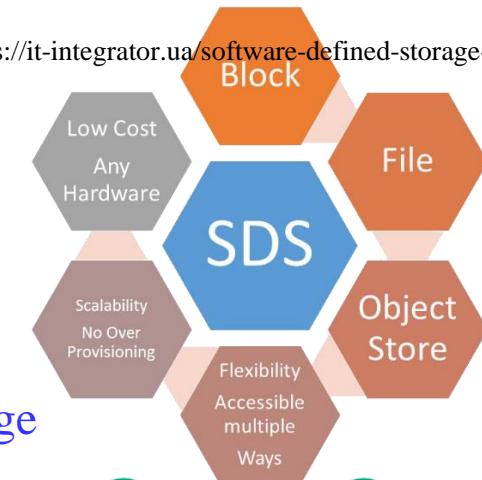


Storage



- Move to All-flash
- Faster protocols – NVMe-oF

<https://it-integrator.ua/software-defined-storage-sd>



## Data Storage Technology Evolution & Software-Defined Storage



### Internal Storage

Distributed Storage  
Shared Network  
File & Block Data

### External Storage

Storage Consolidation  
Dedicated FC SAN  
Block/Structured Data

### Virtualization

vSANs  
Inherit Networking  
Unstructured Data

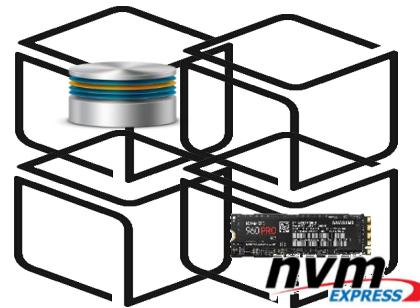
### Secondary Storage

Media/Flash  
Big Data  
Cloud/Analytics

### NVMe

Many Platforms  
Dedicated GbE  
Right Data for Platform

# Storage Clusters & Distributed/Parallel File Systems



Storage Clustering

Local



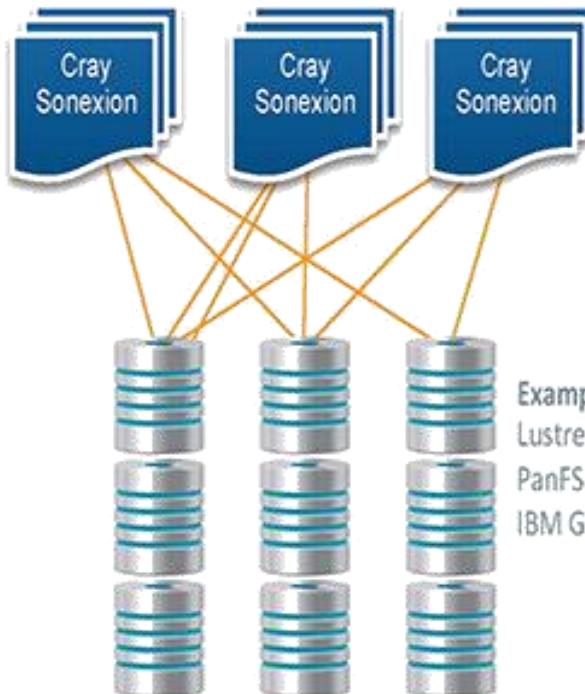
Examples:  
NetApp WAFL  
Sun ZFS  
WinNT  
HDS BlueArc

Distributed File System



Examples:  
Isilon OneFS  
NetApp ONTAP  
Gluster

Parallel File System



Examples:  
Lustre  
PanFS  
IBM GPFS

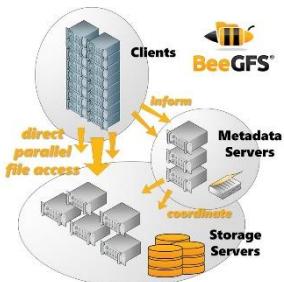
<https://de.wikipedia.org/wiki/BeeGFS-Architecture>



<https://www.cray.com/sites/default/files/SB-Cray-Storage-Hedge-Funds.pdf>

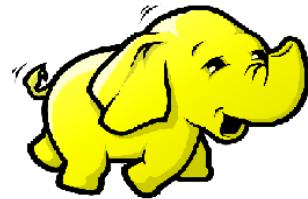
IBM  
GPFS

Lustre®

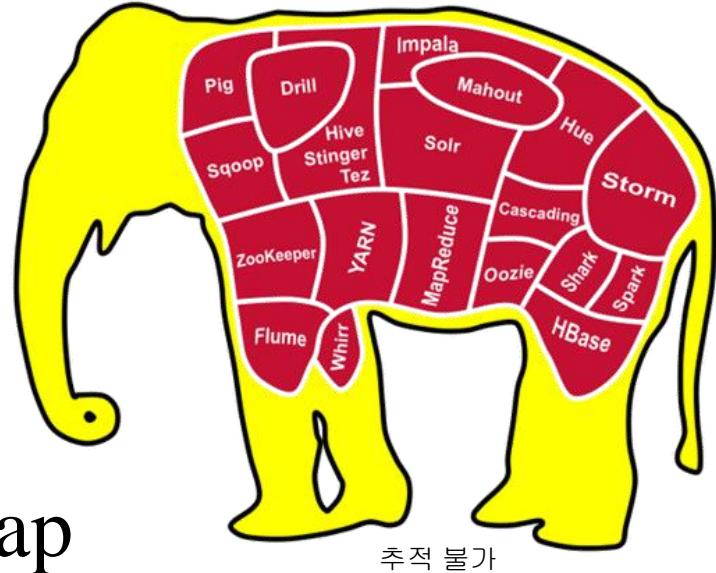




# Open-Source BigData Platform



Hadoop: Distributed Computing Framework



Map  
Reduce

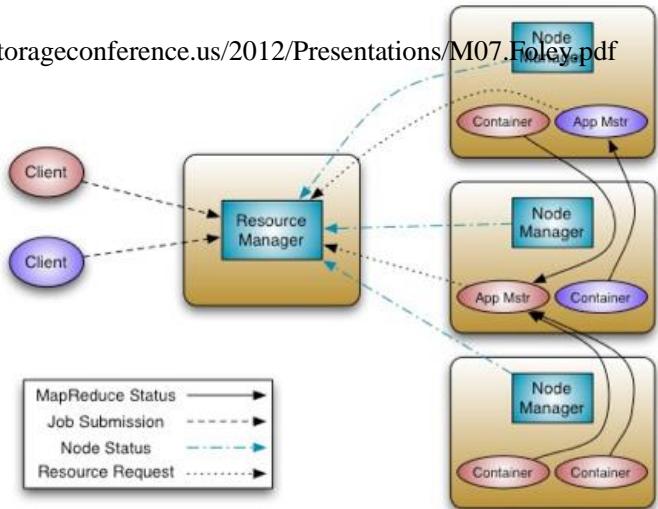


Storage + Compute Clustering

# BigData (HPDA) Cluster Management

## HPDA (High Performance Data Analytics)

<http://storageconference.us/2012/Presentations/M07.Foley.pdf>



## Apache Hadoop YARN

### Applications Run Natively IN Hadoop

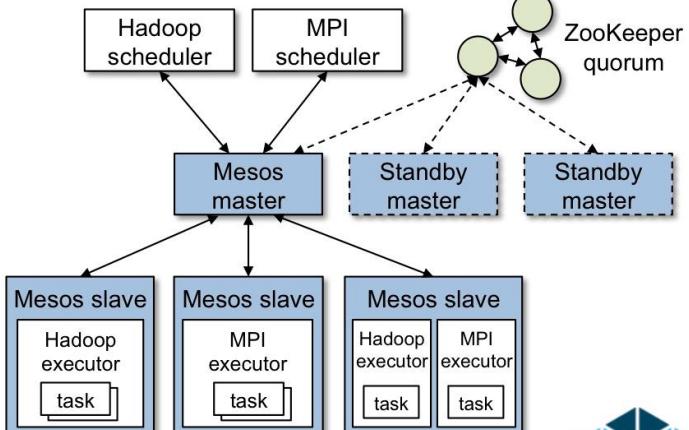


### YARN (Cluster Resource Management)

<http://m.itboth.com/d/qQrYRz/hadoop/>

HDFS2 (Redundant, Reliable Storage)

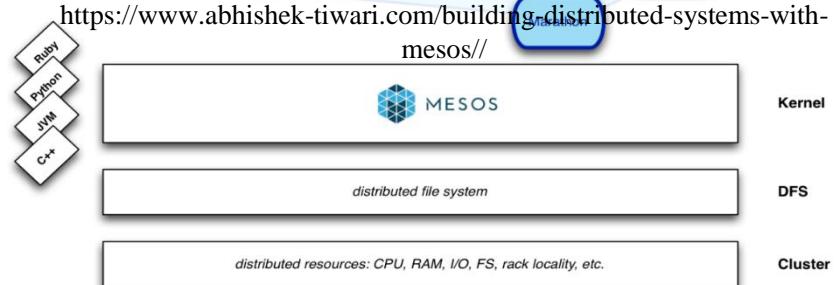
<https://portlandbathrepair.com/hadoop-architecture-ppt/hadoop-architecture-ppt-fresh-10-best-software-architecture-images-on-pinterest/>



## Apache Mesos

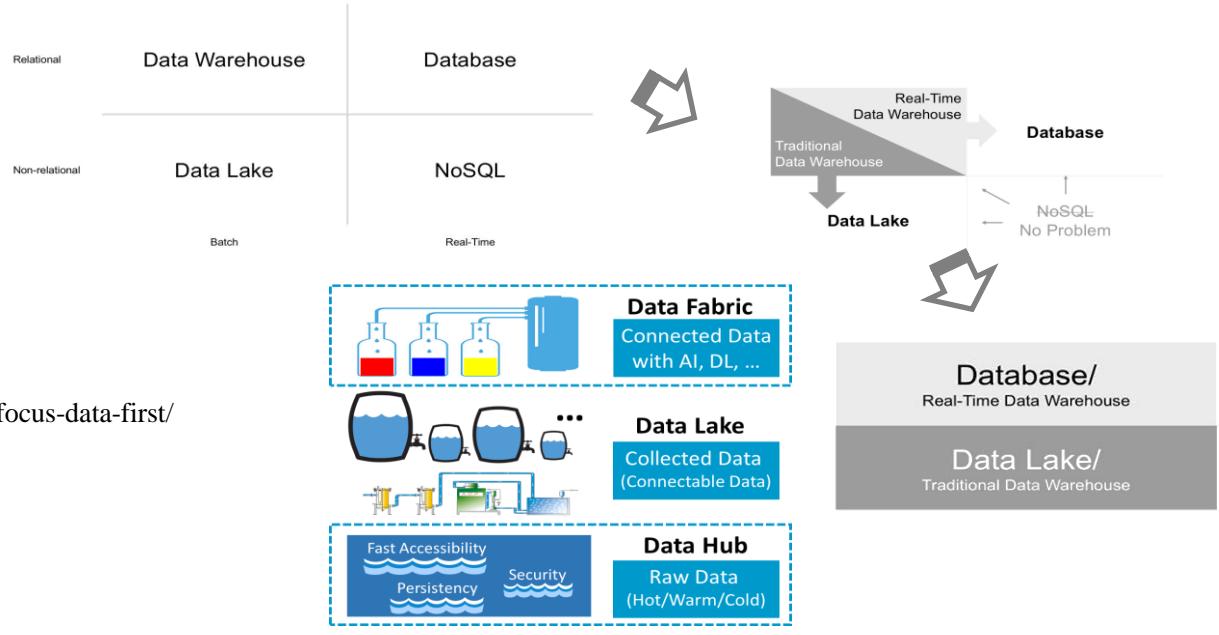
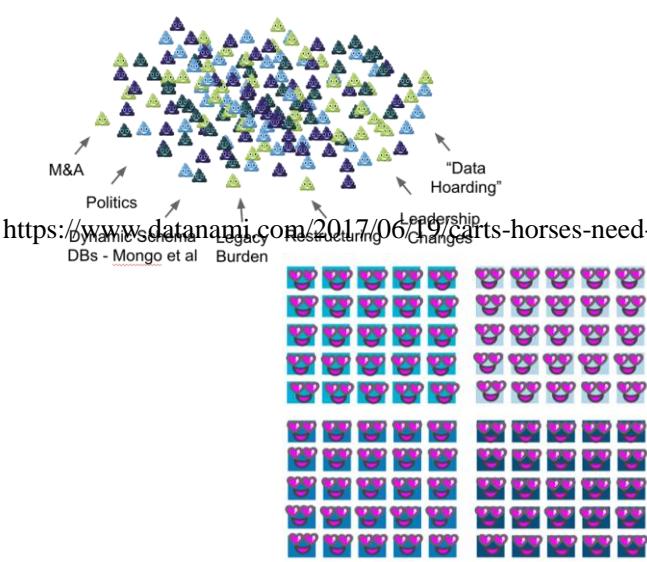


<https://www.abhishek-tiwari.com/building-distributed-systems-with-mesos/>





# ‘New’ Open Data Landscape: Data Fabric vs Data Lake vs Data Hub



→→ Intelligence (Learning & Inference)

(Raw) Data → Information → Knowledge



<https://activerain.com/blogsview/49310/27/unforgettable—that's-what-we-are--->

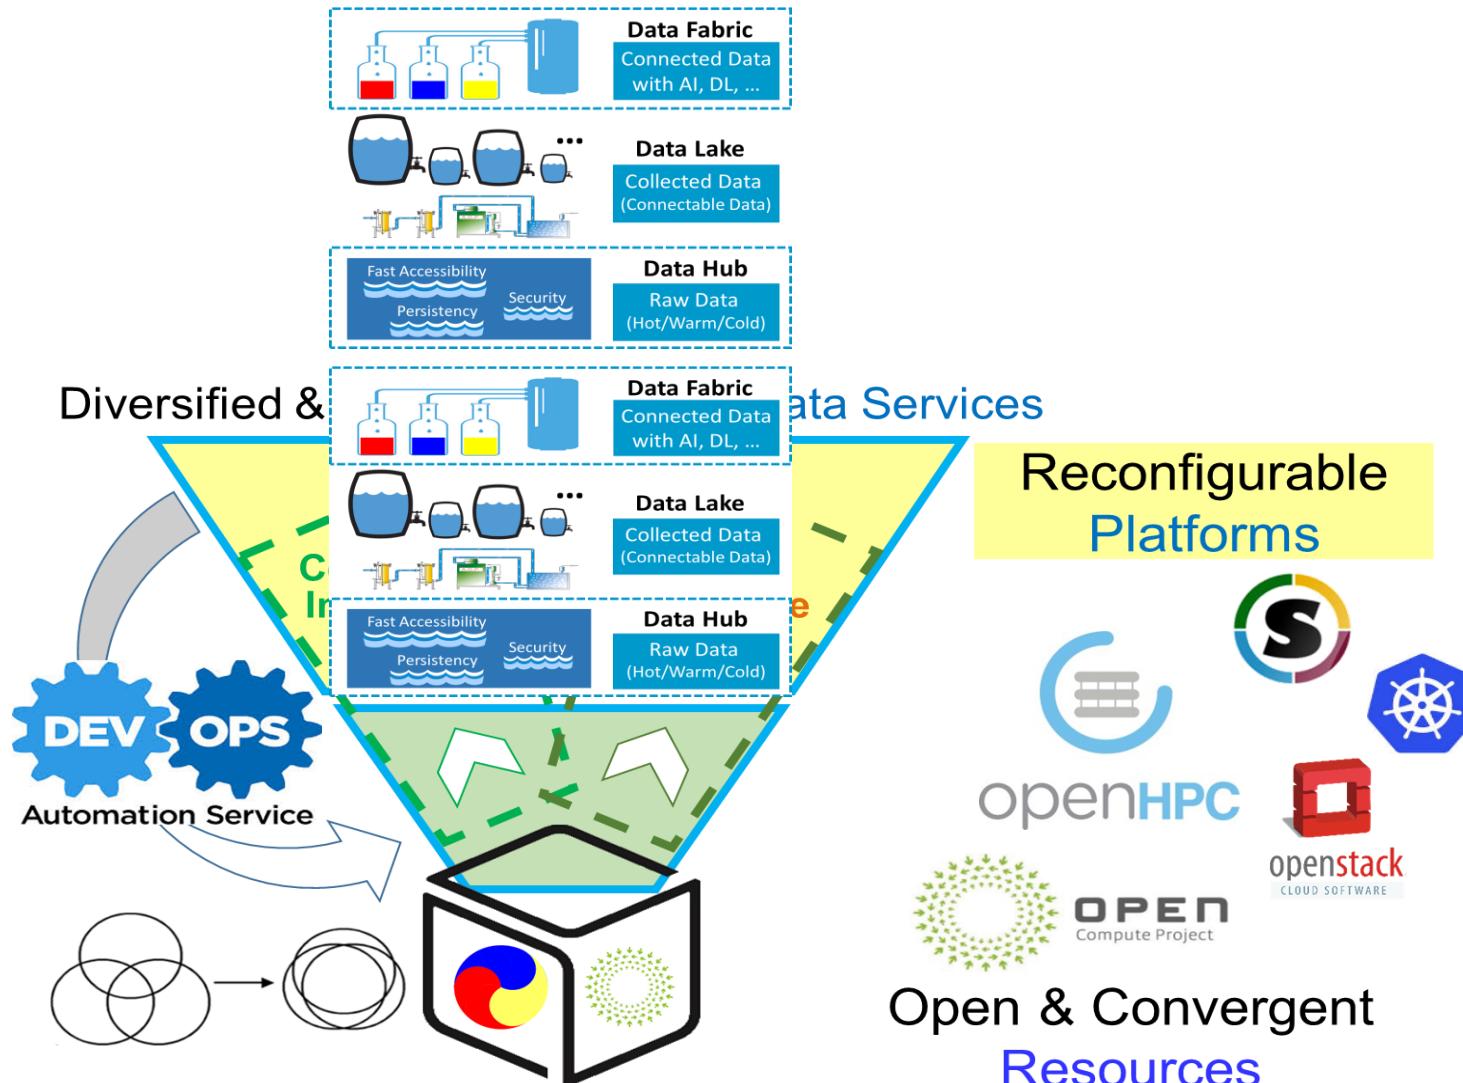
27/unforgettable—that's-what-we-are---



<https://www.quora.com/What-is-the-difference-between-knowledge-and-intelligence>

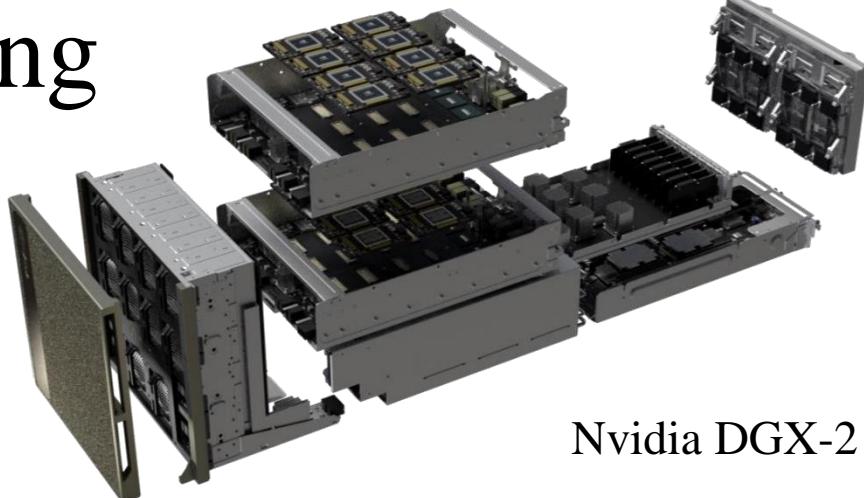


# Configurable Clustering for AI-inspired HPC/BigData (HPDA)



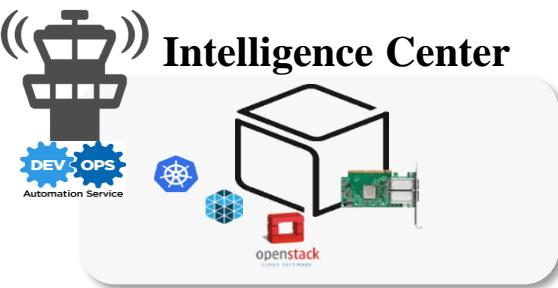
# Cluster for AI Computing

HPC + HPDA (BigData)  
→ AI (ML/DL)



Nvidia DGX-2

## Intelligence Center



<https://www.anandtech.com/show/12587/nvidias-dgx2-sixteen-v100-gpus-30-tb-of-nvme-only-400k>

## Multi-node AI Computing Cluster with Optimized DL Tools

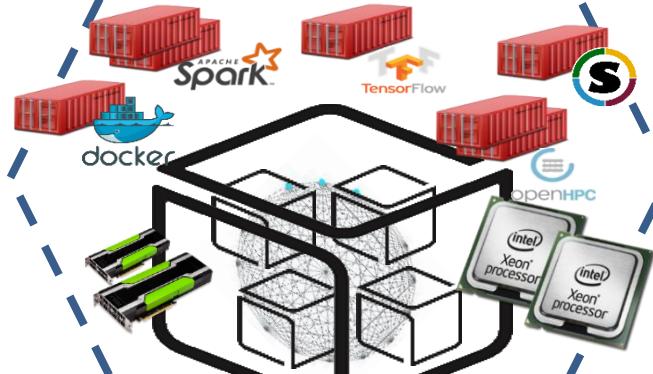
### GPU Acceleration



GPU Cards

### Container Orchestration

kubernetes

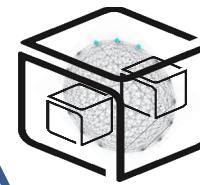


### Parallel File System



NVMe SSDs

ceph



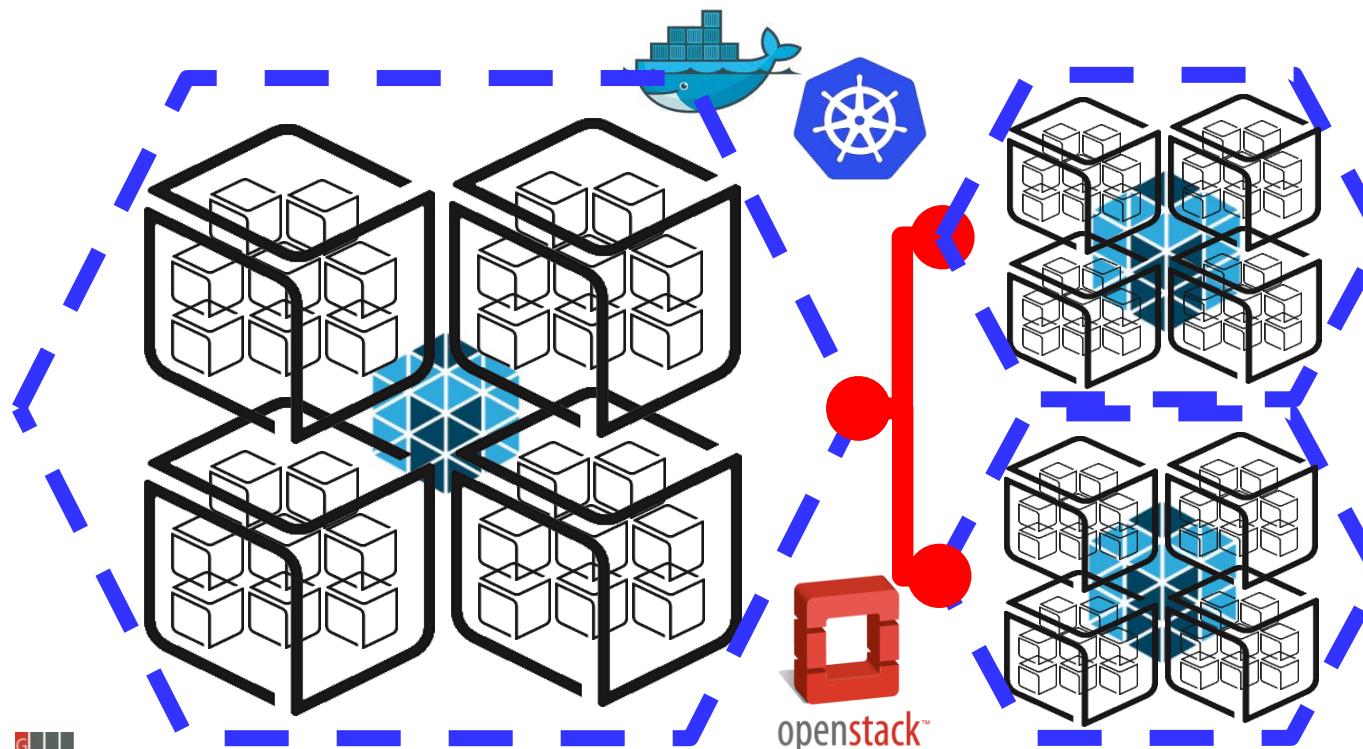
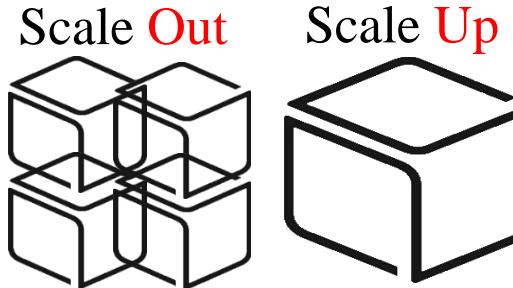
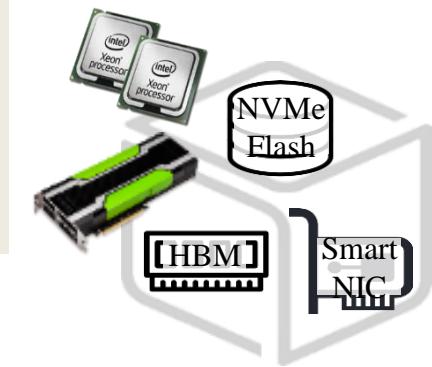
Cloud Storage Cluster

### Cloud Storage



Smart NICs

# Computer System: Resource Scaling/Pooling with Clustering





Gwangju Institute of  
Science & Technology



Thank you!

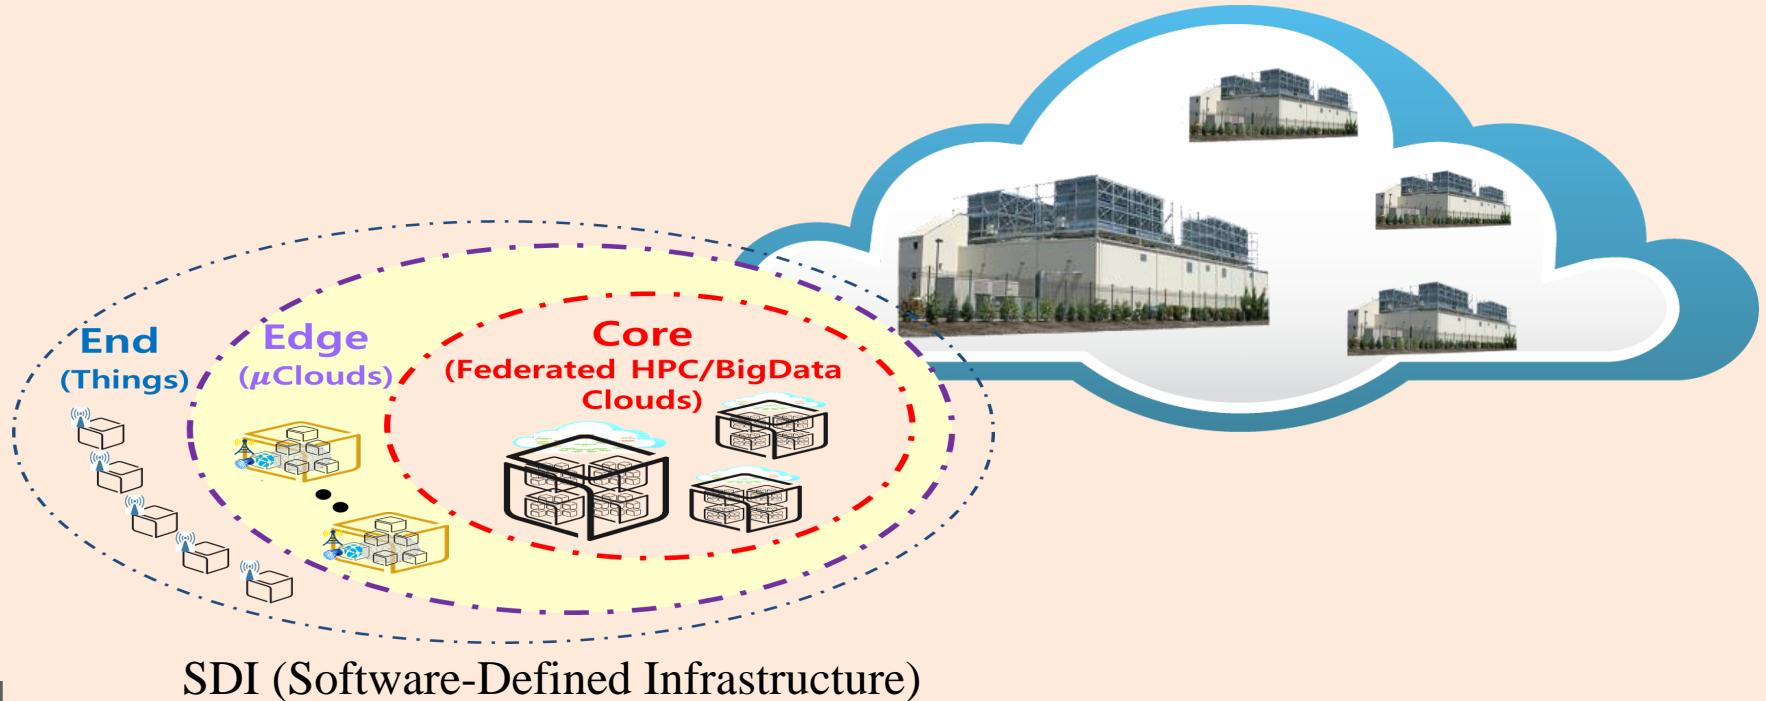
Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>

# Chapter 6:

# Computer Systems

## for Cloud Data Center and SDI



# Chapter Keywords

Data Center Networking

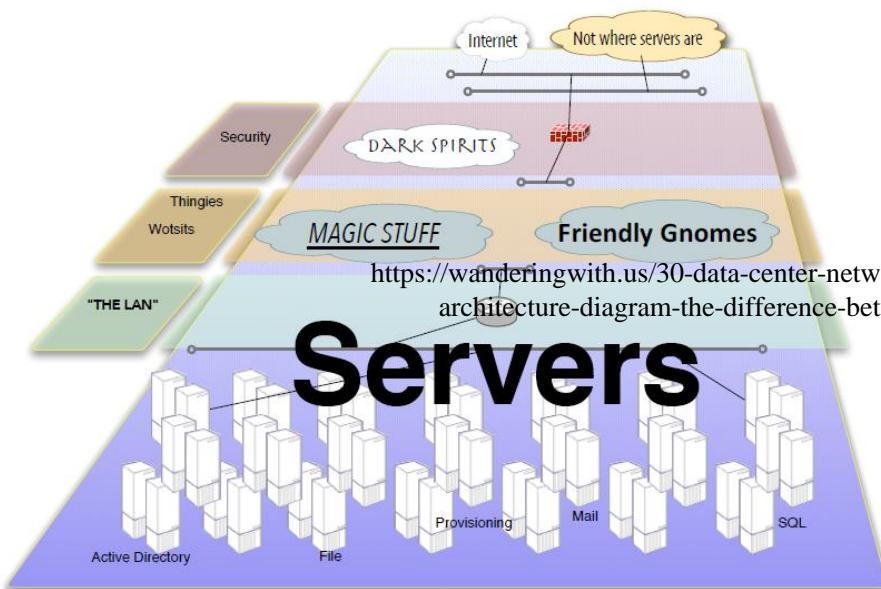
Data Center as a Big Computer

Cloud Computing

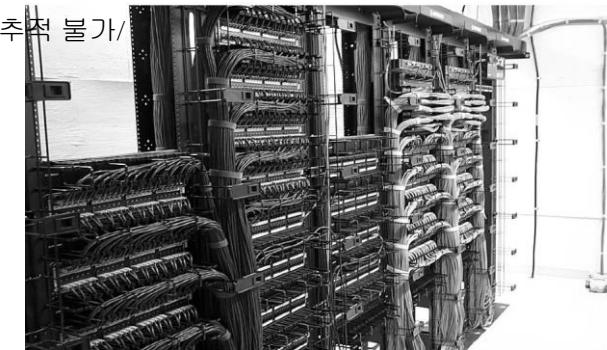
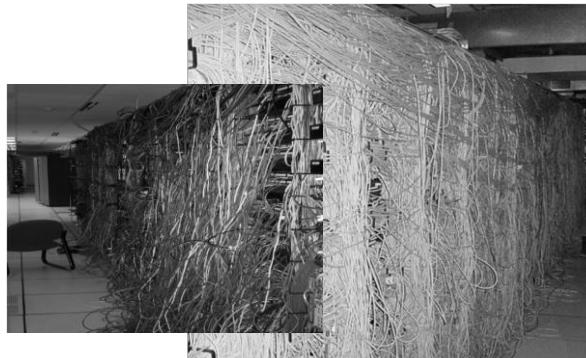
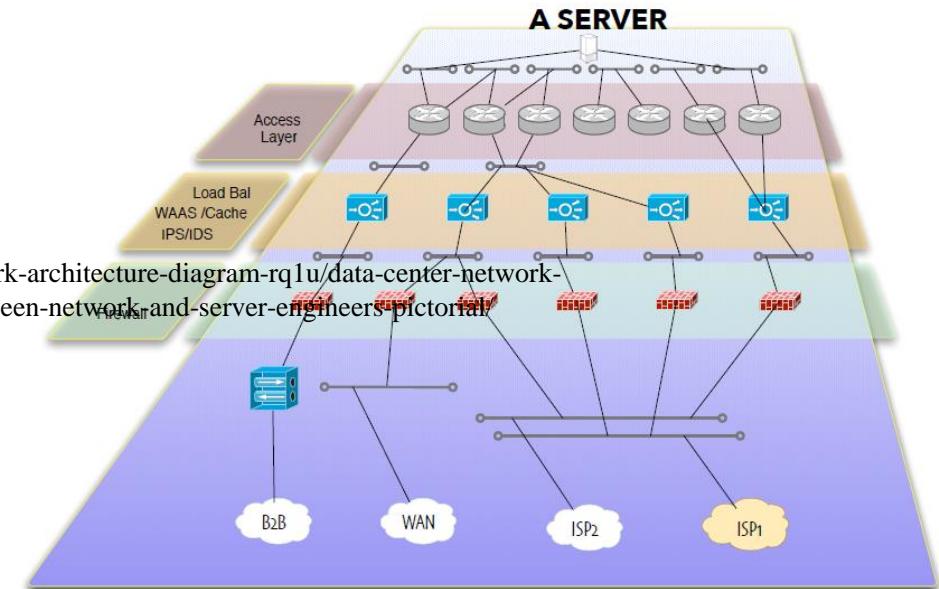
Software-Defined Infrastructure

# Data Center: Separated Operation for Compute (Servers) and Networking

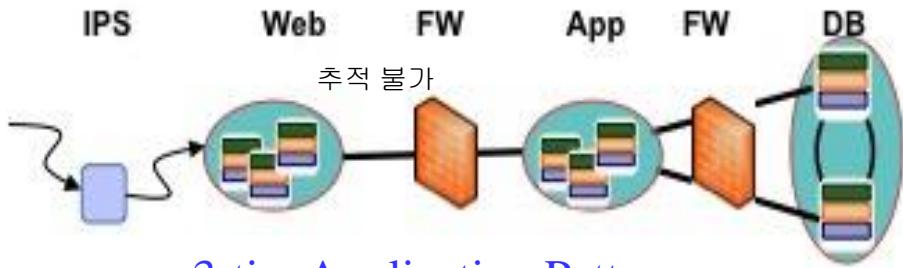
Server Admins See...



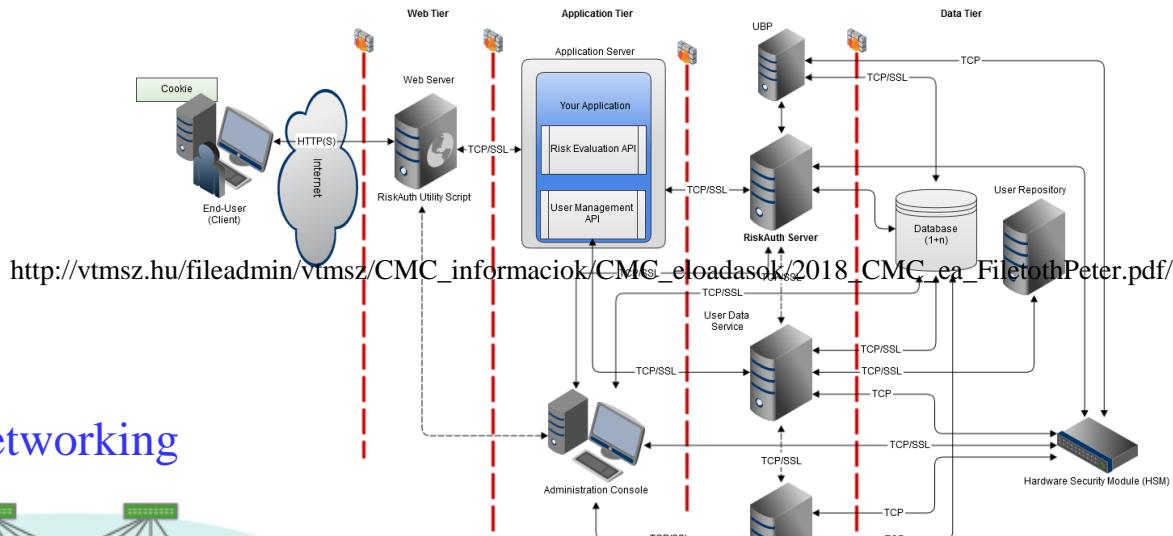
Network Admins see ....



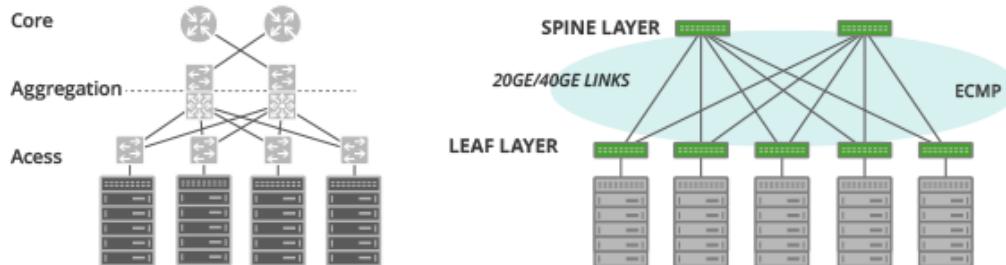
# Data Center 3-Tier Application Pattern & Networking



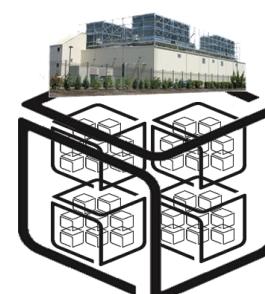
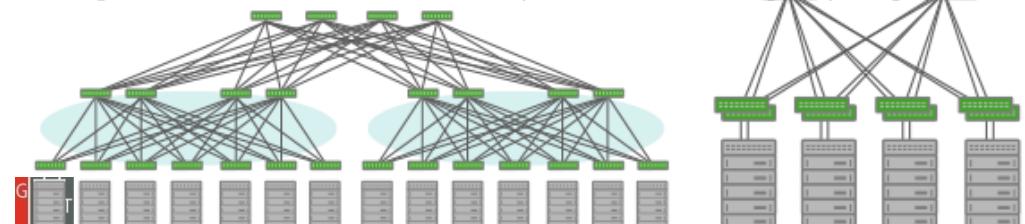
3-tier Application Pattern



Multiple options for Data Center Networking



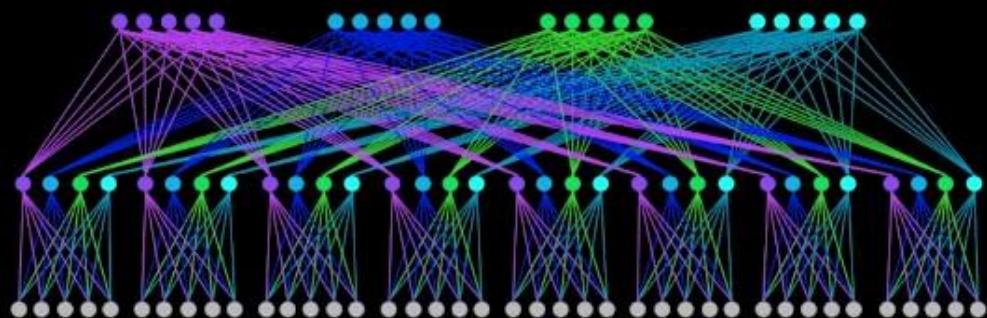
[http://wccclab.cs.nchu.edu.tw/www/images/107-1\\_cloud\\_computing/sdn.pdf](http://wccclab.cs.nchu.edu.tw/www/images/107-1_cloud_computing/sdn.pdf)



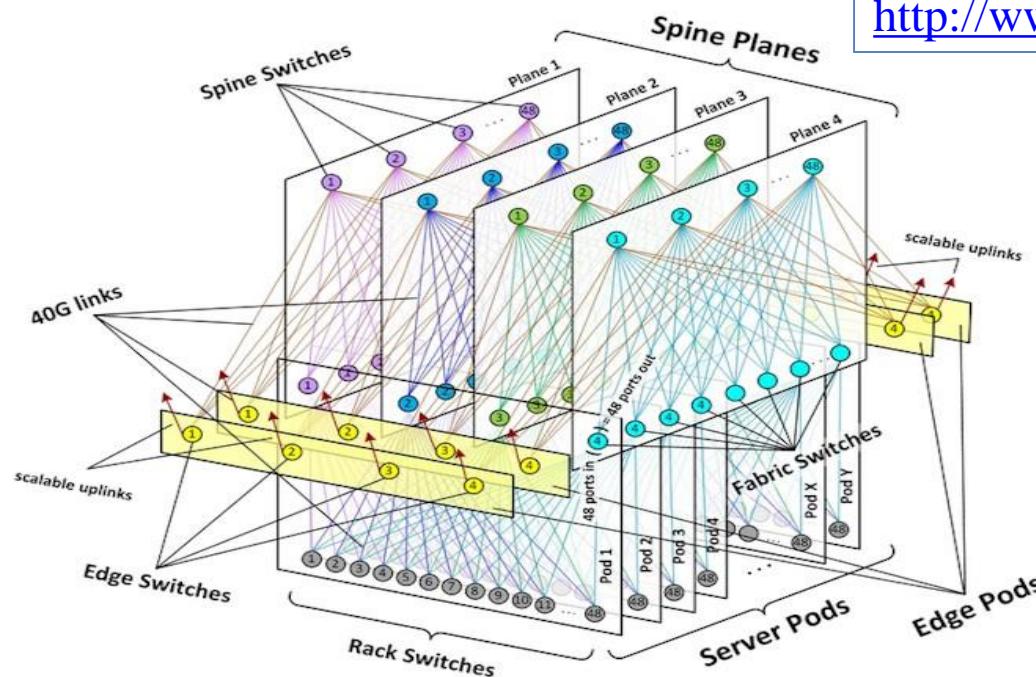
# Data Center Networking: Facebook Fabric



The Fabric: datacenter-wide performance

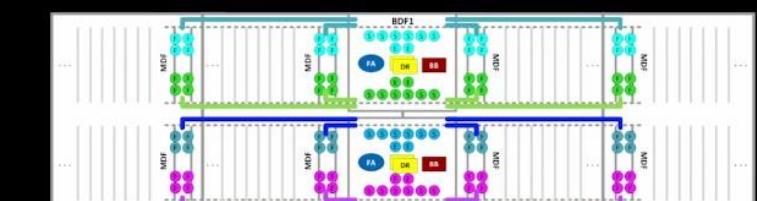


- highly scalable
- smaller units
- all 40G links
- IPv4 + IPv6
- 1 protocol: BGP
- software-driven

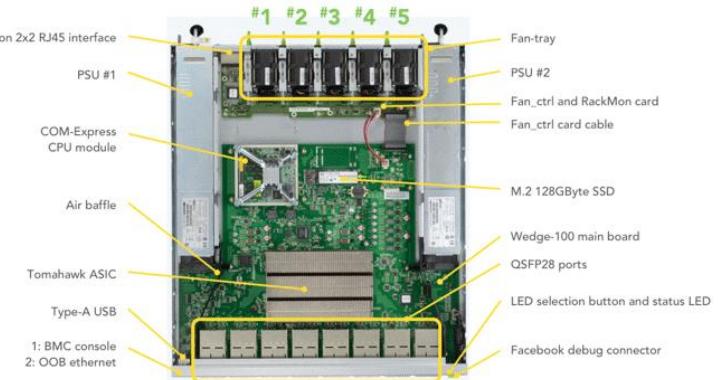
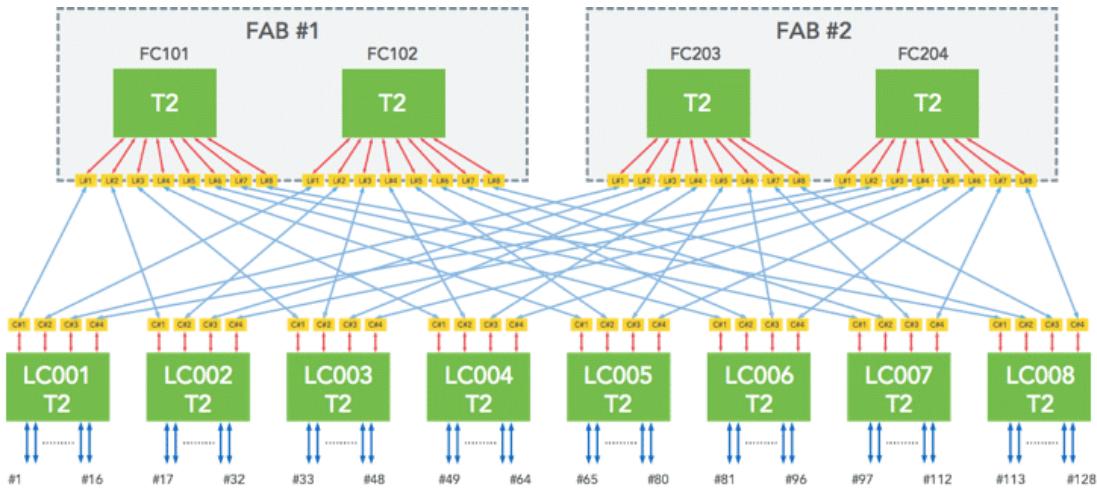


<http://www.youtube.com/watch?v=mLEawo6OzFM>

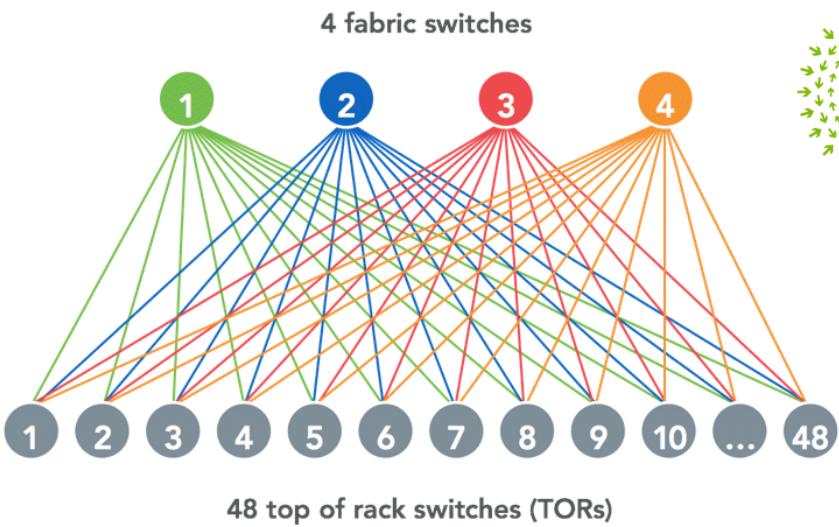
Fabric-optimized datacenter



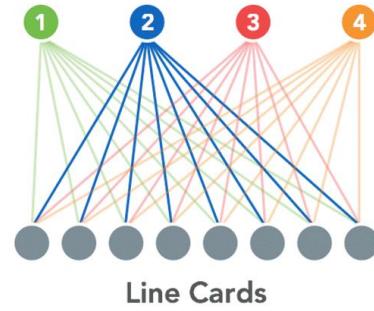
built from center = streamlined deployment  
shorter "straight line" fiber runs



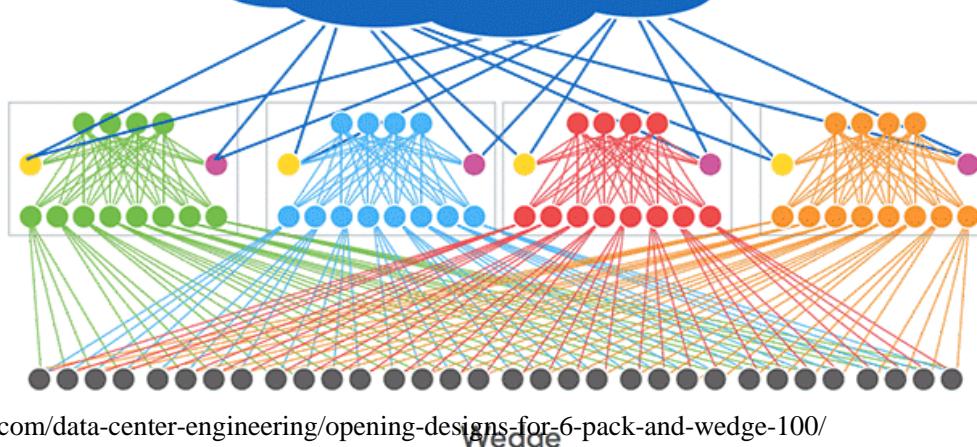
Fabric Cards



**OPEN**  
Compute Project



Facebook Backbone  
(to Internet)

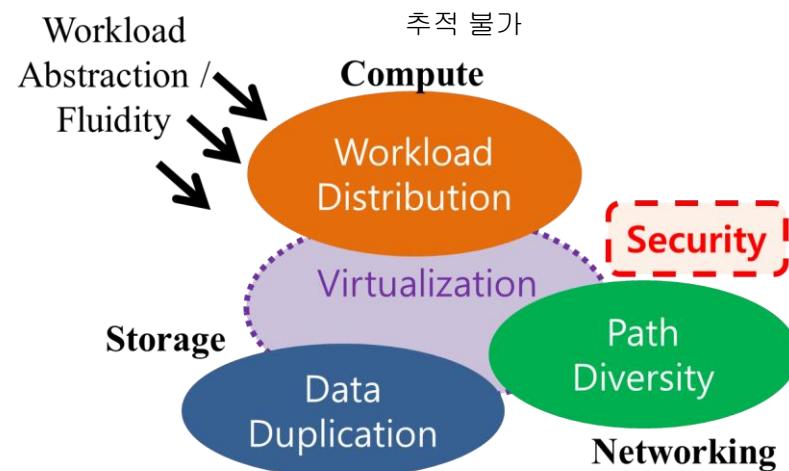
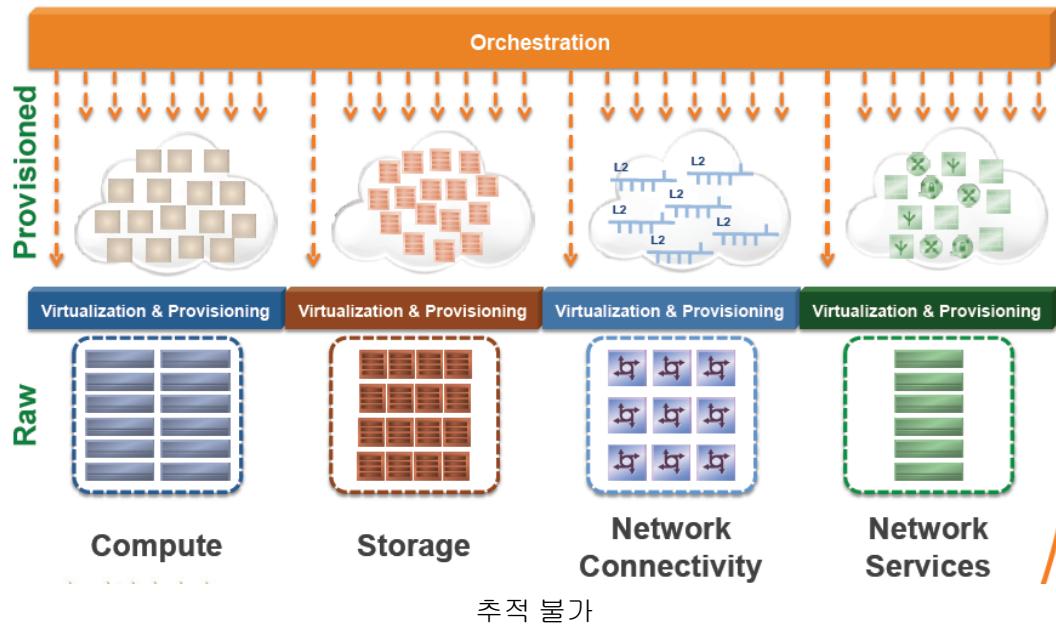
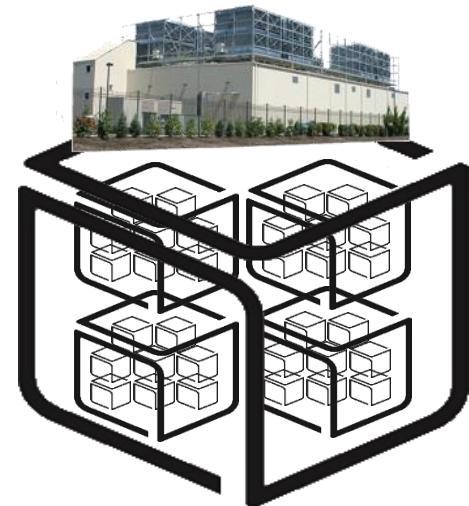


# Data Center Networking: Facebook (Cont.)

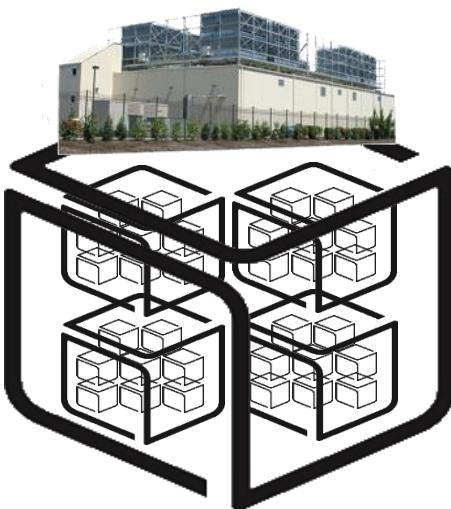
# Software-Defined Data

## Centers

Unified, Programmable &  
Virtualized Resources



# Data Center as a **BIG** Computer



*The Datacenter as a Computer: An Introduction  
to the Design of Warehouse-Scale Machines*

DATACENTER



STATIC PARTITIONING



[https://www.slideshare.net/robdthomas/ibm-big-data-references?next\\_slideshow=1](https://www.slideshare.net/robdthomas/ibm-big-data-references?next_slideshow=1)

DATACENTER



Google

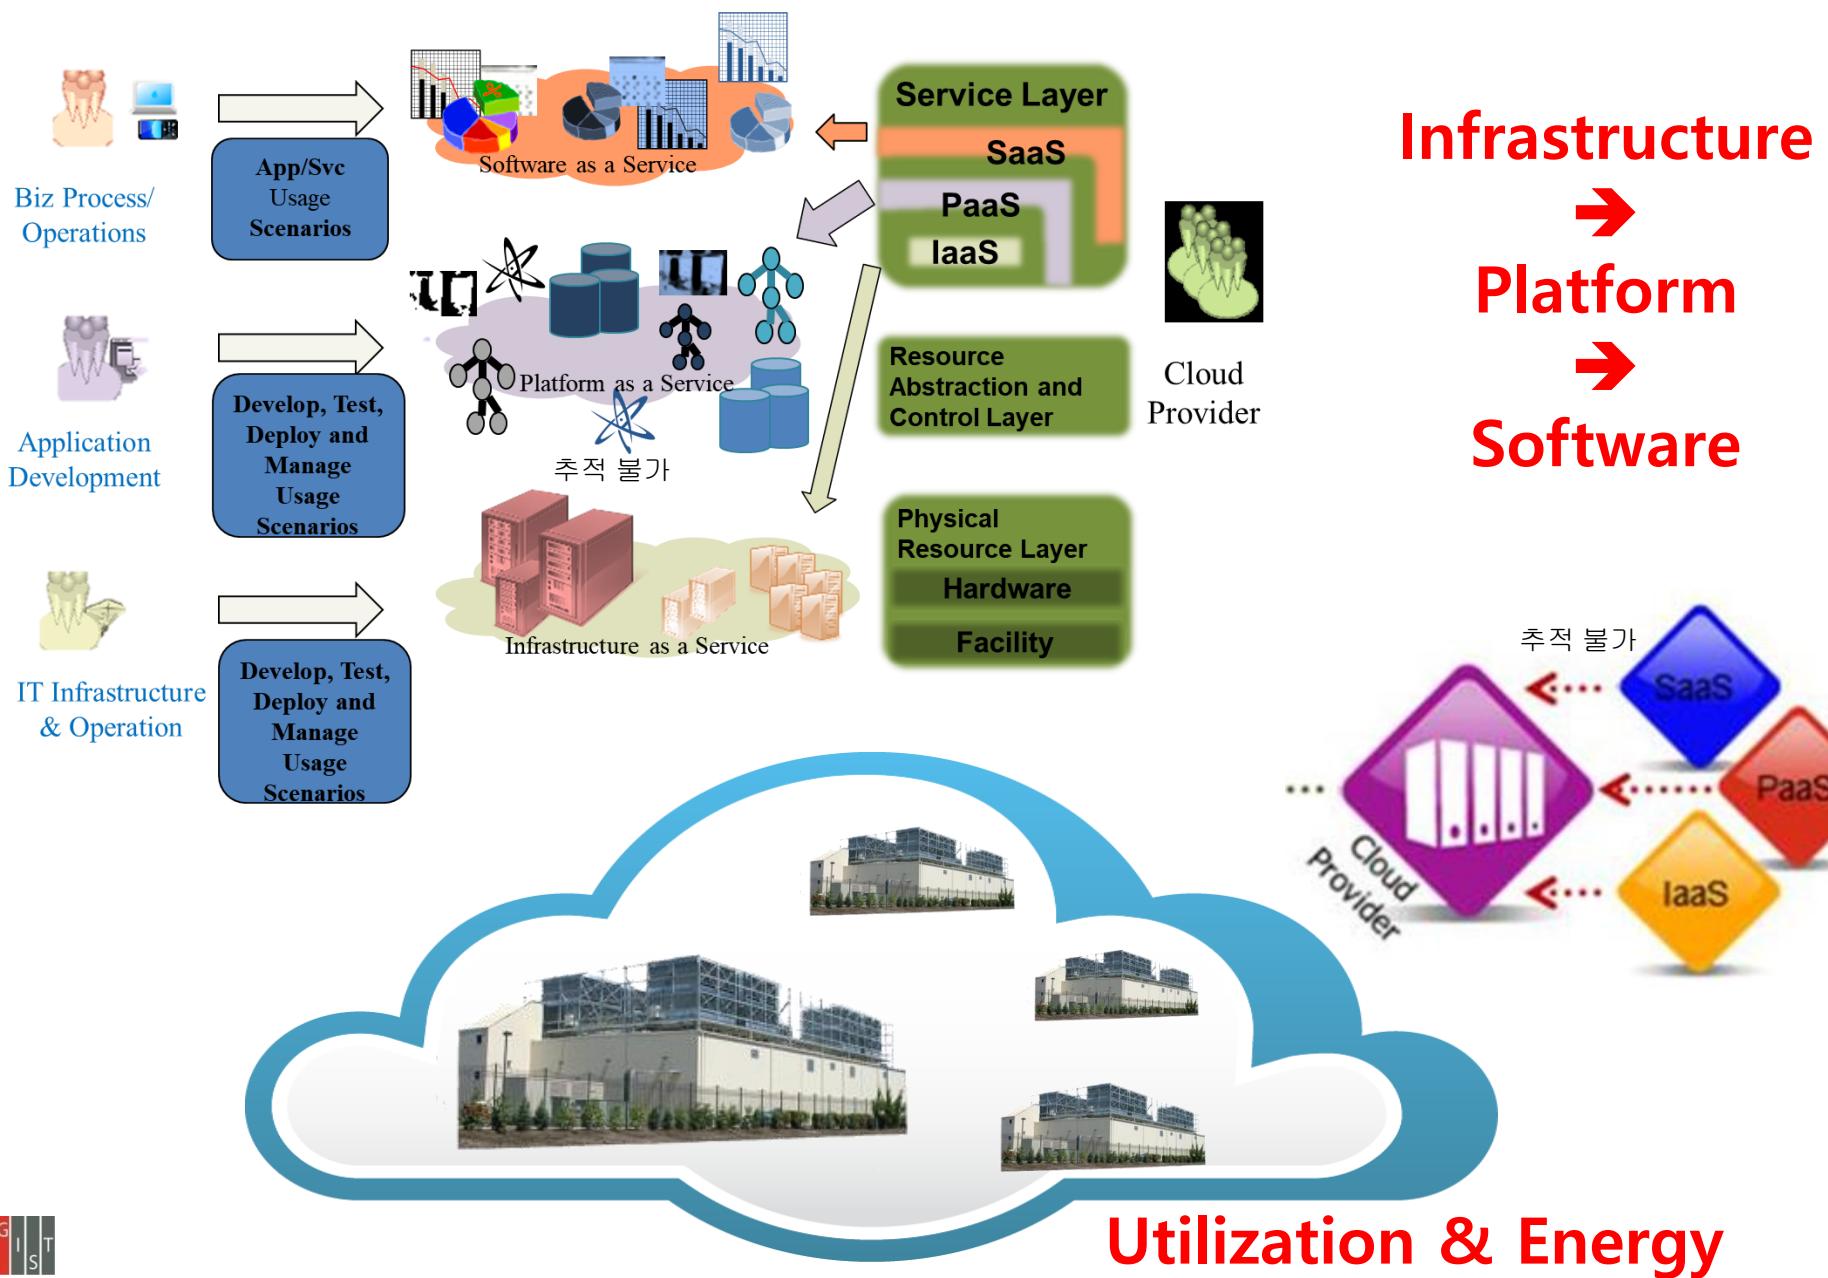


Google Search / Google Maps / MapReduce / Dremel / Gmail



Google Omega:  
flexible, scalable  
Schedulers for large  
compute clusters;  
Treat a data center like  
a single machine

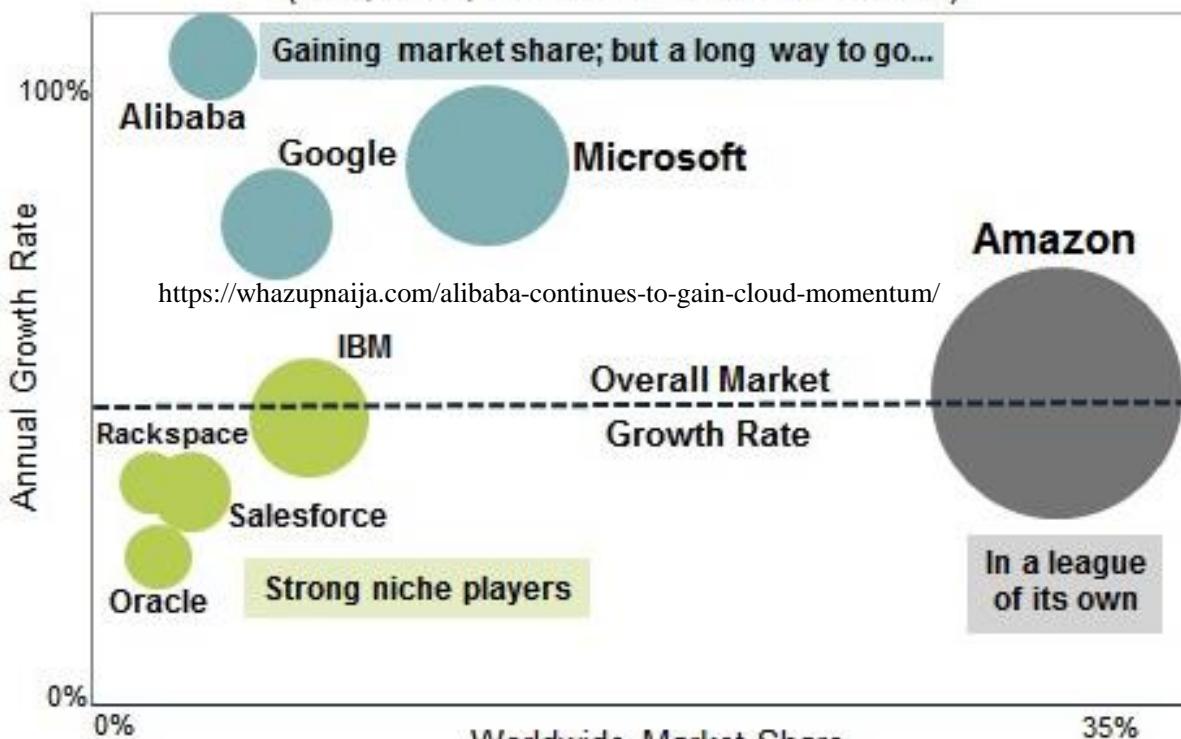
# Cloud Computing: Service Orchestration



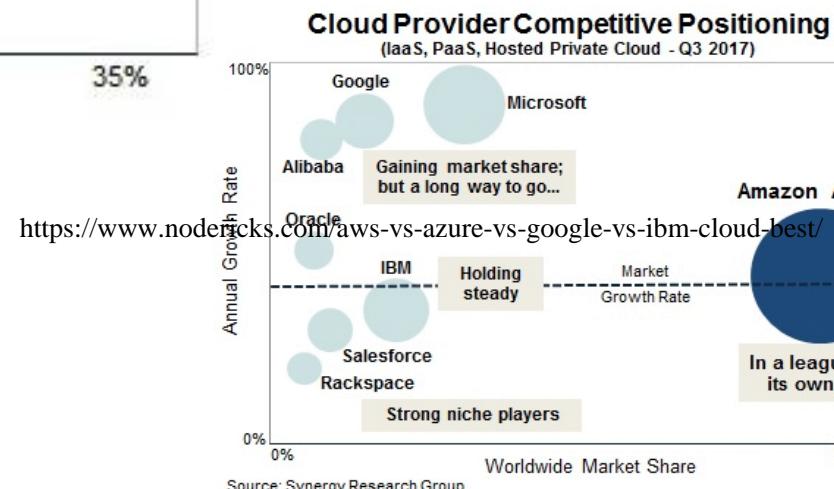
# Market Capacity: Cloud Providers

## Cloud Provider Competitive Positioning

(IaaS, PaaS, Hosted Private Cloud - Q2 2018)



Source: Synergy Research Group



# Market Capacity: Cloud/IT Infrastructure



IDC Cloud Infra  
IT Spending  
(Public+Private):  
57.2bn 2018



Gartner Public  
Cloud Revenue  
186.4bn / 2018

Gartner Public  
Cloud Revenue  
302.5bn / 2021



AWS Revenue  
25bn Year 2018

DC+Cloud Infra  
IT Spending  
124bn Year 2018

Worldwide IT spending  
2.4 Trillion Year 2017  
2.65 Trillion Year 2020

# IoT/HPC/BigData Services from Hyper-scale Cloud (AWS)



Data on road conditions and performance transmitted

AWS IoT

[https://www.google.co.kr/url?sa=i&source=images&cd=&ved=2ahUKEwiM55GX5ufeAhWMWrwKHUsnDcQQjhx6BAgBEAM&url=https%3A%2F%2Fwww.amazonaws.cn%2Fen%2Fiot-platform%2F&psig=AOvVaw1Q49j4AaRJ\\_Sw4X5adQ80L&ust=1542969088950725](https://www.google.co.kr/url?sa=i&source=images&cd=&ved=2ahUKEwiM55GX5ufeAhWMWrwKHUsnDcQQjhx6BAgBEAM&url=https%3A%2F%2Fwww.amazonaws.cn%2Fen%2Fiot-platform%2F&psig=AOvVaw1Q49j4AaRJ_Sw4X5adQ80L&ust=1542969088950725)

Alert cars that moisture levels are high



Alert to driver "Warning: Roads are slippery"

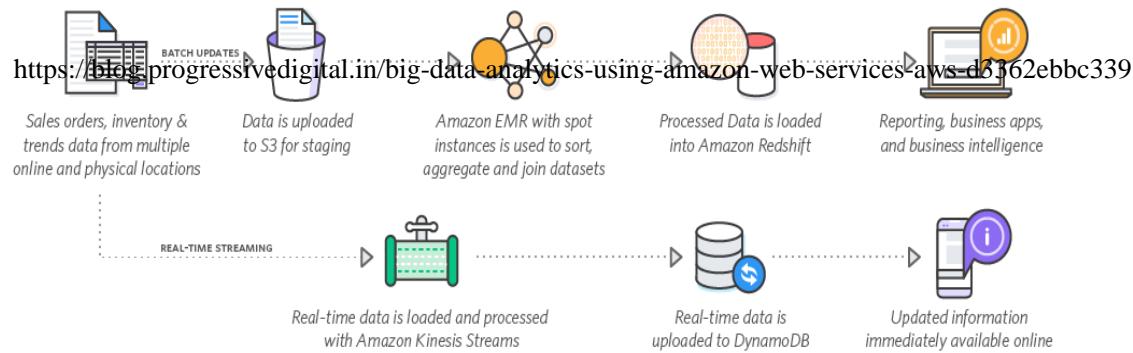
Process data based on rules and interpret that moisture levels are high. Sends alert to nearby cars.

Engine performance data stored in S3 for future analysis

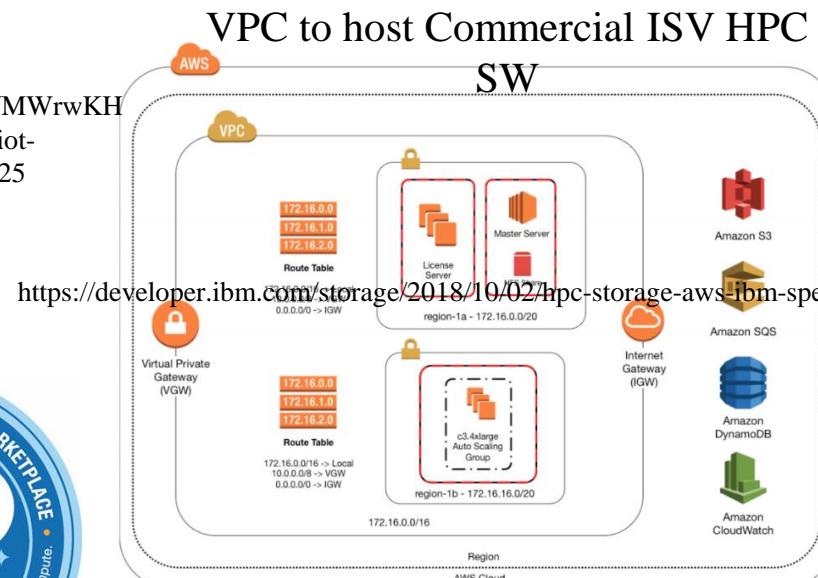
The engine performance data trains an Amazon Machine Learning model for predictions that get more accurate over time



## AWS-IoT



## AWS-BigData



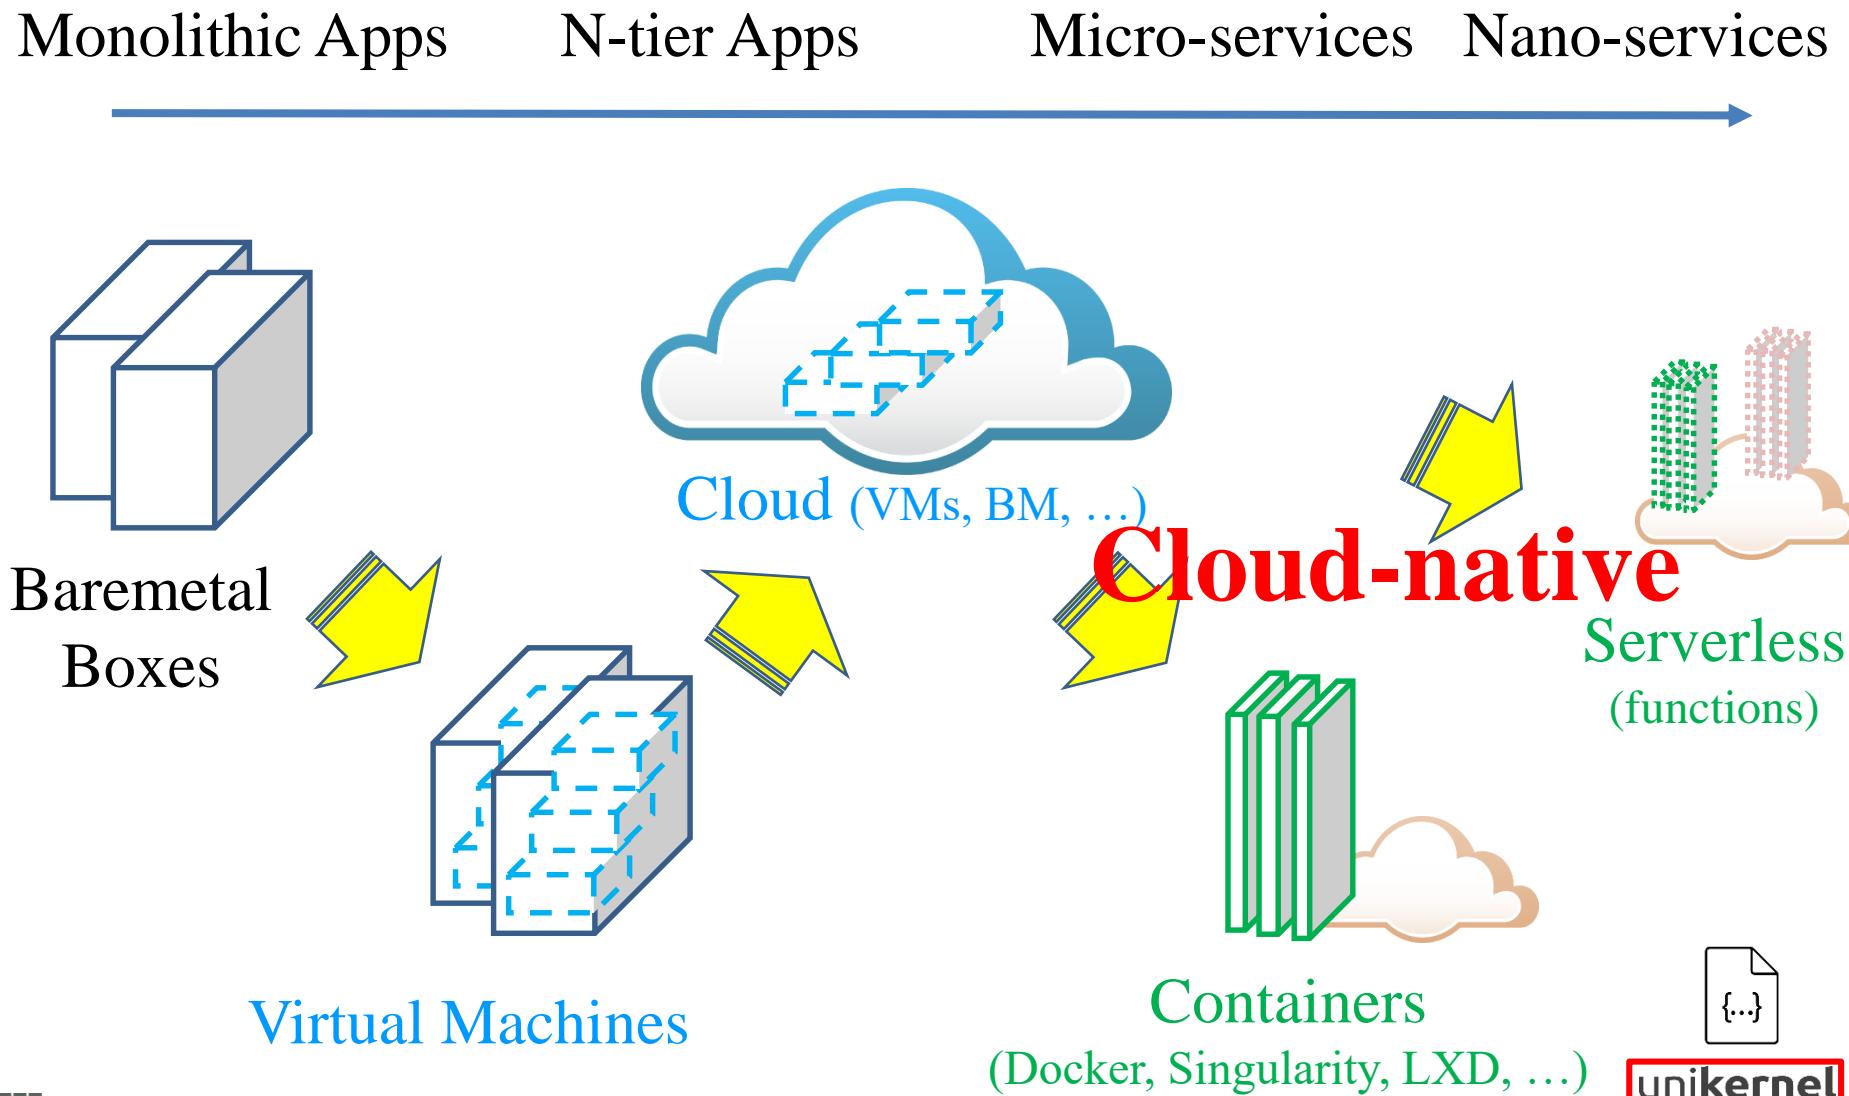
## AWS-HPC



P3 EC2 instance: Nvidia Tesla V100

## AWS-DL

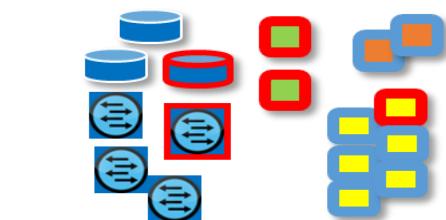
# Recap: From Baremetal to Containers / Serverless



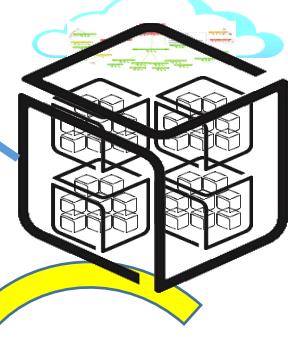
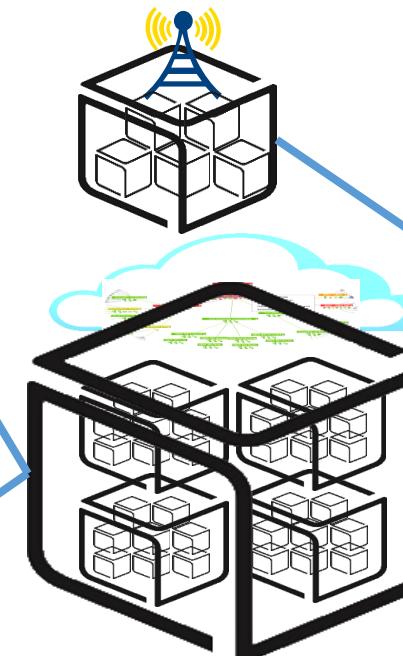
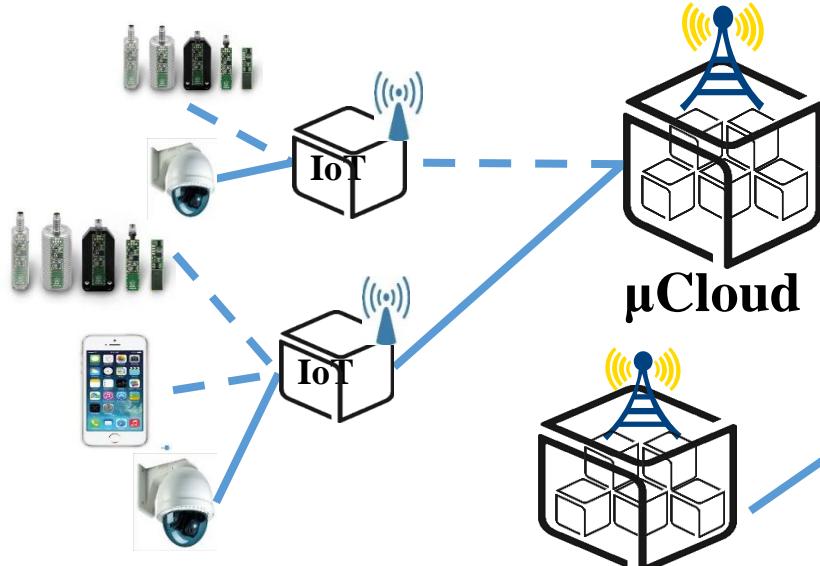
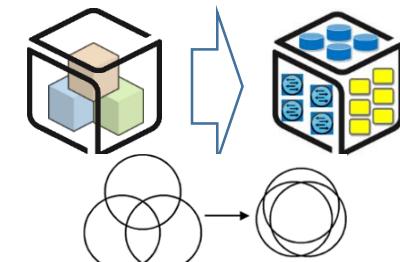


# Convergent SDI for Diversified Services

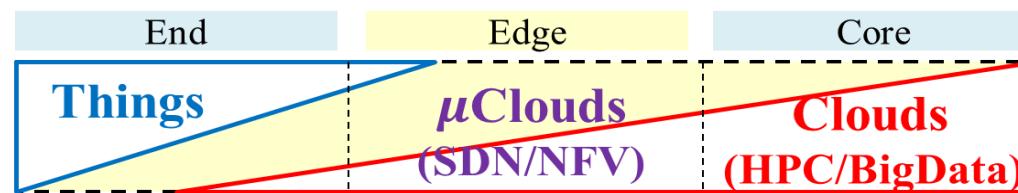
Diversified MicroServices  
Architecture Applications



Resources in  
Hyper-Converged Boxes



Federated  
Cloud DCs



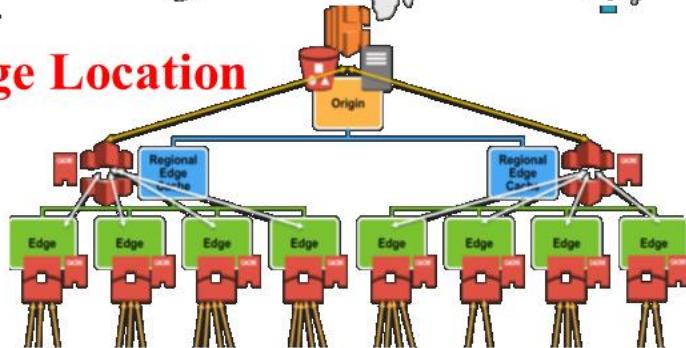
# Converged SDI @ Hyper-scale Cloud



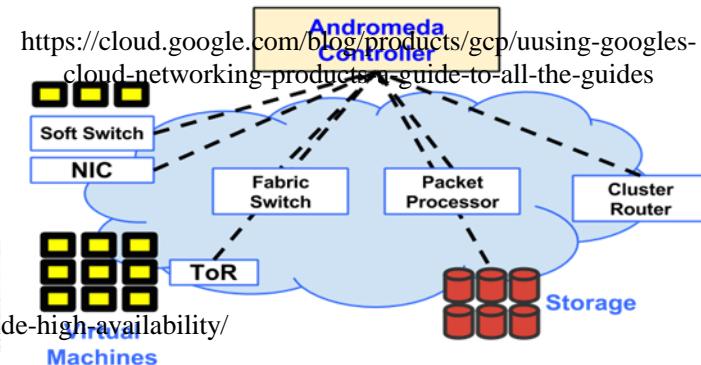
## Regions



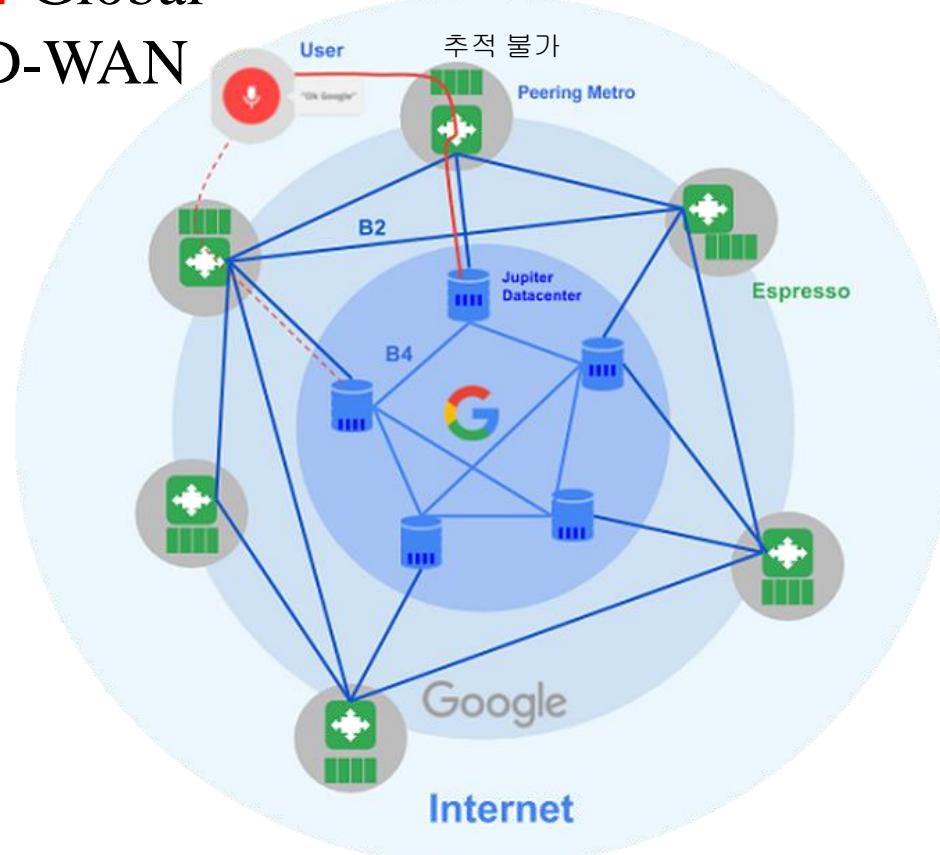
## Edge Location



## B4 Global SD-WAN



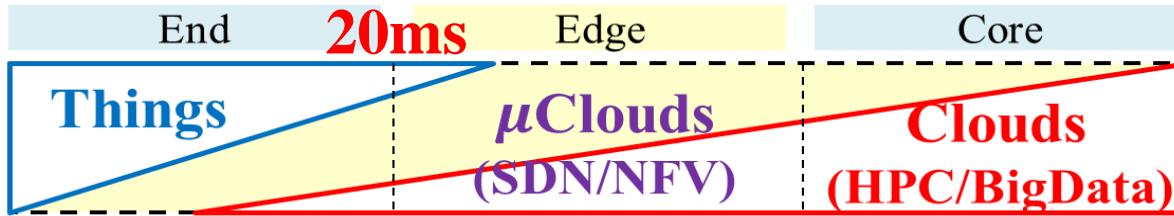
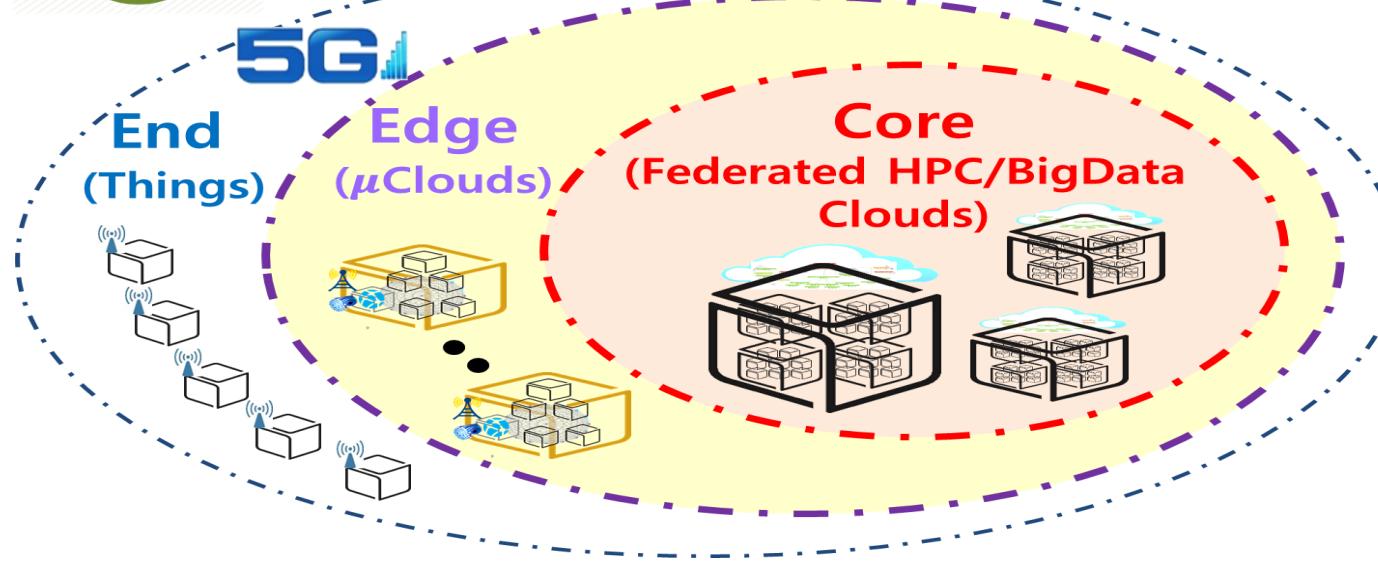
## Andromeda SDN/NFV



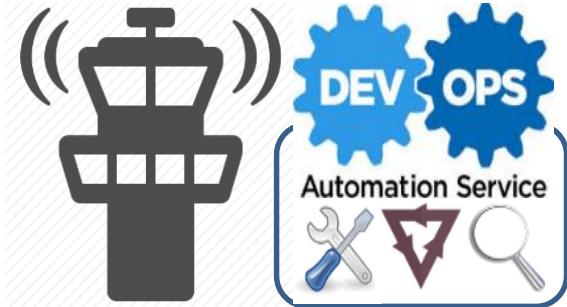
# 5G-leveraged Converged SDI (SDN/NFV/Cloud Integrated)



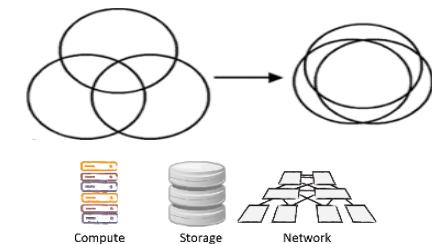
CLOUD NATIVE  
COMPUTING FOUNDATION



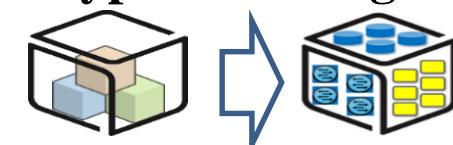
## DevOps Tower



open source  
initiative

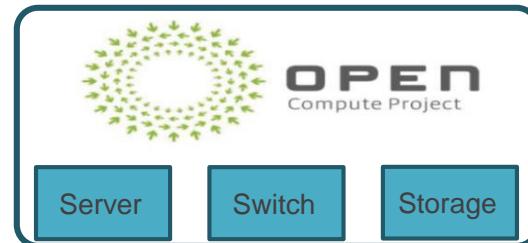
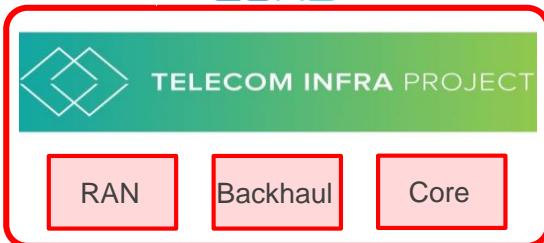
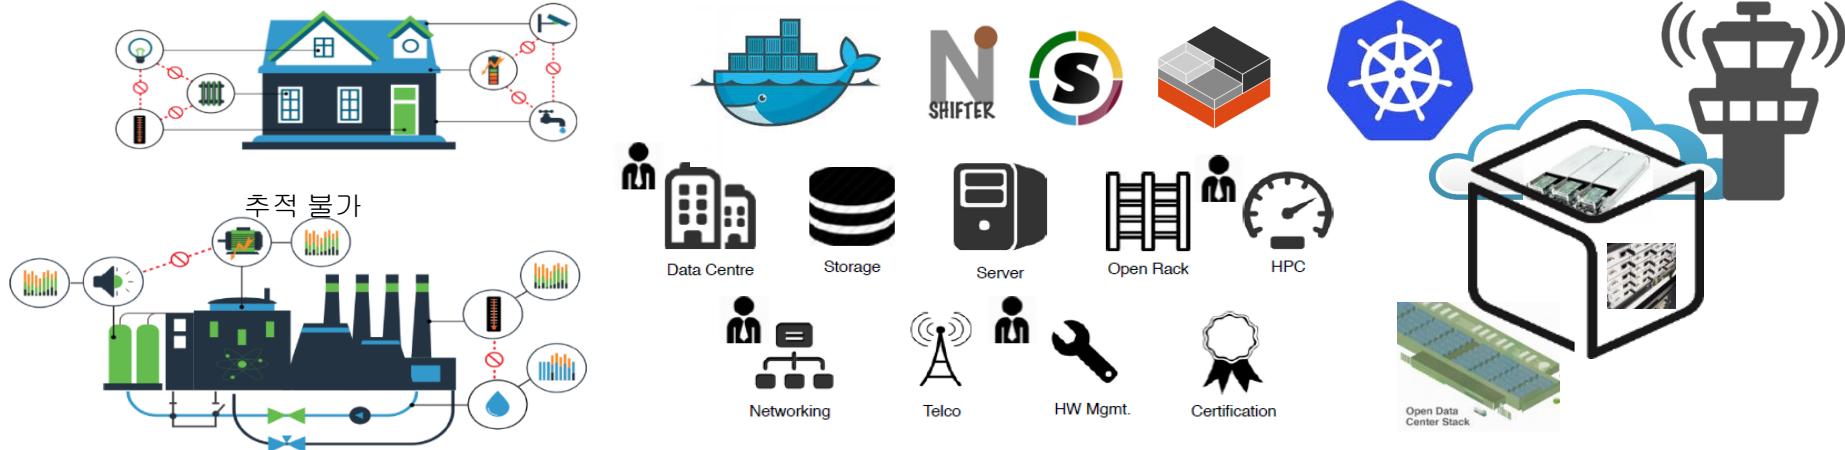


## Hyper-Converged



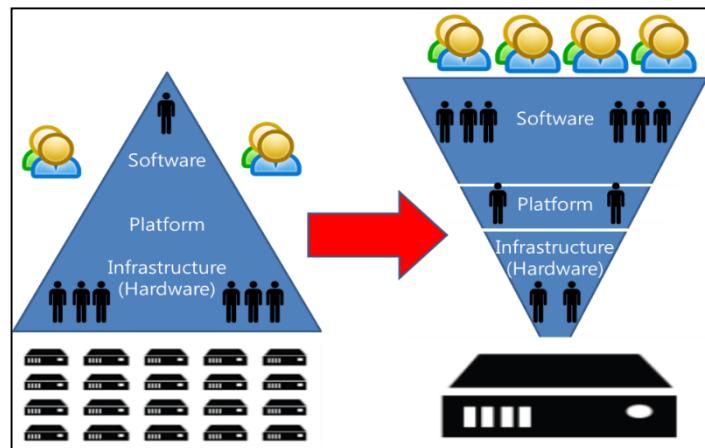
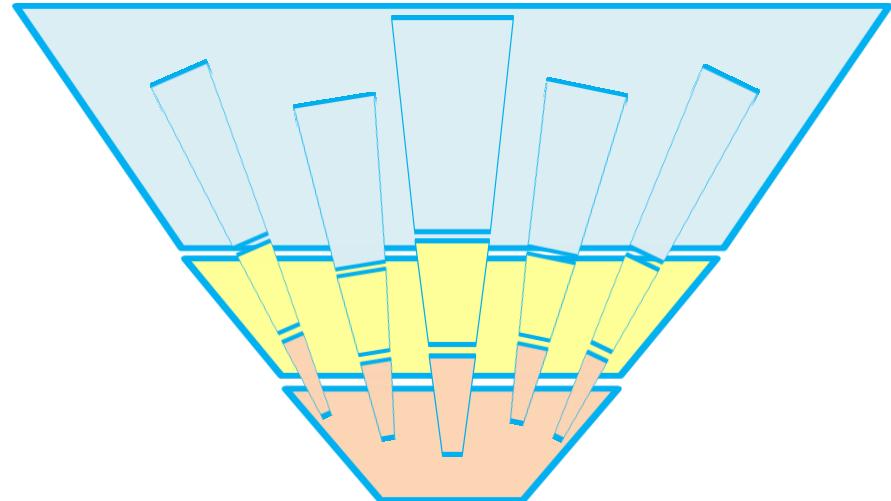
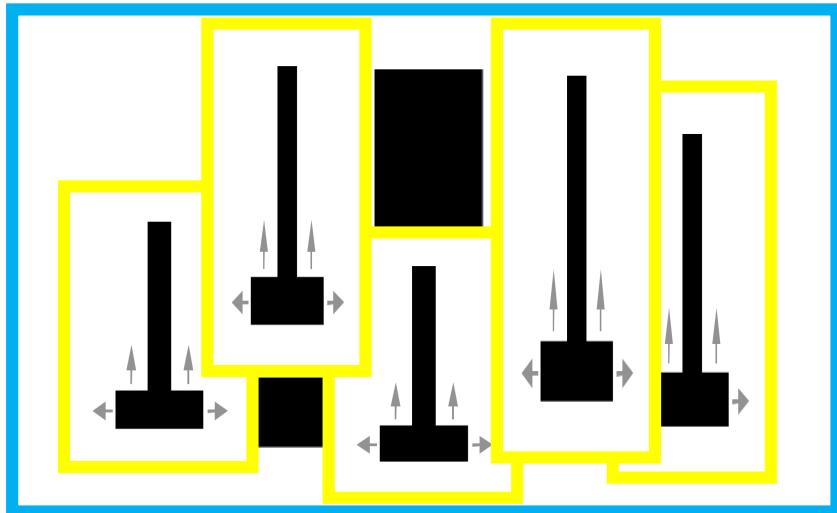
# Open-Source Software/Hardware for Converged SDI

<http://www.apnoms.org/2017/apnoms2017/ppt/tutorial4.pdf>

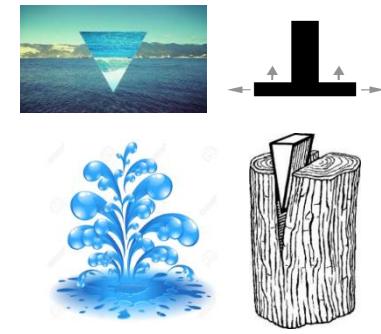
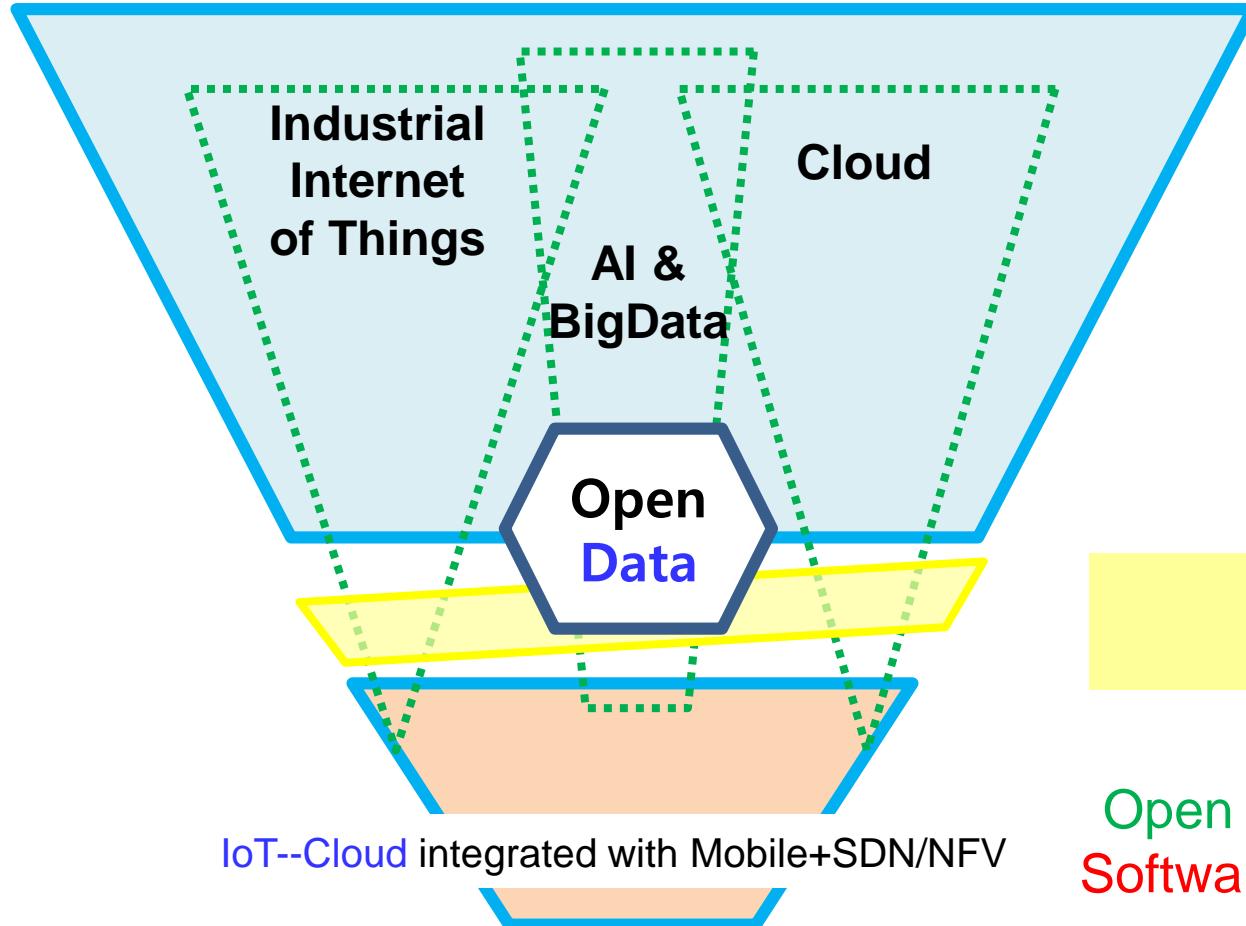


# Infrastructure vs Platform vs Services

Toward Razor-thin Platform that becomes “Composable”



# Diversified Services over Software-Defined Infrastructure



Container-leveraged  
Software as a Service

Open Federated, Composable  
Software-Defined Infrastructure  
(Resources)



Gwangju Institute of  
Science & Technology



# Thank you!

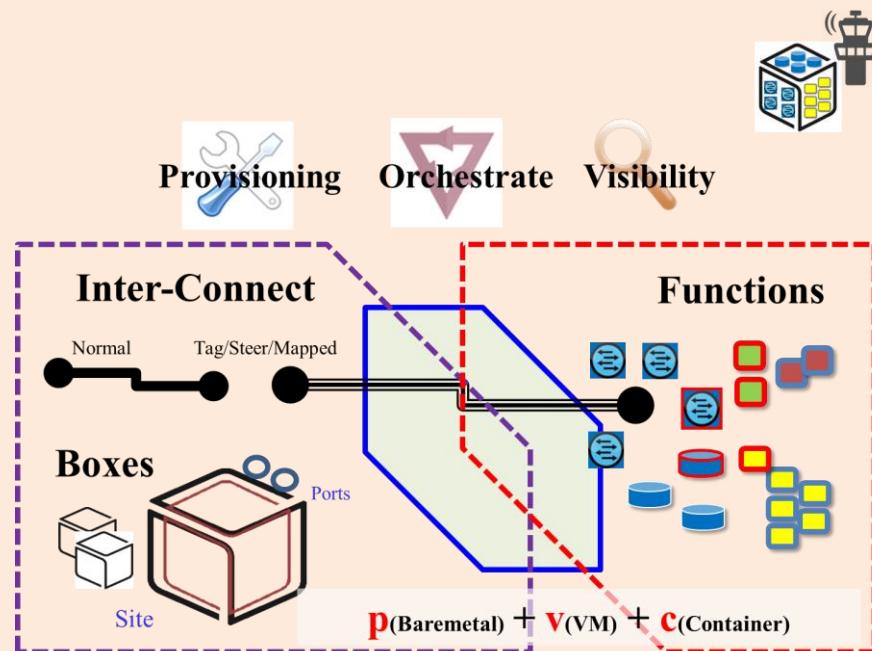
Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>

# Chapter 7:

# Computer System

# SmartX Abstraction



# Chapter Keywords

Inter-Connect

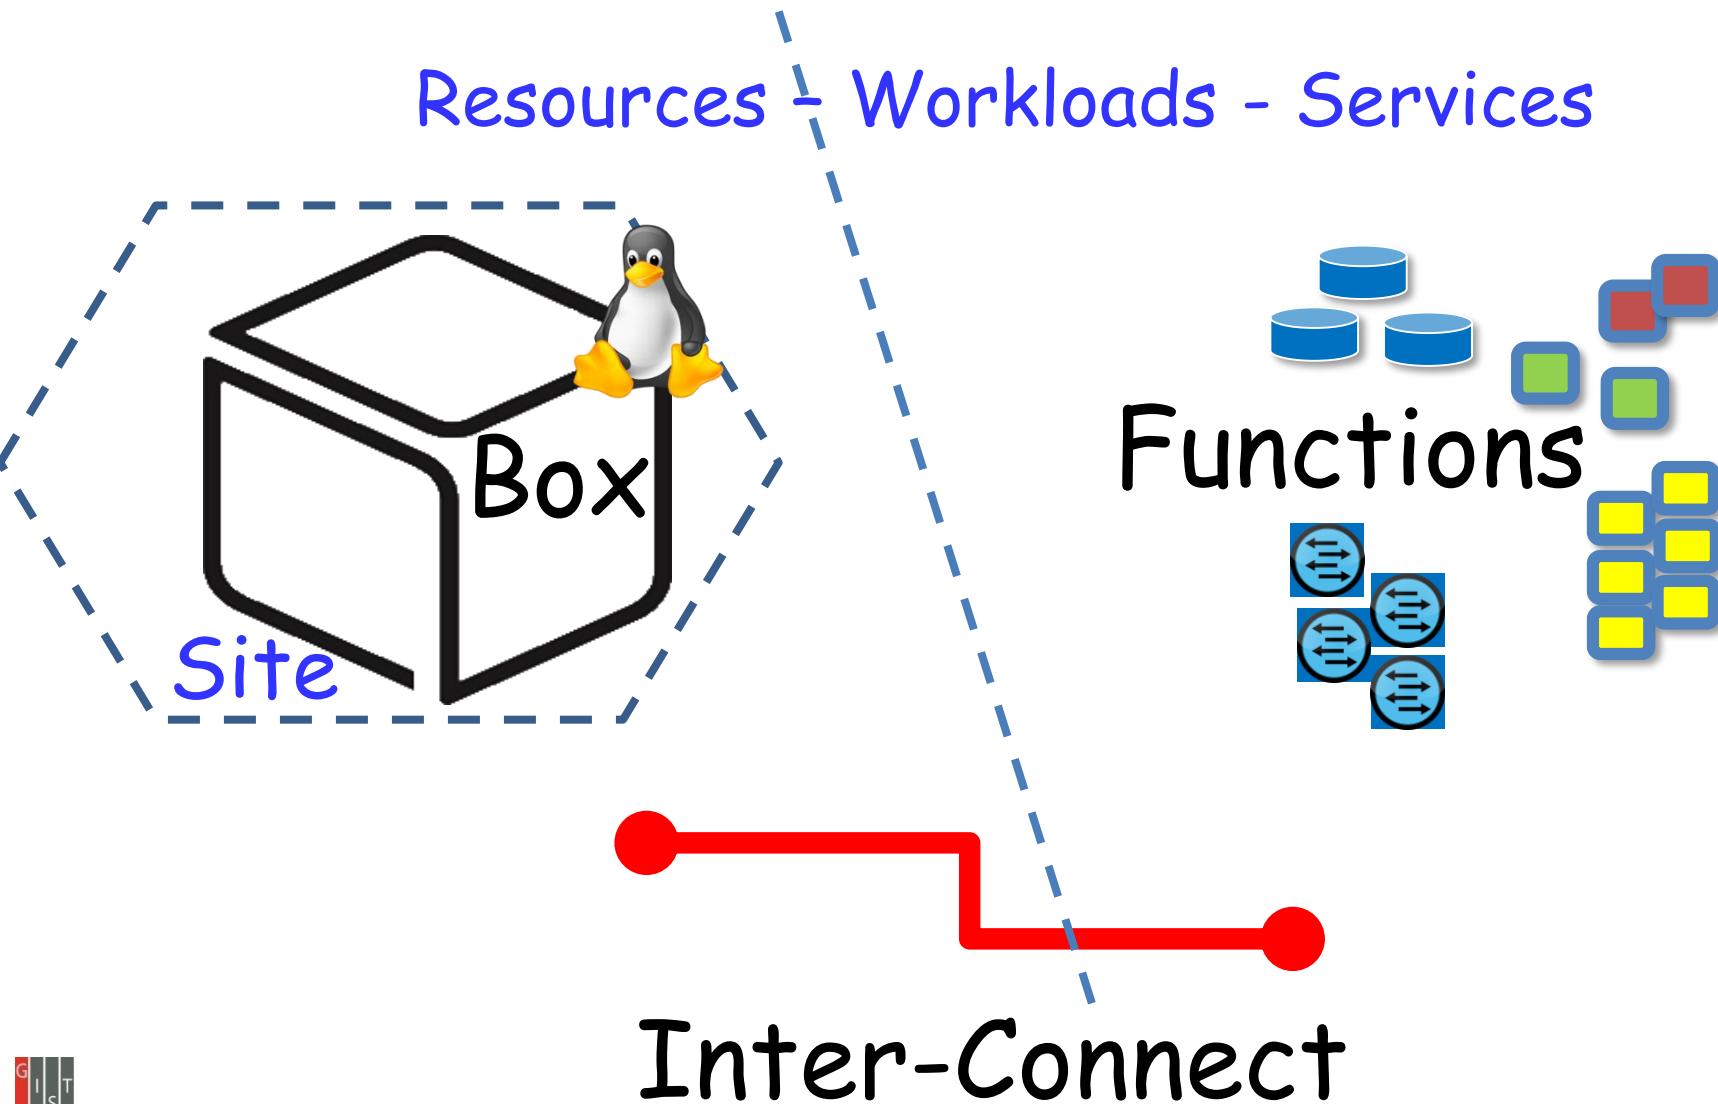
Box

Functions

SmartX Automation

SmartX Playground

# Computer System: Inter-Connected Functions inside/across Boxes/Sites



# Computer System: Inter-Connected Functions inside/across Boxes/Sites



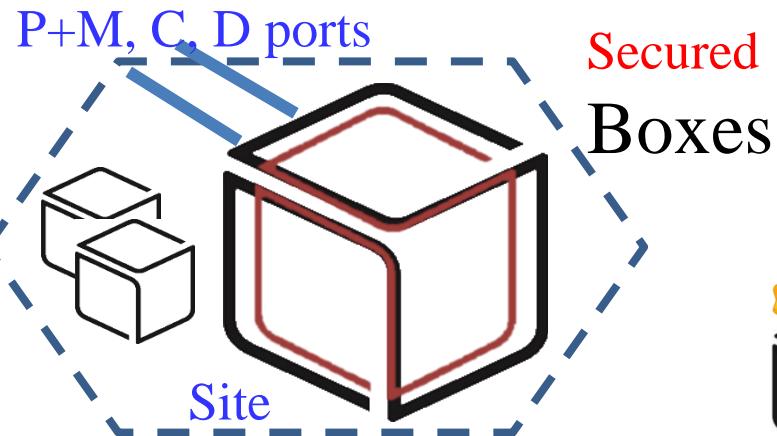
Configuration



Control



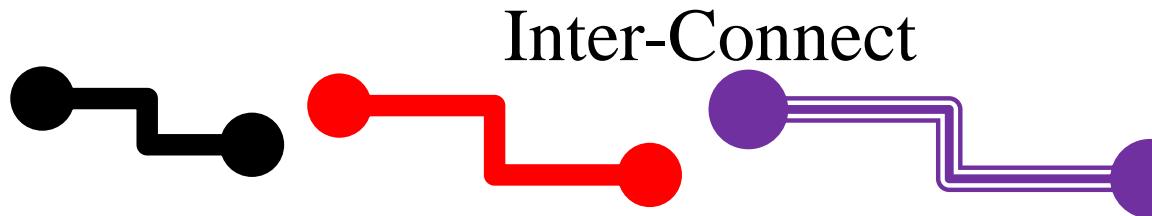
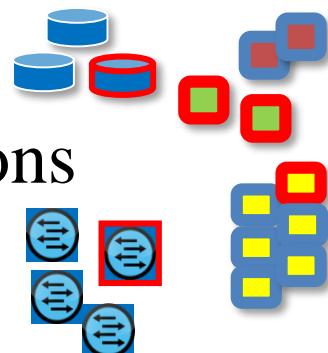
Visibility



Secured BM/Container/VM Functions



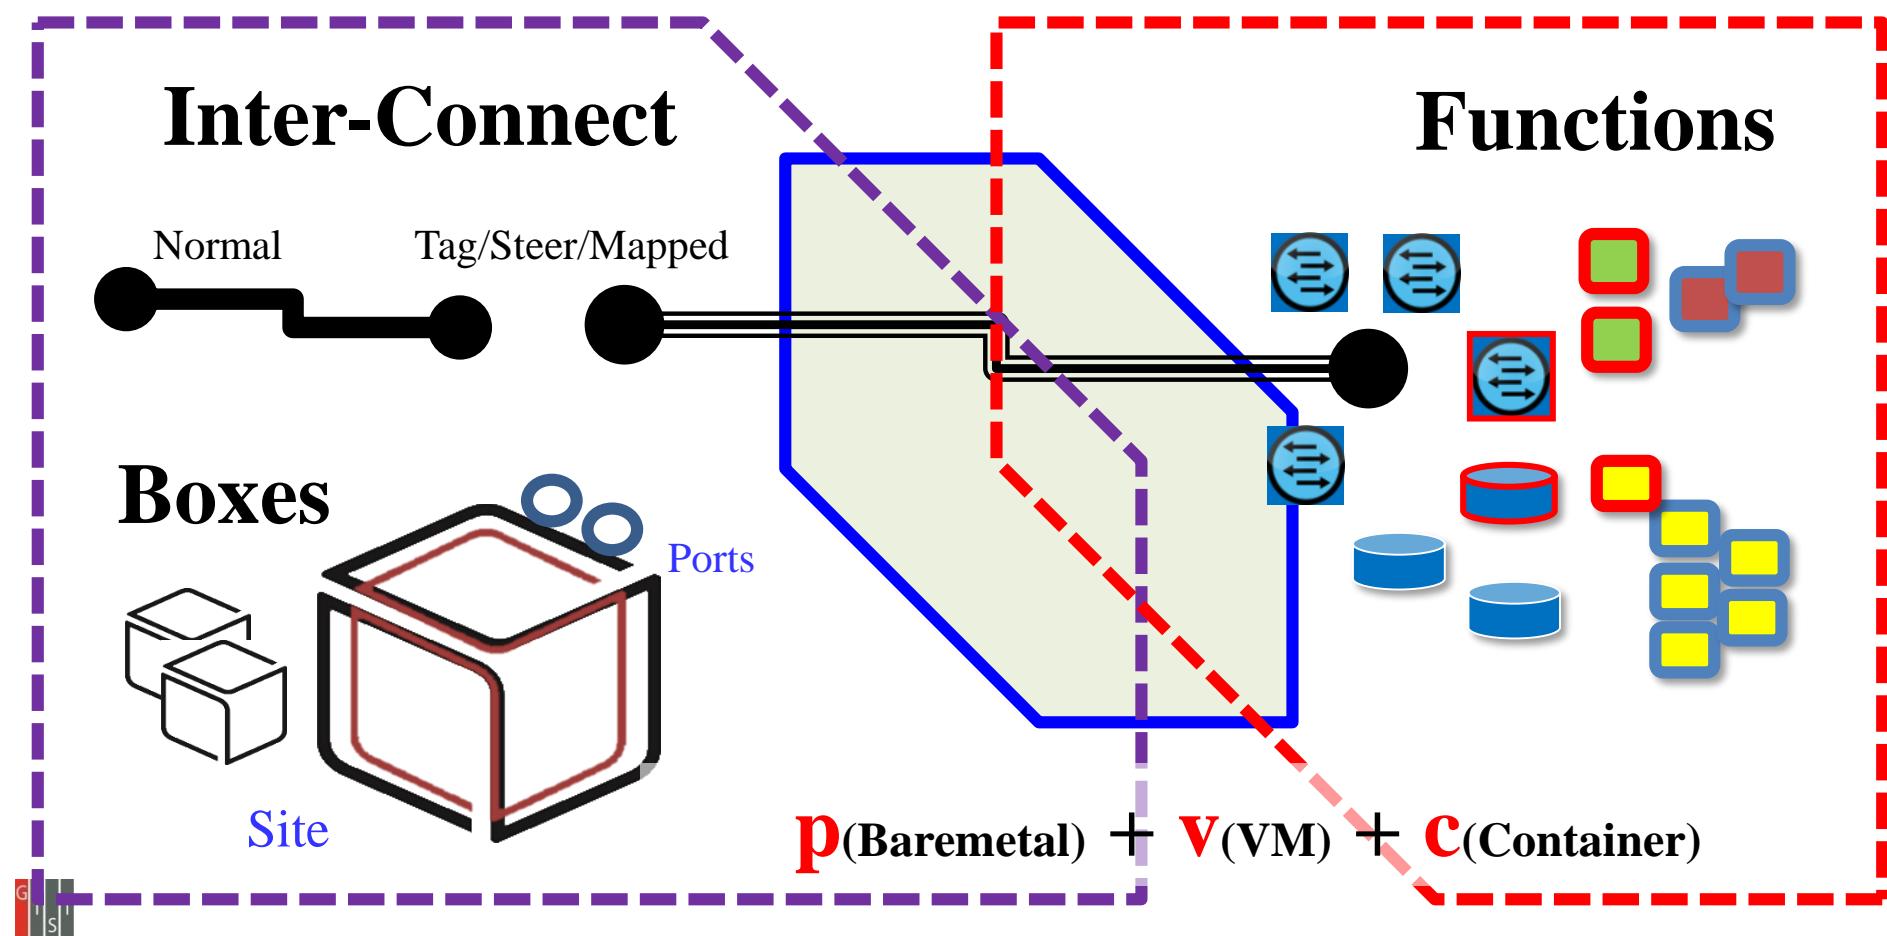
Functions



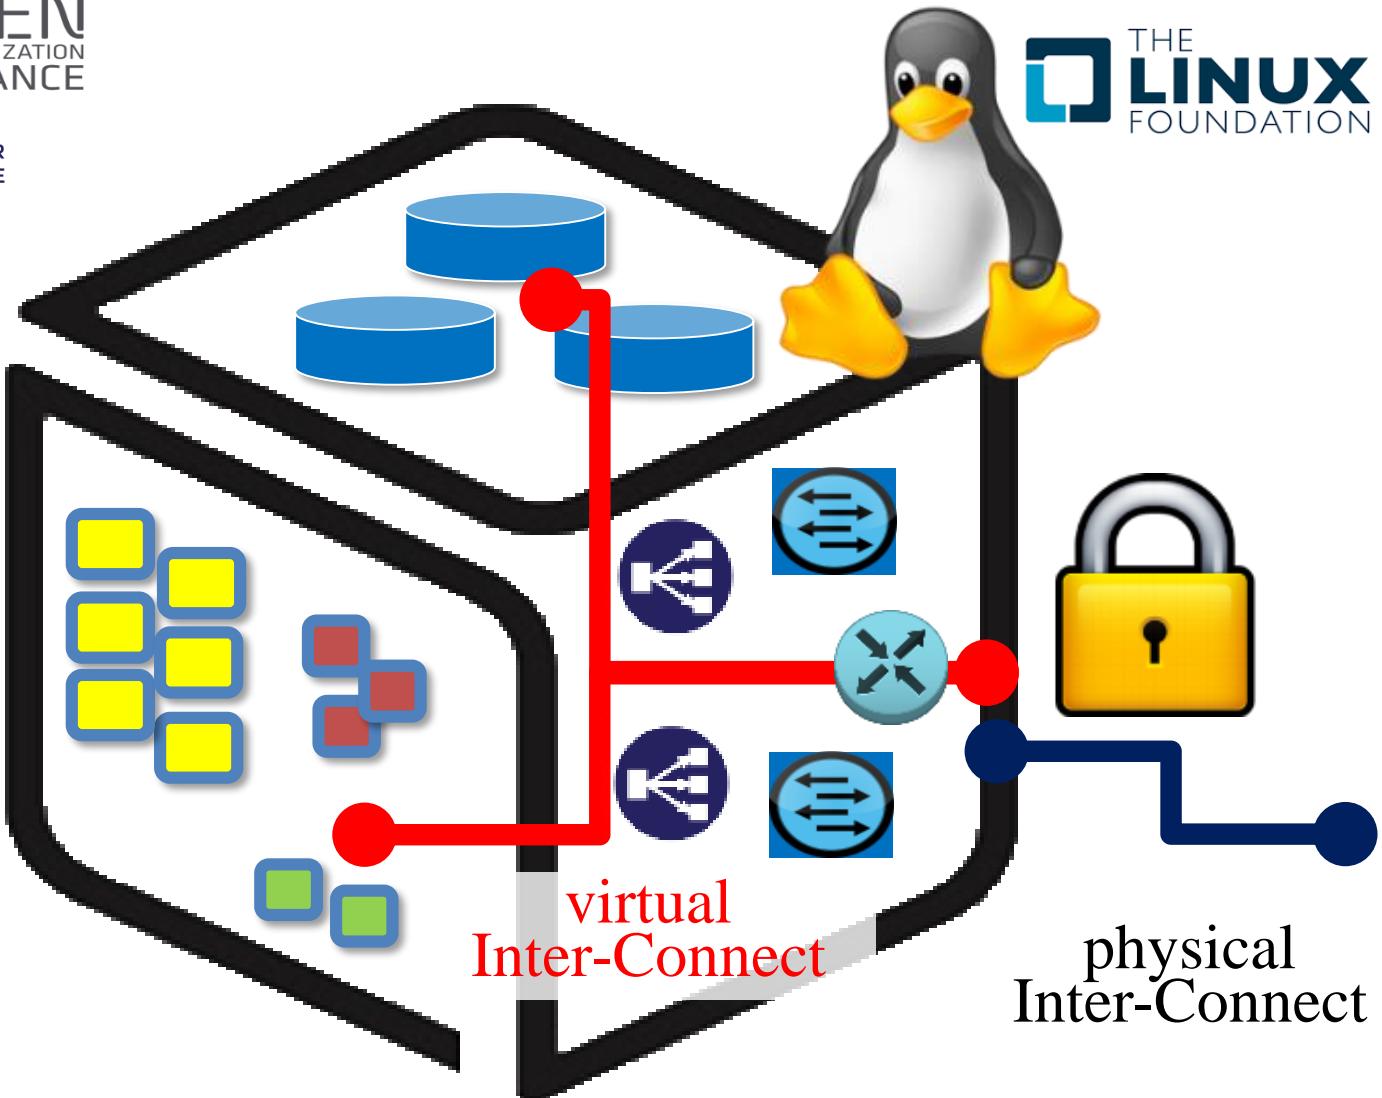
Discover & Connect

Normal → Secured → Tag/Steer/Mapped

# Computer System: Inter-Connected Functions inside/across Boxes/Sites



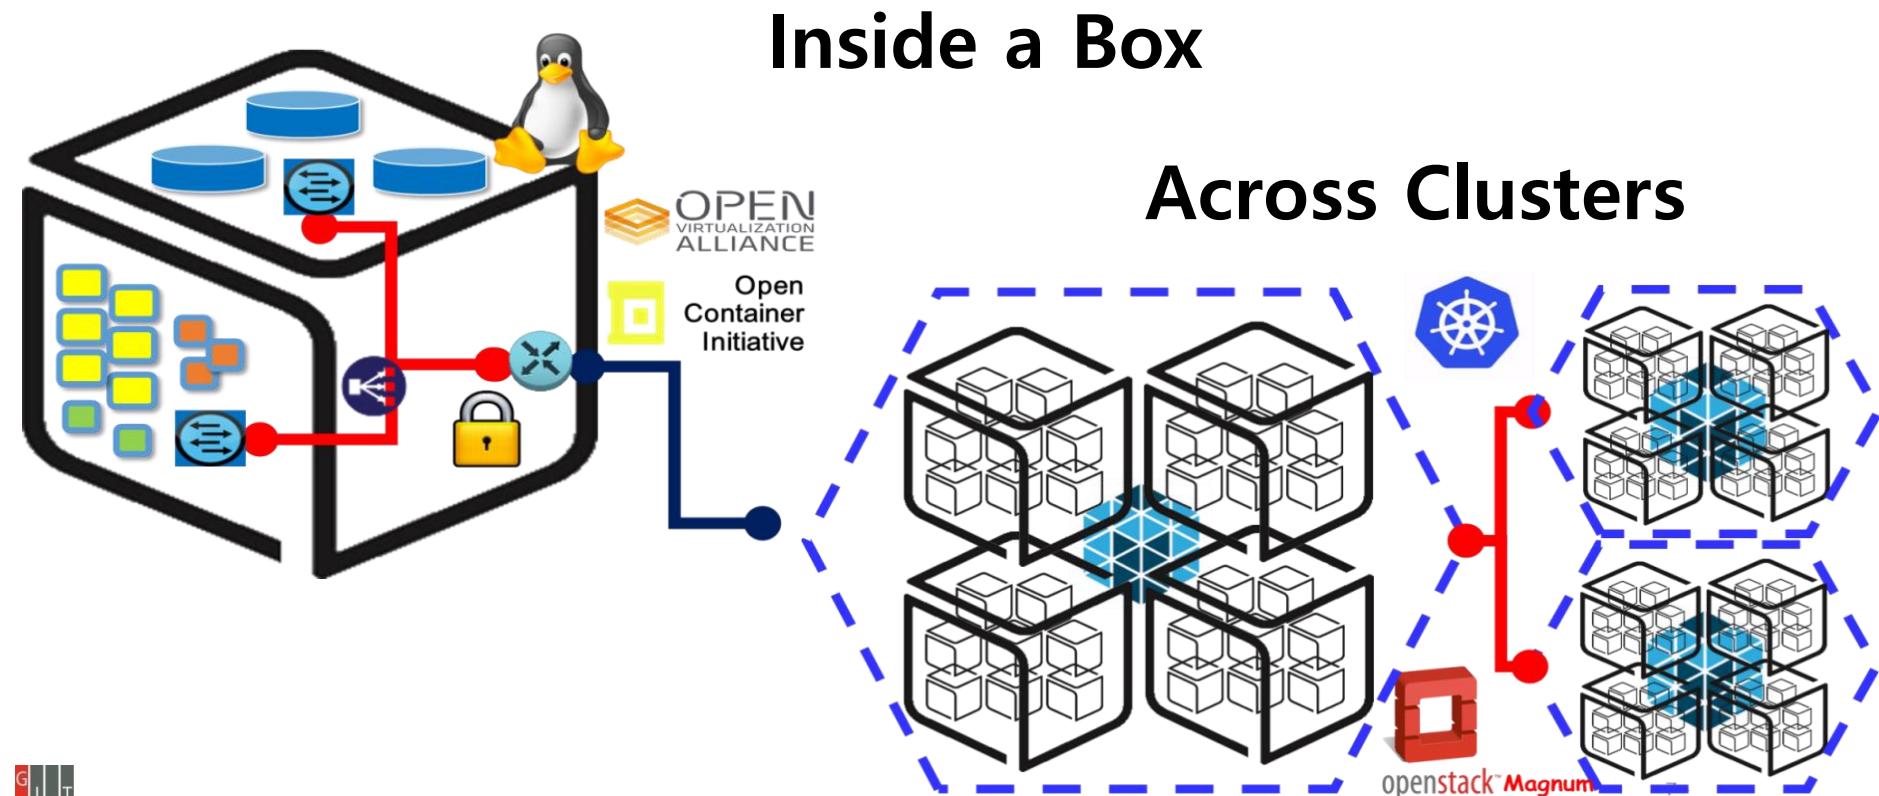
# Computer System: Inter-Connected Functions inside/across Boxes/Sites



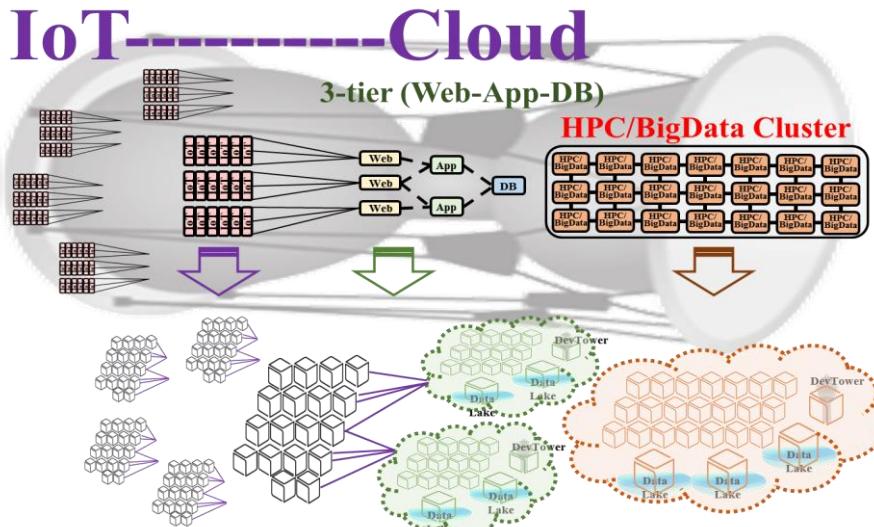


# Computer System: Inter-Connected Functions inside/across Boxes/Sites

**p+v+c Harmonization Challenge:**  
**p(Baremetal) + v(VM) + c(Container)**



# SmartX IoT-Cloud Services with Microservices-based SaaS Applications



**p(Baremetal)**  
+ **c(Container)**  
+ **v(VM)**



Process containers



application containers

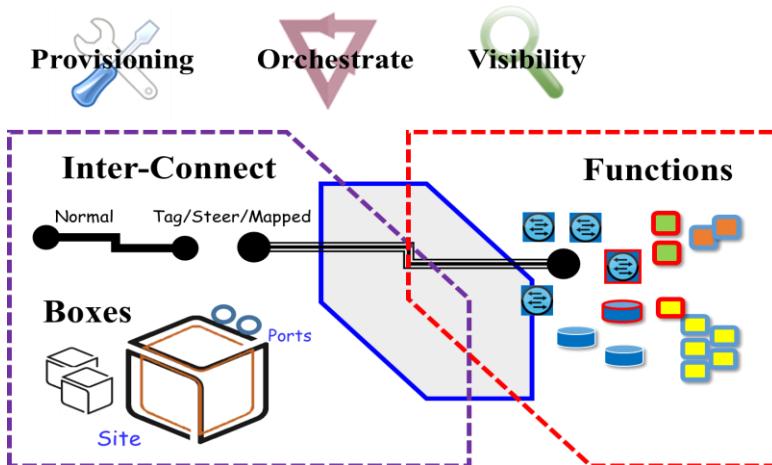


(Monster  
containers)

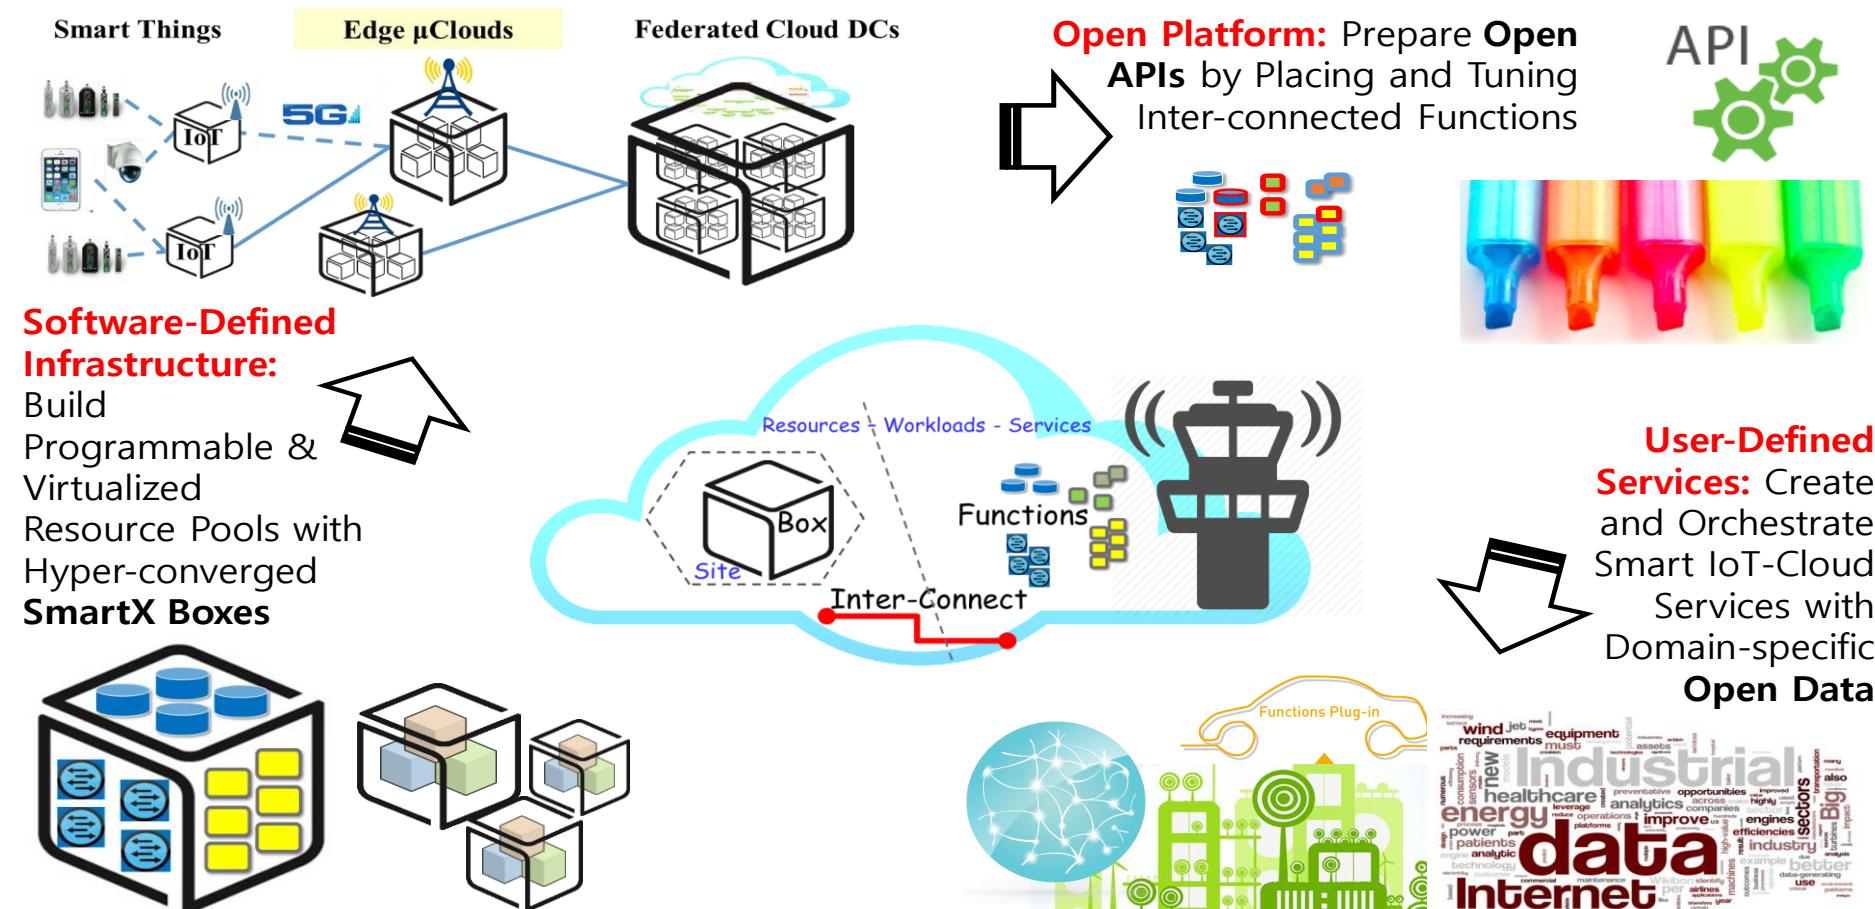
Machine containers



VM hypervisors

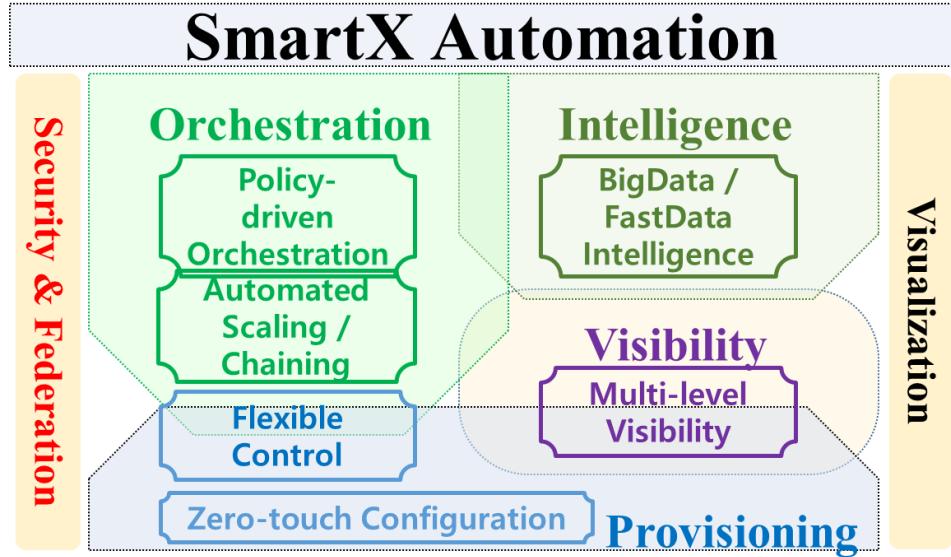
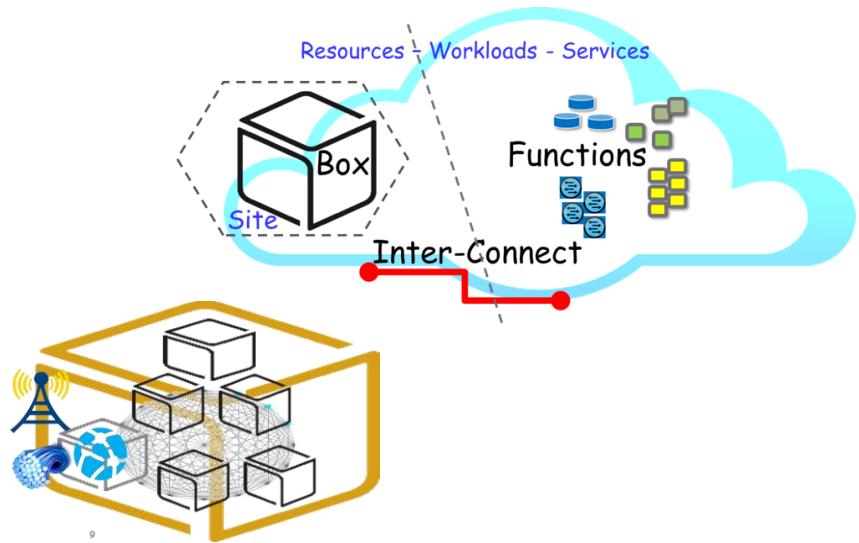
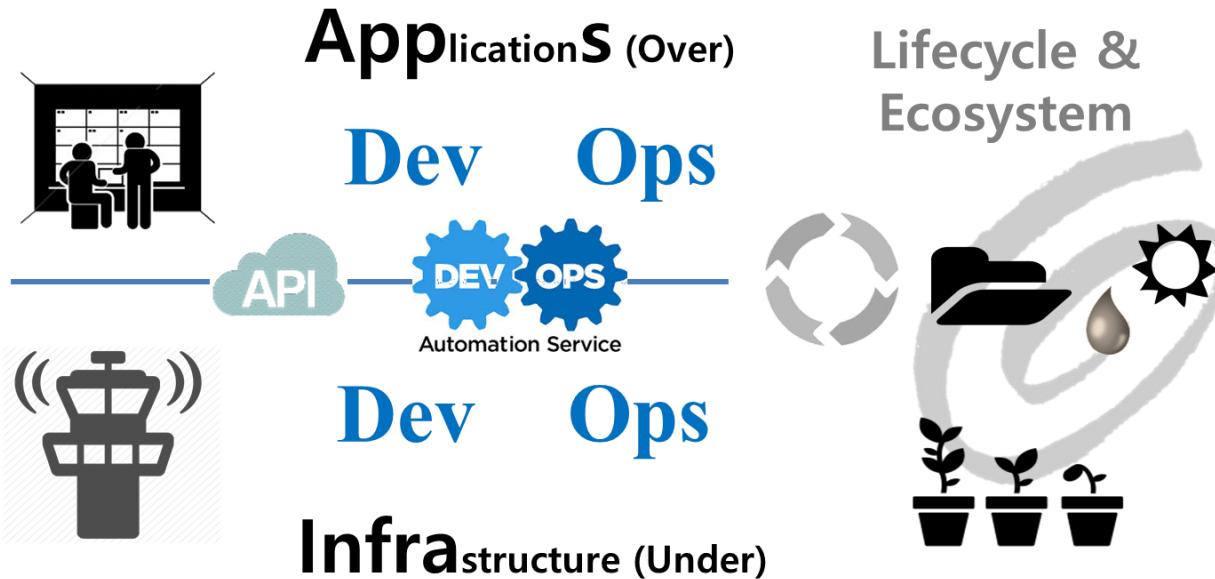


# Realize SmartX Services with Open APIs enabled by SmartX Boxes

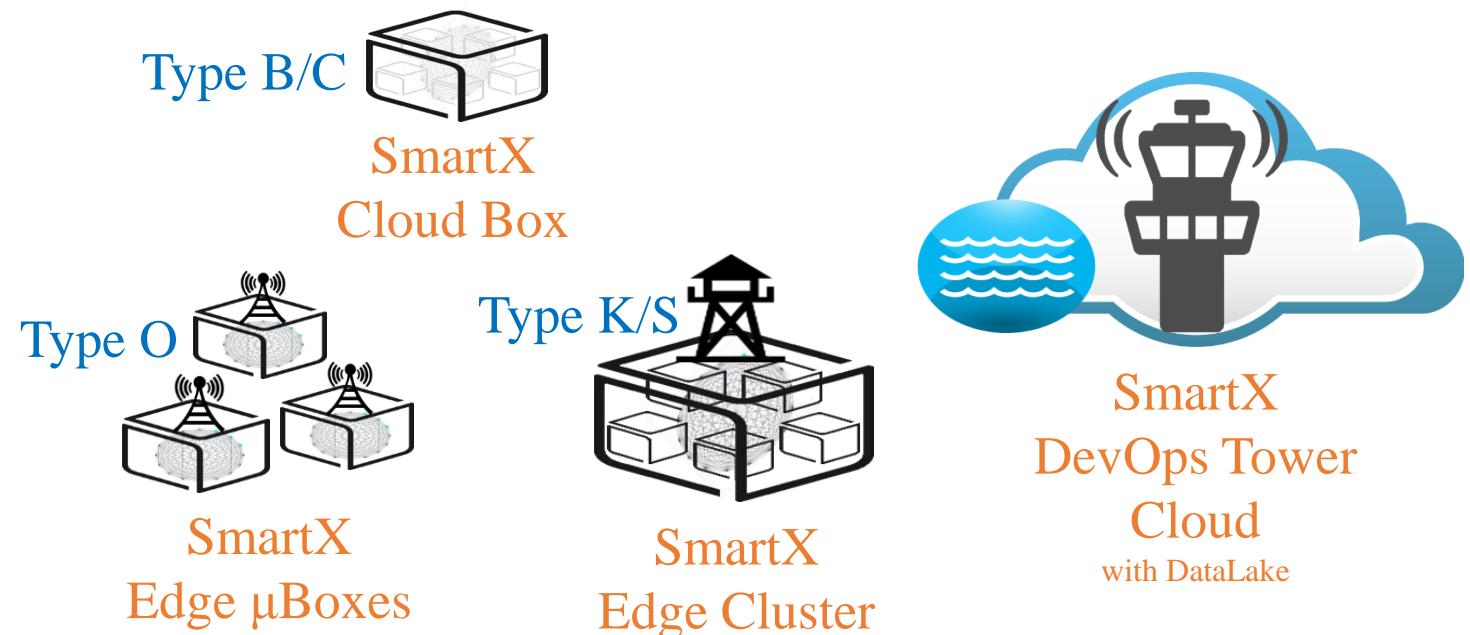
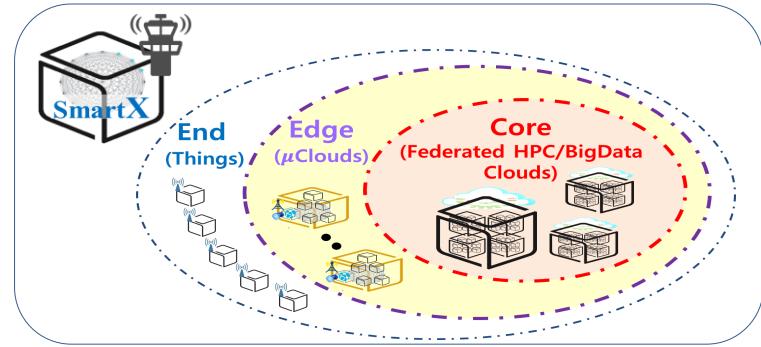




# SmartX Automation Framework



# SmartX Composable Playground & Boxes



Things

$\mu$ Clouds  
(SDN/NFV)

Clouds  
(HPC/BigData)

# Open-Source-leveraged Open Playgrounds

(IoT-SDN/NFV-Cloud (HPC/BigData) Ready)

Container-leveraged SmartX IoT-Cloud Services

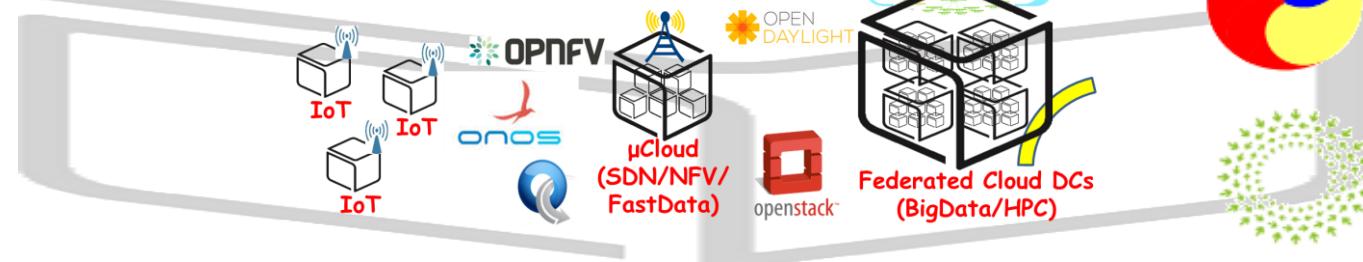


Open APIs  
SmartX Open Platform

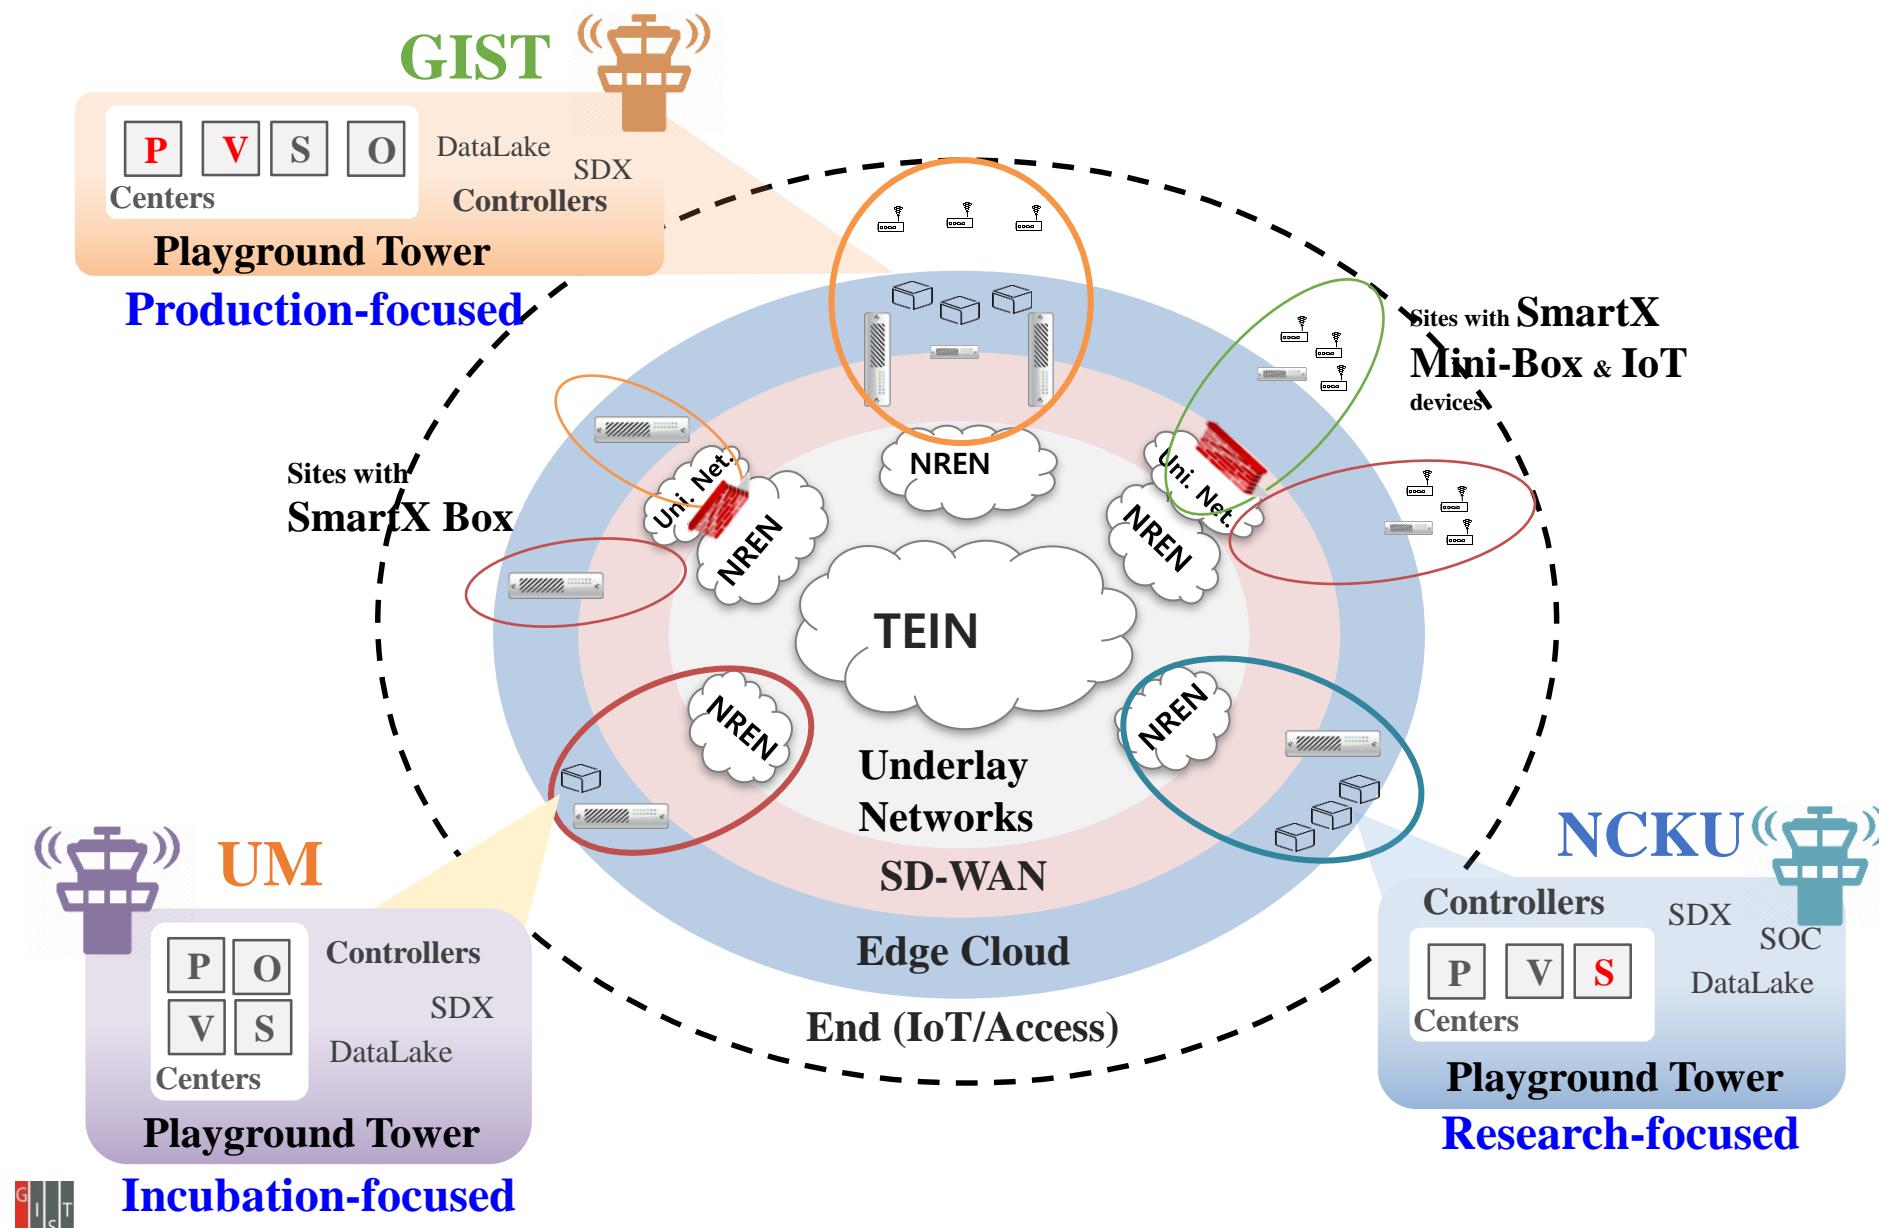
SmartX  
Playground

SmartX  
Open  
Playgrounds

DEV OPS  
Automation Service



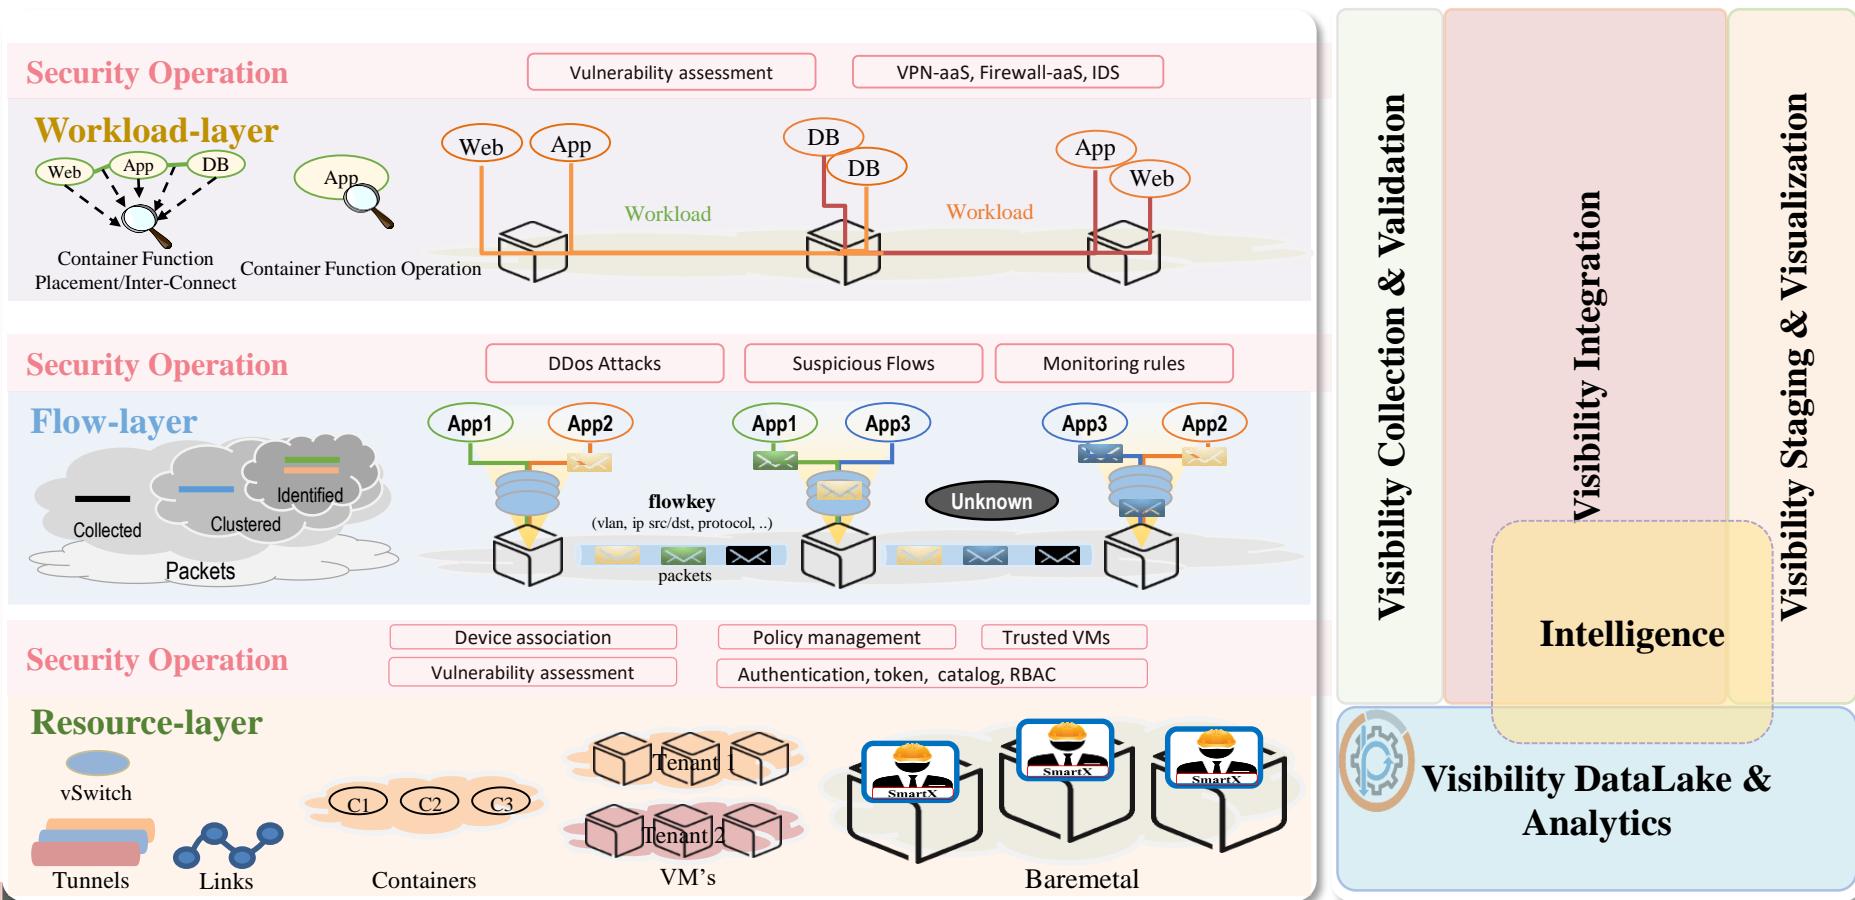
# Federated OF@TEIN+ SmartX Playground



# Visibility & Security Operation for SmartX Playgrounds

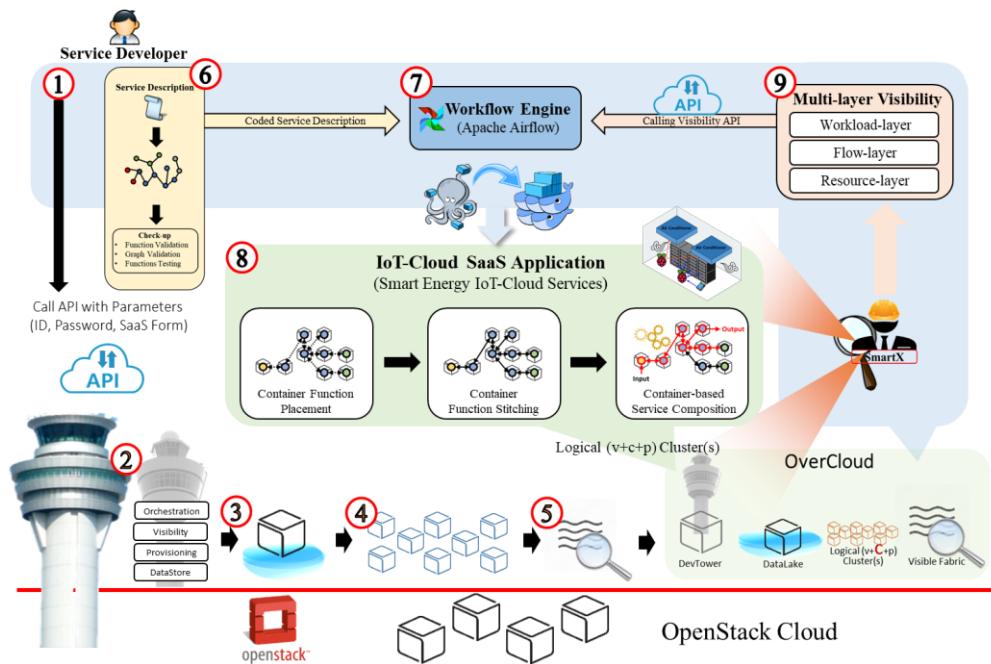
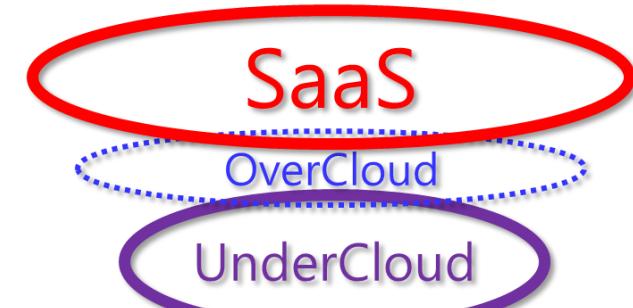
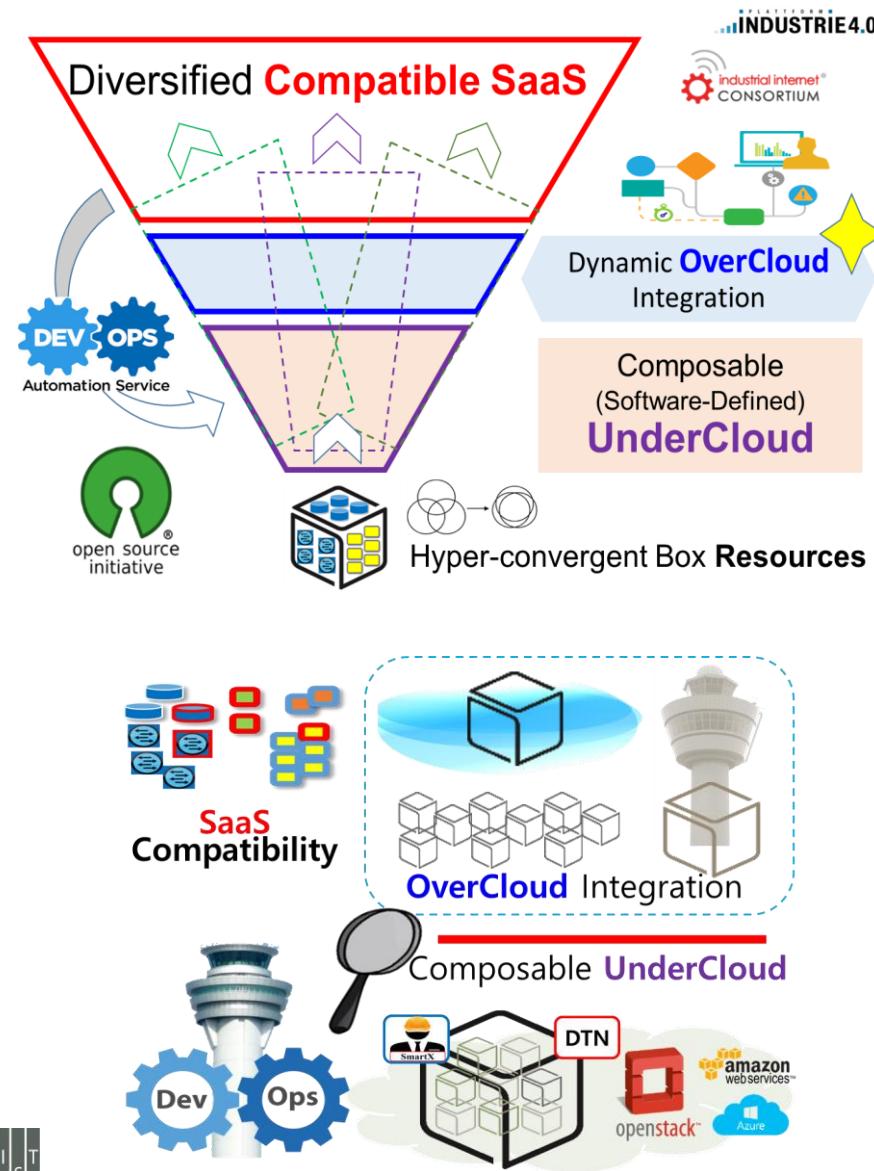
SmartX-  
MultiView

## SmartX MultiView/MultiSec Framework



# Dynamic Cloud-native OverCloud

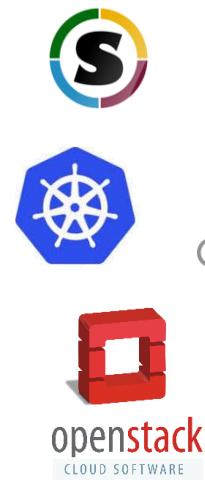
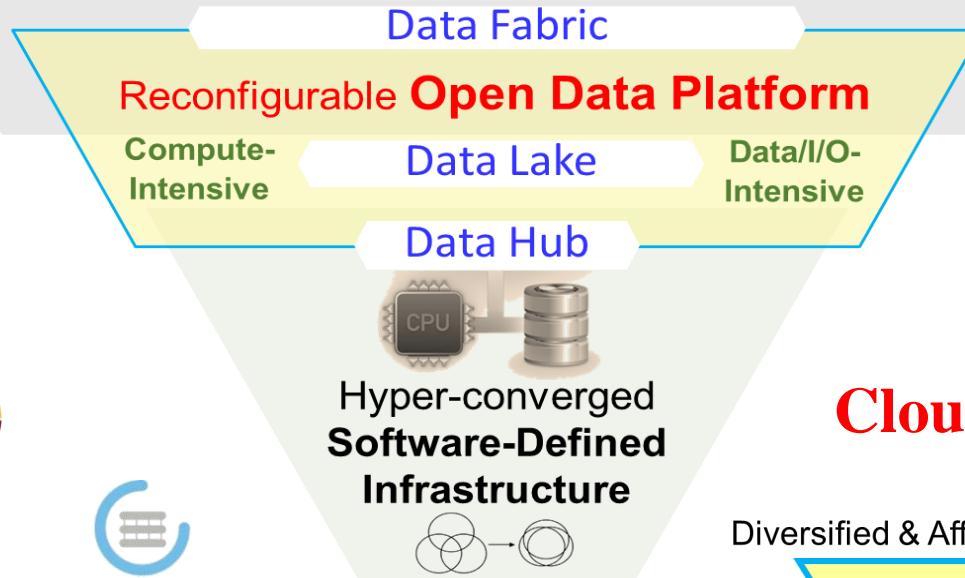
## leveraging Infrastructure Slicing over Converged SDI



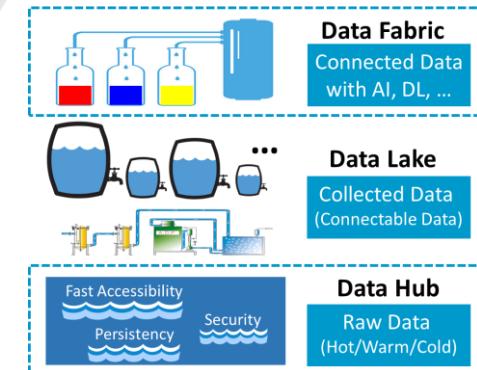
# Open/Reconfigurable Data-Centric Cloud

Cloud-leveraged HPC/AI-inspired Services over Hyper-converged SDI

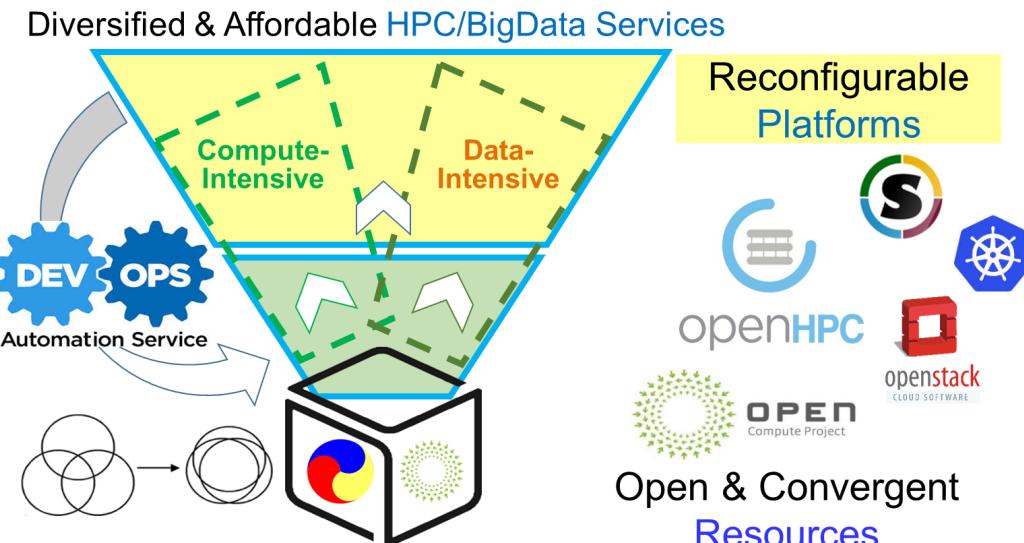
Cloud-based HPC/AI-inspired Services



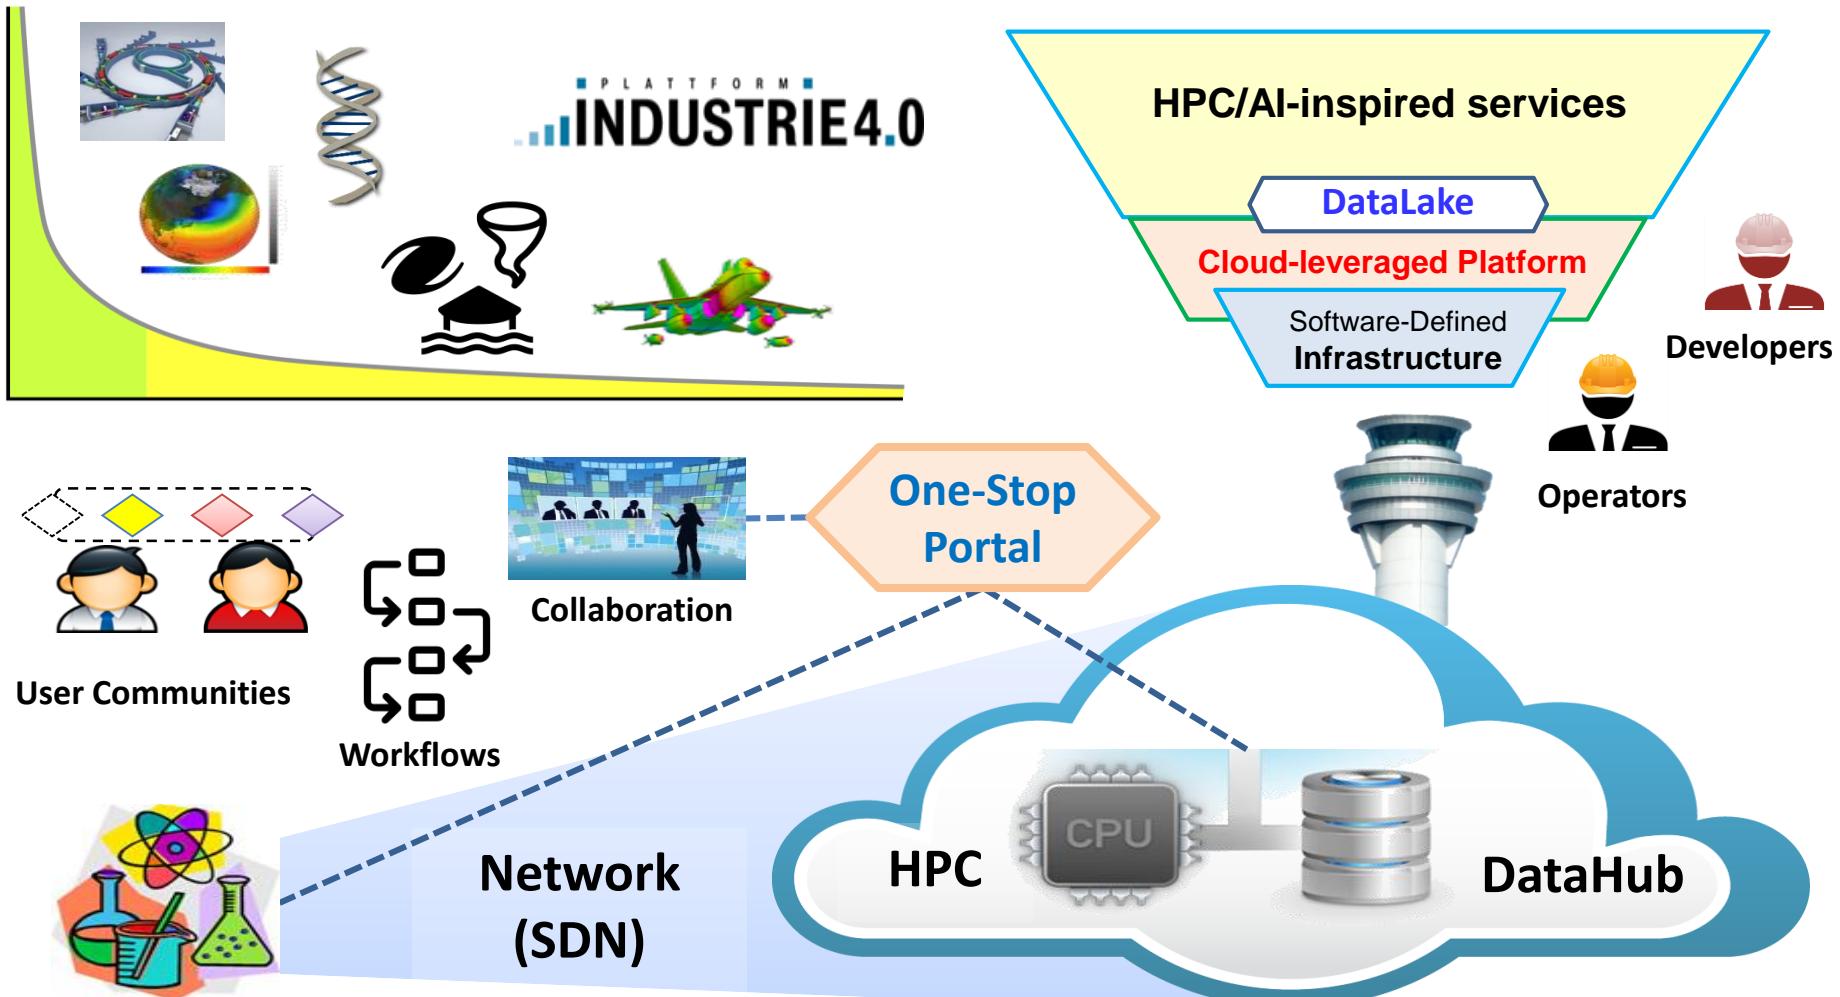
**Flexible HPC/HPDA**



**Cloud-First Integration**



# HPC/AI-inspired Services with Open S&T Ecosystem





Gwangju Institute of  
Science & Technology



# Thank you!

Send Inquiry to [jongwon@gist.ac.kr](mailto:jongwon@gist.ac.kr)

<http://nm.gist.ac.kr>